



RISC-V Model Configuration and Custom Extension Guide

Imperas Software Limited

Imperas Buildings, North Weston,
Thame, Oxfordshire, OX9 2HA, UK
docs@imperas.com



| | |
|-------------|---|
| Author: | Imperas Software Limited |
| Version: | 1.93 |
| Filename: | OVP_RISCV_Model_Custom_Extension_Guide |
| Project: | RISC-V Model Configuration and Custom Extension Guide |
| Last Saved: | Thursday, 08 August 2024 |
| Keywords: | |

Copyright Notice

Copyright © 2024 Imperas Software Limited All rights reserved. This software and documentation contain information that is the property of Imperas Software Limited. The software and documentation are furnished under a license agreement and may be used or copied only in accordance with the terms of the license agreement. No part of the software and documentation may be reproduced, transmitted, or translated, in any form or by any means, electronic, mechanical, manual, optical, or otherwise, without prior written permission of Imperas Software Limited, or as expressly provided by the license agreement.

Right to Copy Documentation

The license agreement with Imperas permits licensee to make copies of the documentation for its internal use only. Each copy shall include all copyrights, trademarks, service marks, and proprietary rights notices, if any.

Destination Control Statement

All technical data contained in this publication is subject to the export control laws of the United States of America. Disclosure to nationals of other countries contrary to United States law is prohibited. It is the reader's responsibility to determine the applicable regulations and to comply with them.

Disclaimer

IMPERAS SOFTWARE LIMITED, AND ITS LICENSORS MAKE NO WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, WITH REGARD TO THIS MATERIAL, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE.

Table of Contents

| | | |
|----------|--|-----------|
| 1 | Introduction..... | 11 |
| 2 | Building a RISC-V Model | 12 |
| 2.1 | INTRODUCTION | 12 |
| 2.2 | USING LOCAL RISC-V MODEL SOURCE DIRECTORY..... | 12 |
| 2.3 | USING VLNV LIBRARY RISC-V MODEL SOURCE..... | 12 |
| 2.4 | USING A LINKED MODEL | 13 |
| 3 | Introduction to Custom Extensions with Linked Model | 14 |
| 3.1 | LINKED MODEL CREATION | 14 |
| 3.2 | CUSTOM CONFIGURATION OPTIONS..... | 14 |
| 3.3 | ADDING CUSTOM INSTRUCTIONS..... | 14 |
| 3.4 | ADDING CUSTOM CSRS | 14 |
| 3.5 | ADDING CUSTOM EXCEPTIONS | 14 |
| 3.6 | ADDING CUSTOM LOCAL INTERRUPTS..... | 14 |
| 3.7 | ADDING CUSTOM FIFO PORTS | 14 |
| 3.8 | ADDING TRANSACTIONAL MEMORY | 14 |
| 3.9 | IMPLEMENTING PMA CONSTRAINTS..... | 14 |
| 4 | Linked Model Creation | 15 |
| 4.1 | INTRODUCTION | 15 |
| 4.2 | CREATING A NEW MODEL LIBRARY | 16 |
| 4.2.1 | <i>Creating a New Model - No Custom Extensions</i> | <i>16</i> |
| 4.2.2 | <i>Creating a New Model – With Custom Extensions.....</i> | <i>18</i> |
| 4.3 | BUILDING THE LINKED MODEL | 20 |
| 4.4 | EXECUTING THE LINKED MODEL | 21 |
| 5 | Custom Configuration Options..... | 22 |
| 5.1 | FUNDAMENTAL CONFIGURATION | 28 |
| 5.2 | INTERRUPTS AND EXCEPTIONS..... | 30 |
| 5.2.1 | <i>Reset</i> | <i>31</i> |
| 5.3 | TIMERS AND TIMER INTERRUPTS | 33 |
| 5.4 | INSTRUCTION AND CSR BEHAVIOR | 34 |
| 5.5 | CLIC | 36 |
| 5.5.1 | <i>CLIC Fundamental Fields.....</i> | <i>36</i> |
| 5.5.2 | <i>CLIC Common Fields.....</i> | <i>37</i> |
| 5.5.3 | <i>CLIC Internal Fields</i> | <i>38</i> |
| 5.5.4 | <i>Custom registers in the CLIC MMIO custom region.....</i> | <i>38</i> |
| 5.6 | CLINT..... | 40 |
| 5.7 | AIA | 41 |
| 5.7.1 | <i>AIA Net Port Interface.....</i> | <i>42</i> |
| 5.7.2 | <i>IMSIC Interrupt Controller Integration.....</i> | <i>42</i> |
| 5.8 | MEMORY SUBSYSTEM..... | 43 |
| 5.8.1 | <i>Memory Access Constraints</i> | <i>44</i> |
| 5.9 | PMP CONFIGURATION | 46 |
| 5.9.1 | <i>Mask Values Available in Integration Registers.....</i> | <i>47</i> |
| 5.10 | FLOATING POINT..... | 48 |
| 5.10.1 | <i>misae CSR D and F Bits.....</i> | <i>49</i> |
| 5.10.2 | <i>mstatus.FS Update Modes.....</i> | <i>49</i> |
| 5.10.3 | <i>Half-Precision Support.....</i> | <i>49</i> |
| 5.10.4 | <i>Zfinx Support</i> | <i>50</i> |
| 5.10.5 | <i>Unimplemented Instructions.....</i> | <i>50</i> |
| 5.10.6 | <i>Flushing Subnormal Values to Zero</i> | <i>50</i> |
| 5.11 | VECTOR EXTENSION | 52 |

| | | |
|-----------|---|------------|
| 5.11.1 | <i>Half-Precision Support</i> | 53 |
| 5.11.2 | <i>Unimplemented Instructions</i> | 53 |
| 5.12 | BIT MANIPULATION EXTENSION | 54 |
| 5.13 | CRYPTOGRAPHIC EXTENSION | 55 |
| 5.14 | HYPERVISOR EXTENSION..... | 57 |
| 5.15 | DSP EXTENSION | 57 |
| 5.16 | CODE SIZE REDUCTION EXTENSION..... | 58 |
| 5.17 | DEBUG MODE | 60 |
| 5.17.1 | <i>Debug Mode Behaviors</i> | 61 |
| 5.17.2 | <i>Debug State Entry</i> | 62 |
| 5.17.3 | <i>Debug State Exit</i> | 62 |
| 5.17.4 | <i>haltreq Signal and debug_priority Field</i> | 62 |
| 5.18 | TRIGGER MODULE | 64 |
| 5.19 | MULTICORE VARIANTS | 66 |
| 5.20 | CSR INDEX NUMBERS | 67 |
| 5.21 | CSR INITIAL VALUES | 68 |
| 5.22 | CSR MASKS..... | 69 |
| 5.23 | UNIMPLEMENTED INSTRUCTIONS..... | 70 |
| 6 | Adding Custom Instructions (addInstructions)..... | 71 |
| 6.1 | INTERCEPT ATTRIBUTES | 71 |
| 6.2 | OBJECT TYPE AND CONSTRUCTOR..... | 72 |
| 6.3 | INSTRUCTION DECODE..... | 72 |
| 6.4 | INSTRUCTION DISASSEMBLY | 75 |
| 6.5 | INSTRUCTION TRANSLATION..... | 76 |
| 6.5.1 | <i>Required VMI Morph-Time Function Knowledge</i> | 80 |
| 6.6 | EXAMPLE EXECUTION..... | 81 |
| 7 | Adding Custom CSRs (addCSRs)..... | 83 |
| 7.1 | CSR TYPE DEFINITIONS..... | 83 |
| 7.2 | EXTENSION-SPECIFIC CONFIGURATION | 84 |
| 7.3 | INTERCEPT ATTRIBUTES | 85 |
| 7.4 | OBJECT TYPE AND CONSTRUCTOR..... | 85 |
| 7.5 | EXAMPLE EXECUTION..... | 95 |
| 7.6 | GENERAL CSR RESET TEMPLATE CODE | 98 |
| 8 | Adding Custom Exceptions (addExceptions) | 100 |
| 8.1 | EXCEPTION CODE | 100 |
| 8.2 | INTERCEPT ATTRIBUTES | 100 |
| 8.3 | OBJECT TYPE AND CONSTRUCTOR..... | 100 |
| 8.4 | INSTRUCTION DECODE..... | 102 |
| 8.5 | INSTRUCTION DISASSEMBLY | 103 |
| 8.6 | INSTRUCTION TRANSLATION..... | 103 |
| 8.7 | EXAMPLE EXECUTION..... | 103 |
| 9 | Adding Custom Local Interrupts (addLocalInterrupts)..... | 107 |
| 9.1 | ENABLING LOCAL INTERRUPT PORTS | 107 |
| 9.2 | INTERRUPT CODES | 107 |
| 9.3 | INTERCEPT ATTRIBUTES | 108 |
| 9.4 | OBJECT TYPE AND CONSTRUCTOR..... | 108 |
| 9.5 | EXAMPLE EXECUTION..... | 111 |
| 10 | Adding Custom FIFOs (fifoExtensions) | 113 |
| 10.1 | INTERCEPT ATTRIBUTES | 113 |
| 10.2 | OBJECT TYPE AND CONSTRUCTOR..... | 114 |

| | | |
|-----------|---|------------|
| 10.3 | EXTENSION CSRS | 115 |
| 10.4 | EXTENSION FIFO PORTS..... | 118 |
| 10.5 | INSTRUCTION DECODE..... | 119 |
| 10.6 | INSTRUCTION DISASSEMBLY..... | 120 |
| 10.7 | INSTRUCTION TRANSLATION..... | 120 |
| 11 | Adding Transactional Memory (tmExtensions) | 123 |
| 11.1 | INTERCEPT ATTRIBUTES | 123 |
| 11.2 | INTERCEPT PARAMETERS..... | 124 |
| 11.3 | OBJECT TYPE, CONSTRUCTOR AND POST-CONSTRUCTOR..... | 125 |
| 11.4 | EXTENSION REGISTERS..... | 127 |
| 11.5 | EXTENSION CSRS | 128 |
| 11.6 | CONTEXT SWITCH MONITOR (RISCVSwitch)..... | 129 |
| 11.7 | TRANSACTIONAL LOAD AND STORE FUNCTIONS (RISCVTLload AND RISCVTLStore) | 130 |
| 11.8 | TRAP AND EXCEPTION RETURN NOTIFIERS (RISCVTrapNotifier AND RISCVRetNotifier)... | 131 |
| 11.9 | INSTRUCTION DECODE..... | 132 |
| 11.10 | INSTRUCTION DISASSEMBLY..... | 133 |
| 11.11 | INSTRUCTION TRANSLATION..... | 133 |
| 11.12 | MEMORY MODEL IMPLEMENTATION GUIDELINES | 137 |
| 12 | Implementing Custom PMA Behavior..... | 138 |
| 12.1 | MEMORY DOMAIN HIERARCHY | 138 |
| 12.2 | STATIC PMA REGION DEFINITIONS | 139 |
| 12.3 | PMA REGION DEFINITIONS USING MODEL COMMANDS | 141 |
| 12.4 | PMA CONTROL BY EXTENSION MODELS..... | 141 |
| 12.4.1 | Example PMA Control Implementation..... | 142 |
| 12.5 | PMA MAPPED PAGE SIZE RESTRICTION | 144 |
| 13 | Appendix: Standard Instruction Patterns | 145 |
| 13.1 | PATTERN RVIP_RD_RS1_RS2 (R-TYPE) | 145 |
| 13.2 | PATTERN RVIP_RD_RS1_SI (I-TYPE)..... | 145 |
| 13.3 | PATTERN RVIP_RD_RS1_SHIFT (I-TYPE - 5 OR 6 BIT SHIFT)..... | 145 |
| 13.4 | PATTERN RVIP_BASE_RS2_OFFSET (S-TYPE) | 145 |
| 13.5 | PATTERN RVIP_RS1_RS2_OFFSET (B-TYPE)..... | 145 |
| 13.6 | PATTERN RVIP_RD_SI (U-TYPE) | 146 |
| 13.7 | PATTERN RVIP_RD_OFFSET (J-TYPE)..... | 146 |
| 13.8 | PATTERN RVIP_RD_RS1_RS2_RS3 (R4-TYPE)..... | 146 |
| 13.9 | PATTERN RVIP_RD_RS1_RS3_SHIFT (NON-STANDARD) | 146 |
| 13.10 | PATTERN RVIP_FD_FS1_FS2 | 146 |
| 13.11 | PATTERN RVIP_FD_FS1_FS2_RM..... | 146 |
| 13.12 | PATTERN RVIP_FD_FS1_FS2_FS3_RM..... | 147 |
| 13.13 | PATTERN RVIP_RD_FS1_FS2 | 147 |
| 13.14 | PATTERN RVIP_VD_VS1_VS2_M..... | 147 |
| 13.15 | PATTERN RVIP_VD_VS1_SI_M..... | 147 |
| 13.16 | PATTERN RVIP_VD_VS1_UI_M..... | 147 |
| 13.17 | PATTERN RVIP_VD_VS1_RS2_M..... | 147 |
| 13.18 | PATTERN RVIP_VD_VS1_FS2_M..... | 148 |
| 13.19 | PATTERN RVIP_RD_VS1_RS2 | 148 |
| 13.20 | PATTERN RVIP_RD_VS1_M..... | 148 |
| 13.21 | PATTERN RVIP_VD_RS2..... | 148 |
| 13.22 | PATTERN RVIP_FD_VS1..... | 148 |
| 13.23 | PATTERN RVIP_VD_FS2..... | 148 |
| 14 | Appendix: Base Model Interface Service Functions | 149 |

| | | |
|-------|-----------------------------------|-----|
| 14.1 | FUNCTION REGISTEREXTCB | 151 |
| 14.2 | FUNCTION GETEXTCLIENTDATA | 152 |
| 14.3 | FUNCTION GETEXTCONFIG | 153 |
| 14.4 | FUNCTION GETXLENMODE | 154 |
| 14.5 | FUNCTION GETXLENARCH | 155 |
| 14.6 | FUNCTION GETXREGNAME | 156 |
| 14.7 | FUNCTION GETFREGNAME | 157 |
| 14.8 | FUNCTION GETVREGNAME | 158 |
| 14.9 | FUNCTION SETTMODE | 159 |
| 14.10 | FUNCTION GETTMODE | 160 |
| 14.11 | FUNCTION GETDATAENDIAN..... | 161 |
| 14.12 | FUNCTION READCSR | 162 |
| 14.13 | FUNCTION WRITECSR..... | 163 |
| 14.14 | FUNCTION READBASECSR | 164 |
| 14.15 | FUNCTION WRITEBASECSR | 165 |
| 14.16 | FUNCTION HALT..... | 166 |
| 14.17 | FUNCTION BLOCK..... | 168 |
| 14.18 | FUNCTION RESTART | 170 |
| 14.19 | FUNCTION UPDATE INTERRUPT | 171 |
| 14.20 | FUNCTION UPDATEDISABLE | 172 |
| 14.21 | FUNCTION UPDATEDISABLENMI | 173 |
| 14.22 | FUNCTION TEST INTERRUPT | 174 |
| 14.23 | FUNCTION RESUMEFROMWFI..... | 175 |
| 14.24 | FUNCTION ILLEGALINSTRUCTION | 176 |
| 14.25 | FUNCTION ILLEGALVERBOSE | 177 |
| 14.26 | FUNCTION VIRTUALINSTRUCTION | 178 |
| 14.27 | FUNCTION VIRTUALVERBOSE | 179 |
| 14.28 | FUNCTION ILLEGALCUSTOM | 180 |
| 14.29 | FUNCTION TAKEEXCEPTION | 181 |
| 14.30 | FUNCTION PENDFETCHEXCEPTION | 182 |
| 14.31 | FUNCTION TAKERESET..... | 183 |
| 14.32 | FUNCTION ACKNOWLEDGECLICINT..... | 184 |
| 14.33 | FUNCTION FETCHINSTRUCTION | 185 |
| 14.34 | FUNCTION DISASSINSTRUCTION | 186 |
| 14.35 | FUNCTION INSTRUCTIONENABLED | 187 |
| 14.36 | FUNCTION MORPHEXTERNAL | 188 |
| 14.37 | FUNCTION MORPHILLEGAL | 189 |
| 14.38 | FUNCTION MORPHVIRTUAL | 190 |
| 14.39 | FUNCTION GETVMIREG | 191 |
| 14.40 | FUNCTION GETVMIREGFS | 192 |
| 14.41 | FUNCTION WRITEREGSIZE | 193 |
| 14.42 | FUNCTION WRITEREG..... | 194 |
| 14.43 | FUNCTION GETFPFLAGSMT | 195 |
| 14.44 | FUNCTION GETDATAENDIANMT | 196 |
| 14.45 | FUNCTION LOADMT..... | 197 |
| 14.46 | FUNCTION STOREMT | 199 |
| 14.47 | FUNCTION CHECKLOADMT | 201 |
| 14.48 | FUNCTION CHECKSTOREMT..... | 203 |
| 14.49 | FUNCTION REQUIREMODEMT..... | 205 |
| 14.50 | FUNCTION REQUIRENOTVMt | 206 |
| 14.51 | FUNCTION CHECKLEGALRMMt..... | 207 |
| 14.52 | FUNCTION MORPHTRAPTVM | 208 |

| | |
|--|------------|
| 14.53 FUNCTION MORPHVOP | 209 |
| 14.54 FUNCTION NEWCSR | 212 |
| 14.55 FUNCTION HPMACCESSVALID..... | 213 |
| 14.56 FUNCTION MAPADDRESS..... | 214 |
| 14.57 FUNCTION UNMAPPMPREGION..... | 215 |
| 14.58 FUNCTION UPDATELdStDOMAIN | 216 |
| 14.59 FUNCTION NEWTLBENTRY | 217 |
| 14.60 FUNCTION FREETLBENTRY | 219 |
| 14.61 FUNCTION NEWEXTREG | 220 |
| 15 Appendix: Extension Object Interface Functions | 221 |
| 15.1 FUNCTION RDFaultCB..... | 223 |
| 15.2 FUNCTION WRFaultCB..... | 224 |
| 15.3 FUNCTION RDSnapCB..... | 225 |
| 15.4 FUNCTION WRSnapCB..... | 226 |
| 15.5 FUNCTION SUPPRESSMEMEXCEPT..... | 227 |
| 15.6 FUNCTION CUSTOMNMI | 228 |
| 15.7 FUNCTION CUSTOMIASSIGN..... | 229 |
| 15.8 FUNCTION TRAPNOTIFIER..... | 231 |
| 15.9 FUNCTION TRAPPRENOTIFIER..... | 232 |
| 15.10 FUNCTION ERETNOTIFIER | 233 |
| 15.11 FUNCTION PRERESETNOTIFIER..... | 235 |
| 15.12 FUNCTION RESETNOTIFIER | 236 |
| 15.13 FUNCTION FIRSTEXCEPTION | 237 |
| 15.14 FUNCTION GETINTERRUPTPRI..... | 238 |
| 15.15 FUNCTION GETHANDLERPC | 239 |
| 15.16 FUNCTION INTUPDATE..... | 241 |
| 15.17 FUNCTION HALTRESTARTNOTIFIER | 242 |
| 15.18 FUNCTION LRSCABORTFN | 243 |
| 15.19 FUNCTION PREMORPH..... | 244 |
| 15.20 FUNCTION POSTMORPH..... | 246 |
| 15.21 FUNCTION AMOCHECK | 248 |
| 15.22 FUNCTION AMOMORPH | 249 |
| 15.23 FUNCTION EMITCSRCHECK | 250 |
| 15.24 FUNCTION UNITSTRIDECHECK..... | 252 |
| 15.25 FUNCTION EMITVFWREDUSUM | 254 |
| 15.26 FUNCTION EMITVFWREDUSUM..... | 257 |
| 15.27 FUNCTION SWITCHCB | 259 |
| 15.28 FUNCTION TLoad | 261 |
| 15.29 FUNCTION TStore | 263 |
| 15.30 FUNCTION DISTINCTPHYSMEM..... | 265 |
| 15.31 FUNCTION INSTALLPHYSMEM..... | 266 |
| 15.32 FUNCTION PMPPRIV | 269 |
| 15.33 FUNCTION PMAENABLE | 270 |
| 15.34 FUNCTION PMACHECK | 271 |
| 15.35 FUNCTION VALIDPTE..... | 272 |
| 15.36 FUNCTION VMTrap | 273 |
| 15.37 FUNCTION SETDOMAINNOTIFIER..... | 275 |
| 15.38 FUNCTION FREEENTRYNOTIFIER..... | 276 |
| 15.39 FUNCTION CLICCUSTOMRD | 277 |
| 15.40 FUNCTION CLICCUSTOMWR | 278 |
| 15.41 FUNCTION CLICUPDATED..... | 279 |

| | |
|---|------------|
| 15.42 FUNCTION RESTRICTIONS | 280 |
| 16 Appendix: riscvInstrInfo Structure | 281 |
| 16.1 COMMON FIELDS | 282 |
| 16.1.1 <i>riscvIType type</i> | 282 |
| 16.1.2 <i>riscvAddr thisPC</i> | 283 |
| 16.1.3 <i>Uns64 instruction</i> | 283 |
| 16.1.4 <i>Uns8 bytes</i> | 283 |
| 16.1.5 <i>riscvRegDesc r[RV_MAX_AREGS]</i> | 283 |
| 16.1.6 <i>riscvAddr tgt</i> | 284 |
| 16.1.7 <i>Uns64 c</i> | 284 |
| 16.1.8 <i>Int32 memBits</i> | 284 |
| 16.1.9 <i>Int32 Bool unsExt</i> | 284 |
| 16.2 CSR INSTRUCTION FIELDS | 284 |
| 16.2.1 <i>riscvCSRUDesc csrUpdate</i> | 284 |
| 16.2.2 <i>Uns32 csr</i> | 285 |
| 16.3 FLOATING POINT INSTRUCTION FIELDS | 285 |
| 16.3.1 <i>riscvRMDesc rm</i> | 285 |
| 16.3.2 <i>Bool useF</i> | 285 |
| 16.4 VECTOR EXTENSION INSTRUCTION FIELDS | 285 |
| 16.4.1 <i>riscvRegDesc mask</i> | 285 |
| 16.4.2 <i>riscvWholeDesc isWhole</i> | 285 |
| 16.4.3 <i>riscvVType vtype</i> | 286 |
| 16.4.4 <i>Uns32 eew</i> | 286 |
| 16.4.5 <i>Uns8 eewDiv</i> | 286 |
| 16.4.6 <i>Uns8 eewIndex</i> | 286 |
| 16.4.7 <i>Uns8 nf</i> | 286 |
| 16.4.8 <i>Bool isFF</i> | 286 |
| 16.5 CODE SIZE REDUCTION EXTENSION INSTRUCTION FIELDS | 286 |
| 16.5.1 <i>Uns32 rlist</i> | 287 |
| 16.5.2 <i>Uns32 alist</i> | 287 |
| 16.5.3 <i>riscvRetValDesc retval</i> | 287 |
| 16.5.4 <i>riscvCompressSet Zc</i> | 287 |
| 16.5.5 <i>Bool doRet</i> | 287 |
| 16.5.6 <i>Bool embedded</i> | 287 |
| 16.6 ZMMUL INSTRUCTION FIELDS | 288 |
| 16.6.1 <i>Bool Zmmul</i> | 288 |
| 16.7 PACKED SIMD EXTENSION INSTRUCTION FIELDS | 288 |
| 16.8 OTHER FIELDS | 288 |
| 17 Appendix: Custom CSR Description | 289 |
| 17.1 RISCVCsrATTRS FIELDS | 289 |
| 17.1.1 <i>Field name</i> | 289 |
| 17.1.2 <i>Field desc</i> | 289 |
| 17.1.3 <i>Field object</i> | 289 |
| 17.1.4 <i>Field csrNum</i> | 289 |
| 17.1.5 <i>Field arch</i> | 290 |
| 17.1.6 <i>Field access</i> | 290 |
| 17.1.7 <i>Field version</i> | 290 |
| 17.1.8 <i>Field presentCB</i> | 290 |
| 17.1.9 <i>Field readCB</i> | 290 |
| 17.1.10 <i>Field readWriteCB</i> | 291 |
| 17.1.11 <i>Field writeCB</i> | 291 |

| | | |
|---------|--|------------|
| 17.1.12 | Field wstateCB..... | 292 |
| 17.1.13 | Field reg | 292 |
| 17.1.14 | Field writeMaskV..... | 293 |
| 17.1.15 | Field writeMaskC32..... | 293 |
| 17.1.16 | Field writeMaskC64..... | 293 |
| 17.1.17 | Field Smstateen..... | 293 |
| 17.1.18 | Field trap..... | 293 |
| 17.1.19 | Field noTraceChange..... | 293 |
| 17.1.20 | Field wEndBlock | 294 |
| 17.1.21 | Field wEndRM | 294 |
| 17.1.22 | Field noSaveRestore..... | 294 |
| 17.1.23 | Field writeRd..... | 294 |
| 17.1.24 | Field aliasV | 294 |
| 17.1.25 | Field undefined..... | 294 |
| 17.1.26 | Field forceRO..... | 294 |
| 17.2 | CSR DEFINITION MACROS..... | 296 |
| 17.2.1 | Macro XCSR_ATTR_UIP | 296 |
| 17.2.2 | Macro XCSR_ATTR_NIP | 296 |
| 17.2.3 | Macro XCSR_ATTR_T | 296 |
| 17.2.4 | Macro XCSR_ATTR_TC | 297 |
| 17.2.5 | Macro XCSR_ATTR_TV | 297 |
| 17.2.6 | Macro XCSR_ATTR_P | 298 |
| 18 | Appendix: Configuring Standard Extensions..... | 299 |
| 18.1 | RATIFIED EXTENSION SUPPORT | 299 |
| 18.1.1 | RISC-V Quality-of-Service (QoS) Identifiers..... | 300 |
| 18.1.2 | Obviating Memory-Management Instructions after Marking PTEs Valid..... | 300 |
| 18.1.3 | Resumable Non-Maskable Interrupts | 300 |
| 18.1.4 | Shadow Stacks and Landing Pads | 300 |
| 18.1.5 | BF16 Extensions..... | 301 |
| 18.1.6 | Zaamo and Zalrsc Extensions..... | 301 |
| 18.1.7 | B Standard Extension for Bit Manipulation Instructions..... | 301 |
| 18.1.8 | Byte and Halfword Atomic memory Operations (Zabha)..... | 301 |
| 18.1.9 | RISC-V Supervisor Counter Delegation (Smcdeleg/Ssccfg)..... | 301 |
| 18.1.10 | May-Be-Operations | 301 |
| 18.1.11 | RISC-V Indirect CSR Access (Smcsrind/Sscsrind)..... | 301 |
| 18.1.12 | RISC-V Pointer Masking Extensions | 301 |
| 18.1.13 | RISC-V Integer Conditional (Zicond) operations extension..... | 301 |
| 18.1.14 | Hardware Updating of PTE A/D Bits (Svadu)..... | 302 |
| 18.1.15 | RISC-V Cycle and Instret Privilege Mode Filtering (Smcntrpmf)..... | 302 |
| 18.1.16 | Atomic Compare-and-Swap (CAS) Instructions (Zacas)..... | 302 |
| 18.1.17 | RISC-V Cryptography Extensions Volume II: Vector Instructions..... | 302 |
| 18.1.18 | "Zfa" Standard Extension for Additional Floating-Point Instructions | 302 |
| 18.1.19 | RISC-V Advanced Interrupt Architecture Extension | 302 |
| 18.1.20 | "Zvfh/Zvfhmin:" Vector Extension for Half-Precision Floating-Point Arithmetic/Vector Extension for Minimal Half-Precision Floating-Point Arithmetic..... | 303 |
| 18.1.21 | "Zihintntl" Non-Temporal Locality Hints..... | 303 |
| 18.1.22 | RISC-V Code Size Reduction..... | 303 |
| 18.1.23 | RISC-V Profiles | 303 |
| 18.1.24 | "Zicntr" and "Zihpm" Counters..... | 303 |
| 18.1.25 | RV32E and RV64E Base Integer Instruction Sets..... | 304 |
| 18.1.26 | "Ztso" Standard Extension for Total Store Ordering | 304 |
| 18.1.27 | RISC-V Wait-on-Reservation-Set (Zawrs) extension | 304 |
| 18.1.28 | Zmmul Extension | 304 |

| | | |
|---------|---|-----|
| 18.1.29 | <i>RISC-V PMP Enhancements for memory access and execution prevention on Machine mode (Smepmp)</i> | 304 |
| 18.1.30 | <i>RISC-V Base Cache Management Operation ISA Extensions</i> | 304 |
| 18.1.31 | <i>RISC-V Bit-Manipulation ISA-extensions</i> | 304 |
| 18.1.32 | <i>RISC-V Count Overflow and Mode-Based Filtering Extension</i> | 305 |
| 18.1.33 | <i>RISC-V Cryptography Extensions Volume I: Scalar & Entropy Source Instructions</i> | 305 |
| 18.1.34 | <i>RISC-V State Enable Extension</i> | 305 |
| 18.1.35 | <i>RISC-V "stimecmp / vstimecmp" Extension</i> | 305 |
| 18.1.36 | <i>RISC-V Vector Extension</i> | 305 |
| 18.1.37 | <i>The RISC-V Instruction Set Manual Volume II: Privileged Architecture</i> | 306 |
| 18.1.38 | <i>"Zfh" and "Zfhmin" Standard Extensions for Half-Precision Floating-Point</i> | 306 |
| 18.1.39 | <i>"Zfinx", "Zdinx", "Zhinx", "Zhinxmin": Standard Extensions for Floating-Point in Integer Registers</i> | 306 |
| 18.1.40 | <i>"Zhintpause" Pause Hint</i> | 307 |
| 18.1.41 | <i>"Zicsr", Control and Status Register (CSR) Instructions</i> | 307 |
| 18.1.42 | <i>"Zifencei" Instruction-Fetch Fence</i> | 307 |
| 18.2 | <i>RISC-V PROFILE EXTENSIONS</i> | 308 |
| 18.3 | <i>OTHER EXTENSIONS</i> | 309 |
| 18.3.1 | <i>"Sdext" ISA Extension</i> | 309 |
| 18.3.2 | <i>"Sdtrig" ISA Extension</i> | 309 |
| 18.3.3 | <i>Core-Local Interrupt Controller (CLIC) RISC-V Privileged Architecture Extensions</i> | 309 |

1 Introduction

The RISC-V architecture is designed to be extensible. Standard features are grouped into subsets which may or may not be present in a particular implementation, and in addition the architecture permits custom extensions (for example, non-standard instructions or control and status registers, CSRs). This presents an issue when a *model* of such a custom processor is required – how can such a model be easily developed and maintained as the architecture evolves?

This document describes a methodology to simplify enhancement of the generic OVP RISC-V processor model with custom extensions. These extensions can add instructions, registers (including CSRs), net ports, bus ports, FIFO ports and documentation to the base model.

Any of the techniques described in the *Imperas Binary Interception Technology User Guide* can be used to enhance a processor model.

In addition, the OVP RISC-V processor model provides some RISC-V specific integration functionality that simplifies addition of many kinds of feature. This document, from chapter 3, describes these model-specific integration facilities, using examples from a template custom RISC-V model implemented in the `vendor.com` library.

Note that many aspects of standard RISC-V model behavior can be controlled using parameters. For processors within the standard envelope, it may be that no custom model is needed and that all required behavior can be specified using parameters. Section 5 of this document describes the available options in detail; the appendix in section 18 presents some of the same information from a RISC-V architectural extension perspective.

2 Building a RISC-V Model

2.1 Introduction

A RISC-V processor model can be built from source provided in a `riscvOVPSim/` `riscvOVPSimPlus` download or within a VLNV library provided in an `OVPSim` or `Imperas` or other installation. This chapter describes how this source can be built and how the subsequent model binary image can be used in a simulation.

2.2 Using Local RISC-V Model Source Directory

To build a source model in a standalone directory an installation of `OVPSim` (www.ovpworld.org) or the Imperas Professional products (www.imperas.com) are required. As part of these installations a `Makefile` build system is provided that can be used in a Linux shell or an `MSYS` shell on Windows.

Below is shown the command line for building the source that is provided within a `riscv-ovpsim` download from GitHub:

```
% cd riscvOVPSim/source
% make -f $IMPERAS_HOME/ImperasLib/buildutils/Makefile.host
```

For example, this may then be executed using `IMPERAS_ISS` by selecting the processor as shown below:

```
% $IMPERAS_ISS --processorfile /d/Imperas/work/riscvOVPSim/source/model.so \
--variant RV32I \
--program application.elf
```

2.3 Using VLNV Library RISC-V Model Source

To build a source model in a VLNV library an installation of `OVPSim` (www.ovpworld.org) or the Imperas Professional products (www.imperas.com) are required. As part of these installations a `Makefile` build system is provided that can be used in a Linux shell or an `MSYS` shell on Windows.

To show the building of a VLNV library, assume that there is a local library source directory at `/home/user1/LocalLib/source` containing the RISC-V processor source in a standard VLNV structure, for example as `vendor/processor/riscv/1.0/model`

The following instructions will build the model into a user-defined output VLNV binary library directory (in this case, `/home/user1/lib/$IMPERAS_ARCH/ImperasLib`):

```
% mkdir -p /home/user1/lib/$IMPERAS_ARCH/ImperasLib
% make -f $IMPERAS_HOME/ImperasLib/buildutils/Makefile.library \
VLNVSRC=/home/user1/LocalLib/source \
VLNVROOT=/home/user1/lib/$IMPERAS_ARCH/ImperasLib
```

To use this model in a simulation platform or with the `IMPERAS_ISS` the following argument should be added to the simulator command line:

```
-vlnvroot lib/$IMPERAS_ARCH/ImperasLib
```

For example, this may then be executed using `IMPERAS_ISS` by selecting the processor as shown below:

```
% $IMPERAS_ISS --processorvendor vendor --processor name riscv \  
--vlnvroot /home/user1/lib/$IMPERAS_ARCH/ImperasLib \  
--variant RV32I \  
--program application.elf
```

2.4 Using a Linked Model

The previous sections have described how a RISC-V model binary image can be created from existing source code. When creating a *custom* or *extended* RISC-V processor model, some important things should be considered:

- 1) How easy will it be to maintain the custom model?
- 2) How will the custom model be adapted to changes to the specifications (either new versions of existing specifications, or entirely new standard extensions)?
- 3) How will the custom model be tested and verified?

The Imperas RISC-V model is designed to address these issues by allowing *linked* models to be created. Linked models use the *unmodified* RISC-V base model source and add custom features by linking in an additional custom shared object. This addresses the issues above in the following way:

- 1) Imperas maintain the base RISC-V model.
- 2) Imperas update the RISC-V model to add new extensions and for new versions or modifications of specifications. As part of doing this configuration parameters are provided that allow any of the current or previous versions to be selected that will configure the processor operations.
- 3) Imperas test and verify all aspects of the RISC-V model so that it can be used as a golden reference

3 Introduction to Custom Extensions with Linked Model

Modeling custom features is done by specifying a *custom configuration* and possibly adding a *custom extension library* to the basic RISC-V processor model, using a *linked model*. A linked model is a RISC-V model that refers to all source files in the base model using *links* rather than copying them. This has the advantage that if the base model is updated (to fix bugs or provide new features) then the linked model will automatically inherit those features – in effect, the linked model implements a skin around the base model. This document discusses these features in turn, referring to the `vendor.com` template, as follows:

3.1 *Linked Model Creation*

Chapter 4 describes what must be done to create a new linked model.

3.2 *Custom Configuration Options*

Chapter 5 describes in detail the custom configuration options that can be specified for a linked model. These options control the availability of standard feature subsets.

3.3 *Adding Custom Instructions*

Chapter 6 describes how custom instructions can be added to a RISC-V model.

3.4 *Adding Custom CSRs*

Chapter 7 describes how custom CSRs can be added to a RISC-V model, and how the behavior of existing CSRs in the base model can be modified.

3.5 *Adding Custom Exceptions*

Chapter 8 describes how custom exceptions can be added to a RISC-V model.

3.6 *Adding Custom Local Interrupts*

Chapter 9 describes how custom local interrupts can be added to a RISC-V model.

3.7 *Adding Custom FIFO Ports*

Chapter 10 is an advanced example describing how custom FIFO ports and supporting instructions can be added to a RISC-V model.

3.8 *Adding Transactional Memory*

Chapter 11 is an advanced example describing how transactional memory and supporting instructions can be added to a RISC-V model.

3.9 *Implementing PMA Constraints*

Chapter 12 describes how custom PMA constraints can be added to a RISC-V model.

4 Linked Model Creation

4.1 Introduction

To create a custom RISC-V model, create a new vendor component VLNV library based on the `vendor.com` template in the Imperas VLNV library. The new VLNV library should use a unique vendor name (don't use `vendor.com`). Beneath the `vendor.com` directory, there are two subdirectories of significance to this process:

1. Directory `processor/riscv/1.0/model` contains files required to extend the base RISC-V model to implement the linked RISC-V model. There are three source files:
 - a. `extensionConfig.h`: this contains an enumeration of the various extensions that can be applied to the model. The RISC-V model in the `vendor.com` directory has six separate extensions, each identified by a member of the `extensionID` enumeration. A custom model may implement either no extensions at all or only a single extension as well – the `vendor.com` model has multiple extensions to provide separate starting points for different types of extension, as described below.
 - b. `riscvConfigList.c`: This defines the RISC-V variants implemented by this linked model, together with a list of extension libraries that implement particular custom features for each variant.
 - c. `riscvInfo.c`: This identifies the available extensions by name and specifies various other standard RISC-V processor features (for example, which debugger to use).

In addition to the three source files, there is a Makefile configured to automatically link to the base RISC-V model.

The combined model is constructed by first taking all source files from the base model (the RISC-V processor model in `riscv.ovpworld.org`) and then copying in any source files from the extension, thereby replacing any files in the base model with files with the same name from the extension.

In theory, *any* base model source file can be replaced in this way, but in practice **only `riscvConfigList.c` and `riscvInfo.c` should ever be replaced in this way** (otherwise the extended model will diverge from base model behavior).

2. Directory `intercept` contains source of each extension library that adds additional custom features to the base RISC-V model. In the `vendor.com` example, there are six extension libraries, each of which adds a particular class of feature so that the features can be described in stand-alone chapters of this document. A typical custom processor will have a single extension library combining features from one or more of these examples; some custom processors may require no extensions at all if they are simply configuring standard features implemented by the base model.

4.2 Creating a New Model Library

A new customized RISC-V linked model must be created by copying files from the `vendor.com` template model into a new vendor directory. The files that must be copied depend on whether the new model has custom extensions or not: if it has no custom extensions but simply requires configuration options to be set correctly, the process is simpler. The steps required in these two cases are described in the following sections.

4.2.1 Creating a New Model - No Custom Extensions

First copy the `vendor.com/processor` directory to a new vendor-specific library directory (called `custom.com` in this section), using commands like this (on Linux):

```
% mkdir -p /home/user1/LocalLib/source/custom.com
% cp -r $IMPERAS_HOME/ImperasLib/source/vendor.com/processor
/home/user1/LocalLib/source/custom.com
```

Once the source has been copied, update the files in the `custom.com/processor/riscv/1.0/model` directory as described below.

4.2.1.1 File `extensionConfig.h`

Delete the `extensionConfig.h` file, which is not required if there are no extensions.

4.2.1.2 File `riscvInfo.c`

This file provides generic information about a processor (not RISC-V specific). Modify it to remove all `vmiVlnvInfo` and `vmiVlnvInfoList` structures, and all references to the `mandatoryExtensions` field in `vmiProcessorInfo` definitions. Then update the `info32` and `info64` `vmiProcessorInfo` definitions to correct the `vendor` and `family` fields to match the new vendor name. For example, the `info32` entry should be modified like this:

```
static const vmiProcessorInfo info32 = {

    .vlnv.vendor      = "custom.com",
    .vlnv.library     = "processor",
    .vlnv.name        = "riscv",
    .vlnv.version     = "1.0",

    .semihost.vendor  = "riscv.ovpworld.org",
    .semihost.library = "semihosting",
    .semihost.name    = "pk",
    .semihost.version = "1.0",

    .helper.vendor    = "imperas.com",
    .helper.library   = "intercept",
    .helper.name      = "riscv32CpuHelper",
    .helper.version   = "1.0",

    .elfCode          = 243,
    .gdbPath          = "$IMPERAS_HOME/lib/$IMPERAS_ARCH/gdb/riscv-none-embed-gdb"
VMI_EXE_SUFFIX,
    .gdbInitCommands  = "set architecture riscv:rv32",
    .family           = "Custom",
    .QLQualified      = True
};
```


Depending on whether the extended model supports only `XLEN` of 32 or 64, either the `info32` or `info64` entries might be deleted, and the return value of `riscvProcInfo()` adjusted accordingly.

4.2.1.3 File `riscvConfigList.c`

This file provides RISC-V-specific information about the processor variants implemented. Each variant is implemented by a single structure of type `riscvConfig`; the multiple possible variants are in a null-terminated list which is returned by function `riscvGetConfigList`. The `vendor.com` template model contains two variants, `RV32X` and `RV64X`, and there is a configuration for each:

```
static const riscvConfig configList[] = {  
    {  
        .name           = "RV32X",  
        .arch           = ISA_U|RV32GC|ISA_X,  
        .user_version    = RVUV_DEFAULT,  
        .priv_version    = RVPV_DEFAULT,  
        .tval_ii_code    = True,  
        .ASID_bits       = 9,  
        .local_int_num   = 7,           // enable local interrupts 16-22  
        .unimp_int_mask  = 0x1f0000,   // int16-int20 absent  
        .extensionConfigs = allExtensions,  
    },  
    {  
        .name           = "RV64X",  
        .arch           = ISA_U|RV64GC|ISA_X,  
        .user_version    = RVUV_DEFAULT,  
        .priv_version    = RVPV_DEFAULT,  
        .tval_ii_code    = True,  
        .ASID_bits       = 9,  
        .local_int_num   = 7,           // enable local interrupts 16-22  
        .unimp_int_mask  = 0x1f0000,   // int16-int20 absent  
        .extensionConfigs = allExtensions,  
    },  
    {0} // null terminator  
};
```

The `riscvConfig` structure allows many aspects of standard RISC-V processor behavior to be configured, as described later in this document. For a new processor with a single variant, modify the null-terminated list to contain a single entry with the correct name with a minimal set of field initializations, removing the `extensionConfigs` field which is not required if there are no extensions:

```
static const riscvConfig configList[] = {  
    {  
        .name           = "variant1",  
        .arch           = ISA_U|RV32GC|ISA_X,  
        .user_version    = RVUV_DEFAULT,  
        .priv_version    = RVPV_DEFAULT,  
    },  
    {0} // null terminator  
};
```

Then refer to sections 5 and 18 in this document to specify more fields in the `riscvConfig` structure to implement the required configuration.

4.2.2 Creating a New Model – With Custom Extensions

To create a new customized RISC-V linked model with custom extensions, first copy the `vendor.com/processor` and `vendor.com/intercept` directories to a new vendor-specific library directory (called `custom.com` in this section), using commands like this (on Linux):

```
% mkdir -p /home/user1/LocalLib/source
% cp -r $IMPERAS_HOME/ImperasLib/source/vendor.com /home/user1/LocalLib/source/custom.com
```

Depending on the nature of the extensions being added, remove all except one of the subdirectories under the `intercept` directory. These directories implement separate extensions, as follows:

1. `addCSRs`: adds custom CSRs;
2. `addInstructions`: adds four instructions operating on general-purpose registers;
3. `addExceptions`: adds custom exceptions and one simple instruction;
4. `addLocalInterrupts`: adds 2 local interrupts with custom priorities;
5. `fifoExtensions`: adds FIFO ports, custom instructions and documentation;
6. `tmExtensions`: adds custom instructions, exceptions, documentation and implements transactional memory.

Once the source has been copied, update the files in the `custom.com/processor/riscv/1.0/model` directory as described below.

4.2.2.1 File `extensionConfig.h`

Modify this to contain a single enumeration entry for the new extension, for example:

```
typedef enum extensionIDE {
    EXTID_CUSTOM = 1234,
} extensionID;
```

When multiple extensions are present on a RISC-V processor, the code here is used to distinguish between them so that context information for an extension can be obtained by client code. The value of the enumeration member is arbitrary but **must be unique amongst all extensions** added to a RISC-V extended processor model.

4.2.2.2 File `riscvInfo.c`

This file provides generic information about a processor (not RISC-V specific). Modify it to contain a single `vmiVlnvInfo` structure for the extension, and a single `vmiVlnvInfoList` structure referencing it, for example:

```
static const vmiVlnvInfo custom = {
    .vendor    = "custom.com",
    .library   = "intercept",
    .name      = "customExt",
```

```

    .version      = "1.0",
};

static const vmiVlnvInfoList customEntry = {
    next : 0,
    info : &custom,
};

```

Then update the `info32` and `info64` `vmiProcessorInfo` definitions to include the new `customEntry` list in the `mandatoryExtensions` field and correct the `vendor` and `family` fields to match the new vendor name. For example, the `info32` entry should be modified like this:

```

static const vmiProcessorInfo info32 = {

    .vlnv.vendor      = "custom.com",
    .vlnv.library     = "processor",
    .vlnv.name        = "riscv",
    .vlnv.version     = "1.0",

    .semihost.vendor  = "riscv.ovpworld.org",
    .semihost.library = "semihosting",
    .semihost.name    = "pk",
    .semihost.version = "1.0",

    .helper.vendor    = "imperas.com",
    .helper.library   = "intercept",
    .helper.name      = "riscv32CpuHelper",
    .helper.version   = "1.0",

    .mandatoryExtensions = &customEntry,

    .elfCode          = 243,
    .gdbPath          = "$IMPERAS_HOME/lib/$IMPERAS_ARCH/gdb/riscv-none-embed-gdb"
    VMI_EXE_SUFFIX,
    .gdbInitCommands  = "set architecture riscv:rv32",
    .family           = "Custom",
    .QLQualified      = True
};

```

The `vendor.com` template example shows how it is possible to add *multiple* extension libraries to a processor, by forming a list of `vmiVlnvInfo` structures, if required.

4.2.2.3 File `riscvConfigList.c`

This file provides RISC-V-specific information about the processor variants implemented. Each variant is implemented by a single structure of type `riscvConfig`; the multiple possible variants are in a null-terminated list which is returned by function `riscvGetConfigList`. The `vendor.com` template model contains two variants, `RV32X` and `RV64X`, and there is a configuration for each:

```

static const riscvConfig configList[] = {

    {
        .name          = "RV32X",
        .arch          = ISA_U|RV32GC|ISA_X,
        .user_version   = RVUV_DEFAULT,
        .priv_version   = RVPV_DEFAULT,
        .tval_ii_code   = True,
        .ASID_bits      = 9,
        .local_int_num  = 7,           // enable local interrupts 16-22
        .unimp_int_mask = 0x1f0000,   // int16-int20 absent
        .extensionConfigs = allExtensions,
    }
};

```

```
    },  
    {  
        .name            = "RV64X",  
        .arch            = ISA_U|RV64GC|ISA_X,  
        .user_version    = RVUV_DEFAULT,  
        .priv_version    = RVPV_DEFAULT,  
        .tval_ii_code    = True,  
        .ASID_bits       = 9,  
        .local_int_num   = 7,           // enable local interrupts 16-22  
        .unimp_int_mask  = 0x1f0000,   // int16-int20 absent  
        .extensionConfigs = allExtensions  
    },  
    {0} // null terminator  
};
```

The `riscvConfig` structure allows many aspects of the RISC-V processor to be configured without additional extension library effort. For a new processor with a single variant, modify the null-terminated list to contain a single entry with the correct name with a minimal set of field initializations:

```
static const riscvConfig configList[] = {  
    {  
        .name            = "variant1",  
        .arch            = ISA_U|RV32GC|ISA_X,  
        .user_version    = RVUV_DEFAULT,  
        .priv_version    = RVPV_DEFAULT,  
        .extensionConfigs = allExtensions,  
    },  
    {0} // null terminator  
};
```

The `extensionConfigs` field is a pointer to a null-terminated list of `riscvExtConfig` structures. Each structure contains an extension id and a pointer to an extension-specific configuration structure (which can often be `NULL`, but see following chapters for more detail). To add a single extension, modify the template code like this:

```
static riscvExtConfigCP allExtensions[] = {  
    &(const riscvExtConfig){  
        .id      = EXTID_CUSTOM,  
        .userData = 0  
    },  
    0 // KEEP LAST: terminator  
};
```

Note that the value `EXTID_CUSTOM` was defined previously in file `extensionConfig.h`.

4.3 Building the linked model

The linked model is built using the standard Makefile system provided in a product installation, located at `$IMPERAS_HOME`. Assuming the linked model source is in a local library directory at `/home/user1/LocalLib/source/custom.com`, the following commands will build the model into an output directory

```
/home/user1/lib/$IMPERAS_ARCH/LocalLib:
```

```
% make -f $IMPERAS_HOME/ImperasLib/buildutils/Makefile.library \  
-C /home/user1/LocalLib/source \  
VLNVSRC=/home/user1/LocalLib/source \  
VLNVROOT=/home/user1/lib/$IMPERAS_ARCH/LocalLib
```

4.4 Executing the linked model

To use this model in a simulation platform the following argument should be added to the simulator command line:

```
-vlnvroot lib/$IMPERAS_ARCH/LocalLib
```

The model can be used with `IMPERAS_ISS` by inclusion on the command line, using:

```
-processorvendor custom.com vlnvroot lib/$IMPERAS_ARCH/LocalLib
```

An example showing the process to create the custom processor model and to run a simple assembler test application is provided in an Imperas release at

```
IMPERAS_HOME/Examples/Models/Processor/FeatureUsage/RISCV_ExtendedModel
```

5 Custom Configuration Options

Many features of a custom RISC-V variant can be defined by fields in the configuration structure of type `riscvExtConfig` for that variant. This section describes the configuration structure and the possible field settings. The defaults specified in this structure can generally be overridden by model parameters if required.

The structure is defined in file `riscvConfig.h` in the base model. It contains configuration options, important CSR defaults and CSR write masks:

```
typedef struct riscvConfigS {

    const char      *name;                // variant name

    // fundamental variant configuration
    riscvArchitecture arch;                // variant architecture
    riscvArchitecture archImplicit;         // implicit feature bits (not in misa)
    riscvArchitecture archMask;            // read/write bits in architecture
    riscvArchitecture archFixed;           // fixed bits in architecture
    riscvUserVer     user_version    : 8;  // user-level ISA version
    riscvPrivVer     priv_version    : 8;  // privileged architecture version
    riscvVectVer     vect_version    : 8;  // vector architecture version
    riscvVectorSet   vect_profile    : 8;  // vector architecture profile
    riscvBitManipVer bitmanip_version: 8;  // bitmanip architecture version
    riscvBitManipSet bitmanip_absent :16;  // bitmanip absent extensions
    riscvCryptoVer   crypto_version  : 8;  // cryptographic architecture version
    riscvVCryptoVer  vcrypto_version : 8;  // vector cryptographic version
    riscvCryptoSet   crypto_absent   :32;  // cryptographic absent extensions
    riscvDSPVer      dsp_version     : 8;  // DSP architecture version
    riscvDSPSet      dsp_absent      : 8;  // DSP absent extensions
    riscvCompressVer compress_version: 8;  // compressed architecture version
    riscvCompressSet compress_present:16;  // compressed present extensions
    riscvHypVer      hyp_version     : 8;  // hypervisor architecture version
    riscvDebugVer    dbg_version     : 8;  // debugger architecture version
    riscvRNMIVer     rnmi_version    : 8;  // rnmi version
    riscvSmepmpVer   smepmp_version  : 8;  // Smepmp version
    riscvCLICVer     CLIC_version    : 8;  // CLIC version
    riscvAIAVer      AIA_version     : 8;  // AIA version
    riscvZfinxVer    Zfinx_version   : 8;  // Zfinx version
    riscvZceaVer     Zcea_version     : 8;  // Zcea version (legacy only)
    riscvZcebVer     Zceb_version     : 8;  // Zceb version (legacy only)
    riscvZceeVer     Zcee_version     : 8;  // Zcee version (legacy only)
    riscvFPL6Ver     fpl6_version    : 8;  // 16-bit floating point version
    riscvFSMode      mstatus_fs_mode : 8;  // mstatus.FS update mode
    riscvDMMode      debug_mode      : 8;  // is Debug mode implemented?
    riscvDERETMode   debug_eret_mode : 8;  // debug mode MRET/SRET/DRET action
    riscvDPPriority  debug_priority  : 8;  // simultaneous debug event priority
    riscvChainTVAL   chain_tval      : 8;  // chained trigger tval to select
    riscvPMPROW1     PMP_ROW1        : 8;  // behavior when PMP R=0 W=1
    riscvPMMMode     Ssnpm            : 8;  // Ssnpm implemented modes
    riscvPMMMode     Ssnpnm           : 8;  // Ssnpnm implemented modes
    riscvPMMMode     Smmpm            : 8;  // Smmpm implemented modes
    const char      **members;         // cluster member variants
    const char      *leaf_hart_prefix; // prefix for hart in cluster
    const char      *mtime_counter;    // mtime counter name

    // memory constraints
    riscvMConstraint amo_constraint    : 2;  // AMO memory constraint
    riscvMConstraint lr_sc_constraint  : 2;  // LR/SC memory constraint
    riscvMConstraint push_pop_constraint : 2;  // PUSH/POP memory constraint
    riscvMConstraint vector_constraint : 2;  // vector load/store constraint
    riscvPMARegionCP pmaStatic;        // static PMA regions

    // unimplemented instructions
    riscvITypeCP unimplementedInstr;    // unimplemented instruction list
}
```

```
// remapped CSR addresses
const char *CSR_remap; // "<csrName>=<number>,..."

// configuration not visible in CSR state
Uns64 reset_address; // reset vector address
Uns64 nmi_address; // NMI address
Uns64 nmiexc_address; // RNMI exception address
Uns64 debug_address; // debug vector address
Uns64 dexc_address; // debug exception address
Uns64 CLINT_address; // internally-implemented CLINT address
Flt64 mtime_Hz; // clock frequency of CLINT mtime
Uns64 unimp_int_mask; // mask of unimplemented interrupts
Uns64 force_mideleg; // always-delegated M-mode interrupts
Uns64 force_sideleg; // always-delegated S-mode interrupts
Uns64 no_ideleg; // non-delegated interrupts
Uns64 no_edeleg; // non-delegated exceptions
Uns64 ecode_mask; // implemented bits in xcause.ecode
Uns64 ecode_nmi; // exception code for NMI
Uns64 ecode_nmi_mask; // implemented bits in mncause.ecode
Uns64 Svnepot_page_mask; // implemented Svnepot page sizes
Uns64 miprio_mask; // writable entries in M-mode iprio array
Uns64 siprio_mask; // writable entries in S-mode iprio array
Uns64 hviprio_mask; // writable entries in VS-mode iprio array
Uns32 counteren_mask; // mcounteren mask
Uns32 scounteren_zero_mask; // zero bits in scounteren
Uns32 hcounteren_zero_mask; // zero bits in hcounteren
Uns32 noinhibit_mask; // counter no-inhibit mask
Uns32 local_int_num; // number of local interrupts
Uns32 lr_sc_grain; // LR/SC region grain size
Uns32 PMP_grain; // PMP region grain size
Uns32 PMP_registers; // number of implemented PMP regions
Uns32 PMP_csrs; // number of implemented PMP CSRs
Uns32 PMP_max_page; // maximum size of PMP page to map
Uns32 Sv_modes; // bit mask of valid Sv modes
Uns32 numHarts; // number of hart contexts if MPCore
Uns32 tvec_align; // trap vector alignment (vectored mode)
Uns32 ELEN; // ELEN (vector extension)
Uns32 SLEN; // SLEN (vector extension)
Uns32 VLEN; // VLEN (vector extension)
Uns32 EEW_index; // maximum index EEW (vector extension)
Uns32 SEW_min; // minimum SEW (vector extension)
Uns32 ASID_cache_size; // ASID cache size
Uns32 TW_time_limit; // mstatus.TW time limit
Uns32 STO_time_limit; // WRS.STO time limit
Uns32 tinfo; // tinfo default value (all triggers)
Uns16 trigger_match; // bitmask of legal trigger match values
Uns16 cmomp_bytes; // cache block bytes (management/prefetch)
Uns16 cmoz_bytes; // cache block bytes (zero)
Uns8 ASID_bits; // number of implemented ASID bits
Uns8 VMID_bits; // number of implemented VMID bits
Uns8 trigger_num; // number of implemented triggers
Uns8 mcontext_bits; // implemented bits in mcontext
Uns8 scontext_bits; // implemented bits in scontext
Uns8 mvalue_bits; // implemented bits in textra.mvalue
Uns8 svalue_bits; // implemented bits in textra.svalue
Uns8 mcontrol_maskmax; // configured value of mcontrol.maskmax
Uns8 dcsr_ebreak_mask; // mask of writable dcsr.ebreak* bits
Uns8 hvictl_IID_bits; // hvictl.IID implemented bits
Uns8 mtime_bits; // bit size of mtime counter
Uns8 IPRIOLEN : 4; // AIA IPRIOLEN value
Uns8 HIPRIOLEN : 4; // AIA HIPRIOLEN value
Bool isPSE : 1; // whether a PSE (internal use only)
Bool enable_expanded : 1; // enable expanded instructions
memEndian endian : 1; // data endianness
Bool endianFixed : 1; // endianness is fixed (UBE/SBE/MBE r/o)
Bool use_hw_reg_names : 1; // use hardware names for X/F registers
Bool no_pseudo_inst : 1; // don't report pseudo-instructions
Bool show_c_prefix : 1; // show opcode compressed prefix
Bool ABI_d : 1; // ABI uses D registers for parameters
Bool misa_B_Zba_Zbb_Zbs : 1; // misa.B if Zba, Zbb and Zbs present
```

```

Bool agnostic_ones      : 1;      // when agnostic elements set to 1
Bool MXL_writable       : 1;      // writable bits in misa.MXL
Bool SXL_writable       : 1;      // writable bits in mstatus.SXL
Bool UXL_writable       : 1;      // writable bits in mstatus.UXL
Bool VSXL_writable      : 1;      // writable bits in mstatus.VSXL
Bool Smstateen         : 1;      // Smstateen implemented?
Bool Smcsrind           : 1;      // Smcsrind implemented?
Bool Sstc               : 1;      // Sstc implemented?
Bool Sscofpmf           : 1;      // Sscofpmf implemented?
Bool Smcntrpmf          : 1;      // Smcntrpmf implemented?
Bool Smcdeleg           : 1;      // Smcdeleg implemented?
Bool Svpbmt             : 1;      // Svpbmt implemented?
Bool Svinval            : 1;      // Svinval implemented?
Bool Svadu              : 1;      // Svadu implemented?
Bool Ssqosid            : 1;      // Ssqosid implemented?
Bool Smaia              : 1;      // Smaia implemented?
Bool IMSIC_present      : 1;      // IMSIC present?
Bool noZaamo            : 1;      // Zaamo not implemented?
Bool noZalrsc           : 1;      // Zalrsc not implemented?
Bool Zacas              : 1;      // Zacas implemented?
Bool Zabha              : 1;      // Zabha implemented?
Bool Zawrs              : 1;      // Zawrs implemented?
Bool Zmmul              : 1;      // Zmmul implemented?
Bool Zfa                : 1;      // Zfa implemented?
Bool Zfhmin             : 1;      // Zfhmin implemented?
Bool Zfbfmin            : 1;      // Zfbfmin implemented?
Bool Zvlsseg            : 1;      // Zvlsseg implemented?
Bool Zvamo              : 1;      // Zvamo implemented?
Bool Zvediv             : 1;      // Zvediv implemented?
Bool Zvqmac             : 1;      // Zvqmac implemented?
Bool Zvfh               : 1;      // Zvfh implemented?
Bool Zvfhmin            : 1;      // Zvfhmin implemented?
Bool Zvfbfmin           : 1;      // Zvfbfmin implemented?
Bool Zvfbfwma           : 1;      // Zvfbfwma implemented?
Bool unitStrideOnly     : 1;      // only unit-stride operations supported
Bool noFaultOnlyFirst   : 1;      // fault-only-first instructions absent?
Bool Zihintntl          : 1;      // whether Zihintntl is present
Bool Zicnd              : 1;      // whether Zicnd is present
Bool Zicbom             : 1;      // whether Zicbom is present
Bool Zicbop             : 1;      // whether Zicbop is present
Bool Zicboz             : 1;      // whether Zicboz is present
Bool Zimop              : 1;      // whether Zimop is present
Bool Zcmop              : 1;      // whether Zcmop is present
Bool Zicfiss            : 1;      // whether Zicfiss is present
Bool Zicfilp            : 1;      // whether Zicfilp is present
Bool noZicsr            : 1;      // whether Zicsr is absent
Bool noZifencei         : 1;      // whether Zifencei is absent
Bool updatePTEA         : 1;      // hardware update of PTE A bit?
Bool updatePTED         : 1;      // hardware update of PTE D bit?
Bool unaligned_low_pri  : 1;      // low-priority unaligned exceptions
Bool unaligned_check_ms : 1;      // whether memory-side checks alignment
Bool unaligned          : 1;      // whether unaligned accesses supported
Bool unalignedAMO       : 1;      // whether AMO supports unaligned
Bool unalignedV         : 1;      // whether vector supports unaligned
Bool wfi_is_nop         : 1;      // whether WFI is treated as NOP
Bool wfi_resume_not_trap : 1;    // whether pending wakeup stops WFI trap
Bool nmi_absent         : 1;      // whether NMI is absent
Bool nmi_is_latched     : 1;      // whether NMI is posedge-latched
Bool nmi_high_priority  : 1;      // whether NMI is high priority
Bool nmi_update_mstatus : 1;      // whether NMI updates mstatus mpie and mie
Bool nmi_update_tcontrol : 1;    // whether NMI updates tcontrol mpte and mte
Bool nmi_zero_mtval     : 1;      // whether NMI zeroes mtval
Bool mtval_is_ro        : 1;      // whether mtval is read-only
Bool mtvec_is_ro        : 1;      // whether mtvec is read-only
Bool mtvec_sext         : 1;      // whether mtvec sign-extended
Bool stvec_sext         : 1;      // whether stvec sign-extended
Bool utvec_sext         : 1;      // whether utvec sign-extended
Bool mtvt_sext          : 1;      // whether mtvec sign-extended
Bool stvt_sext          : 1;      // whether stvec sign-extended
Bool utvt_sext          : 1;      // whether utvec sign-extended
Bool cycle_undefined    : 1;      // whether cycle CSR undefined

```



```

Bool mcycle_undefined : 1; // whether mcycle CSR undefined
Bool time_undefined : 1; // whether time CSR is undefined
Bool instret_undefined : 1; // whether instret CSR undefined
Bool minstret_undefined : 1; // whether minstret CSR undefined
Bool hpmcounter_undefined : 1; // whether hpmcounter* CSRs undefined
Bool mhpmpcounter_undefined : 1; // whether mhpmpcounter* CSRs undefined
Bool tdata2_undefined : 1; // whether tdata2 CSR is undefined
Bool tdata3_undefined : 1; // whether tdata3 CSR is undefined
Bool tinfo_undefined : 1; // whether tinfo CSR is undefined
Bool tcontrol_undefined : 1; // whether tcontrol CSR is undefined
Bool mcontext_undefined : 1; // whether mcontext CSR is undefined
Bool scontext_undefined : 1; // whether scontext CSR is undefined
Bool mscontext_undefined : 1; // whether mscontext CSR is undefined
Bool scontext_0x5A8 : 1; // whether scontext explicitly at 0x5A8
Bool hcontext_undefined : 1; // whether hcontext CSR is undefined
Bool mnoise_undefined : 1; // whether mnoise CSR is undefined
Bool dscratch0_undefined : 1; // whether dscratch0 CSR is undefined
Bool dscratch1_undefined : 1; // whether dscratch1 CSR is undefined
Bool amo_trigger : 1; // whether triggers used with AMO
Bool amo_aborts_lr_sc : 1; // whether AMO aborts active LR/SC
Bool lr_sc_match_size : 1; // whether LR/SC size must match
Bool ignore_non_leaf_DAU : 1; // whether ignore non-leaf PTE D, A, U
Bool no_hit : 1; // whether tdata1.hit* unimplemented
Bool no_sselect_2 : 1; // whether textra.sselect=2 is illegal
Bool d_requires_f : 1; // whether misa D requires F to be set
Bool enable_fflags_i : 1; // whether fflags_i register present
Bool enable_DAZ : 1; // whether floating point DAZ mode
Bool enable_FZ : 1; // whether floating point FZ mode
Bool mstatus_FS_zero : 1; // whether mstatus.FS hardwired to zero
Bool trap_preserves_lr : 1; // whether trap preserves active LR/SC
Bool xret_preserves_lr : 1; // whether xret preserves active LR/SC
Bool fence_g_preserves_vs : 1; // whether G-stage fence preserves VS-stage
Bool vstart0_non_ld_st : 1; // require vstart 0 for non-load-store?
Bool vstart0_ld_st : 1; // require vstart 0 for load-store?
Bool align_whole : 1; // whole register load aligned to hint?
Bool vill_trap : 1; // trap instead of setting vill?
Bool enable_CSR_bus : 1; // enable CSR implementation bus
Bool mcounteren_present : 1; // force mcounteren to be present
Bool PMP_decompose : 1; // decompose unaligned PMP accesses
Bool PMP_undefined : 1; // unimplemented PMP CSR accesses cause
// exceptions
Bool PMP_maskparams : 1; // enable parameters to change PMP CSR
// read-only masks
Bool PMP_active_inM : 1; // PMP is enforced in M-Mode
Bool PMP_initialparams : 1; // enable parameters to change PMP CSR
// reset values
Bool external_int_id : 1; // enable external interrupt ID ports
Bool tval_zero : 1; // whether [smultval are always zero
Bool tval_zero_ebreak : 1; // whether [smultval always zero on ebreak
Bool tval_ii_code : 1; // instruction bits in [smultval for
// illegal instruction exception?
Bool debug_ctrlt_illegal : 1; // whether control transfer illegal in
// debug mode
Bool debug_auiopc_illegal : 1; // whether instructions using PC illegal
// in debug mode
Bool no_resehaltreq : 1; // resehaltreq gives haltreq reason
Bool defer_step_bug : 1; // defer step breakpoint for one
// instruction when interrupt on return
// from debug mode (hardware bug)
Bool mask_tselect : 1; // writes to tselect are masked based on
// trigger_num setting

// CLIC configuration
Uns64 mclicbase; // base of M-mode CLIC region
Uns64 sclicbase; // base of S-mode CLIC region
Uns64 uclicbase; // base of U-mode CLIC region (N extension)
Bool CLICANDBASIC : 1; // whether implements basic mode also
Bool CLICSELHVEC : 1; // selective hardware vectoring?
Bool CLICXNXTI : 1; // *nxti CSRs implemented?
Bool CLICXCSW : 1; // *scratchcs* CSRs implemented?
Bool externalCLIC : 1; // is CLIC externally implemented?
Bool tvt_undefined : 1; // whether *tv* CSRs are undefined

```

```

Bool  intthresh_undefined : 1;    // whether *intthresh CSRs undefined
Bool  mclicbase_undefined : 1;    // whether mclicbase CSR is undefined
Bool  CSIP_present       : 1;    // whether CSIP interrupt is present
Bool  no_clic_tv_align   : 1;    // no CLIC mode alignment restriction
                                   // for xtvec/xtvt

Bool  posedge_other      : 1;    // fixed int[64:N] positive edge
Bool  poslevel_other     : 1;    // fixed int[64:N] positive level
Uns64 posedge_0_63;        // fixed int[63:0] positive edge
Uns64 poslevel_0_63;      // fixed int[63:0] positive level
Uns32 CLICLEVELS;        // number of CLIC interrupt levels
Uns16 nlbits_valid;      // mask of valid cliccfg.nlbits values
Uns8  CLICVERSION;       // CLIC version
Uns8  CLICINTCTLBITS;    // bits implemented in clicintctl[i]
Uns8  CLICCFGMBITS;      // bits implemented for cliccfg.nmbits
Uns8  CLICCFGLBITS;      // bits implemented for cliccfg.nlbits
Uns8  INTTHRESHBITS;     // bits implemented in xintthresh CSRs

// Hypervisor configuration
Uns8  GEILEN;            // number of guest external interrupts
Bool  xtinst_basic;      // only pseudo-instruction in xtinst

// CSR configurable reset register values
struct {
    CSR_REG_DECL (mvendorid);    // mvendorid value
    CSR_REG_DECL (marchid);      // marchid value
    CSR_REG_DECL (mimpid);       // mimpid value
    CSR_REG_DECL (mhartid);      // mhartid value
    CSR_REG_DECL (mconfigptr);   // mconfigptr value
    CSR_REG_DECL (mtvec);        // mtvec value
    CSR_REG_DECL (mstatus);      // mstatus reset value
    CSR_REG_DECL (mseccfg);      // mseccfg value
    CSR_REG_DECL_0_15(pmpcfg);   // pmpcfg values
    CSR_REG_DECL_0_63(pmpaddr);  // pmpaddr values
} csr;

// CSR configurable register masks
struct {
    CSR_REG_DECL (mtvec);        // mtvec mask
    CSR_REG_DECL (stvec);        // stvec mask
    CSR_REG_DECL (utvec);        // utvec mask
    CSR_REG_DECL (mtvt);        // mtvt mask
    CSR_REG_DECL (stvt);        // stvt mask
    CSR_REG_DECL (utvt);        // utvt mask
    CSR_REG_DECL (jvt);         // jvt mask
    CSR_REG_DECL (tdata1);      // tdata1 mask
    CSR_REG_DECL (mip);         // mip mask
    CSR_REG_DECL (sip);         // sip mask
    CSR_REG_DECL (uip);         // uip mask
    CSR_REG_DECL (hip);         // hip mask
    CSR_REG_DECL (mvien);       // mvien mask
    CSR_REG_DECL (mvip);        // mvip mask
    CSR_REG_DECL (hvien);       // hvien mask
    CSR_REG_DECL (hvip);        // hvip mask
    CSR_REG_DECL (envcfg);      // envcfg mask
    CSR_REG_DECL_0_15(romask_pmpcfg); // pmpcfg read-only bit masks
    CSR_REG_DECL_0_63(romask_pmpaddr); // pmpaddr read-only bit masks
} csrMask;

// custom documentation
const char    **specificDocs;    // custom documentation
riscvDocFn    restrictionsCB;    // custom restrictions

// extension configuration information
riscvExtConfigCPP extensionConfigs; // null-terminated list of extension
                                   // configurations
} riscvConfig;

```

Structures of this type should be defined in file `riscvConfigList.c` in the linked model, as shown in section 4. *Always use field initialization by name (as shown in that section) when specifying field values, to avoid any dependency on field order.*

Fields in this structure are described below, in functional groups. Having copied the template model, modify any custom definitions required by referring to these tables.

5.1 Fundamental Configuration

Fields here are applicable to all RISC-V variants. All field defaults can be overridden using a model parameter of the same name unless otherwise specified.

| Field | Type | Description |
|--------------|-------------------|---|
| name | String | This is the name of the variant and must always be specified as a non-null string. It is used to select this configuration when it matches the <code>variant</code> parameter specified for the processor instantiation. |
| arch | riscvArchitecture | This specifies the processor architecture (whether it is 32 or 64 bit, and which standard extensions are present). The value is a bitmask enumeration with values defined in file <code>riscvVariant</code> in the base model. The default value is typically formed by bitwise-or of values from this file, for example, the value: <code>ISA_U RV32GC ISA_X</code> specifies an RV32 processor with I, M, A, C, F, D, U and custom extensions (X). The value should include B and K if the Bit Manipulation or Cryptographic extensions are present, <i>even though those do not appear in the misa CSR</i> . The value in a processor instance can be indirectly modified by parameters <code>misa_MXL</code> , <code>misa_Extensions</code> , <code>add_Extensions</code> and <code>sub_Extensions</code> . |
| archImplicit | riscvArchitecture | This specifies architecture bits for extension features that are implemented and always enabled, but not visible in the <code>misa</code> CSR. Standard extensions that originally were <code>misa</code> -controlled but are now regarded as fundamental (e.g. the B and K extensions) could be specified here, but the model will also set this automatically from equivalent <code>arch</code> bits, based on the specified versions of those extensions, so this should rarely be required to be explicitly set in an extension. The value in a processor instance can be indirectly modified by parameters <code>add_implicit_Extensions</code> and <code>sub_implicit_Extensions</code> . |
| archMask | riscvArchitecture | This specifies writable architecture bits in the <code>misa</code> CSR. Non-writable bits will be fixed to 1 or 0 depending on the value in the <code>arch</code> field. The value in a processor instance can be indirectly modified by parameters <code>misa_MXL_mask</code> , <code>misa_Extensions_mask</code> , <code>add_Extensions_mask</code> and <code>sub_Extensions_mask</code> . |
| archFixed | riscvArchitecture | This specifies bits in the <code>misa</code> CSR that may <i>never</i> be overridden by parameters. For example, a value of <code>ISA_V</code> means that the vector extension may never be configured or deconfigured. It is usually easier to specify as a bitwise-negation of features that <i>may</i> be configured or deconfigured. This field cannot be overridden by a parameter. |
| user_version | riscvUserVer | This specifies the implemented Unprivileged Architecture version, defined by the <code>riscvUserVer</code> enumeration: <pre>typedef enum riscvUserVerE { RVUV_2_2, // 2.2 RVUV_2_3, // 2.3 RVUV_20190305, // 20190305 RVUV_20191213, // 20191213 RVUV_DEFAULT = RVUV_20191213, } riscvUserVer;</pre> |
| priv_version | riscvPrivVer | This specifies the implemented Privileged Architecture version, defined by the <code>riscvPrivVer</code> enumeration: <pre>typedef enum riscvPrivVerE { RVPV_1_10, // 1.10 RVPV_1_11, // 1.11</pre> |

| | | |
|-----------------|-----------|---|
| | | <pre> RVPV_20190405, // 20190405 RVPV_20190608, // 20190608 RVPV_20211203, // 20211203 RVPV_1_12, // 1.12 RVPV_MASTER, RVPV_DEFAULT = RVPV_1_12 } riscvPrivVer; </pre> |
| endian | memEndian | This specifies the initial endianness as reported in <code>mstatus.MBE</code> and equivalent fields for other privilege modes. |
| endianFixed | Bool | This specifies that MBE, SBE and UBE fields in <code>mstatus</code> are read-only (data endianness is fixed). This parameter only has effect for privileged version <code>RVPC_1_12</code> and later. |
| noZicsr | Bool | This specifies whether <code>zicsr</code> is <i>absent</i> (the negation of parameter <code>zicsr</code>). If the field is <code>True</code> , then no CSRs or CSR access instructions are defined, and an alternative privileged scheme must be implemented in the extension. |
| enable_expanded | Bool | The <i>RISC-V Unprivileged ISA</i> specification outlines a scheme to support instructions greater than 32 bits in length. By default, such instructions are Illegal Instructions in the base model, but this field can be set to <code>True</code> to enable support for 48-bit and 64-bit instructions when queried by the <code>vmicxtFetch</code> decoder function. This is considered so fundamental that the field may only be set directly in a processor configuration structure (there is no parameter to override this). Instruction lengths longer than 64 bits are not currently supported. |

5.2 Interrupts and Exceptions

Fields here control interrupt and exception state. All field defaults can be overridden using a model parameter of the same name unless otherwise specified.

| Field | Type | Description |
|---------------------|--------------|--|
| riscvRNMIver | riscvRNMIver | This specifies the implemented Resumable NMI extension version, defined by the <code>riscvRNMIver</code> enumeration: <pre>typedef enum riscvRNMIverE { RNMI_NONE, // RNMI not implemented RNMI_0_2_1, // RNMI 0.2.1 RNMI_0_4_NMIE1, // RNMI version 0.4, but // nmie=1 // at reset RNMI_0_4, // RNMI 0.4 } riscvRNMIver;</pre> Use <code>RNMI_NONE</code> if this extension is not implemented. |
| reset_address | Uns64 | This specifies the address to jump to on reset. |
| nmi_absent | Bool | This specifies that the core does not implement NMI ports or exceptions. |
| nmi_address | Uns64 | This specifies the address to jump to on NMI. This is ignored if <code>nmi_absent</code> is <code>True</code> . |
| nmiexc_address | Uns64 | If the Resumable NMI extension is implemented, this specifies the address to jump to to handle an exception when an NMI is active. |
| nmi_is_latched | Bool | This indicates whether the NMI input is posedge-latched (if <code>True</code>) or level-sensitive (if <code>False</code>). If the NMI is latched, the latch is cleared when the NMI is taken. This is ignored if <code>nmi_absent</code> is <code>True</code> . |
| nmi_high_priority | Bool | This indicates whether the NMI input is higher priority than Debug and Trigger Module events (if <code>True</code>) or lower priority (if <code>False</code>). This is ignored if <code>nmi_absent</code> is <code>True</code> . |
| nmi_update_mstatus | Bool | If the Resumable NMI extension is NOT implemented, setting this to <code>True</code> means that when an NMI exception is taken the <code>mstatus</code> CSR <code>mpie</code> field is set to the value of <code>mie</code> and <code>mie</code> is set to 0. When false, <code>mstatus</code> is not updated upon taking an NMI exception. |
| nmi_update_tcontrol | Bool | If the Resumable NMI extension is NOT implemented, setting this to <code>True</code> means that when an NMI exception is taken the <code>tcontrol</code> CSR <code>mpie</code> field is set to the value of <code>mte</code> and <code>mte</code> is set to 0. When false, <code>tcontrol</code> is not updated upon taking an NMI exception. |
| nmi_zero_mtval | Bool | If the Resumable NMI extension is NOT implemented, and this is <code>True</code> , upon an NMI exception the <code>mtval</code> CSR is set to 0. When false, <code>mtval</code> is not updated upon taking an NMI exception. |
| unimp_int_mask | Uns64 | This is a bitmask specifying interrupts that are <i>not</i> implemented. It determines the net ports created and the writable values of some CSRs (e.g. <code>mie</code> , <code>mideleg</code>). |
| force_mideleg | Uns64 | This is a bitmask specifying interrupts that are always delegated from M-mode to lower execution levels. |
| force_sideleg | Uns64 | This is a bitmask specifying interrupts that are always delegated from S-mode to lower execution levels. |
| no_ideleg | Uns64 | This is a bitmask specifying interrupts that can never be delegated to lower execution levels. |
| no_edeleg | Uns64 | This is a bitmask specifying exceptions that can never be delegated to lower execution levels. |
| ecode_mask | Uns64 | This is a mask specifying writable bits in <code>xcause.ecode</code> . |
| ecode_nmi | Uns64 | This defines the cause reported by an NMI interrupt. This is ignored if <code>nmi_absent</code> is <code>True</code> . |
| local_int_num | Uns32 | This defines the number of local interrupts implemented. For RV32, the maximum number is 16 and for RV64 the maximum number is 48. If the CLIC is implemented, the maximum number |

| | | |
|-------------------|-------|--|
| | | is 4080 in both cases. |
| tvec_align | Uns32 | This specifies the hardware-enforced alignment of <i>xtvec</i> CSRs in <i>non-direct</i> mode (when the least-significant two bits of the <i>xtvec</i> CSR are not 00). For example, a value of 64 implies that <i>mtvec</i> is masked to enforce 64-byte alignment. It has no effect in direct mode (when the least-significant two bits of the <i>xtvec</i> CSR are 00). |
| mtval_is_ro | Bool | This specifies whether <i>mtval</i> is a read-only register (note that when <i>tval_zero</i> is true this does not need to be set, since <i>mtval</i> is hardwired to 0 in that case.) |
| mtvec_is_ro | Bool | This specifies whether <i>mtvec</i> is a read-only register. |
| trap_preserves_lr | Bool | This specifies whether a trap preserves any active LR/SC transaction state (if <i>False</i> , the LR/SC is aborted). |
| xret_preserves_lr | Bool | This specifies whether an <i>xret</i> instruction preserves any active LR/SC transaction state (if <i>False</i> , the LR/SC is aborted). |
| tval_zero | Bool | This specifies whether <i>xtval</i> registers are always zero. |
| tval_ii_code | Bool | This specifies whether <i>xtval</i> registers are filled with the instruction value for Illegal Instruction exceptions. If <i>False</i> , <i>xtval</i> registers are zeroed instead. |
| csrMask.mtvec | Uns64 | This specifies writable bits in the <i>mtvec</i> CSR. Use parameter <i>mtvec_mask</i> to override this value. |
| csrMask.stvec | Uns64 | This specifies writable bits in the <i>stvec</i> CSR (when Supervisor mode is implemented). Use parameter <i>stvec_mask</i> to override this value. |
| csrMask.utvec | Uns64 | This specifies writable bits in the <i>utvec</i> CSR (when N extension is implemented). Use parameter <i>utvec_mask</i> to override this value. |
| csrMask.mip | Uns64 | <i>mip</i> CSR write mask. Use parameter <i>mip_mask</i> to override this value. |
| csrMask.sip | Uns64 | <i>sip</i> CSR write mask (when Supervisor mode is implemented). Use parameter <i>sip_mask</i> to override this value. |
| csrMask.uip | Uns64 | <i>uip</i> CSR write mask (when N extension is implemented). Use parameter <i>uip_mask</i> to override this value. |
| csrMask.hip | Uns64 | <i>hip</i> CSR write mask (when H extension is implemented). Use parameter <i>hip_mask</i> to override this value. |
| mtvec_sext | Bool | This specifies whether the <i>mtvec</i> CSR is sign-extended from the most significant writable bit. |
| stvec_sext | Bool | This specifies whether the <i>stvec</i> CSR is sign-extended from the most significant writable bit (when Supervisor mode is implemented). |
| utvec_sext | Bool | This specifies whether the <i>utvec</i> CSR is sign-extended from the most significant writable bit (when N extension is implemented). |

5.2.1 Reset

The RISC-V Privileged Architecture Manual specifies a small state set that must be updated on a reset event. The base RISC-V model will automatically perform these actions at reset:

1. The hart will be restarted if it is in a halted or WFI state.
2. The hart will be forced into Machine mode.
3. If PMP registers are present, these will be reset to the state specified in the configuration (see section 5.9).
4. If *Smstateen* CSRs are present, these will be reset as described in that extension.

5. The `mcause` register will be zeroed.
6. If the vector extension is present, the `vtype` and `vl` CSRs will be reset.
7. The `misra` CSR will be reset to indicate all configured extensions are enabled.
8. If the Trigger Module is implemented, all triggers are reset to their initial state.
9. If Debug mode is implemented, the `dcsr` CSR is reset to its initial state.
10. Any LR/SC exclusive access is cleared.
11. If the AIA extension is present, all `miprio`, `siprio` and `vsiprio` state is reset.
12. If a CLINT is present and internally implemented, the `msip` CSR will be cleared.
13. If a CLIC is present and implemented by the model, all interrupts will be configured as M-mode with priority 255 and `mintstatus` will be cleared.
14. If implemented, timers will be zeroed and pending timer interrupts cleared.
15. The hart will start executing at the configured reset address (`reset_address`).

No other state changes are made at reset. Extensions may specify additional behavior using reset notifiers (see sections 15.11 and 15.12).

5.3 Timers and Timer Interrupts

Fields here define model aspects related to timers and timer interrupts. All field defaults can be overridden using a model parameter of the same name.

| Field | Type | Description |
|----------------|-------|--|
| time_undefined | Bool | This specifies whether the <code>time</code> CSR is undefined (its behavior must be emulated by a trap). |
| mtime_Hz | Flt64 | If the <code>time</code> CSR is implemented (<code>time_undefined</code> is <code>False</code>), this specifies the frequency of the timer (in Hertz). This parameter is also used if the internal CLINT model is implemented (see section 5.5.4). |
| Sstc | Bool | This specifies whether the <code>stimecmp/stimecmph</code> CSRs are implemented and, if Hypervisor mode is present, whether the <code>vstimecmp/vstimecmph</code> CSRs are implemented. |

5.4 Instruction and CSR Behavior

Fields here define instruction and CSR behavior. All field defaults can be overridden using a model parameter of the same name unless otherwise specified.

| Field | Type | Description |
|---------------------|-------|---|
| wfi_is_nop | Bool | This specifies whether the WFI instruction is treated as a NOP (if False, it will cause execution of the current core to suspend waiting for an interrupt). |
| noZaamo | Bool | When the atomic extension (A) is implemented, this specifies whether the Zaamo instruction subset is <i>absent</i> . This is the negation of the Zaamo parameter value. |
| noZalrsc | Bool | When the atomic extension (A) is implemented, this specifies whether the Zalrsc instruction subset is <i>absent</i> . This is the negation of the Zalrsc parameter value. |
| Zacas | Bool | When the atomic extension (A) is implemented, this specifies whether the Zacas instruction subset is implemented. |
| Zabha | Bool | When the atomic extension (A) is implemented, this specifies whether the Zabha instruction subset is implemented. |
| Zawrs | Bool | This specifies whether the Zawrs extension is implemented. |
| Zimop | Bool | This specifies whether the Zimop extension is implemented. |
| Zcmop | Bool | This specifies whether the Zcmop extension is implemented. |
| Zicfiss | Bool | This specifies whether the Zicfiss extension is implemented. |
| Zicfilp | Bool | This specifies whether the Zicfilp extension is implemented. |
| Zmmul | Bool | If the M extension is present, this specifies that only multiply instructions are implemented (not divide or remainder). |
| Zihintntl | Bool | This specifies whether Zihintntl instructions are decoded (treated as NOPs by the model) |
| Zicond | Bool | This specifies whether Zicond is present. |
| noZifencei | Bool | This specifies whether Zifencei is <i>absent</i> (the negation of parameter Zifencei). If the field is True, instruction fence.i is undefined. |
| Smstateen | Bool | This specifies whether xstateen state enable CSRs are implemented. |
| Smcsrind | Bool | This specifies whether Smcsrind is present (and also Sscsrind if Supervisor mode is implemented). |
| Smcdeleg | Bool | This specifies whether the Smcdeleg extension is implemented (and also Ssccfg). |
| Sscofpmf | Bool | This specifies whether the Sscofpmf extension is implemented. |
| Smcntrpmf | Bool | This specifies whether the Smcntrpmf extension is implemented. |
| Ssqosid | Bool | This specifies whether the Ssqosid extension is implemented. |
| TW_time_limit | Uns32 | If wfi_is_nop is False or Zawrs is True, this specifies the number of simulated cycles that will pass when a hart is stalled by WFI or WRS.NTO before an Illegal Instruction or Virtual Instruction trap is taken. |
| wfi_resume_not_trap | Bool | If wfi_is_nop is True or TW_time_limit is zero and a WFI restart event is pending when the WFI is executed, this specifies that the WFI should be treated as a NOP if it would otherwise trap because mstatus.TW=1. |
| STO_time_limit | Uns32 | If Zawrs is True, this specifies the number of simulated cycles that will pass when a hart is stalled by WRS.STO before the hart resumes execution. |
| cycle_undefined | Bool | This specifies whether the cycle CSR is undefined (its behavior must be emulated by a trap). |
| mcycle_undefined | Bool | This specifies whether the mcycle CSR is undefined (its behavior must be emulated by a trap). |
| time_undefined | Bool | This specifies whether the time CSR is undefined (its behavior must be emulated by a trap). If present, the frequency of the timer (in Hertz) is given by mtime_Hz (see section 5.3). |

| | | |
|-----------------------|-------|---|
| instret_undefined | Bool | This specifies whether the instret CSR is undefined (its behavior must be emulated by a trap). |
| minstret_undefined | Bool | This specifies whether the minstret CSR is undefined (its behavior must be emulated by a trap). |
| hpmcounter_undefined | Bool | This specifies whether the hpmcounter* CSRs are undefined (their behavior must be emulated by a trap). |
| mhpmcounter_undefined | Bool | This specifies whether the mhpmcounter* and mhpmevent* CSRs are undefined (their behavior must be emulated by a trap). |
| counteren_mask | Uns32 | This is a bitmask specifying writable bits in mcounteren, scounteren and hcounteren registers. Note that writable bits in scounteren and hcounteren may be further modified by scounteren_zero_mask and hcounteren_zero_mask (see below). |
| scounteren_zero_mask | Uns32 | This is a bitmask specifying bits in scounteren that are always zero even if they are indicated as writable in mcounteren by counteren_mask (see above). |
| hcounteren_zero_mask | Uns32 | This is a bitmask specifying bits in hcounteren that are always zero even if they are indicated as writable in mcounteren by counteren_mask (see above). |
| noinhibit_mask | Uns32 | This is a bitmask specifying counters that may <i>not</i> be inhibited using mcountinhibit (in other words, 1 bits in this mask are always 0 in mcountinhibit). |

5.5 CLIC

Fields here define the *Core Local Interrupt Controller* (CLIC) configuration. All field defaults can be overridden using a model parameter of the same name if the field is applicable – see the description below of when each field is used.

The CLIC is implemented if field `CLICLEVELS` is non-zero. When implemented, the behavior can either be modeled by an *internal* memory-mapped component or *externally*, depending on the value of field `externalCLIC`. Depending on the settings of these two fields, more fields are used and parameters made available to further configure the CLIC behavior; these are described in separate tables below.

5.5.1 CLIC Fundamental Fields

| Field | Type | Description |
|---------------------------|-------|--|
| <code>CLICLEVELS</code> | Uns32 | If zero, this specifies that the CLIC is unimplemented; otherwise it must be a number in the range 2-256, specifying the number of levels implemented. |
| <code>externalCLIC</code> | Bool | If the CLIC is implemented (<code>CLICLEVELS != 0</code>), this specifies whether the <i>internal</i> CLIC model is used or whether the CLIC is modeled by an <i>external</i> memory-mapped component. If implemented externally, five new net ports are created that must be connected to the external CLIC model: <ol style="list-style-type: none">1. <code>irq_id_i</code>: input port written with highest-priority pending interrupt;2. <code>irq_lev_i</code>: input port written with level of highest-priority pending interrupt;3. <code>irq_sec_i</code>: input port written with execution level of highest-priority pending interrupt;4. <code>irq_shv_i</code>: input port written with indication of whether the highest-priority pending interrupt uses selective hardware vectoring;5. <code>irq_i</code>: active-high input indicating that an external CLIC interrupt is pending. |

5.5.2 CLIC Common Fields

These fields are used when the CLIC is implemented (`CLICLEVELS!=0`), whether internally or externally. All field defaults can be overridden using a model parameter of the same name unless otherwise specified.

| Field | Type | Description |
|----------------------------|-------|---|
| CLICVERSION | Uns32 | This specifies the CLIC version. |
| CLICXNXTI | Bool | This specifies whether <code>xnxti</code> CSRs are implemented. |
| CLICXCSW | Bool | This specifies whether <code>xscratchcsx</code> CSRs are implemented. |
| tv _t _undefined | Bool | This specifies whether <code>xtvt</code> CSRs are implemented – if <code>False</code> , then <code>xtvec</code> registers are used instead. |
| intthresh_undefined | Bool | This specifies whether <code>xintthresh</code> CSRs are undefined. |
| INTTHRESHBITS | Uns8 | This specifies the number of writable bits in <code>xintthresh</code> CSRs (if defined). |
| mclicbase_undefined | Bool | This specifies whether the <code>mclicbase</code> CSR is undefined. |
| csrMask.mtv _t | Uns64 | This specifies writable bits in the <code>mtvt</code> CSR. Use parameter <code>mtvt_mask</code> to override this value. |
| csrMask.stv _t | Uns64 | This specifies writable bits in the <code>stvt</code> CSR (if Supervisor mode is implemented). Use parameter <code>stvt_mask</code> to override this value. |
| csrMask.utv _t | Uns64 | This specifies writable bits in the <code>utvt</code> CSR (if the N extension is implemented). Use parameter <code>utvt_mask</code> to override this value. |
| mtvt_sext | Bool | This specifies whether the <code>mtvt</code> CSR is sign-extended from the most significant writable bit. |
| stvt_sext | Bool | This specifies whether the <code>stvt</code> CSR is sign-extended from the most significant writable bit (if Supervisor mode is implemented). |
| utvt_sext | Bool | This specifies whether the <code>utvt</code> CSR is sign-extended from the most significant writable bit (if the N extension is implemented). |

5.5.3 CLIC Internal Fields

These fields are used when the CLIC is implemented *internally* (CLICLEVELS!=0 and externalCLIC=False).

| Field | Type | Description |
|------------------|--------------|---|
| CLIC_version | riscvCLICVer | This specifies the implemented CLIC version, defined by the <code>riscvCLICVer</code> enumeration: <pre>typedef enum riscvCLICVerE { RVCLC_20180831, // 20180831 RVCLC_0_9_20191208, // 0.9-draft-20191208 RVCLC_0_9_20220315, // 0.9-draft-20221315 RVCLC_0_9_20221108, // 0.9-draft-20221108 RVCLC_0_9_20230801, // 0.9-draft-20230801 RVCLC_MASTER, RVCLC_DEFAULT = RVCLC_0_9_20230801 } riscvCLICVer;</pre> |
| CLICANDBASIC | Bool | This specifies whether basic interrupt control is also implemented. |
| CLICINTCTLBITS | Uns32 | This specifies the number of bits implemented in <code>clicintctl[i]</code> . |
| CLICCFGMBITS | Uns32 | This specifies the maximum value to which <code>cliccfg.nmbits</code> may be set. Attempts to set larger values are clamped to this maximum. |
| CLICCFGLBITS | Uns32 | This specifies the maximum value to which <code>cliccfg.nlbits</code> may be set. Attempts to set larger values are clamped to this maximum. |
| mclicbase | Uns64 | mclicbase value |
| sclicbase | Uns64 | sclicbase value (if Supervisor mode is implemented) |
| uclicbase | Uns64 | uclicbase value (if the N extension is implemented) |
| nlbits_valid | Uns16 | This is a bitmask of valid values for <code>cliccfg.nlbits</code> . For example, <code>cliccfg.nlbits</code> may only hold the value 8 if bit 8 is set in <code>nlbits_valid</code> . Attempts to set <code>cliccfg.nlbits</code> to values not selected by this bitmask will be rejected, and the previous value of the field will be retained. |
| CLICSELHVEC | Bool | This specifies whether <i>selective hardware vectoring</i> is supported. |
| CSIP_present | Bool | This specifies whether the optional positive-edge-triggered CSIP interrupt is present. Depending on the CLIC version selected, this may be either ID 12 or ID 16. |
| no_clic_tv_align | Bool | This specifies whether the alignment rules for <code>xtvec.base</code> and <code>xtvt</code> required by the <code>sxclic</code> extensions are implemented. When <code>False</code> the compliant behavior of forcing at least 64 byte alignment when CLIC mode is active for <code>xtvec.base</code> and <code>xtvt</code> is implemented. |
| posedge_0_63 | Uns64 | This mask specifies interrupts in the range 0 to 63 that have fixed positive edge-sensitive configuration. |
| poslevel_0_63 | Uns64 | This mask specifies interrupts in the range 0 to 63 that have fixed positive level-sensitive configuration. |
| posedge_other | Bool | This specifies whether all interrupts with index 64 and higher have fixed positive edge-sensitive configuration. |
| poslevel_other | Bool | This specifies whether all interrupts with index 64 and higher have fixed positive level-sensitive configuration. |

5.5.4 Custom registers in the CLIC MMIO custom region

Behavior for registers in the CLIC MMIO custom region may be implemented by an extension by registering a notifier function in the constructor with the `clicCustomRd/clicCustomWr` member of the `riscvExtCB` struct. These notifiers will be

called when the custom region of the CLIC MMIO address space, at offset 0x800 to 0xFFF, is read or written.

See Appendix sections 15.39 and 15.40 for details.

5.6 CLINT

Fields here define the *Core Local Interruptor* (CLINT) configuration. This block models a legacy SiFive-specific component and is not intended for general use. All field defaults can be overridden using a model parameter of the same name.

| Field | Type | Description |
|---------------|-------|---|
| CLINT_address | Uns64 | If zero, this specifies that the CLINT is unimplemented, otherwise it specifies the address of the CLINT block. |
| mtime_Hz | Flt64 | If the CLINT is implemented (CLINT_address!=0), this specifies the frequency of the CLINT <code>mtime</code> counter. |

5.7 AIA

Fields here define the *Advanced Interrupt Architecture* (AIA) configuration. All field defaults can be overridden using a model parameter of the same name unless otherwise specified.

Note that the model does not implement an interrupt controller component (e.g. APLIC or IMSIC): if such components are required, they must be implemented in the platform and integrated using the interface as described below.

| Field | Type | Description |
|------------------|-------------|--|
| Smaia | Bool | This specifies that the AIA is implemented. Subsequent fields in this table are ignored unless Smaia is True. |
| IMSIC_present | Bool | This indicates whether an <i>Incoming MSI Controller</i> (IMSIC) is implemented in the platform. If True, then xstopei CSRs are enabled and an IMSIC bus port is present which must be connected to the externally-implemented IMSIC. See below for information about connecting such a model. |
| AIA_version | riscvAIAVer | This specifies the implemented Smaia extension version, defined by the riscvAIAVerVer enumeration: <pre>typedef enum riscvAIAVerE { RVAIA_1_0_RC1, // 1.0-RC1 RVAIA_1_0_RC3, // 1.0-RC3 RVAIA_1_0_RC5, // 1.0-RC5 RVAIA_1_0, // 1.0 ratified version RVAIA_MASTER, RVAIA_DEFAULT = RVAIA_1_0 } riscvAIAVer;</pre> |
| miprio_mask | Uns64 | This specifies which M-mode IPRIO array entries are writable. An IPRIO entry with index <i>i</i> is writable only if bit <i>i</i> of miprio_mask is set. |
| siprio_mask | Uns64 | If Supervisor mode is present, this specifies which S-mode IPRIO array entries are writable. An IPRIO entry with index <i>i</i> is writable only if bit <i>i</i> of siprio_mask is set. |
| hviprio_mask | Uns64 | If the Hypervisor extension is present, this specifies which entries in hviprio CSRs are writable. An hviprio CSR entry corresponding to interrupt <i>i</i> is writable only if bit <i>i</i> of hviprio_mask is set. |
| IPRIOLEN | Uns8 | This specifies the number of priority bits implemented for external interrupts (in the range 1 to 8). |
| HIPRIOLEN | Uns8 | If the Hypervisor extension is present, this specifies the number of priority bits implemented for virtual interrupts in hviprio registers (in the range 6 to 8). |
| hviectl_IID_bits | Uns8 | If the Hypervisor extension is present, this specifies the number of bits implemented in the hviectl.IID field (in the range 6 to 12). |
| csrMask.mvip | Uns64 | This specifies writable bits in the mvip CSR. Bits 12:0 are ignored and always zero. Use parameter mvip_mask to override this value. |
| csrMask.mvien | Uns64 | This specifies writable bits in the mvien CSR. Bits 12:0 are ignored and always zero. Bits that are set in mvip_mask but clear in mvien_mask will be hard-wired to 1 in mvien. Use parameter mvien_mask to override this value. |
| csrMask.hvip | Uns64 | hvip CSR write mask (only bits 63:13 used). Bits 12:0 are ignored and always zero. Use parameter hvip_mask to override this value. |
| csrMask.hvien | Uns64 | This specifies writable bits in the hvien CSR. Bits 12:0 are ignored and always zero. Bits that are set in hvip_mask but clear in hvien_mask will be hard-wired to 1 in hvien. Use parameter hvien_mask to override this value. |

5.7.1 AIA Net Port Interface

When the AIA is implemented, net ports are present allowing the external interrupt controller model to supply the highest-priority pending interrupt. These are:

`miprio`: this input should be written with the id of the highest-priority pending M-mode interrupt from the external interrupt controller model.

`siprio`: this input is present if Supervisor mode is implemented. It should be written with the id of the highest-priority pending S-mode interrupt from the external interrupt controller model.

`vsiprio`: this input is present if the Hypervisor extension is implemented. It should be written with the id of the highest-priority pending VS-mode interrupt from the external interrupt controller model.

5.7.2 IMSIC Interrupt Controller Integration

This model does not implement an interrupt controller to drive the AIA interface: this must be implemented as a platform component and connected to the processor model. For an APLIC, use the standard hart interrupt net ports plus the priority ports described in the previous section. When an IMSIC is implemented (`IMSIC_present` is `True`), additional work is required, as follows:

1. Connect a 16-bit bus to the artifact `CSR` bus of the hart (this bus port is present when `IMSIC_present` is `True`). This will be used to observe and react to changes to IMSIC-related CSRs.
2. Connect a second 16-bit bus to the artifact `IMSIC` bus of the hart (also present when `IMSIC_present` is `True`). This bus is used to implement IMSIC registers.
3. Install watchpoints on the `CSR` bus to observe writes to `mtopei`, `stopei` and `vstopei` CSRs. Writes made by the hart must cause the highest-priority pending interrupt of the indicated type to be acknowledged by the IMSIC. These CSRs have indices `0x35C`, `0x15C` and `0x25C` respectively, meaning write callbacks must be installed at addresses `0x35C0`, `0x15C0` and `0x25C0` on the `CSR` artifact bus.
4. If the Hypervisor extension is implemented, also install a watchpoint on the `CSR` bus to observe writes to the `hstatus` CSR (`0x600`). The value written to `hstatus.VGIEN` selects the guest interrupt file that should be reported by the IMSIC via the `vsiprio` input and also the IMSIC registers mapped using `vsiselect` and `vsireg` CSRs.
5. Expose IMSIC indirectly-accessed interrupt-file registers to the hart by installing read and write callbacks on the `IMSIC` artifact bus. An IMSIC register with index `0xAB` is mapped on the bus at address `0xMAB0`, where `M` indicates the register mode (`S=1`, `VS=2`, `M=3`); as a concrete example, implementing Machine-mode IMSIC register `0x72` requires installation of 4-byte or 8-byte callbacks at address `0x3720` on the `IMSIC` artifact bus.

5.8 Memory Subsystem

Fields here define memory model features. All field defaults can be overridden using a model parameter of the same name unless otherwise specified.

| Field | Type | Description |
|---------------------|-------------|---|
| lr_sc_grain | Uns32 | This defines the lock granularity for LR and SC instructions. It must be a power of 2 in the range 1..(1<<16). |
| Sv_modes | Uns32 | This is a bitmask specifying supported virtual memory modes in the standard format (for example 1<<8 is Sv39). |
| ASID_bits | Uns32 | This defines the number of bits implemented in the ASID (maximum of 9 for RV32 and 16 for RV64). A value of 0 indicates that ASID is not implemented. |
| updatePTEA | Bool | This indicates whether hardware update of page table entry A bit is always active. See also parameter Svadu below. |
| updatePTED | Bool | This indicates whether hardware update of page table entry D bit is implemented. See also parameter Svadu below. |
| Svadu | Bool | This indicates whether the Svadu extension is implemented (this extension adds CSR controls for hardware update of PTE A/D bits). If Svadu is True, updatePTEA and updatePTED are both forced to be False. |
| ignore_non_leaf_DAU | Bool | This indicates whether the value of D, A and U bits in non-leaf page table entries should be ignored. If False, any of these bits being non-zero will cause a Page Fault trap. |
| Svnapot_page_mask | Uns64 | If non-zero, this indicates that the Svnapot extension is implemented, with contiguous pages of sizes specified by the mask. For example, a value of 0x10000 indicates that 64kB NAPOT contiguous pages are supported. |
| Svpbmt | Bool | This indicates whether the Svpbmt extension is implemented (page-based memory types, specified by bits 62:61 of a page table entry). Note that except for their effect on Page Faults, the encoded memory types do not alter the behavior of this model, which always implements strongly-ordered non-cacheable semantics. |
| Svinval | Bool | This indicates whether the Svinval extension is implemented (fine-grained address-translation cache invalidation instructions <code>sinal.vma</code> , <code>sfence.w.inval</code> and <code>sfence.inval.ir</code> , together with <code>hinal.vvma</code> and <code>hinal.gvma</code> if Hypervisor mode is also present). |
| Ssnpm | riscvPMMode | This indicates whether the Ssnpm extension is implemented (Supervisor mode next-level pointer masking). Supported options are defined by the <code>riscvPMMode</code> enumeration: <pre>typedef enum riscvPMModeE { RVPMD_NONE, // not implemented RVPMD_48, // XLEN-48 RVPMD_57, // XLEN-57 RVPMD_48_57, // XLEN-48 and XLEN-57 } riscvPMMode;</pre> |
| Smnpm | riscvPMMode | This indicates whether the Smnpm extension is implemented (Machine mode next-level pointer masking). Supported options are defined by the <code>riscvPMMode</code> enumeration (see above). |
| Smmnpm | riscvPMMode | This indicates whether the Smmnpm extension is implemented (Machine mode pointer masking). Supported options are defined by the <code>riscvPMMode</code> enumeration (see above). |
| unaligned | Bool | This indicates whether unaligned accesses are supported. |
| unaligned_low_pri | Bool | If an access is both unaligned and would also cause a page fault or access fault, this specifies that the unaligned address fault has higher priority (if False) or lower priority (if |

| | | |
|---------------------|------------------|--|
| | | True) than the page fault or access fault. |
| unaligned_check_ms | Bool | For processors that normally always permit unaligned accesses, this specifies that load address misaligned and store/AMO address misaligned exceptions can still be generated by memory components. |
| unalignedAMO | Bool | This indicates whether unaligned accesses are supported for AMO operations. Use parameter <code>Zam</code> to modify this value. |
| amo_aborts_lr_sc | Bool | This indicates whether an AMO operation aborts any active <code>lr/sc</code> pair. |
| lr_sc_match_size | Bool | This indicates <code>lr/sc</code> data size must match for <code>sc</code> instruction to succeed (RV64 only). |
| Zicbom | Bool | This indicates whether the <code>Zicbom</code> extension is implemented (instructions <code>cbo.clean</code> , <code>cbo.flush</code> and <code>cbo.inval</code>). The instructions may cause traps if used illegally but otherwise are NOPs in this model. |
| cmomp_bytes | Unsl6 | If the <code>Zicbom</code> extension is implemented, this specifies the cache block size for those instructions. |
| Zicbop | Bool | This indicates whether the <code>Zicbop</code> extension is implemented (instructions <code>prefetch.i</code> , <code>prefetch.r</code> and <code>prefetch.w</code>). If implemented, the instructions behave as NOPs in this model. |
| Zicboz | Bool | This indicates whether the <code>Zicboz</code> extension is implemented (instruction <code>cbo.zero</code>). |
| cmoz_bytes | Unsl6 | If the <code>Zicboz</code> extension is implemented, this specifies the cache block size for <code>cbo.zero</code> . |
| amo_constraint | riscvMConstraint | This specifies constraints on memory accesses performed by AMO operations, defined by the <code>riscvMConstraint</code> enumeration (see below for more information about these constraints). <pre>typedef enum riscvMConstraintE { RVMC_NONE, RVMC_USER1, RVMC_USER2, } riscvMConstraint;</pre> |
| lr_sc_constraint | riscvMConstraint | This specifies constraints on memory accesses performed by LR/SC operations, defined by the <code>riscvMConstraint</code> enumeration (see below for more information about these constraints). <pre>typedef enum riscvMConstraintE { RVMC_NONE, RVMC_USER1, RVMC_USER2, } riscvMConstraint;</pre> |
| push_pop_constraint | riscvMConstraint | This specifies constraints on memory accesses performed by PUSH/POP operations, defined by the <code>riscvMConstraint</code> enumeration (see below for more information about these constraints). <pre>typedef enum riscvMConstraintE { RVMC_NONE, RVMC_USER1, RVMC_USER2, } riscvMConstraint;</pre> |

5.8.1 Memory Access Constraints

Memory access constraints specified by `amo_constraint`, `lr_sc_constraint` and `push_pop_constraint` fields allow access for AMO, LR/SC and PUSH/POP instructions to be denied to certain memory regions, respectively. The enumeration members have the following meanings:

1. RVMC_NONE: no access restriction.
2. RVMC_USER1: deny access to memory regions including permission MEM_PRIV_USER1.
3. RVMC_USER2: deny access to memory regions including permission MEM_PRIV_USER2.

Region privileges MEM_PRIV_USER1 and MEM_PRIV_USER2 will typically be defined on certain regions by custom PMA code in a processor extension. For example, in the Andes processor, custom PMA code disallows access by AMO instructions if the PMA region is configured as a NAMO (no AMO) region. Function `refinePMARegionRange` specifies the required privilege as follows:

```
// refine privilege
if(e.MTYP==PMAMT_Black_Hole) {
    *privP = MEM_PRIV_NONE;
} else if(!e.NAMO) {
    *privP = MEM_PRIV_RWX;
} else {
    *privP = MEM_PRIV_RWX|MEM_PRIV_USER1;
}
```

This privilege is then used in function `setDomainPriv` to enforce PMA region permissions:

```
vmirtProtectMemory(dataDomain, low, high, priv, MEM_PRIV_SET);
```

In Andes processor configurations, values of `amo_constraint`, `lr_sc_constraint` and `push_pop_constraint` fields are all set to RVMC_USER1, meaning access to NAMO regions will be denied for all AMO, LR, SC, PUSH and POP instructions.

5.9 PMP Configuration

Fields here define PMP features. All field defaults can be overridden using a model parameter. Most parameters have the same name as the field, with the exception of the `romask_...` fields, which are overridden with a `mask_...` parameter.

| Field | Type | Description |
|----------------------------------|-----------------------------|--|
| <code>PMP_registers</code> | Uns32 | This defines the number of PMP regions that are implemented (maximum of 64). A value of 0 indicates the PMP feature is absent. |
| <code>PMP_crs</code> | Uns32 | This defines the number of PMP address registers (and implies the number of <code>PMP_cfg</code> registers) that are present, including those that are RAZ/WI (maximum of 64). A value of 0 indicates the number of registers defaults to 64 (for <code>priv_version</code> 1.12 and later, and 16 for earlier versions). |
| <code>PMP_decompose</code> | Bool | This indicates whether unaligned PMP accesses are decomposed into individual aligned accesses (thus allowing accesses that straddle region boundaries). |
| <code>PMP_undefined</code> | Bool | This indicates whether accesses to unimplemented PMP registers cause Illegal Instruction traps. If <code>False</code> , then such accesses are ignored. |
| <code>PMP_grain</code> | Uns32 | This defines the minimum granularity for PMP regions in the standard format (0: 4 bytes, 1: 8 bytes, etc). |
| <code>PMP_ROW1</code> | <code>riscvPMPROW1</code> | This defines the behavior of <code>pmpcfg</code> R/W/X fields when an illegal value with R=0 and W=1 is written. Choices are described by the <code>riscvPMPROW1</code> enumeration: <pre>typedef enum riscvPMPROW1E { RVPMP1_RWX_00X, // set R=0, W=0, use new X RVPMP1_RWX_11X, // set R=1, W=1, use new X RVPMP1_RWX_PPX, // revert R and W, use new X RVPMP1_RWX_PPP, // revert R, W and X RVPMP1_RWX_000, // set RWX=000 } riscvPMPROW1;</pre> |
| <code>PMP_active_inM</code> | Bool | When true all PMP checks are enforced in M mode when any PMP region is active. Normally this is only true for regions with the lock bit set. |
| <code>pmpaddr0..63</code> | XLEN | This defines the reset value for the <code>pmpaddr</code> registers. There are 64 separate fields. |
| <code>pmpcfg0..15</code> | Uns64 (even) Uns32 (odd) | This defines the reset value for the <code>pmpcfg</code> registers. There are 15 separate fields. The odd numbered fields are ignored when XLEN is 64. The high order word of even numbered fields is ignored when XLEN is 32. |
| <code>romask_pmpaddr0..63</code> | XLEN | This defines the read-only bit masks for the <code>pmpaddr</code> registers. These masks are applied along with other architecturally defined masks (they do not override other masks.) |
| <code>romask_pmpcfg0..15</code> | XLEN | This defines the read-only bit masks for the <code>pmpcfg</code> registers. These masks are applied along with other architecturally defined masks (they do not override other masks.) |
| <code>Smepmp_version</code> | <code>riscvSmepmpVer</code> | This specifies the implemented <code>Smepmp</code> extension version, defined by the <code>riscvSmepmpVer</code> enumeration: <pre>typedef enum riscvSmepmpVerE { RVSP_NONE, // Smepmp not implemented RVSP_0_9_5, // 0.9.5 RVSP_1_0, // 1.0 RVSP_DEFAULT = RVSP_1_0 } riscvSmepmpVer;</pre> |

The `PMP_registers` field defines the number of PMP regions implemented. Fields for specifying reset and mask values for unimplemented region indexes are ignored.

By default, the maximum number of PMP regions is 64 when `priv_version` is 1.12 or later, and 16 when `priv_version` is earlier than 1.12 and it is a fatal error if the value of `PMP_registers` exceeds this. This can be overridden by the value of `PMP_crs`. When `PMP_crs` is not zero it species the number of PMP address registers that are implemented, which may be any value between 1 and 64, and the value of `PMP_registers` may be set to any value between 0 and `PMP_crs`.

The `pmpcfg` and `romask_pmpcfg` fields behave differently depending on the XLEN that is implemented:

- Fields corresponding to even numbered registers are 64 bits. When XLEN is 32 the upper word is ignored.
- Fields corresponding to odd numbered registers are 32 bits. When XLEN is 64 the corresponding register does not exist and the field is ignored.

The parameters to override the `pmpaddr`, `pmpcfg`, `romask_pmpaddr` and `romask_pmpcfg` fields are not available by default. The parameters to enable them are:

- **PMP_initialparams**
Enables the *reset* value related parameters `pmpaddr<x>` and `pmpcfg<x>`
- **PMP_maskparams**
Enables the *mask* related parameters `mask_pmpaddr<x>` and `mask_pmpcfg<x>`.

Note the parameters to override the `romask_pmpaddr` and `romask_pmpcfg` fields have a slightly different name because they specify a *write* mask, not a read-only mask like the field values. The parameter values are negated before overriding the configuration field. (The field values must be read-only masks to maintain backwards compatibility with existing models that have a default values of 0 for this field.)

5.9.1 Mask Values Available in Integration Registers

The values of the write masks cannot be determined by the usual technique of saving a register's value, writing all 0's followed by all 1's, then reading and restoring the register because the `pmpcfg` entries become locked if the L field is written with a 1.

Therefore the mask values have been made available in integration support registers named `mask_pmpcfg0..15` and `mask_pmpaddr0..63`.

Note that these are write mask values, so are inverted versions of the `romask` config values.

Note that these are simulator-only, not architecturally-defined, registers, and can be accessed only through the debug interface, not from a program executing on the model. But they can be useful for a harness or intercept library to discover which PMP register bits are configured as not writable.

5.10 Floating Point

Fields here are applicable when floating point is implemented (D or F extensions indicated as implemented by the processor `arch` field). All field defaults can be overridden using a model parameter of the same name. Field `fp16_version` may also be set to `RVFP16_IEEE754` by setting parameter `Zfh` to `True` (if parameter `fp16_version` is not also set).

| Field | Type | Description |
|------------------------------|----------------------------|---|
| <code>mstatus_fs_mode</code> | <code>riscvFSMode</code> | This field specifies how the implementation-dependent <code>mstatus.FS</code> field is updated by the model. The possible behaviors are specified by the <code>riscvFSMode</code> enumeration (see section 5.10.2 below): <pre>typedef enum riscvFSModeE { RVFS_WRITE_NZ, RVFS_WRITE_ANY, RVFS_EXECUTE_NOT_S, RVFS_EXECUTE_ANY, RVFS_ALWAYS_DIRTY, RVFS_FORCE_DIRTY } riscvFSMode;</pre> |
| <code>d_requires_f</code> | <code>Bool</code> | This field specifies whether <code>misa.D</code> can only be set if <code>misa.F</code> is also set. |
| <code>Zfinx_version</code> | <code>riscvZfinxVer</code> | This field specifies whether, the <code>Zfinx</code> option is specified, and if so with what version. The possible values are specified by the <code>riscvZfinxVer</code> enumeration: <pre>typedef enum riscvZfinxVerE { RVZFINX_NA, // Zfinx not implemented RVZFINX_0_4, // Zfinx 0.4 RVZFINX_0_41, // Zfinx 0.41 RVZFINX_1_0, // Zfinx 1.0 RVZFINX_DEFAULT = RVZFINX_1_0, } riscvZfinxVer;</pre> |
| <code>fp16_version</code> | <code>riscvFP16Ver</code> | This field allows a format to be selected for 16-bit floating point. Values are defined by the <code>riscvFP16Ver</code> enumeration: <pre>typedef enum riscvFP16VerE { RVFP16_NA, // 16-bit FP not supported RVFP16_IEEE754, // IEEE half-precision format RVFP16_BFLOAT16, // BFLOAT16 format RVFP16_DYNAMIC // dynamic (IEEE 754 or BF16) } riscvFP16Ver;</pre> If <code>RVFP16_DYNAMIC</code> is selected, the format to be used for 16-bit floating point can be dynamically changed at runtime using artifact register <code>fp16Format</code> . |
| <code>Zfa</code> | <code>Bool</code> | This field specifies whether additional floating point instructions in the <code>Zfa</code> extension are implemented. |
| <code>Zfhmin</code> | <code>Bool</code> | If half-precision floating point is present (with format indicated by <code>fp16_version</code>), this field indicates whether only a minimal half-precision subset is implemented. If <code>False</code> , half-precision instructions are fully supported. |
| <code>Zfbfmin</code> | <code>Bool</code> | This field indicates whether the <code>Zfbfmin</code> extension is supported (minimal <code>BFLOAT16</code> support) |
| <code>mstatus_FS_zero</code> | <code>Bool</code> | If floating point is <i>not</i> implemented but <i>Supervisor mode is present</i> , this field specifies whether <code>mstatus.FS</code> is forced to zero (normally, it is writable to enable floating point emulation). |
| <code>enable_fflags_i</code> | <code>Bool</code> | If <code>True</code> , this field causes an 8-bit artifact register, <code>fflags_i</code> , to be present. This register reports floating point flags updated by each instruction (unlike the standard <code>fflags</code> CSR, which reports cumulative flags). |
| <code>enable_DAZ</code> | <code>Bool</code> | If <code>True</code> , this field enables the floating point <i>Denormals-Are-Zero</i> mode where subnormal results are flushed to +/- zero. See 5.10.6 for |

| | | |
|-----------|------|--|
| | | more detailed information. <i>Note that such behavior is not compliant with the RISC-V or IEEE 2008 specifications.</i> |
| enable_FZ | Bool | If True, this field enables the floating point <i>Flush-to-Zero</i> mode where subnormal input values to certain floating point instruction are flushed to +/- zero before. See 5.10.6 for more detailed information. <i>Note that such behavior is not compliant with the RISC-V or IEEE 2008 specifications.</i> |

5.10.1 `misa` CSR `D` and `F` Bits

When both single and double precision floating point are implemented, `D` and `F` bits in the `misa` CSR often have implementation-specific dependencies. The model allows four different behaviors to be specified:

1. **default:** `D` and `F` bits can be set independently;
2. **`d_requires_f=1`:** `D` and `F` bits can be set independently, but `D` bit cannot be set to 1 if `F` bit is 0;
3. **`archMask` bit 5 (`F`) is 1 and bit 3 (`D`) is 0:** only bit 3 can be set in `misa`, and bit 5 is a read-only shadow of this.
4. **`archMask` bit 5 (`F`) is 0 and bit 3 (`D`) is 1:** only bit 5 can be set in `misa`, and bit 3 is a read-only shadow of this.

In both cases 3 and 4, single and double precision floating point cannot be enabled separately.

5.10.2 `mstatus.FS` Update Modes

The RISC-V Privileged Specification allows the behavior of the `mstatus.FS` field to be implementation defined. The model allows the behavior of this CSR field to be configured using the `mstatus_fs_mode` configuration option as follows:

1. **`RVFS_WRITE_NZ`:** `mstatus.FS` will be set to *dirty* (3) state whenever an FPR is written or a floating point operation signals an exception. Floating point operations that do not write an FPR or cause an exception will not set `mstatus.FS`. As an example, floating point comparison operations with `Normal` operands will *not* set `mstatus.FS` because these do not signal an exception.
2. **`RVFS_WRITE_ANY`:** `mstatus.FS` will be set whenever an FPR is written or a floating point operation could *potentially* signal an exception. As an example, floating point comparison operations with `Normal` operands *will* set `mstatus.FS`, because those operations could signal an exception with some inputs (e.g. `SNaN` operands).
3. **`RVFS_EXECUTE_NOT_S`:** `mstatus.FS` will be set to *dirty* (3) state when any floating point instruction retires, except for floating point stores, which do not modify `mstatus.FS`.
4. **`RVFS_EXECUTE_ANY`:** `mstatus.FS` will be set to *dirty* (3) state when any floating point instruction retires.
5. **`RVFS_ALWAYS_DIRTY`:** `mstatus.FS` can only hold either *off* (0) or *dirty* (3) states.
6. **`RVFS_FORCE_DIRTY`:** `mstatus.FS` is forced to *dirty* (3) state.

5.10.3 Half-Precision Support

Scalar half-precision support may be configured as follows:

1. **No half-precision support**
Set `fp16_version=RVFP16_NA` and `Zfhmin=False` in the configuration structure or set parameter `Zfh=False` and parameter `Zfhmin=False` at simulation time.
2. **Standard Zfh extension**
Set `fp16_version=RVFP16_IEEE754` and `Zfhmin=False` in the configuration structure or set parameter `Zfh=True` at simulation time (in which case parameter `Zfhmin` is ignored).
3. **Standard Zfhmin extension**
Set `fp16_version=RVFP16_IEEE754` and `Zfhmin=True` in the configuration structure or set parameter `Zfh=False` and parameter `Zfhmin=True` at simulation time.
4. **Non-Standard half-precision support**
Set `fp16_version` accordingly

5.10.4 Zfinx Support

The `Zfinx` extension replaces some standard floating point behavior. The options defined in this specification are configured as follows:

1. **Zfinx**
Enable `F` extension and set `Zfinx_version` to a non-zero value.
2. **Zdinx**
Enable `F` and `D` extensions and set `Zfinx_version` to a non-zero value.
3. **Zhinx**
Enable `F` extension, set `Zfh=True` and set `Zfinx_version` to a non-zero value.
4. **Zhinxmin**
Enable `F` extension, set `Zfhmin=True` and set `Zfinx_version` to a non-zero value.

5.10.5 Unimplemented Instructions

If a variant implements only a *subset* of the floating-point extension (for example, all instructions except `DIV/SQRT/FMAC`) then this can be specified using the `unimplementedInstr` field, described in section 5.23).

5.10.6 Flushing Subnormal Values to Zero

The configuration parameters `enable_DAZ` and `enable_FZ` may be set to `True` to cause subnormal results or inputs, respectively, to be flushed to zero for certain instructions.

Note that this behavior is not compliant with the RISC-V arc

Enabling `DAZ/FZ` in the RiscV model only affects "real" floating point operations and not operations that are may be implemented by simple bit manipulation. This means these parameters have no effect on these RISC-V floating point instructions:

```
fabs.[sd]  
fneg.[sd]  
fsgnjn.[sd]  
fsgnjx.[sd]  
fclass.[sd]
```

They also won't affect moves to and from floating point registers, or conversions from integral to floating point types (note that these can never generate subnormal results).

The `enable_DAZ` and `enable_FZ` settings will affect these instructions:

```
fsqrt.[sd]
fadd.[sd]
fdiv.[sd]
fmax.[sd]
fmin.[sd]
fmul.[sd]
fsub.[sd]
fmadd.[sd]
fmsub.[sd]
fnmadd.[sd]
fnmsub.[sd]
fcvt.l.[sd]
fcvt.lu.[sd]
fcvt.w.[sd]
fcvt.wu.[sd]
feq.[sd]
fle.[sd]
flt.[sd]
```

In all these cases, floating-point inputs and results will be flushed to correctly-signed zeros, without setting any floating point flags as a consequence. Floating point flags may be set by the operation itself with its freshly flushed-to-zero inputs, of course.

Additional information may be found in section **10 Floating Point Operations** in the document *OVP_VMI_Morph_Time_Function_Reference.pdf*.

5.11 Vector Extension

Fields here are applicable when the Vector Extension (V) is indicated as implemented by the processor `arch` field. All field defaults can be overridden using a model parameter of the same name, unless otherwise indicated.

| Field | Type | Description |
|-----------------------------|-----------------------------|---|
| <code>vect_version</code> | <code>riscvVectVer</code> | <p>This specifies the implemented Vector Extension version, defined by the <code>riscvVectVer</code> enumeration:</p> <pre>typedef enum riscvVectVerE { RVVV_0_7_1, // 0.7.1-draft-20190605 RVVV_0_7_1_P, // 0.7.1+ RVVV_0_8_20190906, // 0.8-draft-20190906 RVVV_0_8_20191004, // 0.8-draft-20191004 RVVV_0_8_20191117, // 0.8-draft-20191117 RVVV_0_8_20191118, // 0.8-draft-20191118 RVVV_0_8, // 0.8 RVVV_0_9, // 0.9 RVVV_1_0_20210130, // 1.0-draft-20210130 RVVV_1_0_20210608, // 1.0-draft-20210608 RVVV_1_0, // 1.0 RVVV_MASTER, RVVV_DEFAULT = RVVV_1_0, } riscvVectVer;</pre> <p>Version <code>RVVV_1_0_20210608</code> corresponds to the 1.0-rc1-20210608 release candidate.</p> |
| <code>vect_profile</code> | <code>riscvVectorSet</code> | <p>If non-zero, this specifies an applicable embedded profile, and should be one of:</p> <pre>RVVS_Application (zero value) RVVS_Zve32x RVVS_Zve32f RVVS_Zve64x RVVS_Zve64f RVVS_Zve64d</pre> <p>Any one of these values can be specified using parameters <code>Zve32x</code>, <code>Zve32f</code>, <code>Zve64x</code>, <code>Zve64f</code> or <code>Zve64d</code> when the model is instantiated, if required.</p> |
| <code>ELEN</code> | <code>Uns32</code> | This specifies the maximum vector element size, in bits (32 or 64). |
| <code>SLLEN</code> | <code>Uns32</code> | This specifies the vector stride, in bits. From Vector Extension version 1.0, this must match <code>VLEN</code> . |
| <code>VLEN</code> | <code>Uns32</code> | This specifies the vector size, in bits (a power of two in the range 32 to 65536). |
| <code>EEW_index</code> | <code>Uns32</code> | If non-zero, this specifies the maximum EEW that may be used for offset elements in indexed load and store instructions. If offsets larger than this are used, an Illegal Instruction exception is raised. |
| <code>SEW_min</code> | <code>Uns32</code> | This specifies the minimum vector element size, in bits. If zero, a value of 8 is assumed. |
| <code>Zvlssseg</code> | <code>Bool</code> | This specifies whether the <code>Zvlssseg</code> extension is implemented. |
| <code>Zvamo</code> | <code>Bool</code> | This specifies whether the <code>Zvamo</code> extension is implemented. |
| <code>Zvediv</code> | <code>Bool</code> | <p>This specifies whether the <code>Zvediv</code> extension is implemented.</p> <p><i>Note: this must currently always be False because the base model does not implement this extension.</i></p> |
| <code>Zvqmac</code> | <code>Bool</code> | This specifies whether the <code>Zvqmac</code> extension is implemented. |
| <code>Zvfh</code> | <code>Bool</code> | This specifies whether the <code>Zvfh</code> extension is implemented. |
| <code>Zvfhmin</code> | <code>Bool</code> | This specifies whether the <code>Zvfhmin</code> extension is implemented. |
| <code>Zvfbfmin</code> | <code>Bool</code> | This specifies whether the <code>Zvfbfmin</code> extension is implemented. |
| <code>Zvfbfwma</code> | <code>Bool</code> | This specifies whether the <code>Zvfbfwma</code> extension is implemented. |
| <code>unitStrideOnly</code> | <code>Bool</code> | This specifies whether only unit-stride vector loads and stores are supported. If <code>True</code> , then any other vector load/store instruction will |

| | | |
|-------------------|------|---|
| | | be reported as an Illegal Instruction. <i>Note that such behavior is not compliant with the Vector Extension specification.</i> |
| noFaultOnlyFirst | Bool | This specifies whether <i>fault-only-first</i> vector loads are unimplemented. If <code>True</code> , then any vector fault-only-first instruction will be reported as an Illegal Instruction. <i>Note that such behavior is not compliant with the Vector Extension specification.</i> |
| vstart0_non_ld_st | Bool | This specifies whether vector instructions that are not load/store instructions require the <code>vstart</code> CSR to be zero. If <code>True</code> , then attempting to execute such instructions with <code>vstart!=0</code> will cause an Illegal Instruction trap. |
| vstart0_ld_st | Bool | This specifies whether vector load/store instructions require the <code>vstart</code> CSR to be zero. If <code>True</code> , then attempting to execute such instructions with <code>vstart!=0</code> will cause an Illegal Instruction trap. |
| align_whole | Bool | This specifies whether whole-register load addresses must be aligned using the encoded EEW. |
| vill_trap | Bool | This specifies whether illegal <code>vtype</code> values cause a trap. |
| agnostic_ones | Bool | When tail/mask agnostic behavior is implemented, this specifies whether agnostic elements are filled with ones (if <code>True</code>) or preserved (if <code>False</code>). |
| unalignedV | Bool | This specifies whether vector load/store instructions support unaligned memory accesses. |

5.11.1 Half-Precision Support

Vector half-precision floating point requires either `zfh` or `zfhmin` to be also configured (see section 5.10.3). Given this prerequisite, vector half-precision support may then be configured in these three ways:

1. **`zvfh=False, zvfhmin=False`**
Set `zvfh=False` and `zfhmin=False` in the configuration structure or set parameter `zvfh=False` and parameter `zvfhmin=False` at simulation time.
2. **`zvfh=True, zvfhmin=True`**
Set `zvfh=True` in the configuration structure or set parameter `zvfh=True` at simulation time (`zvfhmin` is ignored because it is implied by `zvfh`).
3. **`zvfh=False, zvfhmin=True`**
Set `zvfh=False` and `zfhmin=True` in the configuration structure or set parameter `zvfh=False` and parameter `zvfhmin=True` at simulation time.

5.11.2 Unimplemented Instructions

If a variant implements only a *subset* of the vector extension then this can be specified using the `unimplementedInstr` field, described in section 5.23).

5.12 Bit Manipulation Extension

Fields here are applicable when the Bit Manipulation Extension (B) is indicated as implemented by the processor `arch` field¹. All field defaults can be overridden using a model parameter of the same name, unless otherwise indicated.

| Field | Type | Description |
|---------------------------------|-------------------------------|---|
| <code>bitmanip_version</code> | <code>riscvBitManipVer</code> | <p>This specifies the implemented Bit Manipulation Extension version, defined by the <code>riscvBitManipVer</code> enumeration:</p> <pre>typedef enum riscvBitManipVerE { RVBV_0_90, // 0.90 RVBV_0_91, // 0.91 RVBV_0_92, // 0.92 RVBV_0_93_DRAFT, // 0.93 (intermediate) RVBV_0_93, // 0.93 RVBV_0_94, // 0.94 RVBV_1_0_0, // 1.0.0 RVBV_MASTER, RVBV_DEFAULT = RVBV_0_92, } riscvBitManipVer;</pre> |
| <code>bitmanip_absent</code> | <code>riscvBitManipSet</code> | <p>By default, all Bit Manipulation Extension instruction subsets are implemented. This field is a bitmask allowing <i>unimplemented</i> subsets to be specified:</p> <pre>typedef enum riscvBitManipSetE { RVBS_Zba = (1<<0), RVBS_Zbb = (1<<1), RVBS_Zbc = (1<<2), RVBS_Zbe = (1<<3), RVBS_Zbf = (1<<4), RVBS_Zbm = (1<<5), RVBS_Zbp = (1<<6), RVBS_Zbr = (1<<7), RVBS_Zbs = (1<<8), RVBS_Zbt = (1<<9), } riscvBitManipSet;</pre> <p>The presence or absence of each of these subsets can be individually specified using parameters <code>Zba</code>, <code>Zbb</code>, <code>Zbc</code>, <code>Zbe</code>, <code>Zbf</code>, <code>Zbm</code>, <code>Zbp</code>, <code>Zbr</code>, <code>Zbs</code> and <code>Zbt</code>.</p> |
| <code>misa_B_Zba_Zbb_Zbs</code> | <code>Bool</code> | <p>For bit manipulation extension versions 1.0.0 and later, this specifies that when all of the <code>Zba</code>, <code>Zbb</code> and <code>Zbs</code> sub-extensions are present then <code>misa.B</code> is set. If <code>False</code>, then <code>misa.B</code> is always unset for these extension versions.</p> |

¹ The B bit must be specified in the `arch` field *even when not visible in the `misa` register* for this extension to be active.

5.13 Cryptographic Extension

Fields here are applicable when the Cryptographic Extension (K) is indicated as implemented by the processor `arch` field², except for `vcrypto_version`, which requires both the Cryptographic Extension *and* the Vector Extension (V). All field defaults can be overridden using a model parameter of the same name, unless otherwise indicated.

| Field | Type | Description |
|------------------------------|------------------------------|--|
| <code>crypto_version</code> | <code>riscvCryptoVer</code> | <p>This specifies the implemented Scalar Cryptographic Extension version, defined by the <code>riscvCryptoVer</code> enumeration:</p> <pre>typedef enum riscvCryptoVerE { RVKV_0_7_2, // 0.7.2 RVKV_0_8_1, // 0.8.1 RVKV_0_9_0, // 0.9.0 RVKV_0_9_2, // 0.9.2 RVKV_1_0_0_RC1, // 1.0.0-RC1 RVKV_1_0_0_RC5, // 1.0.0-RC5 RVKV_DEFAULT = RVKV_1_0_0_RC5 } riscvCryptoVer;</pre> |
| <code>vcrypto_version</code> | <code>riscvVCryptoVer</code> | <p>This specifies the implemented Vector Cryptographic Extension version, defined by the <code>riscvVCryptoVer</code> enumeration:</p> <pre>typedef enum riscvVCryptoVerE { RVKVV_0_3_0, // version 0.3.0 RVKVV_0_5_2, // version 0.5.2 RVKVV_1_0_0_RC1, // version 1.0.0-rc1 RVKVV_MASTER, // version master RVKVV_DEFAULT = RVKVV_1_0_0_RC1, } riscvVCryptoVer;</pre> |
| <code>crypto_absent</code> | <code>riscvCryptoSet</code> | <p>By default, all Cryptographic Extension instruction subsets are implemented. This field is a bitmask allowing <i>unimplemented</i> subsets to be specified:</p> <pre>typedef enum riscvCryptoSetE { // SCALAR SUBSETS RVKS_Zbkb = (1<<0), RVKS_Zbkc = (1<<1), RVKS_Zbkx = (1<<2), RVKS_Zkr = (1<<3), RVKS_Zknd = (1<<4), RVKS_Zkne = (1<<5), RVKS_Zknh = (1<<6), RVKS_Zksed = (1<<7), RVKS_Zksh = (1<<8), RVKS_Zkb = RVKS_Zbkb, RVKS_Zkg = RVKS_Zbkc, // VECTOR SUBSETS RVKS_Zvbb = (1<<9), RVKS_Zvbc = (1<<10), RVKS_Zvkg = (1<<11), RVKS_Zvknha = (1<<12), RVKS_Zvknhb = (1<<13), RVKS_Zvkned = (1<<14), RVKS_Zvksed = (1<<15), RVKS_Zvksh = (1<<16), RVKS_Zvkb = RVKS_Zvbb } riscvCryptoSet;</pre> <p>The presence or absence of each <i>scalar</i> subset can be individually specified using parameters <code>Zbkb</code>, <code>Zbkc</code>, <code>Zbkx</code>, <code>Zkr</code>, <code>Zknd</code>, <code>Zkne</code>, <code>Zknh</code>, <code>Zksed</code> and <code>Zksh</code>. Deprecated parameters <code>Zkb</code> and <code>Zkg</code> are equivalent to <code>Zbkb</code> and <code>Zbkc</code>, respectively.</p> |

² The K bit must be specified in the `arch` field *even when not visible in the `misra` register* for this extension to be active

| | | |
|-------------------------------|------|--|
| | | The presence or absence of each <i>vector</i> subset (when V extension is present) can be individually specified using parameters <code>Zvbb</code> , <code>Zvbc</code> , <code>Zvkg</code> , <code>Zvknha</code> , <code>Zvknhb</code> , <code>Zvkned</code> , <code>Zvksed</code> and <code>Zvksh</code> . Parameter <code>Zvkb</code> is a deprecated alias for <code>Zvkb</code> . |
| <code>mnoise_undefined</code> | Bool | This specifies that the <code>mnoise</code> CSR is undefined, and that accesses to it will trap to Machine mode. If defined, it has address <code>0x7A9</code> . |

5.14 Hypervisor Extension

Fields here are applicable when the Hypervisor Extension (H) is indicated as implemented by the processor `arch` field. All field defaults can be overridden using a model parameter of the same name unless otherwise indicated.

| Field | Type | Description |
|-----------------------------------|--------------------------|---|
| <code>hyp_version</code> | <code>riscvHypVer</code> | This specifies the implemented Hypervisor Extension version, defined by the <code>riscvHypVer</code> enumeration: <pre>typedef enum riscvHypVerE { RVHV_0_6_1, // 0.6.1 RVHV_1_0, // 1.0 RVHV_DEFAULT = RVHV_1_0, } riscvHypVer;</pre> |
| <code>VMID_bits</code> | <code>Uns32</code> | This defines the number of bits implemented in the VMID (maximum of 7 for RV32 and 14 for RV64). |
| <code>GEILEN</code> | <code>Uns32</code> | This specifies the number of guest external interrupts implemented (maximum of 31 for RV32 and 63 for RV64). |
| <code>xtinst_basic</code> | <code>Bool</code> | If <code>True</code> , this specifies that only pseudo-instructions are reported by <code>htinst/mtinst</code> ; if <code>False</code> , then true instructions can be reported as well. |
| <code>fence_g_preserves_vs</code> | <code>Bool</code> | If <code>True</code> , this specifies that <code>HFENCE.GVMA</code> preserves cached VS-stage address translations; if <code>False</code> , then <code>HFENCE.GVMA</code> will invalidate all cached VS-stage address translations. |
| <code>csrMask.hvip</code> | <code>Uns64</code> | <code>hvip</code> CSR write mask (only bits 63:13 used). Use parameter <code>hvip_mask</code> to override this value. |

5.15 DSP Extension

Fields here are applicable when the DSP Extension (P) is indicated as implemented by the processor `arch` field. All field defaults can be overridden using a model parameter of the same name, unless otherwise indicated.

| Field | Type | Description |
|--------------------------|--------------------------|--|
| <code>dsp_version</code> | <code>riscvDSPVer</code> | This specifies the implemented DSP Extension version, defined by the <code>riscvDSPVer</code> enumeration: <pre>typedef enum riscvDSPVerE { RVDSPV_0_5_2, // 0.5.2 RVDSPV_0_9_6, // 0.9.6 RVDSPV_DEFAULT = RVDSPV_0_5_2, } riscvDSPVer;</pre> |
| <code>dsp_absent</code> | <code>riscvDSPSet</code> | By default, all DSP Extension instruction subsets are implemented. This field is a bitmask allowing <i>unimplemented</i> subsets to be specified: <pre>typedef enum riscvDSPSetE { RVPS_Zpsfoperand = (1<<0), } riscvDSPSet;</pre> <p>The single entry <code>Zpsfoperand</code> is valid for RV32 only, and indicates whether instructions operating on register <i>pairs</i> are implemented. Its presence can be modified using a parameter of the same name.</p> |

5.16 Code Size Reduction Extension

Fields here are applicable when the Code Size Reduction Extension is implemented. All field defaults can be overridden using a model parameter of the same name, unless otherwise indicated. The parameters available for configuration of this extension depend on whether a *legacy* or *recent* specification version is configured.

| Field | Type | Description |
|------------------|------------------|---|
| compress_version | riscvCompressVer | <p>This specifies the implemented Code Size Reduction version, defined by the <code>riscvCompressVer</code> enumeration:</p> <pre>typedef enum riscvCompressVerE { RVCV_NA_LEGACY, RVCV_0_70_1, // 0.70.1 RVCV_0_70_5, // 0.70.5 RVCV_1_0_0_RC57, // 1.0.0-RC5.7 RVCV_1_0, // 1.0 (ratified) RVCV_DEFAULT = RVCV_1_0, } riscvCompressVer;</pre> <p>A value of <code>RVCV_NA_LEGACY</code> indicates that either the Code Size Reduction extension is absent or that a legacy version is in use, in which case subsets are separately-versioned (see below). Other versions do not have separately-versioned subsets.</p> |
| zcea_version | riscvZceaVer | <p>This specifies the implemented legacy Code Size Reduction zcea subset Extension version, defined by the <code>riscvZceaVer</code> enumeration:</p> <pre>typedef enum riscvZceaVerE { RVZCEA_NA, RVZCEA_0_50_1, // 0.50.1 RVZCEA_DEFAULT = RVZCEA_0_50_1, } riscvZceaVer;</pre> <p>This is ignored unless the legacy Code Size Reduction extension is configured.</p> |
| zceb_version | riscvZcebVer | <p>This specifies the implemented legacy Code Size Reduction zceb subset Extension version, defined by the <code>riscvZcebVer</code> enumeration:</p> <pre>typedef enum riscvZcebVerE { RVZCEB_NA, RVZCEB_0_50_1, // 0.50.1 RVZCEB_DEFAULT = RVZCEB_0_50_1, } riscvZcebVer;</pre> <p>This is ignored unless the legacy Code Size Reduction extension is configured.</p> |
| zcee_version | riscvZceeVer | <p>This specifies the implemented legacy Code Size Reduction zcee subset Extension version, defined by the <code>riscvZceeVer</code> enumeration:</p> <pre>typedef enum riscvZceeVerE { RVZCEE_NA, RVZCEE_1_0_0_RC, // 1.0.0-RC RVZCEE_DEFAULT = RVZCEE_1_0_0_RC, } riscvZceeVer;</pre> <p>This is ignored unless the legacy Code Size Reduction extension is configured.</p> |
| compress_present | riscvCompressSet | <p>This field is a bitmask allowing <i>implemented</i> Code Size Reduction Extension subsets to be specified:</p> <pre>typedef enum riscvCompressSetE { // legacy values RVCS_Zcea = (1<<0), RVCS_Zceb = (1<<1), RVCS_Zcee = (1<<2), // new values RVCS_Zca = (1<<3), RVCS_Zcb = (1<<4), RVCS_Zcd = (1<<5), }</pre> |

| | | |
|-------------|-------|--|
| | | <pre>RVCS_Zcf = (1<<6), RVCS_Zcmb = (1<<7), RVCS_Zcmp = (1<<8), RVCS_Zcmpe = (1<<9), RVCS_Zcmt = (1<<10), } riscvCompressSet;</pre> <p>Values for subsets Zcea, Zceb and Zcee are derived from the legacy version parameters above, if the legacy Code Size Reduction Extension is configured.</p> <p>Subset Zcd is an artifact used to indicate presence of compressed extension double-precision load/store instructions; its value is automatically maintained by the model.</p> <p>For other subsets, presence or absence can be individually specified using parameters Zca, Zcb, Zcf, Zcmb, Zcmp, Zcmpe and Zcmt, when the legacy Code Size Reduction extension is not configured.</p> |
| csrMask.jvt | Uns64 | jvt CSR write mask (Zcmt extension). Use parameter jvt_mask to override this value. |

5.17 Debug Mode

Fields here are applicable when Debug mode is required. All field defaults can be overridden using a model parameter of the same name unless otherwise specified.

| Field | Type | Description |
|-----------------|----------------|--|
| debug_version | riscvDebugVer | <p>This specifies the implemented Debug version, defined by the <code>riscvDebugVer</code> enumeration:</p> <pre>typedef enum riscvDebugVerE { RVDBG_0_13_2, // 0.13.2-DRAFT RVDBG_0_14_0, // 0.14.0-DRAFT RVDBG_1_0_0, // 1.0.0-STABLE RVDBG_1_0, // 1.0-STABLE RVDBG_DEFAULT = RVDBG_1_0 } riscvDebugVer;</pre> <p>Versions <code>RVDBG_1_0_0</code> and <code>RVDBG_1_0</code> correspond to two versions both called <i>1.0-STABLE</i>; They differ because in version <code>RVDBG_1_0</code>, the <code>nmi</code> field has moved to the <code>itrigger</code> view of <code>tdata1</code>, and behavior of that field has changed.</p> |
| debug_mode | riscvDMMode | <p>This field specifies how Debug mode is implemented. The possible behaviors are specified by the <code>riscvDMMode</code> enumeration:</p> <pre>typedef enum riscvDMModeE { RVDM_NONE, RVDM_VECTOR, RVDM_INTERRUPT, RVDM_HALT, RVDM_INJECT } riscvDMMode;</pre> <p>See below for a detailed description of how this affects behavior.</p> |
| debug_address | Uns64 | This specifies the address to jump to on a debug event (when <code>debug_mode</code> is <code>RVDM_VECTOR</code>). |
| dexc_address | Uns64 | This specifies the address to jump to on a debug exception (when <code>debug_mode</code> is <code>RVDM_VECTOR</code>). |
| debug_priority | riscvDPriority | <p>This specifies relative priority of simultaneous trigger-after, step, execute-address, <code>resethaltreq</code> and <code>haltreq</code> breakpoint events, defined by the <code>riscvDPriority</code> enumeration:</p> <pre>typedef enum riscvDPriorityE { RVDP_A_S_X_H, RVDP_A_S_H_X, RVDP_A_H_S_X, RVDP_H_A_S_X, RVDP_DEFAULT = RVDP_H_A_S_X, } riscvDPriority;</pre> <p>Value <code>RVDP_A_S_X_H</code> indicates the priority order: trigger-after, step, execute-address, <code>resethaltreq</code>, <code>haltreq</code> (in highest to lowest order).</p> <p>Value <code>RVDP_A_S_H_X</code> indicates the priority order: trigger-after, step, <code>resethaltreq</code>, <code>haltreq</code>, execute-address.</p> <p>Value <code>RVDP_A_H_S_X</code> indicates the priority order: trigger-after, <code>resethaltreq</code>, <code>haltreq</code>, step, execute-address.</p> <p>Value <code>RVDP_H_A_S_X</code> indicates the priority order: <code>resethaltreq</code>, <code>haltreq</code>, trigger-after, step, execute-address. See section 5.17.4 for more information.</p> |
| debug_eret_mode | riscvDERETMode | <p>This specifies the required behavior when <code>MRET</code>, <code>SRET</code> or <code>URET</code> instructions are executed in Debug mode. The possible behaviors are specified by the <code>riscvDERETMode</code> enumeration:</p> <pre>typedef enum riscvDERETModeE { RVDRM_NOP, // treat as NOP RVDRM_JUMP, // jump to dexc_address RVDRM_TRAP, // trap to dexc_address }</pre> |

| | | |
|----------------------|------|--|
| | | <code>} riscvDERETMode</code> For RVDRM_JUMP, the instruction is considered to have retired. For RVDRM_TRAP, the instruction is <i>not</i> considered to have retired. |
| debug_ctrlt_illegal | Bool | This specifies that control transfer instructions will cause Illegal Instruction traps in Debug mode. If <code>False</code> , the instructions are executed normally. |
| debug_auipec_illegal | Bool | This specifies that <code>auipec</code> instructions will cause Illegal Instruction traps in Debug mode. If <code>False</code> , the instructions are executed normally. |
| dscratch0_undefined | Bool | This specifies that the <code>dscratch0</code> register is undefined, and that accesses to it will trap to Machine mode. |
| dscratch1_undefined | Bool | This specifies that the <code>dscratch1</code> register is undefined, and that accesses to it will trap to Machine mode. |
| no_resethaltreq | Bool | If <code>False</code> , report code 5 in <code>dcsr</code> when <code>resethaltreq</code> is applied at reset, otherwise (if <code>True</code>) report code 3. |
| defer_step_bug | Bool | <p>The 0.13.2 Debug specification was ambiguous about required behavior when a debug single step is attempted while there is a pending exception that will be taken before the non-debug-mode instruction can execute.</p> <p>By default, the model will take the step breakpoint <i>before</i> the first trap handler instruction is executed in this case, which was the intention of the specification and explicitly stated in the 0.14.0 draft. Setting this configuration option to <code>True</code> enables an alternative behavior where the step breakpoint is instead taken <i>after the first trap handler instruction has executed</i>. This may be required for compatibility with legacy hardware where the specification was misinterpreted. There is no parameter to override this configuration option and it should be explicitly set in the configuration of any variant that requires it.</p> |

5.17.1 Debug Mode Behaviors

Field `debug_mode` can be used to specify four different simulation behaviors, as follows:

1. If set to value `RVDM_VECTOR`, then operations that would cause entry to Debug mode result in the processor jumping to the address specified by the `debug_address` field. It will execute at this address, in Debug mode, until a `dret` instruction causes return to non-Debug mode. Any exception generated during this execution will cause a jump to the address specified by the `dexc_address` field.
2. If set to value `RVDM_INTERRUPT`, then operations that would cause entry to Debug mode result in the processor simulation call (e.g. `opProcessorSimulate`) returning, with a stop reason of `OP_SR_INTERRUPT`. In this usage scenario, the Debug Module is implemented in the simulation harness.
3. If set to value `RVDM_HALT`, then operations that would cause entry to Debug mode result in the processor halting. Depending on the simulation environment, this might cause a return from the simulation call with a stop reason of `OP_SR_HALT`, or debug mode might be implemented by another platform component which then restarts the debugged processor again.
4. If set to value `RVDM_INJECT`, then operations that would cause entry to Debug mode result in the processor continuing to execute from the current address in Debug mode. The harness should detect that Debug mode has been entered by monitoring the `DM` integration support register, and inject Debug-mode

instructions one at a time using function `opProcessorSimulateInstruction`. Debug mode is exited by either an explicit write of `False` to the `DM` register or by execution of an injected `dret` instruction, as described in section 5.17.3 below.

5.17.2 Debug State Entry

The specification does not define how Debug mode is implemented. In this model, Debug mode is enabled by a Boolean pseudo-register, `DM`. When `DM` is `True`, the processor is in Debug mode. When `DM` is `False`, mode is defined by `mstatus` in the usual way. Entry to Debug mode can be performed in any of these ways:

1. By writing `True` to register `DM` (e.g. using `opProcessorRegWrite`) followed by simulation of at least one cycle (e.g. using `opProcessorSimulate`) - in this case, `dcsr.cause` will report a cause of `trigger` (2);
2. By writing a 1 then 0 to net `haltreq` (using `opNetWrite`) followed by simulation of at least one cycle (e.g. using `opProcessorSimulate`) - in this case, `dcsr.cause` will report a cause of `haltreq` (3);
3. By writing a 1 to net `resethaltreq` (using `opNetWrite`) while the `reset` signal undergoes a negative edge transition, followed by simulation of at least one cycle (e.g. using `opProcessorSimulate`) - in this case, `dcsr.cause` will report a cause of `resethaltreq` (5) or `haltreq` (3), depending on the value of field `no_resethaltreq`;
4. By executing an `ebreak` instruction when Debug mode entry for the current processor mode is enabled by `dcsr.ebreakm`, `dcsr.ebreaks` or `dcsr.ebreaku` - in this case, `dcsr.cause` will report a cause of `ebreak` (1);
5. By executing a single instruction when Debug mode entry for the current processor mode is enabled by `dcsr.step` - in this case, `dcsr.cause` will report a cause of `step` (4);
6. By a Trigger Module trigger, when that trigger is configured to enter Debug mode - in this case, `dcsr.cause` will report a cause of `trigger` (2).

In all cases, the processor will save required state in `dpc` and `dcsr` and then perform actions described above, depending in the value of the `debug_mode` field.

5.17.3 Debug State Exit

Exit from Debug mode can be performed in either of these ways:

1. By writing `False` to register `DM` (e.g. using `opProcessorRegWrite`) followed by simulation of at least one cycle (e.g. using `opProcessorSimulate`);
2. By executing a `dret` instruction when Debug mode.

In both cases, the processor will perform the steps described in section 4.6 (Resume) of the Debug Specification.

5.17.4 `haltreq` Signal and `debug_priority` Field

There are various events that can cause entry into Debug mode:

1. The hart might execute a special instruction, `EBREAK`, which implements a software breakpoint;
2. The hart might enter Debug mode because of a Trigger Module event;
3. The processor might be forced into Debug mode by a special external signal, `hltreq`.
4. The hart might be configured to perform a single step by Debug mode. When single-stepping, the hart returns from Debug mode, executes a single instruction in non-debug mode, and then re-enters Debug mode automatically.

Apart from the `hltreq` signal, these events are all synchronous. The type of event causing Debug mode entry is reported to Debug mode.

The model prioritizes *synchronous* events in this order:

1. Any Trigger Module pre-trigger is handled first (this will typically be an execute address trigger);
2. Explicit entry by `EBREAK` has next highest priority;
3. Any Trigger Module post-trigger is handled next;
4. If the instruction was single-stepped, that is handled last.

This means, for example, that if single-stepping an instruction which also causes a Trigger Module post-trigger, it is the post-trigger event that is reported to Debug mode, not the step event.

The Debug Specification is not clear about the priority of the asynchronous `hltreq` signal with respect to the other priorities above; for example, it doesn't define what happens if (for example) the `hltreq` signal is raised while single-stepping an instruction which also causes a Trigger Module post-trigger. The `debug_priority` field allows the model to be configured so that the `hltreq` event priority can be programmed with any one of four priorities:

`RVDP_A_S_X_H`

`hltreq` lowest priority.

`RVDP_A_S_H_X`

`hltreq` priority above execute address pre-trigger on the next instruction, but lower priority than post-triggers and single-step.

`RVDP_A_H_S_X`

`hltreq` lower priority than post-triggers, but higher priority than execute address pre-trigger on the next instruction and single-step.

`RVDP_H_A_S_X`

`hltreq` highest priority.

5.18 Trigger Module

Fields here are applicable when the Trigger Module is configured. All field defaults can be overridden using a model parameter of the same name unless otherwise indicated.

| Field | Type | Description |
|---------------------|-------|---|
| trigger_num | Uns32 | This field specifies how many trigger registers are implemented. If <code>trigger_num</code> is 0, the trigger module is not configured. |
| mask_tselect | Bool | When true, writes to <code>tselect</code> will be masked based on the value of <code>trigger_num</code> , which determines the number of bits implemented for the <code>tselect</code> register. Attempts to write values greater than the <code>trigger_num</code> setting will be ignored, regardless of whether they have been masked. |
| tinfo | Uns32 | This specifies the implemented trigger types (bits 15:0) and <code>sdtrig</code> version (bits 31:24). For example, if trigger types 5 and 6 are supported and <code>sdtrig</code> version is 1, the value should be <code>0x01000060</code> . |
| trigger_match | Uns32 | This is bitmask specifying legal values for the <code>match</code> field for address triggers (types 2 and 6). For example, if <code>match</code> values 0, 2 and 3 are supported, <code>trigger_match</code> should be specified as <code>0x000d</code> . |
| mcontext_bits | Uns32 | This specifies the number of implemented bits in the <code>mcontext</code> register. |
| scontext_bits | Uns32 | This specifies the number of implemented bits in the <code>scontext</code> register. |
| mvalue_bits | Uns32 | This specifies the number of implemented bits in the <code>textra32/textra64 mvalue</code> field. If zero, the <code>mselect</code> field in the same register is tied to zero. |
| svalue_bits | Uns32 | If Supervisor mode is present, this specifies the number of implemented bits in the <code>textra32/textra64 svalue</code> field. If zero, the <code>sselect</code> field in the same register is tied to zero. |
| mcontrol_maskmax | Uns32 | This specifies the value of the <code>mcontrol.maskmax</code> field. |
| tdata2_undefined | Bool | This specifies that the <code>tdata2</code> register is undefined, and that accesses to it will trap to Machine mode. |
| tdata3_undefined | Bool | This specifies that the <code>tdata3</code> register is undefined, and that accesses to it will trap to Machine mode. |
| tinfo_undefined | Bool | This specifies that the <code>tinfo</code> register is undefined, and that accesses to it will trap to Machine mode. |
| tcontrol_undefined | Bool | This specifies that the <code>tcontrol</code> register is undefined, and that accesses to it will trap to Machine mode. |
| mcontext_undefined | Bool | This specifies that the <code>mcontext</code> register is undefined, and that accesses to it will trap to Machine mode. |
| scontext_undefined | Bool | This specifies that the <code>scontext</code> register is undefined, and that accesses to it will trap to Machine mode. |
| mscontext_undefined | Bool | This specifies that the <code>mscontext</code> register is undefined, and that accesses to it will trap to Machine mode (Debug Version 0.14.0 and later). |
| hcontext_undefined | Bool | This specifies that the <code>hcontext</code> register is undefined, and that accesses to it will trap to Machine mode (when Hypervisor extension is present). |
| amo_trigger | Bool | This specifies whether AMO operations cause load/store triggers to be activated. |
| no_hit | Bool | This specifies whether the optional <code>hit</code> bits in <code>tdata1</code> are unimplemented. |
| no_sselect_2 | Bool | If Supervisor mode is present, this specifies whether the |

| | | |
|-----------------------------|-----------------------------|--|
| | | sselect field in <code>textra32/textra64</code> registers is unable to hold value 2 (indicating match by ASID is not allowed). |
| <code>csrMask.tdata1</code> | Uns64 | <code>tdata1</code> CSR write mask. Use parameter <code>tdata1_mask</code> to override this value. |
| <code>chain_tval</code> | <code>riscvChainTVAL</code> | <p>When chained triggers are implemented, the Debug Specification does not state which trigger supplies the value reported in <code>xtval</code> when the trigger causes a breakpoint exception. This field allows one of four possible behaviors to be specified: <code>RVCT_FIRST</code>, <code>RVCT_LAST</code>, <code>RVCT_FIRST_NON_EPC</code> and <code>RVCT_LAST_NON_EPC</code>. These have the following meanings:</p> <p><code>RVCT_FIRST</code>: report value from the first trigger in the chain that matches.</p> <p><code>RVCT_LAST</code>: report value from the last trigger in the chain that matches.</p> <p><code>RVCT_FIRST_NON_EPC</code>: report value from the first trigger in the chain that matches, but prefer values that are not the current PC (because this would be redundant with <code>xepc</code>).</p> <p><code>RVCT_LAST_NON_EPC</code>: report value from the last trigger in the chain that matches, but prefer values that are not the current PC (because this would be redundant with <code>xepc</code>).</p> |

5.19 Multicore variants

Fields here are applicable only for multicore processor variants (SMP or AMP). All field defaults can be overridden using a model parameter of the same name.

| Field | Type | Description |
|----------|---------------|--|
| numHarts | Uns32 | This field specifies how many processors are implemented beneath the root level in a multicore variant. A value of 0 specifies that the processor is not a multicore variant. |
| members | const char ** | For a multicore variant (<code>numHarts!=0</code>), a null value for this parameter specifies that the processor is an SMP processor. If <code>numHarts!=0</code> and <code>members</code> is non-null, then <code>members</code> must be a list of <code>numHarts</code> strings, each string specifying the variant name of a <i>cluster member</i> . The cluster member name must be the name of another variant in the current processor configuration list. The processor hierarchy will then be constructed with a heterogeneous set of variants, where the <i>nth</i> hart is of the type corresponding to the <i>nth</i> element of <code>members</code> . |

5.20 CSR Index Numbers

CSR numbers are defined by the Privileged Specification and other extension specifications. Sometimes, CSR numbers are reallocated as specifications evolve; for example, in the RNMI specification (which is not yet ratified), CSR `mnstatus` was originally assigned number `0x353`, but in more recent versions of the specification it has been allocated number `0x744` instead. These changes in CSR numbers are automatically handled by the base model if the correct extension version number is selected.

Occasionally, variants are created that use *non-standard* CSR numbers. For example, some SiFive variants with interim RNMI implementations use number `0x743` for `mnstatus`, which does not correspond with any specification version. To implement this, the `CSR_remap` field in the configuration can be used. If non-NULL, this field is a string containing comma-separated entries of the form:

`<csrName>=<number>`

For SiFive cores implementing an interim RNMI extension, this field is set to:

`"mnstatus=0x743"`

This indicates that number `0x743` must be used for `mnstatus` instead of the default number (`0x744`). CSR number remapping can be overridden using string parameter `CSR_remap` if required.

5.21 CSR Initial Values

The `csr` field contains default values for some CSRs that are typically read-only or not otherwise fully configurable by other options mentioned above. All field defaults can be overridden using a model parameter of the same name unless otherwise stated.

| Field | Type | Description |
|-----------------------------|-------|---|
| <code>csr.mvendorid</code> | Uns64 | <code>mvendorid</code> CSR value. |
| <code>csr.marchid</code> | Uns64 | <code>marchid</code> CSR value. |
| <code>csr.mimpid</code> | Uns64 | <code>mimpid</code> CSR value. |
| <code>csr.mhartid</code> | Uns64 | <code>mhartid</code> CSR value. |
| <code>csr.mconfigptr</code> | Uns64 | <code>mconfigptr</code> CSR value. Not used for Privileged Architecture versions 1.11 and earlier. |
| <code>csr.mtvec</code> | Uns64 | <code>mtvec</code> CSR initial value. |
| <code>csr.mstatus</code> | Uns64 | <code>mstatus</code> CSR initial value. Fields <code>mstatus.FS</code> and <code>mstatus.VS</code> can be overridden using parameters <code>mstatus_FS</code> and <code>mstatus_VS</code> , respectively. |

5.22 CSR Masks

The `csrMask` field defines write masks for some CSRs. If the fields are zero, default values are used as described in the *Default if Zero* column below. All field defaults can be overridden using parameters with `_mask` suffix; for example use parameter `mtvec_mask` to override the value of `csrMask.mtvec`.

| Field | Type | Description | Default if Zero |
|-----------------------------|-------|---|---|
| <code>csrMask.mtvec</code> | Uns64 | mtvec CSR write mask. | all bits writable except <code>mtvec[1:0]</code> , which depend on basic IC/CLIC presence |
| <code>csrMask.stvec</code> | Uns64 | stvec CSR write mask. | all bits writable except <code>stvec[1:0]</code> , which depend on basic IC/CLIC presence |
| <code>csrMask.utvec</code> | Uns64 | utvec CSR write mask (N extension). | all bits writable except <code>utvec[1:0]</code> , which depend on basic IC/CLIC presence |
| <code>csrMask.mtvt</code> | Uns64 | mtvt CSR write mask (CLIC). | all bits writable except <code>mtvt[5:0]</code> |
| <code>csrMask.stvt</code> | Uns64 | stvt CSR write mask (CLIC). | all bits writable except <code>stvt[5:0]</code> |
| <code>csrMask.utvt</code> | Uns64 | utvt CSR write mask (CLIC and N extension). | all bits writable except <code>utvt[5:0]</code> |
| <code>csrMask.jvt</code> | Uns64 | jvt CSR write mask (Zcmt extension). | all bits writable except <code>jvt[5:0]</code> |
| <code>csrMask.tdata1</code> | Uns64 | tdata1 CSR write mask (Trigger Module). | all bits writable |
| <code>csrMask.mip</code> | Uns64 | mip CSR write mask | 0x337 |
| <code>csrMask.sip</code> | Uns64 | sip CSR write mask | 0x103 |
| <code>csrMask.uip</code> | Uns64 | uip CSR write mask (N extension) | 0x001 |
| <code>csrMask.hip</code> | Uns64 | hip CSR write mask (H extension) | 0x004 |
| <code>csrMask.mvien</code> | Uns64 | mvien CSR write mask (Smaia extension, only bits 63:13 used) | 0x0 |
| <code>csrMask.mvip</code> | Uns64 | mvip CSR write mask (Smaia extension, only bits 63:13 used) | 0x0 |
| <code>csrMask.hvien</code> | Uns64 | hvien CSR write mask (Smaia extension with H extension, only bits 63:13 used) | 0x0 |
| <code>csrMask.hvip</code> | Uns64 | hvip CSR write mask (H extension, only bits 63:13 used) | 0x0 |
| <code>csrMask.envcfg</code> | Uns64 | menvcfg/henvcfg/senvcfg write mask. Not used for Privileged Architecture versions 1.11 and earlier. | 0x80000000000000f1 |

5.23 Unimplemented Instructions

Sometimes a RISC-V processor variant will implement only a *subset* of the instructions in a standard extension. For example, a processor might implement all floating-point instructions except divide instructions, or only the LR/SC instructions from the atomic extension. This behavior could be implemented in an extended model by decoding the unimplemented instructions and reimplementing them as Illegal Instructions (see section 6). An easier alternative is to use the `unimplementedInstr` field in the configuration structure.

If this field is non-NULL, it must be a null-terminated list of `riscvIType` enumeration members for instructions that are unimplemented. This enumeration is declared in file `riscvDecodeTypes.h` and has one entry for each fundamental instruction type recognized by the instruction decoder. As an example, if a variant implements the atomic (A) extension but only the LR/SC instructions from that, and the floating-point extensions (F and D) but not DIV/SQRT/FMAC instructions, then this could be specified as follows:

```
//
// Definition of unimplemented instructions
//
static const riscvIType unimplementedInstructions[] = {

    // Unimplemented atomics
    RV_IT_AMOADD_R,
    RV_IT_AMOAND_R,
    RV_IT_AMOMAX_R,
    RV_IT_AMOMAXU_R,
    RV_IT_AMOMIN_R,
    RV_IT_AMOMINU_R,
    RV_IT_AMOOR_R,
    RV_IT_AMOSWAP_R,
    RV_IT_AMOXOR_R,

    // Unimplemented floating-point instructions
    RV_IT_FDIV_R,
    RV_IT_FSQRT_R,
    RV_IT_FMADD_R4,
    RV_IT_FMSUB_R4,
    RV_IT_FNMADD_R4,
    RV_IT_FNMSUB_R4,

    // list terminator
    0
};

static const riscvConfig configList[] = {

    {
        .name          = "RV32X",
        .arch          = ISA_U|RV32GC|ISA_X,
        . . . fields omitted . . .
        .unimplementedInstr = &unimplementedInstructions[0]
    },

    {0} // null terminator
};
```

6 Adding Custom Instructions (addInstructions)

The `addInstructions` extension demonstrates how to add custom instructions to a RISC-V model. The extension adds four instructions in the custom space of a RISC-V processor. Each instruction takes two 32-bit GPR inputs (`rs1` and `rs2`) and writes a result to a target GPR (`rd`), as follows:

chacha20qr1

```
rd = ((rs1 ^ rs2) << 16) | ((rs1 ^ rs2) >> (32-16))
```

chacha20qr2

```
rd = ((rs1 ^ rs2) << 12) | ((rs1 ^ rs2) >> (32-12))
```

chacha20qr3

```
rd = ((rs1 ^ rs2) << 8) | ((rs1 ^ rs2) >> (32-8))
```

chacha20qr4

```
rd = ((rs1 ^ rs2) << 7) | ((rs1 ^ rs2) >> (32-7))
```

All behavior of this extension object is implemented in file `addInstructionsExtensions.c`. Sections will be discussed in turn below.

6.1 Intercept Attributes

The behavior of the extension is defined using the standard `vmiosAttr` structure:

```
vmiosAttr modelAttrs = {
    ///////////////////////////////////////////////////////////////////
    // VERSION
    ///////////////////////////////////////////////////////////////////

    .versionString = VMI_VERSION,           // version string
    .modelType     = VMI_INTERCEPT_LIBRARY, // shared object type
    .interceptType = VMI_IT_PROC_EXTENSION,  // intercept type
    .packageName   = "addCSRs",             // description
    .objectSize    = sizeof(vmiosObject),    // size in bytes of OSS object

    ///////////////////////////////////////////////////////////////////
    // CONSTRUCTOR/DESTRUCTOR ROUTINES
    ///////////////////////////////////////////////////////////////////

    .constructorCB = addInstructionsConstructor, // object constructor

    ///////////////////////////////////////////////////////////////////
    // INSTRUCTION INTERCEPT ROUTINES
    ///////////////////////////////////////////////////////////////////

    .morphCB      = addInstructionsMorph,      // instruction morph callback
    .disCB        = addInstructionsDisassemble, // disassemble instruction

    ///////////////////////////////////////////////////////////////////
    // ADDRESS INTERCEPT DEFINITIONS
    ///////////////////////////////////////////////////////////////////

    .intercepts   = {{0}}
}
```

In this extension, there is a constructor, a JIT translation (morpher) function and a disassembly function.

6.2 Object Type and Constructor

The object type is defined as follows:

```
typedef struct vmiosObjectS {  
    // Info for associated processor  
    riscvP      riscv;  
  
    // extended instruction decode table  
    vmidDecodeTableP decode32;  
  
    // extension callbacks  
    riscvExtCB    extCB;  
} vmiosObject;
```

To be compatible with the integration facilities described in this document, an extension object must contain the `riscv`, `decode32` and `extCB` fields shown here. It may also contain other instance-specific fields. The constructor initializes the fields as follows:

```
static VMIOS_CONSTRUCTOR_FN(addInstructionsConstructor) {  
    riscvP riscv = (riscvP)processor;  
    object->riscv = riscv;  
  
    // prepare client data  
    object->extCB.clientData = object;  
  
    // register extension with base model using unique ID  
    riscv->cb.registerExtCB(riscv, &object->extCB, EXTID_ADDINST);  
}
```

The `extCB` field defines one side of the interface between the extension object and the base RISC-V model. It is filled by the extension object constructor with callback function pointers and other data that are used by the base model to communicate with the extension object. To complement this, the base model itself implements a set of *interface functions* that can be called from the extension object to query and modify the base model state. These are discussed in following sections.

The constructor must initialize the `clientData` field within the `extCB` structure with the extension object and then call the `registerExtCB` interface function to register this extension object with the model. The last argument to this function is an identification number that should uniquely identify this extension object in the case that multiple libraries are installed on one RISC-V processor.

6.3 Instruction Decode

Because this extension object is implementing new instructions, it needs to define a utility function to decode those new instructions. Provided that the instructions follow the standard RISC-V pattern in terms of operand locations and types, an interface function is available to simplify this process.

The first step is to define an enumeration with an entry for each new instruction:

```
typedef enum riscvExtITypeE {
    // extension instructions
    EXT_IT_CHACHA20QR1,
    EXT_IT_CHACHA20QR2,
    EXT_IT_CHACHA20QR3,
    EXT_IT_CHACHA20QR4,

    // KEEP LAST
    EXT_IT_LAST
} riscvExtIType;
```

Then, information about each instruction is given in a table of `riscvExtInstrAttrs` entries, with one entry for each instruction. Members of the table should be initialized using the `EXT_INSTRUCTION` macro, defined (like all interface function types and macros) in file `riscvModelCallbackTypes.h` in the RISC-V model source:

```
const static riscvExtInstrAttrs attrsArray32[] = {
    EXT_INSTRUCTION(
        EXT_IT_CHACHA20QR1, "chacha20qr1", RVANY, RVIP_RD_RS1_RS2, FMT_R1_R2_R3,
        "|0000000|.....|.....|000|.....|0001011|"
    ),
    EXT_INSTRUCTION(
        EXT_IT_CHACHA20QR2, "chacha20qr2", RVANY, RVIP_RD_RS1_RS2, FMT_R1_R2_R3,
        "|0000000|.....|.....|001|.....|0001011|"
    ),
    EXT_INSTRUCTION(
        EXT_IT_CHACHA20QR3, "chacha20qr3", RVANY, RVIP_RD_RS1_RS2, FMT_R1_R2_R3,
        "|0000000|.....|.....|010|.....|0001011|"
    ),
    EXT_INSTRUCTION(
        EXT_IT_CHACHA20QR4, "chacha20qr4", RVANY, RVIP_RD_RS1_RS2, FMT_R1_R2_R3,
        "|0000000|.....|.....|011|.....|0001011|"
    )
};
```

Each instruction is described by 6 arguments to the `EXT_INSTRUCTION` macro:

1. The instruction enumeration member name (e.g. `EXT_IT_CHACHA20QR1`);
2. The instruction opcode, as a string (e.g. `"chacha20qr1"`);
3. The applicable architecture (using enumeration values defined in file `riscvVariant.h` in the base model). Value `RVANY` means no constraint; other values can constrain an instruction to a particular `XLEN` (e.g. `RV32` or `RV64`) or to a particular extension code (e.g. `RVANYB`) or a combination of the two (e.g. `RV32B`). Given this constraint, the base model will automatically check that the constraint is valid and generate an Illegal Instruction exception if it is not.
4. The instruction operands. In this case, the value `RVIP_RD_RS1_RS2` means an instruction targeting GPR `Rd` and taking GPRs `Rs1` and `Rs2` as arguments, with all registers in the standard positions in a RISC-V instruction. Other operand layouts can be specified instead, as defined by the following enumeration in `riscvModelCallbackTypes.h`:

```
typedef enum riscvExtInstrPatternE {
```

```

// GPR INSTRUCTIONS
RVIP_RD_RS1_RS2,      // op  xd, xs1, xs2      (R-Type)
RVIP_RD_RS1_SI,       // op  xd, xs1, imm      (I-Type)
RVIP_RD_RS1_SHIFT,    // op  xd, xs1, shift    (I-Type - 5 or 6 bit shift)
RVIP_BASE_RS2_OFFSET, // op  base, xs2, offset (S-Type)
RVIP_RS1_RS2_OFFSET,  // op  xs1, xs2, offset  (B-Type)
RVIP_RD_SI,           // op  xd, imm           (U-Type)
RVIP_RD_OFFSET,       // op  xd, offset        (J-Type)
RVIP_RD_RS1_RS2_RS3,  // op  xd, xs1, xs2, xs3 (R4-Type)
RVIP_RD_RS1_RS3_SHIFT, // op  xd, xs1, xs2, shift (Non-Standard)

// FPR INSTRUCTIONS
RVIP_FD_FS1_FS2,      // op  fd, fs1, fs2
RVIP_FD_FS1_FS2_RM,   // op  fd, fs1, fs2, rm
RVIP_FD_FS1_FS2_FS3_RM, // op  fd, fs1, fs2, fs3, rm
RVIP_RD_FS1_FS2,      // op  xd, fs1, fs2

// VECTOR INSTRUCTIONS
RVIP_VD_VS1_VS2_M,    // op  vd, vs1, vs2, vm
RVIP_VD_VS1_SI_M,     // op  vd, vs1, simm, vm
RVIP_VD_VS1_UI_M,     // op  vd, vs1, uimm, vm
RVIP_VD_VS1_RS2_M,    // op  vd, vs1, rs2, vm
RVIP_VD_VS1_FS2_M,    // op  vd, vs1, fs2, vm
RVIP_RD_VS1_RS2,      // op  xd, vs1, vs2
RVIP_RD_VS1_M,        // op  xd, vs1, vm
RVIP_VD_RS2,          // op  vd, xs2
RVIP_FD_VS1,          // op  fd, vs1
RVIP_VD_FS2,          // op  vd, fs2

RVIP_LAST             // KEEP LAST: for sizing
} riscvExtInstrPattern;

```

See section 13 for a detailed description of each of these.

- How arguments should be shown when the instruction is disassembled, as described by a format string defined in `riscvDisassembleFormats.h` in the base model. In this case, the value `FMT_R1_R2_R3` specifies that the three GPR arguments should be shown as a comma-separated list.
- The instruction pattern, in terms of ones, zeros and dot (don't care) bits, separated by vertical bar characters which are ignored. For example, the `chacha20qr1` instruction is defined using this pattern string (the comment identifies standard RISC-V instruction fields):

```

| dec | rs2 | rs1 | fn3 | rd | dec |
"0000000|....|....|000|....|0001011"

```

The table of instructions is used by the `fetchInstruction` interface function, as follows:

```

static riscvExtIType decode(
    riscvP          riscv,
    vmiosObjectP    object,
    riscvAddr       thisPC,
    riscvExtInstrInfoP info
) {
    return riscv->cb.fetchInstruction(
        riscv, thisPC, info, &object->decode32, attrsArray32, EXT_IT_LAST, 32
    );
}

```

Arguments to `cb.fetchInstruction` are:

1. The RISC-V processor object;
2. The instruction address;
3. A pointer to an object of type `riscvExtInstrInfo` which is filled by `cb.fetchInstruction` with details of the decoded instruction;
4. A pointer to a decode table structure used by `cb.fetchInstruction`;
5. A pointer to the instruction table, defined above;
6. An indication of the number of entries in the instruction table;
7. The size (in bits) of instructions in the table. In this case, instructions are 32 bits.

Function `cb.fetchInstruction` returns either the index number of the decoded instruction or `EXT_IT_LAST`, if the instruction is not specified in this instruction table.

The `riscvExtInstrInfo` type which is filled by `cb.fetchInstruction` has this definition:

```
typedef struct riscvExtInstrInfoS {
    riscvAddr      thisPC;           // instruction address
    Uns32          instruction;      // instruction word
    Uns8           bytes;           // instruction bytes
    const char     *opcode;          // opcode name
    const char     *format;          // disassembly format string
    riscvArchitecture arch;         // architecture requirements
    riscvRegDesc   r[4];             // argument registers
    riscvRegDesc   mask;             // mask register (vector instructions)
    riscvRMDesc    rm;              // rounding mode
    Uns64          c;               // constant value
    void           *userData;        // client-specific data
} riscvExtInstrInfo;
```

Fields `thisPC`, `instruction`, `bytes`, `arch`, `opcode` and `format` are always filled. Fields `r`, `mask`, `rm` and `c` are filled if applicable to the operand pattern, otherwise they are zeroed. Field `userData` is available for the extension object to use if it requires to pass further data to other stages (disassembly or JIT instruction translation).

All instructions added by this extension have operands specified by `RVIP_RD_RS1_RS2`, so in this case, `r[0]`, `r[1]` and `r[2]` are filled with register descriptions extracted from the instruction, and other fields are zero.

Typically, the decode function will be used unmodified in a new extension object, except for changing its name: only the instruction type enumeration and `attrsArray32` entries should change.

6.4 Instruction Disassembly

The decode routine described in the previous section can be used in combination with the `disassInstruction` interface function to implement instruction disassembly in standard form:

```
static VMIO_DISASSEMBLE_FN(addInstructionsDisassemble) {
```

```

riscvP      riscv = (riscvP)processor;
const char  *result = 0;
riscvExtInstrInfo info;

// action is only required if the instruction is implemented by this
// extension
if(decode(riscv, object, thisPC, &info) != EXT_IT_LAST) {
    result = riscv->cb.disassInstruction(riscv, &info, attrs);
}

return result;
}

```

This function calls `decode`. If the result is not `EXT_IT_LAST`, then the instruction was successfully decoded by this extension object, and interface function

`disassInstruction` is called to produce the disassembly string. This takes three arguments:

1. The RISC-V processor instance;
2. The `riscvExtInstrInfo` argument block (filled by function `decode`);
3. The disassembly attributes passed to `chachaDisassemble` (whether normal or uncooked disassembly is required).

This will produce disassembly for instructions in this extension that conforms exactly to the disassembly generated by base model instructions. For example:

```

Info `iss/cpu0`, 0x00000000000102b4(processWord+8): 01010413 addi    s0,sp,16
Info `iss/cpu0`, 0x00000000000102b8(processWord+c): 00050513 mv      a0,a0
Info `iss/cpu0`, 0x00000000000102bc(processWord+10): 00058593 mv      a1,a1
Info `iss/cpu0`, 0x00000000000102c0(processWord+14): 00b5050b chacha20qr1 a0,a0,a1
Info `iss/cpu0`, 0x00000000000102c4(processWord+18): 00b5150b chacha20qr2 a0,a0,a1
Info `iss/cpu0`, 0x00000000000102c8(processWord+1c): 00b5250b chacha20qr3 a0,a0,a1
Info `iss/cpu0`, 0x00000000000102cc(processWord+20): 00b5350b chacha20qr4 a0,a0,a1
Info `iss/cpu0`, 0x00000000000102d0(processWord+24): 00b5050b chacha20qr1 a0,a0,a1
Info `iss/cpu0`, 0x00000000000102d4(processWord+28): 00b5150b chacha20qr2 a0,a0,a1
Info `iss/cpu0`, 0x00000000000102d8(processWord+2c): 00b5250b chacha20qr3 a0,a0,a1
Info `iss/cpu0`, 0x00000000000102dc(processWord+30): 00b5350b chacha20qr4 a0,a0,a1
Info `iss/cpu0`, 0x00000000000102e0(processWord+34): 00050513 mv      a0,a0
Info `iss/cpu0`, 0x00000000000102e4(processWord+38): 00c12403 lw      s0,12(sp)
Info `iss/cpu0`, 0x00000000000102e8(processWord+3c): 01010113 addi    sp,sp,16

```

Typically, the disassembly function will be used unmodified in a new extension object, except for changing its name.

6.5 Instruction Translation

Information about each instruction is given in a table of `riscvExtMorphAttr` entries, with one entry for each instruction. This type is defined in `riscvModelCallbackTypes.h` as follows:

```

typedef struct riscvExtMorphAttrS {
    extMorphFn      morph; // function to translate one instruction
    octiaInstructionClass iClass; // supplemental instruction class
    Uns32           variant; // required variant
    void            *userData; // client-specific data
} riscvExtMorphAttr;

```

Field `morph` specifies a function to be called to translate the instruction. Field `iClass` is used to specify an instruction *class*; this information is used by some supplemental tools

(such as timing estimators) but is not required if those tools are not used. Field `variant` is a model-specific variant feature code: this can be used to specify whether an instruction is implemented in a particular instantiation of this extension object (useful if the library can be configured to support multiple supplementary instructions whose presence is determined by a configuration register, for example). Field `userData` can hold any extension-specific data.

In this example, the table is specified like this:

```
const static riscvExtMorphAttr dispatchTable[] = {
    [EXT_IT_CHACHA20QR1] = {morph:emitCHACHA20QR, userData:(void *)16},
    [EXT_IT_CHACHA20QR2] = {morph:emitCHACHA20QR, userData:(void *)12},
    [EXT_IT_CHACHA20QR3] = {morph:emitCHACHA20QR, userData:(void *) 8},
    [EXT_IT_CHACHA20QR4] = {morph:emitCHACHA20QR, userData:(void *) 7},
};
```

Here, the same callback function (`emitCHACHA20QR`) is used for all four extension instructions, with behavior controlled using an integer rotate passed using the `userData` field.

The JIT translation function is specified like this:

```
static VMIO_MORPH_FN(addInstructionsMorph) {

    riscvP          riscv = (riscvP)processor;
    riscvExtMorphState state = {riscv:riscv, object:object};

    // decode instruction
    riscvExtITType type = decode(riscv, object, thisPC, &state.info);

    // action is only required if the instruction is implemented by this
    // extension
    if(type != EXT_IT_LAST) {

        // fill translation attributes
        state.attrs = &dispatchTable[type];

        // translate instruction
        riscv->cb.morphExternal(&state, 0, opaque);
    }

    // no callback function is required
    return 0;
}
```

This function defines a structure of type `riscvExtMorphState`, which is used to encapsulate all information required to translate a single instruction:

```
typedef struct riscvExtMorphStateS {
    riscvExtInstrInfo  info;          // decoded instruction information
    riscvExtMorphAttrCP attrs;        // instruction attributes
    riscvP             riscv;         // current processor
    vmiosObjectP       object;        // current extension object
} riscvExtMorphState;
```

The structure is declared with `riscv` and `object` fields initialized:

```
riscvExtMorphState state = {riscv:riscv, object:object};
```

Then, the `info` field is filled with information about the current instruction:

```
riscvExtIType type = decode(riscv, object, thisPC, &state.info);
```

If the instruction is implemented by this extension object, the `attrs` field is filled with the relevant entry from table `dispatchTable`, and the interface function `morphExternal` is called to perform the translation:

```
// fill translation attributes
state.attrs = &dispatchTable[type];

// translate instruction
riscv->cb.morphExternal(&state, 0, opaque);
```

Function `morphExternal` takes three arguments:

1. The `riscvExtMorphState` structure;
2. A *disable reason string*. If this string is non-NULL, then the instruction is not translated, but instead an Illegal Instruction exception is triggered, with the given string reported as a reason in verbose mode. In this example, the string is NULL, so the instruction is always translated, but typically code at this point would validate the variant in the `riscvExtMorphAttr` entry against current configuration in the extension object, and return a non-NULL disable reason string if required.
3. The `opaque` function argument passed to `chachaMorph`.

Interface function `morphExternal` performs standard checks for instruction validity (for example, if the instruction architecture is defined as RV64 then an Illegal Instruction exception is raised if the current XLEN is 32). If the instruction is valid, the appropriate callback function from the `dispatchTable` is called to generate JIT code to implement the instruction.

Typically, the JIT translation function will be used unmodified in a new extension object, except for changing its name.

Each translation callback function is passed a single argument of type `riscvExtMorphStateP`, which holds all information required to generate code for the current instruction. In this case, function `emitCHACHA20QR` is implemented like this:

```
static EXT_MORPH_FN(emitCHACHA20QR) {

    riscvP riscv = state->riscv;

    // get abstract register operands
    riscvRegDesc rd = getRVReg(state, 0);
    riscvRegDesc rs1 = getRVReg(state, 1);
    riscvRegDesc rs2 = getRVReg(state, 2);

    // get VMI registers for abstract operands
    vmiReg rdA = getVMIReg(riscv, rd);
    vmiReg rs1A = getVMIReg(riscv, rs1);
    vmiReg rs2A = getVMIReg(riscv, rs2);

    // emit embedded call to perform operation
    UnsPS rotl = (UnsPS)state->attrs->userData;
```

```
Uns32 bits = 32;
vmimtArgReg(bits, rs1A);
vmimtArgReg(bits, rs2A);
vmimtArgUns32(rotl);
vmimtCallResult((vmiCallFn)qrN_c, bits, rdA);

// handle extension of result if 64-bit XLEN
writeRegSize(riscv, rd, bits, True);
}
```

The first part of this function extracts RISC-V *abstract register definitions* from the passed state object for registers `r[0]`, `r[1]` and `r[2]`:

```
riscvRegDesc rd = getRVReg(state, 0);
riscvRegDesc rs1 = getRVReg(state, 1);
riscvRegDesc rs2 = getRVReg(state, 2);
```

The abstract register description combines register index (0-31), type information (X, F or V register) with size (8, 16, 32 or 64 bits). Function `getRVReg` is a simple utility function:

```
inline static riscvRegDesc getRVReg(riscvExtMorphStateP state, Uns32 argNum) {
    return state->info.r[argNum];
}
```

The next part of the function translates abstract register definitions into VMI register definitions:

```
vmiReg rdA = getVMIReg(riscv, rd);
vmiReg rs1A = getVMIReg(riscv, rs1);
vmiReg rs2A = getVMIReg(riscv, rs2);
```

Function `getVMIReg` is a simple wrapper function around an interface function of the same name:

```
inline static vmiReg getVMIReg(riscvP riscv, riscvRegDesc r) {
    return riscv->cb.getVMIReg(riscv, r);
}
```

Once the registers have been converted to VMI register descriptions, all available VMI morph-time operations can be used (see the *VMI Morph Time Function Reference Manual* and *OVP Processor Modeling Guide*). For simplicity, it is usually easiest to implement extension instructions as *embedded calls*, in which each extension instruction is implemented by a simple function. To do this, first define a utility function to implement the instruction operation. In this case, the utility function is this:

```
static Uns32 qrN_c(Uns32 rs1, Uns32 rs2, Uns32 rotl) {
    return ((rs1 ^ rs2) << rotl) | ((rs1 ^ rs2) >> (32-rotl));
}
```

This function takes two 32-bit register inputs (`rs1` and `rs2`) and constant rotation. It returns operation results as described at the start of this chapter. Within the translation callback function, a call to this utility function is emitted. First, the constant rotation amount is extracted from the `userData` field in the `riscvExtMorphAttr` entry:

```
UnsPS rotl = (UnsPS)state->attrs->userData;
```

The two GPR sources and the constant amount are passed as arguments to the utility function, and a call to that emitted, and the result is assigned to register `rd`:

```
Uns32 bits = 32;
vmimtArgReg(bits, rs1A);
vmimtArgReg(bits, rs2A);
vmimtArgUns32(rotl);
vmimtCallResult((vmiCallFn)qrN_c, bits, rdA);
```

The instructions operate on data of fixed width (32 bits). If run on a RISC-V with XLEN of 64, the 32-bit result must be extended to 64 bits. This extension is specified as follows:

```
writeRegSize(riscv, rd, bits, True);
```

Function `writeRegSize` is again a simple wrapper round an interface function of the same name:

```
inline static void writeRegSize(
    riscvP      riscv,
    riscvRegDesc r,
    Uns32       srcBits,
    Bool        signExtend
) {
    riscv->cb.writeRegSize(riscv, r, srcBits, signExtend);
}
```

The `writeRegSize` interface function will extend the value in abstract register `r` from the `srcBits` to the full register size encoded in the abstract register. The extension can either be zero-extension or sign-extension (in this case, sign-extension).

6.5.1 Required VMI Morph-Time Function Knowledge

When using embedded functions to implement extension functions, knowledge of only a small subset of the VMI Morph-Time Function function API is required. The necessary functions are those that specify function arguments, together with `vmimtCallResultAttrs` and its aliases:

```
//
// Add various argument types to the stack frame
//
void vmimtArgProcessor(void);
void vmimtArgUns32(Uns32 arg);
void vmimtArgUns64(Uns64 arg);
void vmimtArgFlt64(Flt64 arg);
void vmimtArgReg(vmiRegArgType argType, vmiReg r);
void vmimtArgRegSimAddress(Uns32 bits, vmiReg r);
void vmimtArgSimAddress(Addr arg);
void vmimtArgSimPC(Uns32 bits);
void vmimtArgNatAddress(const void *arg);

//
// Deprecated name for argument type function
//
#define vmimtArgDouble vmimtArgFlt64

//
// Make a call with all current stack frame arguments. If 'rd' is not VMI_NOREG,
```



```
// the function result (of size bits) is assigned to this register. Argument
// 'attrs' is used to define optimization attributes of a called function
// (see the comment preceding the definition of that type).
//
void vmimtCallResultAttrs(
    vmiCallFn    arg,
    Uns32        bits,
    vmiReg        rd,
    vmiCallAttrs attrs
);

//
// Backwards-compatible vmimtCall
//
#define vmimtCall(_ARG) \
    vmimtCallResultAttrs(_ARG, 0, VMI_NOREG, VMCA_NA)

//
// Backwards-compatible vmimtCallResult
//
#define vmimtCallResult(_ARG, _BITS, _RD) \
    vmimtCallResultAttrs(_ARG, _BITS, _RD, VMCA_NA)

//
// Backwards-compatible vmimtCallAttrs
//
#define vmimtCallAttrs(_ARG, _ATTRS) \
    vmimtCallResultAttrs(_ARG, 0, VMI_NOREG, _ATTRS)
```

Refer to the *VMI MorphTime Function Reference* manual for detailed information about these.

6.6 Example Execution

This example can be found in

Examples/Models/Processor/FeatureUsage/RISCV_ExtendedProcessor

Using the assembler test program

asmtest/test_ex_ch6.S

The code below shows the main part of the simple assembler program that can be used to exercise the extension:

```
//
// CHACHA20QRN <N>, <rd>, <rs1>, <rs2>
//
#define CHACHA20QRN(_N, _RD, _RS1, _RS2) .word ( \
    (0x0b << 0) | \
    (_RD << 7) | \
    ((_N-1) << 12) | \
    (_RS1 << 15) | \
    (_RS2 << 20) | \
    (0x01 << 25) \
)

//
// CHACHA20QR1-4 <rd> <rs1>, <rs2>
//
#define CHACHA20QR1(_RD, _RS1, _RS2) CHACHA20QRN(1, _RD, _RS1, _RS2)
#define CHACHA20QR2(_RD, _RS1, _RS2) CHACHA20QRN(2, _RD, _RS1, _RS2)
#define CHACHA20QR3(_RD, _RS1, _RS2) CHACHA20QRN(3, _RD, _RS1, _RS2)
#define CHACHA20QR4(_RD, _RS1, _RS2) CHACHA20QRN(4, _RD, _RS1, _RS2)
```

```
START_TEST:

    li      s0, 0x12345678
    li      s1, 0xaabbccdd

    // validate CHACHA instructions
    CHACHA20QR1(r_s2, r_s0, r_s1)
    CHACHA20QR2(r_s2, r_s0, r_s1)
    CHACHA20QR3(r_s2, r_s0, r_s1)
    CHACHA20QR4(r_s2, r_s0, r_s1)

EXIT_TEST
```

This can be run using the `iss.exe` simulator like this:

```
iss.exe \
--trace \
--tracechange \
--tracemode \
--traceshowicount \
--addressbits 32 \
--processorvendor vendor.com \
--processorname riscv \
--variant RV64X \
--program test.elf \
--extlib iss/cpu0=riscv.ovpworld.org/intercept/customControl/1.0 \
```

Which produces this output:

```
Info 1: 'iss/cpu0', 0x0000000080000000(_start): Machine 4000206f j      80002400
Info 2: 'iss/cpu0', 0x0000000080002400(START_TEST): Machine 12345437 lui      s0,0x12345
Info   s0 0000000000000000 -> 0000000012345000
Info 3: 'iss/cpu0', 0x0000000080002404(START_TEST+4): Machine 6784041b addiw   s0,s0,1656
Info   s0 0000000012345000 -> 0000000012345678
Info 4: 'iss/cpu0', 0x0000000080002408(START_TEST+8): Machine 000ab4b7 lui      s1,0xab
Info   s1 0000000000000000 -> 00000000000ab000
Info 5: 'iss/cpu0', 0x000000008000240c(START_TEST+c): Machine bbd4849b addiw   s1,s1,-
1091
Info   s1 00000000000ab000 -> 00000000000aabbd
Info 6: 'iss/cpu0', 0x0000000080002410(START_TEST+10): Machine 00c49493 slli    s1,s1,0xc
Info   s1 00000000000aabbd -> 00000000aabbd000
Info 7: 'iss/cpu0', 0x0000000080002414(START_TEST+14): Machine cdd48493 addi    s1,s1,-
803
Info   s1 00000000aabbd000 -> 00000000aabbccdd
Info 8: 'iss/cpu0', 0x0000000080002418(START_TEST+18): Machine 0294090b chacha20qr1
s2,s0,s1
Info   s2 0000000000000000 -> ffffffff9aa5b88f
Info 9: 'iss/cpu0', 0x000000008000241c(START_TEST+1c): Machine 0294190b chacha20qr2
s2,s0,s1
Info   s2 ffffffff9aa5b88f -> ffffffff9aa5b88
Info 10: 'iss/cpu0', 0x0000000080002420(START_TEST+20): Machine 0294290b chacha20qr3
s2,s0,s1
Info   s2 ffffffff9aa5b88 -> ffffffff8f9aa5b8
Info 11: 'iss/cpu0', 0x0000000080002424(START_TEST+24): Machine 0294390b chacha20qr4
s2,s0,s1
Info   s2 ffffffff8f9aa5b8 -> 0000000047cd52dc
Info 12: 'iss/cpu0', 0x0000000080002428(START_TEST+28): Machine 4501      li      a0,0
Info 13: 'iss/cpu0', 0x000000008000242a(START_TEST+2a): Machine custom0
```

Note the four custom instructions being executed and that result values are sign-extended from bit 31.

7 Adding Custom CSRs (addCSRs)

The `addCSRs` extension demonstrates how to add custom CSRs to a RISC-V model. The extension adds five CSRs in the custom space of a RISC-V processor. Three of the CSRs are implemented using plain registers, and one of these is read-only. The other two CSRs are implemented using callback functions. In detail, the CSRs are:

`custom_rw1_32`

This is a 32-bit M-mode CSR implemented as a plain register with a write mask (some bits are not writable). When accessed with XLEN 64, the value is zero extended from 32 to 64 bits.

`custom_rw1_64`

This is a 64-bit M-mode CSR implemented as a plain register with a write mask (some bits are not writable). When accessed with XLEN 32, the most-significant 32 bits are zero.

`custom_ro1`

This is a 64-bit M-mode read-only CSR implemented as a plain register.

`custom_rw3_32`

This is a 32-bit M-mode CSR implemented using callback functions.

`custom_rw4_64`

This is a 64-bit M-mode CSR implemented using callback functions.

In addition, the example shows how to modify the behavior of the existing `mstatus` CSR to add an extra field to it.

All behavior of this extension object is implemented in file `addCSRsExtensions.c` and header files `addCSRsCSR.h` and `addCSRsConfig.h`. Sections will be discussed in turn below.

7.1 CSR Type Definitions

Utility structures are defined for each of the CSRs implemented by this extension, in file `addCSRsCSR.h`. The structures use macros from file `riscvCSR.h` in the base model. 32-bit CSRs are defined using macros `CSR_REG_TYPE_32` (to define the structure type name) and `CSR_REG_STRUCT_DECL_32` (to define a union allowing the CSR to be accessed either using fields or as an entire 32-bit value). For example, here is the definition of custom CSR `custom_rw1_32` from file `addCSRsCSR.h`:

```
// 32-bit view
typedef struct {
    Uns32 F1 : 8;
    Uns32 _u1 : 8;
    Uns32 F3 : 8;
    Uns32 _u2 : 8;
} CSR_REG_TYPE_32(custom_rw1_32);

// define 32 bit type
```

```
CSR_REG_STRUCT_DECL_32(custom_rw1_32);

// define write masks
#define WM32_custom_rw1_32 0x00ff00ff
#define WM64_custom_rw1_32 0x00ff00ff
```

This CSR has two true fields (F1 and F3) and two unused field that are always zero. The last two `#define` lines specify *write masks* for the CSR, when it is written with XLEN of 32 and 64. In this case, the values are the same and allow change only to fields F1 and F3.

For 64-bit CSRs, equivalent macros `CSR_REG_TYPE_64` and `CSR_REG_STRUCT_DECL_64` are used instead. For example, here is the definition of custom CSR `custom_rw2_64` from file `addCSRsCSR.h`:

```
// 64-bit view
typedef struct {
    Uns32 F1 : 8;
    Uns32 _u1 : 8;
    Uns32 F3 : 8;
    Uns32 _u2 : 8;
    Uns32 F5 : 8;
    Uns32 _u3 : 8;
    Uns32 F7 : 8;
    Uns32 _u4 : 8;
} CSR_REG_TYPE_64(custom_rw2_64);

// define 32 bit type
CSR_REG_STRUCT_DECL_64(custom_rw2_64);

// define write masks
#define WM32_custom_rw2_64 0x00ff00ff
#define WM64_custom_rw2_64 0xff00ff0000ff00ffULL
```

In this case the write masks for XLEN of 32 and 64 differ, because only when XLEN is 64 are the most-significant bits accessible.

7.2 Extension-Specific Configuration

File `addCSRsConfig.h` defines a structure specifying extension-specific configuration information. In this case, the extension allows one CSR default value to be configurable:

```
typedef struct addCSRsConfigS {
    // extension CSR register values in configuration
    struct {
        CSR_REG_DECL(custom_rol);
    } csr;
} addCSRsConfig;

DEFINE_S (addCSRsConfig);
DEFINE_CS(addCSRsConfig);
```

The extension-specific configuration structure is used in the definition of variant configurations in file `riscvConfigList.h` of the linked model:

```
static riscvExtConfigCP allExtensions[] = {

    // example adding CSRs
    &(const riscvExtConfig){
```

```

        .id      = EXTID_ADDCSR,
        .userData = &(const addCSRsConfig){
            .csr = {
                .custom_rol = {u32 : {bits : 0x12345678}}
            }
        },
        . . . fields omitted . . .
    };

```

This specifies that the default value for custom CSR `custom_rol` should be 0x12345678.

7.3 Intercept Attributes

The behavior of the extension is defined using the standard `vmiosAttr` structure in `addCSRSExtensions.c`:

```

vmiosAttr modelAttrs = {

    //////////////////////////////////////
    // VERSION
    //////////////////////////////////////

    .versionString = VMI_VERSION,          // version string
    .modelType     = VMI_INTERCEPT_LIBRARY, // type
    .interceptType = VMI_IT_PROC_EXTENSION, // intercept type
    .packageName  = "addCSRs",             // description
    .objectSize   = sizeof(vmiosObject),    // size in bytes of OSS object

    //////////////////////////////////////
    // CONSTRUCTOR/DESTRUCTOR ROUTINES
    //////////////////////////////////////

    .constructorCB = addCSRsConstructor,    // object constructor

    //////////////////////////////////////
    // ADDRESS INTERCEPT DEFINITIONS
    //////////////////////////////////////

    .intercepts    = {{0}}
}

```

In this extension, there is a constructor to register the custom CSRs with the base model.

7.4 Object Type and Constructor

The object type is defined as follows:

```

typedef struct vmiosObjectS {

    // Info for associated processor
    riscvP      riscv;

    // configuration (including CSR reset values)
    addCSRsConfig config;

    // extension CSR info
    addCSRsCSRs  csr;                // extension CSR values
    riscvCSRAttrs csrs[XCSR_ID(LAST)]; // extension CSR definitions
    Bool         mstatus30;           // modified mstatus bit 30

    // extension callbacks
    riscvExtCB   extCB;
}

```

```
} vmiosObject;
```

To be compatible with the integration facilities described in this document, an extension object must contain the `riscv`, `config`, `csr`, `csrs` and `extCB` fields shown here. It may also contain other instance-specific fields: in this case, a field `mstatus30` is added, to hold the value of bit 30 of the `mstatus` register, which is a custom field added by this extension.

The `config` field holds configuration-specific information (used, for example, to hold initial values for read-only CSR fields when these are configuration dependent (see section 7.2).

The `csr` field holds current values of each CSR defined by this extension. The `addCSRSCSRs` value container type is defined in file `addCSRSCSR.h` like this:

```
typedef struct addCSRSCSRs {
    CSR_REG_DECL (custom_rw1_32);    // 0xBC0
    CSR_REG_DECL (custom_rw2_64);    // 0xBC1
    CSR_REG_DECL (custom_rw3_32);    // 0xBC2
    CSR_REG_DECL (custom_rw4_64);    // 0xBC3
    CSR_REG_DECL (custom_rol);       // 0xFC0
} addCSRSCSRs;
```

The `csrs` field holds description information for each CSR defined by this extension (required by the debugger and when tracing register value changes, for example). The `riscvCSRAttrs` type is defined in the base model:

```
typedef struct riscvCSRAttrs {
    const char      *name;           // register name
    const char      *desc;           // register description
    vmiosObjectP    object;          // custom extension
    Uns32           csrNum;           // CSR number (includes privilege and r/w access)
    riscvArchitecture arch;          // required architecture (presence)
    riscvArchitecture access;        // required architecture (access)
    riscvPrivVer    version;         // minimum specification version
    riscvCSRPresentFn presentCB;     // CSR present callback
    riscvCSRReadFn  readCB;          // read callback
    riscvCSRReadFn  readWriteCB;     // read callback (in r/w context)
    riscvCSRWriteFn writeCB;         // write callback
    riscvCSRWStateFn wstateCB;       // adjust JIT code generator state
    vmiReg          reg;             // register
    vmiReg          writeMaskV;      // configuration-dependent write mask
    Uns32           writeMaskC32;    // constant 32-bit write mask
    Uns64           writeMaskC64;    // constant 64-bit write mask

    riscvCSRStateenBit Smstateen :8; // whether xstateen-controlled access
    riscvCSRTrap      trap       :2; // whether trapped
    riscvCSRTrace     noTraceChange:2; // trace mode
    Bool              wEndBlock  :1; // whether write terminates this block
    Bool              wEndRM     :1; // whether write invalidates RM assumption
    Bool              noSaveRestore:1; // whether to exclude from save/restore
    Bool              writeRd    :1; // whether write updates Rd
    Bool              aliasV     :1; // whether CSR has virtual alias
    Bool              undefined  :1; // whether CSR is undefined
    Bool              forceRO    :1; // type is RO even if write CB defined
} riscvCSRAttrs;
```

Usage of this type is described in detail in Appendix 0; this chapter describes some common use cases.

Field `csrs` is an array of these structures, of size `XCSR_ID(LAST)`, which is a member of the `extCSRId` enumeration in `addCSRsExtensions.c`:

```
typedef enum extCSRIdE {  
  
    // custom CSRs, plain registers  
    XCSR_ID (custom_rw1_32),    // 0xBC0  
    XCSR_ID (custom_rw2_64),    // 0xBC1  
    XCSR_ID (custom_ro1),       // 0xFC0  
  
    // custom CSRs, implemented by callbacks  
    XCSR_ID (custom_rw3_32),    // 0xBC3  
    XCSR_ID (custom_rw4_64),    // 0xBC4  
  
    // base CSR with additional fields  
    XCSR_ID (mstatus),          // 0x300  
  
    // keep last (used to define size of the enumeration)  
    XCSR_ID (LAST)  
  
} extCSRId;
```

This enumeration contains one member for each custom CSR added by this extension, an entry for the standard `mstatus` CSR (whose behavior is being modified) and a final `LAST` member for sizing.

The constructor initializes the fields in the `vmiosObject` structure as follows:

```
static VMIOS_CONSTRUCTOR_FN(addCSRsConstructor) {  
  
    riscvP riscv = (riscvP)processor;  
  
    object->riscv = riscv;  
  
    // prepare client data  
    object->extCB.clientData = object;  
    object->extCB.resetNotifier = CSRReset;  
  
    // register extension with base model using unique ID  
    riscv->cb.registerExtCB(riscv, &object->extCB, EXTID_ADDCSR);  
  
    // copy configuration from template  
    object->config = *getExtConfig(riscv);  
  
    // initialize CSRs  
    addCSRsCSRInit(object);  
}
```

The `extCB` field defines the interface between the extension object and the base RISC-V model. The constructor must initialize the `clientData` field within that structure with the extension object and then call the `registerExtCB` interface function to register this extension object with the model. The last argument to this function is an identification number that should uniquely identify this extension object in the case that multiple libraries are installed on one RISC-V processor.

The `resetNotifier` field is initialized with a notifier function that is called whenever a processor reset occurs. This callback is called *after* all base model reset behavior has been performed, so it can update standard CSR values to reflect custom behavior if required. In this case, the notifier writes reset values to various custom CSRs:

```
static RISC_V_RESET_NOTIFIER_FN(CSRReset) {  
  
    vmiosObjectP object = clientData;  
  
    // reset custom mstatus field  
    object->mstatus30 = 0;  
  
    // reset custom CSRs by index  
    riscv->cb.writeCSR(riscv, 0xBC0, 0);  
    riscv->cb.writeCSR(riscv, 0xBC1, 0x1234);  
    riscv->cb.writeCSR(riscv, 0xBC2, 0);  
    riscv->cb.writeCSR(riscv, 0xBC3, 0);  
}
```

The reset notifier sets the `mstatus30` field to 0. It then calls the `writeCSR` interface function to reset all custom CSRs added by this extension to initial values. This interface function will write a CSR value given the index number of the CSR:

```
#define RISC_V_WRITE_CSR_NUM_FN(_NAME) Uns64 _NAME( \  
    riscvP riscv, \  
    Uns32  csrNum, \  
    Uns64  newValue \  
)  
typedef RISC_V_WRITE_CSR_NUM_FN((*riscvWriteCSRNumFn));  
  
typedef struct riscvModelCBS {  
    . . . fields omitted . . .  
    riscvWriteCSRNumFn writeCSR;  
    . . . fields omitted . . .  
} riscvModelCB;
```

The `config` field is initialized by copying default values from the *extension configuration* for this variant. The purpose of this is to allow different variants to be defined which have different default values for the extension registers. Function `getExtConfig` retrieves configuration information for the current processor:

```
static addCSRsConfigCP getExtConfig(riscvP riscv) {  
  
    riscvExtConfigCP cfg = riscv->cb.getExtConfig(riscv, EXTID_ADDCSR);  
  
    VMI_ASSERT(cfg, "ADDCSR config not found");  
  
    return cfg->userData;  
}
```

Function `addCSRsCSRInit` initializes the custom CSRs:

```
static void addCSRsCSRInit(vmiosObjectP object) {  
  
    riscvP  riscv = object->riscv;  
    extCSRId id;  
  
    // initialize CSR values that have configuration values defined  
    WR_XCSR(object, custom_rol, object->config.csr.custom_rol.u64.bits);  
  
    // register each CSR with the base model using the newCSR interface
```



```

// function
for(id=0; id<XCSR_ID(LAST); id++) {

    extCSRAttrsCP src = &csrs[id];
    riscvCSRAttrs *dst = &object->csrs[id];

    riscv->cb.newCSR(dst, &src->baseAttrs, riscv, object);
}

// perform initial CSR reset
CSRReset(riscv, object);
}

```

This function first initializes the *current* value of read-only CSR `custom_ro1` using the *default* value from the configuration:

```
WR_XCSR(object, custom_ro1, object->config.csr.custom_ro1.u64.bits);
```

The macro `WR_XCSR` is defined in file `riscvModelCallbackTypes.h` in the base model and is used to write the value of an entire extension CSR. The `for` loop iterates over all members of the `extCSRId` enumeration, filling one entry in the `csrs` array in the extension object from a matching entry in a static configuration template structure, `csrs`, using the `newCSR` interface function:

```

for(id=0; id<XCSR_ID(LAST); id++) {

    extCSRAttrsCP src = &csrs[id];
    riscvCSRAttrs *dst = &object->csrs[id];

    riscv->cb.newCSR(dst, &src->baseAttrs, riscv, object);
}

```

Copying the CSR definitions from a template into the `vmiosObject` structure in this way allows the definitions to be modified on an instance-specific basis if required (for example, using model parameters). The `csrs` template structure is defined like this:

```

static const extCSRAttrs csrs[XCSR_ID(LAST)] = {

    //
    // -----
    // CSRs IMPLEMENTED AS PLAIN REGISTERS
    // -----
    //
    //          name          num    arch extension attrs    description
rCB rwCB wCB
    XCSR_ATTR_TC(custom_rw1_32, 0xBC0, 0, 0, 0,0,0,0, "32-bit R/W CSR (plain)",
0, 0, 0),
    XCSR_ATTR_TC(custom_rw2_64, 0xBC1, 0, 0, 0,0,0,0, "XLEN R/W CSR (plain)",
0, 0, 0),
    XCSR_ATTR_TC(custom_ro1, 0xFC0, 0, 0, 0,0,0,0, "R/O CSR (plain)",
0, 0, 0),

    //
    // -----
    // CSRs IMPLEMENTED WITH CALLBACKS
    // -----
    //
    //          name          num    arch extension attrs    description
rCB          rwCB wCB
    XCSR_ATTR_TC(custom_rw3_32, 0xBC2, 0, 0, 0,0,0,0, "32-bit R/W CSR (cb)",
custom_rw3_32R, 0, custom_rw3_32W),

```

```

XCSR_ATTR_TC(custom_rw4_64, 0xBC3, 0, 0, 0,0,0,0, "XLEN R/W CSR (cb)",
custom_rw4_64R, 0, custom_rw4_64W),

//
// -----
// MODIFIED BEHAVIOR OF BASE CSR
// -----
//
//          name          num      arch extension attrs      description
rCB          rwCB wCB
XCSR_ATTR_P__(mstatus, 0x300, 0, 0, 0,0,0,0, "Machine Status",
mstatusR, 0, mstatusW ),
};

```

Each entry for a *new* CSR in the template is filled with CSR name, number, extension requirements, attributes, description and callbacks using macro `XCSR_ATTR_TC_` defined in file `riscvModelCallbackTypes.h` in the base model:

```

#define XCSR_ATTR_TC( \
_ID, _NUM, _ARCH, _EXT, _ENDB, _ENDRM, _NOTR, _TRAP, _DESC, _RCB, _RWCB, _WCB \
) [XCSR_ID(_ID)] = { \
    .extension = _EXT, \
    .baseAttrs = { \
        name      : #_ID, \
        desc      : _DESC, \
        csrNum     : _NUM, \
        arch      : _ARCH, \
        wEndBlock  : _ENDB, \
        wEndRM     : _ENDRM, \
        noTraceChange : _NOTR, \
        trap      : _TRAP, \
        readCB     : _RCB, \
        readWriteCB : _RWCB, \
        writeCB    : _WCB, \
        reg       : XCSR_REG_MT(_ID), \
        writeMaskC32 : WM32_##_ID, \
        writeMaskC64 : WM64_##_ID \
    } \
}

```

This macro defines a CSR with a constant write mask and optional callbacks. There are other similar macro variants for CSRs with no write mask, or with configurable write masks: see Appendix 0 for full details. The macro takes the following arguments:

1. The CSR *identifier*. This is used to construct both the CSR enumeration member name and the CSR name string (for reporting).
2. The CSR number, using standard RISC-V CSR numbering conventions.
3. Any architectural restrictions for the CSR, specified in the same way as architectural restrictions on instructions, discussed previously.
4. The extension identifier (see above).
5. *End-block* attribute: whether writes to the CSR should terminate a code block.
6. *End-rounding* attribute: whether writes to the CSR modify rounding mode.
7. *No-trace* attribute: specifies whether changes to the CSR are reported when trace change is enabled (RCSRT_YES indicates always reported, RCSRT_NO indicates never reported, RCSRT_VOLATILE indicates only reported if traceVolatile parameter is set in the base model).
8. A *trap* attribute: whether accesses to the CSR are trapped by `mstatus.TVM=1` (if *trap* is CSRT_TVM) or by `hvipcl.VTI=1` (if *trap* is CSRT_VTI).

9. A CSR description string (used in documentation generation).
10. An optional read callback function.
11. An optional read-modify-write callback function (only for CSRs such as `mip` that have special behavior in this case).
12. An optional write callback function.

Often, CSRs can be implemented as plain registers with no other associated behavior. In such cases, the callback fields in the template structure can be null. In this example, CSRs `custom_rw1_32`, `custom_rw2_64` and `custom_ro1` are all implemented as plain registers with defined write masks (see section 7.1).

When a CSR must have behavior associated with it, it must be implemented using callbacks. In this example, CSRs `custom_rw3_32` and `custom_rw4_64` are implemented in this way with read and write callbacks.

A *read* callback is called whenever the CSR is read (either by a true model access or in another way, for example by the debugger or when tracing CSR values). Read callbacks are defined using the `RISCV_CSR_READFN` macro, defined in `riscvCSRtypes.h` in the base model like this:

```
#define RISCV_CSR_READFN(_NAME) Uns64 _NAME( \
    riscvCSRAttrsCP attrs,      \
    riscvP                      riscv    \
)
typedef RISCV_CSR_READFN((*riscvCSRReadFn));
```

A CSR read callback is passed a description of the CSR being read (a pointer of type `riscvCSRAttrsCP`) and the RISC-V processor doing the read. The example read function for CSR `custom_rw3_32` is:

```
static RISCV_CSR_READFN(custom_rw3_32R) {
    vmiosObjectP object = attrs->object;
    Int32 result = RD_XCSR(object, custom_rw3_32);

    return result;
}
```

This function reads the entire value of the CSR using the `RD_XCSR` macro and returns it. In a real case, the callback would do more than this, because the same effect can be achieved with a plain register read.

Within a read (or write) callback function, the current extension object of type `vmiosObjectP` can be found using the expression `attrs->object`. This means that the callback can easily refer to any custom structures in the extension object. Whether this is a true access (by a processor) or an artifact access (by a debugger or for tracing) is indicated by the `artifactAccess` flag on the base processor, allowing the callback to modify its behavior in these cases. For example, to perform an operation only when a non-artifact access, the callback could contain an `if` statement:

```
if(!riscv->artifactAccess) {
```

```

    // behavior only if a true processor access
}

```

CSR *write* callbacks are defined using the `RISCV_CSR_WRITEFN` macro, defined in `riscvCSRtypes.h` like this:

```

#define RISCV_CSR_WRITEFN(_NAME) Uns64 _NAME( \
    riscvCSRAttrsCP attrs, \
    riscvP riscv, \
    Uns64 newValue \
)
typedef RISCV_CSR_WRITEFN(*riscvCSRWriteFn);

```

A CSR write callback is passes a description of the CSR being written, the RISC-V processor doing the read, and the new CSR value. The example write function for CSR `custom_rw3_32` is:

```

static RISCV_CSR_WRITEFN(custom_rw3_32W) {
    vmiosObjectP object = attrs->object;

    WR_XCSR(object, custom_rw3_32, newValue);

    return newValue;
}

```

This function writes the entire value of the CSR using the `WR_XCSR` macro and returns it. In a real case, the callback would do more than this, because the same effect can be achieved with a plain register write.

The CSR installation process automatically handles CSR access constraints based on address. For example, custom CSR `custom_rw3_32` is known to allow read/write access because its address (`0xBC0`) is in the custom read/write range, whereas CSR `custom_ro1` is known to be read-only because its address (`0xFC0`) is in the custom read-only range.

If a CSR is defined in an extension with the same number as a standard CSR in the base model, then the extension implementation will *override* that in the base. This allows extension objects to modify the behavior of standard CSRs in custom ways. In this example, CSR `mstatus` is redefined so that an extra one-bit field can be added to it (bit 30). The CSR is redefined using the `XCSR_ATTR_P__` macro, defined in file `riscvModelCallbackTypes.h` in the base model:

```

#define XCSR_ATTR_P__( \
    _ID, _NUM, _ARCH, _EXT, _ENDB, _ENDRM, _NOTR, _TRAP, _DESC, _RCB, _RWC, _WCB \
) [XCSR_ID(_ID)] = { \
    .extension = _EXT, \
    .baseAttrs = { \
        name      : #_ID, \
        desc      : _DESC, \
        csrNum    : _NUM, \
        arch      : _ARCH, \
        wEndBlock : _ENDB, \
        wEndRM    : _ENDRM, \
        noTraceChange : _NOTR, \
        trap      : _TRAP, \
        readCB    : _RCB, \
        readWriteCB : _RWC, \
        writeCB   : _WCB, \
    } \
}

```

```
        writeCB      : _WCB,                \
        writeMaskC32 : -1,                  \
        writeMaskC64 : -1,                  \
    }                                           \
}
```

This macro defines a CSR which is implemented using callbacks only (there is no field to hold the value for it in the `csr` structure in the extension object). In this case, `mstatus` is reimplemented using a read callback function (`mstatusR`) and write callback function (`mstatusW`). Function `mstatusR` is defined like this:

```
static RISCY_CSR_READFN(mstatusR) {
    vmiosObjectP object = attrs->object;

    // get value from base model
    mstatusU result = {u64 : riscv->cb.readBaseCSR(riscv, CSR_ID(mstatus))};

    // fill custom field from extension object
    result.f.custom1 = object->mstatus30;

    // return composed result
    return result.u64;
}
```

This function first uses the interface function `readBaseCSR` to read the value of `mstatus` from the base model into a union of type `mstatusU`:

```
mstatusU result = {u64 : riscv->cb.readBaseCSR(riscv, CSR_ID(mstatus))};
```

The `readBaseCSR` interface function has this prototype:

```
#define RISCY_READ_BASE_CSR_FN(_NAME) Uns64 _NAME( \
    riscvP      riscv,                \
    riscvCSRId id,                    \
)
typedef RISCY_READ_BASE_CSR_FN((*riscvReadBaseCSRFn));

typedef struct riscvModelCBS {
    . . . fields omitted . . .
    riscvReadBaseCSRFn      readBaseCSR;
    . . . fields omitted . . .
} riscvModelCB;
```

The `riscvCSRId` type defines the CSRs known to the base model. Type `mstatusU` is defined in the extension object like this:

```
typedef union {
    Uns64 u64;
    struct {
        Uns64 standard1 : 30;
        Uns64 custom1   : 1;
        Uns64 standard2 : 33;
    } f;
} mstatusU;
```

This bitfield structure defines the location of the new `custom1` field within the (otherwise opaque) `mstatus` CSR. Having read the `mstatus` value from the base model, function

`mstatusR` inserts the value of the `custom1` field from the master value held in the extension object, and returns the composed value:

```
// fill custom field from extension object
result.f.custom1 = object->mstatus30;

// return composed result
return result.u64;
```

Function `mstatusW` is defined like this:

```
static RISCY_CSR_WRITEFN(mstatusW) {

    vmiosObjectP object = attrs->object;

    // assign value to mstatusU for field extraction
    mstatusU result = {u64 : newValue};

    // extract custom field
    object->mstatus30 = result.f.custom1;

    // set value in base model
    result.u64 = riscv->cb.writeBaseCSR(riscv, CSR_ID(mstatus), newValue);

    // fill custom field from extension object
    result.f.custom1 = object->mstatus30;

    // return composed result
    return result.u64;
}
```

This function first saves the value being written to a union of type `mstatusU` and extracts the `custom1` field from that into the extension object field:

```
// assign value to mstatusU for field extraction
mstatusU result = {u64 : newValue};

// extract custom field
object->mstatus30 = result.f.custom1;
```

It then calls the interface function `writeBaseCSR` to write the value of `mstatus` in the base model:

```
result.u64 = riscv->cb.writeBaseCSR(riscv, CSR_ID(mstatus), newValue);
```

The `writeBaseCSR` interface function has this prototype:

```
#define RISCY_WRITE_BASE_CSR_FN(_NAME) Uns64 _NAME( \
    riscvP      riscv,          \
    riscvCSRId id,              \
    Uns64       newValue        \
)
typedef RISCY_WRITE_BASE_CSR_FN((*riscvWriteBaseCSRFn));

typedef struct riscvModelCBS {
    . . . fields omitted . . .
    riscvWriteBaseCSRFn      writeBaseCSR;
    . . . fields omitted . . .
} riscvModelCB;
```

Interface function `writeBaseCSR` returns the new value of the base model `mstatus` CSR, allowing for non-writable bits. To compose a result from the `mstatusw` function, the `mstatus30` bit is inserted into this return value:

```
// fill custom field from extension object
result.f.custom1 = object->mstatus30;

// return composed result
return result.u64;
```

7.5 Example Execution

This example can be found in

`Examples/Models/Processor/FeatureUsage/RISCV_ExtendedProcessor`

Using the assembler test program

`asmtest/test_ex_ch7.S`

The code below shows the main part of the simple assembler program that can be used to exercise the extension:

```
// 64-bit processor load mnemonics
#define SX sd
#define LX ld

.macro SETUP_M_HANDLER _BASE=defaultMHandler, _SCRATCH=defaultMScratch

    la        t0, \_BASE
    csrwr     mtvec, t0
    la        t0, \_SCRATCH
    csrwr     mscratch, t0
.endm

START_TEST:

    // set up default machine-mode exception handler
    SETUP_M_HANDLER customMHandler

    // read CSR initial values
    csrr      s1, 0xBC0
    csrr      s1, 0xBC1
    csrr      s1, 0xBC2
    csrr      s1, 0xBC3
    csrr      s1, 0xFC0

    // write CSRs
    li        s1, -1
    csrwr     0xBC0, s1
    csrwr     0xBC1, s1
    csrwr     0xBC2, s1
    csrwr     0xBC3, s1
    csrwr     0xFC0, s1

    // test mstatus with custom field at bit 30
    csrr      s1, mstatus
    li        s1, -1
    csrwr     mstatus, s1
    csrr      s1, mstatus
    EXIT_TEST

.align 6

customMHandler:
```

```

// save gp, a0, t0 (gp in scratch)
csrrw    gp, mscratch, gp
SX       a0, 0(gp)
SX       t0, 8(gp)

// calculate faulting instruction size in t0
csrr     a0, mepc
lhu      a0, 0(a0)
andi     a0, a0, 3
addi     a0, a0, -3
li       t0, 2
bnez     a0, 1f
addi     t0, t0, 2
1:
csrr     a0, mepc          // skip instruction
add      a0, a0, t0
csrw     mepc, a0

// restore registers and return
LX       a0, 0(gp)
LX       t0, 8(gp)
csrrw    gp, mscratch, gp
mret

```

This program attempts to read and write each custom CSR and the modified `mstatus` CSR. There is a simple exception handler to trap illegal accesses. This can be run using the `iss.exe` simulator like this:

```

iss.exe \
--trace                               \
--tracechange                         \
--tracemode                           \
--traceshowicount                     \
--addressbits 32                      \
--processorvendor vendor.com          \
--processorname riscv                 \
--variant RV64X                       \
--program test.elf                    \
--extlib iss/cpu0=riscv.ovpworld.org/intercept/customControl/1.0 \

```

It produces the following output:

```

Info 1: 'iss/cpu0', 0x0000000080000000(_start): Machine 4000206f j      80002400
Info 2: 'iss/cpu0', 0x0000000080002400(START_TEST): Machine 00000297 auipc  t0,0x0
Info  t0 0000000000000000 -> 0000000080002400
Info 3: 'iss/cpu0', 0x0000000080002404(START_TEST+4): Machine 08028293 addi  t0,t0,128
Info  t0 0000000080002400 -> 0000000080002480
Info 4: 'iss/cpu0', 0x0000000080002408(START_TEST+8): Machine 30529073 csrw   mtvec,t0
Info  mtvec 0000000000000000 -> 0000000080002480
Info 5: 'iss/cpu0', 0x000000008000240c(START_TEST+c): Machine fffff297 auipc  t0,0xfffff
Info  t0 0000000080002480 -> 000000008000140c
Info 6: 'iss/cpu0', 0x0000000080002410(START_TEST+10): Machine 5f428293 addi  t0,t0,1524
Info  t0 000000008000140c -> 0000000080001a00
Info 7: 'iss/cpu0', 0x0000000080002414(START_TEST+14): Machine 34029073 csrw   mscratch,t0
Info  mscratch 0000000000000000 -> 0000000080001a00
Info 8: 'iss/cpu0', 0x0000000080002418(START_TEST+18): Machine bc0024f3 csrr   s1,custom_rw1_32
Info 9: 'iss/cpu0', 0x000000008000241c(START_TEST+1c): Machine bc1024f3 csrr   s1,custom_rw2_64
Info  s1 0000000000000000 -> 0000000000000034
Info 10: 'iss/cpu0', 0x0000000080002420(START_TEST+20): Machine bc2024f3 csrr   s1,custom_rw3_32
Info  s1 0000000000000034 -> 0000000000000000

```



```

Info 11: 'iss/cpu0', 0x0000000080002424(START_TEST+24): Machine bc3024f3 csrr
s1,custom_rw4_64
Info 12: 'iss/cpu0', 0x0000000080002428(START_TEST+28): Machine fc0024f3 csrr
s1,custom_rol
Info s1 0000000000000000 -> 0000000012345678
Info 13: 'iss/cpu0', 0x000000008000242c(START_TEST+2c): Machine 54fd      li      s1,-1
Info s1 0000000012345678 -> ffffffff
Info 14: 'iss/cpu0', 0x000000008000242e(START_TEST+2e): Machine bc049073 csrw
custom_rw1_32,s1
Info custom_rw1_32 0000000000000000 -> 0000000000ff00ff
Info 15: 'iss/cpu0', 0x0000000080002432(START_TEST+32): Machine bc149073 csrw
custom_rw2_64,s1
Info custom_rw2_64 0000000000000034 -> ff00ff0000ff00ff
Info 16: 'iss/cpu0', 0x0000000080002436(START_TEST+36): Machine bc249073 csrw
custom_rw3_32,s1
Info custom_rw3_32 0000000000000000 -> ffffffff
Info 17: 'iss/cpu0', 0x000000008000243a(START_TEST+3a): Machine bc349073 csrw
custom_rw4_64,s1
Info custom_rw4_64 0000000000000000 -> ffffffff
Info 18: 'iss/cpu0', 0x000000008000243e(START_TEST+3e): Machine fc049073 csrw
custom_rol,s1
Info mstatus 0000000200000000 -> 0000000200001800
Info mepc 0000000000000000 -> 000000008000243e
Info mcause 0000000000000000 -> 0000000000000002
Info mtval 0000000000000000 -> 00000000fc049073
Info 19: 'iss/cpu0', 0x0000000080002480(customMHandler): Machine 340191f3 csrrw
gp,mscratch,gp
Info gp 0000000000000000 -> 0000000080001a00
Info mscratch 0000000080001a00 -> 0000000000000000
Info 20: 'iss/cpu0', 0x0000000080002484(customMHandler+4): Machine 00a1b023 sd
a0,0(gp)
Info 21: 'iss/cpu0', 0x0000000080002488(customMHandler+8): Machine 0051b423 sd
t0,8(gp)
Info 22: 'iss/cpu0', 0x000000008000248c(customMHandler+c): Machine 34102573 csrr
a0,mepc
Info a0 0000000000000000 -> 000000008000243e
Info 23: 'iss/cpu0', 0x0000000080002490(customMHandler+10): Machine 00055503 lhu
a0,0(a0)
Info a0 000000008000243e -> 0000000000009073
Info 24: 'iss/cpu0', 0x0000000080002494(customMHandler+14): Machine 890d      andi
a0,a0,3
Info a0 0000000000009073 -> 0000000000000003
Info 25: 'iss/cpu0', 0x0000000080002496(customMHandler+16): Machine 1575      addi
a0,a0,-3
Info a0 0000000000000003 -> 0000000000000000
Info 26: 'iss/cpu0', 0x0000000080002498(customMHandler+18): Machine 4289      li      t0,2
Info t0 0000000080001a00 -> 0000000000000002
Info 27: 'iss/cpu0', 0x000000008000249a(customMHandler+1a): Machine e111      bnez
a0,8000249e
Info 28: 'iss/cpu0', 0x000000008000249c(customMHandler+1c): Machine 0289      addi
t0,t0,2
Info t0 0000000000000002 -> 0000000000000004
Info 29: 'iss/cpu0', 0x000000008000249e(customMHandler+1e): Machine 34102573 csrr
a0,mepc
Info a0 0000000000000000 -> 000000008000243e
Info 30: 'iss/cpu0', 0x00000000800024a2(customMHandler+22): Machine 9516      add
a0,a0,t0
Info a0 000000008000243e -> 0000000080002442
Info 31: 'iss/cpu0', 0x00000000800024a4(customMHandler+24): Machine 34151073 csrw
mepc,a0
Info mepc 000000008000243e -> 0000000080002442
Info 32: 'iss/cpu0', 0x00000000800024a8(customMHandler+28): Machine 0001b503 ld
a0,0(gp)
Info a0 0000000080002442 -> 0000000000000000
Info 33: 'iss/cpu0', 0x00000000800024ac(customMHandler+2c): Machine 0081b283 ld
t0,8(gp)
Info t0 0000000000000004 -> 0000000080001a00
Info 34: 'iss/cpu0', 0x00000000800024b0(customMHandler+30): Machine 340191f3 csrrw
gp,mscratch,gp
Info gp 0000000080001a00 -> 0000000000000000
Info mscratch 0000000000000000 -> 0000000080001a00

```

```

Info 35: 'iss/cpu0', 0x00000000800024b4(customMHandler+34): Machine 30200073 mret
Info  mstatus 0000000200001800 -> 0000000200000080
Info 36: 'iss/cpu0', 0x0000000080002442(START_TEST+42): Machine 300024f3 csrr
s1,mstatus
Info  s1 ffffffffffffffff -> 0000000200000080
Info 37: 'iss/cpu0', 0x0000000080002446(START_TEST+46): Machine 54fd      li      s1,-1
Info  s1 0000000200000080 -> ffffffffffffffff
Info 38: 'iss/cpu0', 0x0000000080002448(START_TEST+48): Machine 30049073 csrw
mstatus,s1
Info  mstatus 0000000200000080 -> 8000000240227888
Info 39: 'iss/cpu0', 0x000000008000244c(START_TEST+4c): Machine 300024f3 csrr
s1,mstatus
Info  s1 ffffffffffffffff -> 8000000240227888
Info 40: 'iss/cpu0', 0x0000000080002450(START_TEST+50): Machine 4501      li      a0,0
Info 41: 'iss/cpu0', 0x0000000080002452(START_TEST+52): Machine custom0

```

In the trace output, note that:

1. CSR `custom_rw2_64` has initial value `0x34`, defined by the reset notifier (line 9). Although the reset notifier attempted to write `0x1234` to this CSR, the CSR write mask prevented read-only bits from being reset to non-zero values.
2. Attempting to write the read-only custom CSR `custom_ro1` causes an exception (line 18).
3. Standard CSR `mstatus` now has a writable custom field at bit 30 (lines 38 and 39).

7.6 General CSR Reset Template Code

Section 7.4 showed how the reset notifier can be used to reset extension model CSRs. This notifier can also be used to reset *standard* CSRs if required, using one of two base model interface service functions, depending on requirements:

writeCSR

This will write a CSR using the extension model view (if there is one) or the base model view (if not). The CSR is identified by *number*.

writeBaseCSR

This will write a CSR using the base model view, bypassing any extension model view. The CSR is identified by `riscvCSRId` value (defined in file `riscvCSR.h` in the base model).

Both these functions apply write masking and any CSR write callbacks, so they won't normally modify fields that wouldn't be writable by a CSR update instruction. Some CSRs have fields that can only be modified by hardware and are *read-only* for normal CSR access; these CSR fields can be unlocked by setting the `riscv->artifactAccess` Boolean to `True` while performing the write.

This example is a template showing how the standard Trigger Module CSRs can be reset in the reset notifier:

```

static RISC_V_RESET_NOTIFIER_FN(CSRReset) {

    Uns32 trigger_num = riscv->configInfo.trigger_num;
    Uns32 i;

```

```
// enable full write access to CSRs
riscv->artifactAccess = True;

// reset trigger module CSRs
for(i=0; i<trigger_num; i++) {
    riscv->cb.writeBaseCSR(riscv, CSR_ID(tselect), i);
    riscv->cb.writeBaseCSR(riscv, CSR_ID(tdata1), 0);
    riscv->cb.writeBaseCSR(riscv, CSR_ID(tdata2), 0);
    riscv->cb.writeBaseCSR(riscv, CSR_ID(tdata3), 0);
}

// reset tselect and tcontrol (not banked)
riscv->cb.writeBaseCSR(riscv, CSR_ID(tselect), 0);
riscv->cb.writeBaseCSR(riscv, CSR_ID(tcontrol), 0);

// disable full write access to CSRs
riscv->artifactAccess = False;
}
```

8 Adding Custom Exceptions (addExceptions)

The addExceptions extension demonstrates how to add custom exceptions to a RISC-V model. The extension adds a custom exception with cause 24, and also a custom instruction that triggers the exception.

All behavior of this extension object is implemented in file addExceptionExtensions.c. Sections will be discussed in turn below.

8.1 Exception Code

The new exception code is defined by the riscvExtException enumeration:

```
typedef enum riscvExtExceptionE {  
    EXT_E_EXCEPT24 = 24,  
} riscvExtException;
```

8.2 Intercept Attributes

The behavior of the extension is defined using the standard vmiosAttr structure:

```
vmiosAttr modelAttrs = {  
  
    ///////////////////////////////////////////  
    // VERSION  
    ///////////////////////////////////////////  
  
    .versionString = VMI_VERSION,          // version string  
    .modelType     = VMI_INTERCEPT_LIBRARY, // type  
    .interceptType = VMI_IT_PROC_EXTENSION, // intercept type  
    .packageName  = "addCSRs",            // description  
    .objectSize   = sizeof(vmiosObject),   // size in bytes of OSS object  
  
    ///////////////////////////////////////////  
    // CONSTRUCTOR/DESTRUCTOR ROUTINES  
    ///////////////////////////////////////////  
  
    .constructorCB = addExceptionsConstructor, // object constructor  
  
    ///////////////////////////////////////////  
    // INSTRUCTION INTERCEPT ROUTINES  
    ///////////////////////////////////////////  
  
    .morphCB      = addExceptionsMorph,      // instruction morph callback  
    .disCB        = addExceptionsDisassemble, // disassemble instruction  
  
    ///////////////////////////////////////////  
    // ADDRESS INTERCEPT DEFINITIONS  
    ///////////////////////////////////////////  
  
    .intercepts   = {{0}}  
}
```

In this library, there is a constructor, a JIT translation (morpher) function and a disassembly function.

8.3 Object Type and Constructor

The object type is defined as follows:

```
typedef struct vmiosObjectS {  
  
    // Info for associated processor  
    riscvP      riscv;  
  
    // extended instruction decode table  
    vmidDecodeTableP decode32;  
  
    // extension callbacks  
    riscvExtCB    extCB;  
  
} vmiosObject;
```

See chapter 6 for a detailed description of these fields, which are required when instructions are being added by an extension. The constructor initializes the fields as follows:

```
static VMIOS_CONSTRUCTOR_FN(addExceptionsConstructor) {  
  
    riscvP riscv = (riscvP)processor;  
  
    object->riscv = riscv;  
  
    // prepare client data  
    object->extCB.clientData = object;  
  
    // install custom exceptions  
    object->extCB.firstException = firstException;  
  
    // install notifier when trap is taken  
    object->extCB.trapNotifier = takeTrap;  
  
    // register extension with base model using unique ID  
    riscv->cb.registerExtCB(riscv, &object->extCB, EXTID_ADDEXCEPT);  
}
```

In addition to initializations that were explained in chapter 6, the constructor also initializes the `firstException` and `trapNotifier` fields in the interface object.

Function `firstException` returns the first exception description in a null-terminated list of exceptions implemented by this extension. It is defined like this:

```
//  
// Fill one member of exceptions  
//  
#define EXT_EXCEPTION(_NAME, _DESC) { \br/>    name: #_NAME, code: EXT_E_##_NAME, description: _DESC, \br/>}  
  
//  
// Table of exception descriptions  
//  
static const vmiExceptionInfo exceptions[] = {  
    EXT_EXCEPTION (EXCEPT24, "Custom Exception 24"),  
    {0}  
};  
  
//  
// Return the first exception implemented by the derived model  
//  
static RISCV_FIRST_EXCEPTION_FN(firstException) {  
    return exceptions;  
}
```

Each exception has a name, an index number (the cause number) and a description. In general, any number of exceptions can be specified in the list.

Function `takeTrap` is called whenever the RISC-V processor takes a trap of any kind (interrupt or exception). It gives the extension object a chance to update state that is dependent on that trap. In this case, `takeTrap` simply reports whenever the new exception is taken:

```
static RISC_V_TRAP_NOTIFIER_FN(takeTrap) {  
  
    // vmiosObjectP object = clientData;  
    Uns64      pc      = getPC(riscv);  
    Uns32      code    = RD_CSR_FIELD(riscv, mcause, ExceptionCode);  
  
    if(code==EXT_E_EXCEPT24) {  
        vmiMessage("I", CPU_PREFIX "_TRAP",  
                  SRCREF_FMT "TRAP:%u MODE:%u INTERRUPT:%u",  
                  SRCREF_ARGS(riscv, pc),  
                  code,  
                  mode,  
                  RD_CSR_FIELD(riscv, mcause, Interrupt)  
        );  
    }  
}
```

Macro `RD_CSR_FIELD` is defined in `riscvCSR.h` in the base model. It returns the value of a field within a standard CSR structure.

Function `takeTrap` is defined using the `RISC_V_TRAP_NOTIFIER_FN` macro defined in `riscvModelCallbacks.h` in the base model:

```
#define RISC_V_TRAP_NOTIFIER_FN(_NAME) void _NAME( \  
    riscvP      riscv,          \  
    riscvMode mode,            \  
    void        *clientData     \  
)  
typedef RISC_V_TRAP_NOTIFIER_FN((*riscvTrapNotifierFn));
```

The trap notifier is passed the executing RISC-V processor, the mode to which the trap is being taken and a `clientData` opaque pointer. This third argument is the `vmiosObjectP` pointer for the extension object; the commented-out line should be included if access to the intercept object is required in the notifier (in this case, it is not).

8.4 Instruction Decode

This extension implements a single new instruction that triggers the new custom exception. Adding new instructions is discussed in detail in chapter 6, so only brief details are given here.

The instruction type enumeration is:

```
typedef enum riscvExtITypeE {  
  
    // extension instructions  
    EXT_IT_EXCEPT24,  
  
    // KEEP LAST
```

```
EXT_IT_LAST  
} riscvExtITType;
```

The instruction table is this:

```
const static riscvExtInstrAttrs attrsArray32[] = {  
    // |  
    dec | rs2 | rs1 | fn3 | rd | dec |  
    EXT_INSTRUCTION(EXT_IT_EXCEPT24, "except24", RVANY, RVIP_RD_RS1_RS2, FMT_NONE,  
    "0000001|00000|00000|100|00000|0001011|")  
};
```

The instruction disassembly uses the value `FMT_NONE` to specify the instruction should be disassembled without arguments.

8.5 Instruction Disassembly

Disassembly uses an identical function to that described in chapter 6, so it is not described here.

8.6 Instruction Translation

The translation attribute table is specified like this:

```
const static riscvExtMorphAttr dispatchTable[] = {  
    [EXT_IT_EXCEPT24] = {morph:emitEXCEPT24},  
};
```

In this case, JIT translation function `emitEXCEPT24` is implemented like this:

```
static EXT_MORPH_FN(emitEXCEPT24) {  
    vmimtArgProcessor();  
    vmimtCall((vmiCallFn)takeExcept24);  
}
```

This emits code to call function `takeExcept24`, passing the current processor as an argument. Function `takeExcept24` calls interface function `takeException`, which causes the processor to take a standard exception with the numeric cause passed as the second argument:

```
static void takeExcept24(riscvP riscv) {  
    riscv->cb.takeException(riscv, EXT_E_EXCEPT24, 0);  
}
```

8.7 Example Execution

This example can be found in

Examples/Models/Processor/FeatureUsage/RISCV_ExtendedProcessor

Using the assembler test program

asmtest/test_ex_ch8.S

The code below shows the main part of the simple assembler program that can be used to exercise the extension:

```

#define EXCEPT24 .word ( \
    (0x0b << 0) | \
    (0 << 7) | \
    (4 << 12) | \
    (0 << 15) | \
    (0 << 20) | \
    (0x01 << 25) \
)

START_TEST:

    // set up default machine-mode exception handler
    SETUP_M_HANDLER customMHandler

    // validate EXCEPT24 instruction
    EXCEPT24

    // check medeleg
    li      s0, -1
    csrwr   medeleg, s0

    EXIT_TEST

.align 6

customMHandler:

    // save gp, a0, t0 (gp in scratch)
    csrrw   gp, mscratch, gp
    SX      a0, 0(gp)
    SX      t0, 8(gp)

    // calculate faulting instruction size in t0
    csrr     a0, mepc
    lhu      a0, 0(a0)
    andi     a0, a0, 3
    addi     a0, a0, -3
    li       t0, 2
    bnez     a0, 1f
    addi     t0, t0, 2
1:
    csrr     a0, mepc          // skip instruction
    add      a0, a0, t0
    csrwr   mepc, a0

    // restore registers and return
    LX      a0, 0(gp)
    LX      t0, 8(gp)
    csrrw   gp, mscratch, gp
    mret

```

This can be run using the `iss.exe` simulator like this:

```

iss.exe \
    --trace \
    --tracechange \
    --tracemode \
    --traceshowicount \
    --addressbits 32 \
    --processorvendor vendor.com \
    --processorname riscv \
    --variant RV64X \
    --program test.elf \
    --extlib iss/cpu0=riscv.ovpworld.org/intercept/customControl/1.0 \

```


Which produces this output:

```

Info 1: 'iss/cpu0', 0x0000000080000000(_start): Machine 4000206f j      80002400
Info 2: 'iss/cpu0', 0x0000000080002400(START_TEST): Machine 00000297 auipc  t0,0x0
Info  t0 0000000000000000 -> 0000000080002400
Info 3: 'iss/cpu0', 0x0000000080002404(START_TEST+4): Machine 04028293 addi  t0,t0,64
Info  t0 0000000080002400 -> 0000000080002440
Info 4: 'iss/cpu0', 0x0000000080002408(START_TEST+8): Machine 30529073 csrw  mtvec,t0
Info  mtvec 0000000000000000 -> 0000000080002440
Info 5: 'iss/cpu0', 0x000000008000240c(START_TEST+c): Machine fffff297 auipc  t0,0xfffff
Info  t0 0000000080002440 -> 000000008000140c
Info 6: 'iss/cpu0', 0x0000000080002410(START_TEST+10): Machine 5f428293 addi  t0,t0,1524
Info  t0 000000008000140c -> 0000000080001a00
Info 7: 'iss/cpu0', 0x0000000080002414(START_TEST+14): Machine 34029073 csrw  mscratch,t0
Info  mscratch 0000000000000000 -> 0000000080001a00
Info 8: 'iss/cpu0', 0x0000000080002418(START_TEST+18): Machine 0200400b except24
Info (ADD_EXCEPT_TRAP) CPU 'iss/cpu0' 0x80002440 340191f3 csrrw  gp,mscratch,gp: TRAP:24
MODE:3 INTERRUPT:0
Info  mstatus 0000000a00000000 -> 0000000a00001800
Info  mepc 0000000000000000 -> 0000000080002418
Info  mcause 0000000000000000 -> 0000000000000018
Info 9: 'iss/cpu0', 0x0000000080002440(customMHandler): Machine 340191f3 csrrw  gp,mscratch,gp
Info  gp 0000000000000000 -> 0000000080001a00
Info  mscratch 0000000080001a00 -> 0000000000000000
Info 10: 'iss/cpu0', 0x0000000080002444(customMHandler+4): Machine 00a1b023 sd  a0,0(gp)
Info 11: 'iss/cpu0', 0x0000000080002448(customMHandler+8): Machine 0051b423 sd  t0,8(gp)
Info 12: 'iss/cpu0', 0x000000008000244c(customMHandler+c): Machine 34102573 csrr  a0,mepc
Info  a0 0000000000000000 -> 0000000080002418
Info 13: 'iss/cpu0', 0x0000000080002450(customMHandler+10): Machine 00055503 lhu  a0,0(a0)
Info  a0 0000000080002418 -> 000000000000400b
Info 14: 'iss/cpu0', 0x0000000080002454(customMHandler+14): Machine 890d      andi  a0,a0,3
Info  a0 000000000000400b -> 0000000000000003
Info 15: 'iss/cpu0', 0x0000000080002456(customMHandler+16): Machine 1575      addi  a0,a0,-3
Info  a0 0000000000000003 -> 0000000000000000
Info 16: 'iss/cpu0', 0x0000000080002458(customMHandler+18): Machine 4289      li    t0,2
Info  t0 0000000080001a00 -> 0000000000000002
Info 17: 'iss/cpu0', 0x000000008000245a(customMHandler+1a): Machine e111      bnez  a0,8000245e
Info 18: 'iss/cpu0', 0x000000008000245c(customMHandler+1c): Machine 0289      addi  t0,t0,2
Info  t0 0000000000000002 -> 0000000000000004
Info 19: 'iss/cpu0', 0x000000008000245e(customMHandler+1e): Machine 34102573 csrr  a0,mepc
Info  a0 0000000000000000 -> 0000000080002418
Info 20: 'iss/cpu0', 0x0000000080002462(customMHandler+22): Machine 9516      add  a0,a0,t0
Info  a0 0000000080002418 -> 000000008000241c
Info 21: 'iss/cpu0', 0x0000000080002464(customMHandler+24): Machine 34151073 csrw  mepc,a0
Info  mepc 0000000080002418 -> 000000008000241c
Info 22: 'iss/cpu0', 0x0000000080002468(customMHandler+28): Machine 0001b503 ld  a0,0(gp)
Info  a0 000000008000241c -> 0000000000000000
Info 23: 'iss/cpu0', 0x000000008000246c(customMHandler+2c): Machine 0081b283 ld  t0,8(gp)
Info  t0 0000000000000004 -> 0000000080001a00
Info 24: 'iss/cpu0', 0x0000000080002470(customMHandler+30): Machine 340191f3 csrrw  gp,mscratch,gp
Info  gp 0000000080001a00 -> 0000000000000000
Info  mscratch 0000000000000000 -> 0000000080001a00
Info 25: 'iss/cpu0', 0x0000000080002474(customMHandler+34): Machine 30200073 mret

```

```
Info  mstatus 0000000a00001800 -> 0000000a00000080
Info 26: 'iss/cpu0', 0x000000008000241c(START_TEST+1c): Machine 547d      li      s0,-1
Info  s0 0000000000000000 -> ffffffff
Info 27: 'iss/cpu0', 0x000000008000241e(START_TEST+1e): Machine 30241073 csr
medeleg,s0
Info  medeleg 0000000000000000 -> 000000000000b3ff
Info 28: 'iss/cpu0', 0x0000000080002422(START_TEST+22): Machine 4501      li      a0,0
Info 29: 'iss/cpu0', 0x0000000080002424(START_TEST+24): Machine custom0
```

At instruction 8, the `except24` extension instruction is executed, causing a Machine mode exception with cause 0x18 (24).

9 Adding Custom Local Interrupts (addLocalInterrupts)

The RISC-V architecture allows a processor to add custom interrupts, with numbers 16-31 (for RV32) and 16-63 (for RV64). The `addLocalInterrupts` extension demonstrates how to add two such local interrupts to a RISC-V model, with numbers 21 and 22.

Most behavior of this extension object is implemented in file `addLocalInterruptsExtensions.c`, but the processor configuration in file `riscvConfigList.c` of the linked processor model must also be modified to enable the local interrupt ports. Sections will be discussed in turn below.

9.1 Enabling Local Interrupt Ports

Local interrupt ports 21 and 22 are enabled by two lines in the configuration structure for each processor variant (in file `riscvConfigList.c` of the linked processor model):

```
static const riscvConfig configList[] = {  
    {  
        .name           = "RV32X",  
        .arch           = ISA_U|RV32GC|ISA_X,  
        .user_version    = RVUV_DEFAULT,  
        .priv_version    = RVPV_DEFAULT,  
        .tval_ii_code    = True,  
        .ASID_bits       = 9,  
        .local_int_num   = 7,           // enable local interrupts 16-22  
        .unimp_int_mask  = 0x1f0000,   // int16-int20 absent  
        .extensionConfigs = allExtensions,  
    },  
    {  
        .name           = "RV64X",  
        .arch           = ISA_U|RV64GC|ISA_X,  
        .user_version    = RVUV_DEFAULT,  
        .priv_version    = RVPV_DEFAULT,  
        .tval_ii_code    = True,  
        .ASID_bits       = 9,  
        .local_int_num   = 7,           // enable local interrupts 16-22  
        .unimp_int_mask  = 0x1f0000,   // int16-int20 absent  
        .extensionConfigs = allExtensions,  
    },  
    {0} // null terminator  
};
```

Specifying `local_int_num` of 7 indicates that local interrupts 16-22 are potentially implemented. Then, the specification of `unimp_int_mask` of `0x1f0000` indicates that local interrupts 16-20 are *not* implemented, leaving local interrupts 21 and 22 as the only implemented local interrupts. Using a combination of `local_int_num` and `unimp_int_mask` in this way allows any subset of the defined local interrupts to be specified as implemented.

9.2 Interrupt Codes

The new local interrupt codes are defined by the `riscvExtInt` enumeration in `addLocalInterruptsExtensions.c`:

```
typedef enum riscvExtIntE {
    EXT_I_INT21 = 21,
    EXT_I_INT22 = 22
} riscvExtInt;
```

9.3 Intercept Attributes

The behavior of the extension is defined using the standard `vmiosAttr` structure:

```
vmiosAttr modelAttrs = {

    //////////////////////////////////////////
    // VERSION
    //////////////////////////////////////////

    .versionString = VMI_VERSION,          // version string
    .modelType     = VMI_INTERCEPT_LIBRARY, // type
    .interceptType = VMI_IT_PROC_EXTENSION, // intercept type
    .packageName  = "addCSRs",             // description
    .objectSize   = sizeof(vmiosObject),    // size in bytes of OSS object

    //////////////////////////////////////////
    // CONSTRUCTOR/DESTRUCTOR ROUTINES
    //////////////////////////////////////////

    .constructorCB = addLocalInterruptsConstructor, // object constructor

    //////////////////////////////////////////
    // ADDRESS INTERCEPT DEFINITIONS
    //////////////////////////////////////////

    .intercepts    = {{0}}
};
```

In this library, there is a constructor that implements installation of the local interrupts.

9.4 Object Type and Constructor

The object type is defined as follows:

```
typedef struct vmiosObjects {

    // Info for associated processor
    riscvP      riscv;

    // extension callbacks
    riscvExtCB  extCB;

} vmiosObject;
```

The constructor initializes the fields as follows:

```
static VMIO_CONSTRUCTOR_FN(addLocalInterruptsConstructor) {

    riscvP riscv = (riscvP)processor;

    object->riscv = riscv;

    // prepare client data
    object->extCB.clientData = object;

    // install notifier for suppression of memory exceptions
    object->extCB.getInterruptPri = getInterruptPriority;

    // install notifier when trap is taken
```

```
object->extCB.trapNotifier = takeTrap;

// register extension with base model using unique ID
riscv->cb.registerExtCB(riscv, &object->extCB, EXTID_ADDLOCALINT);
}
```

In addition to initializations that were explained in previous chapters, the constructor also initializes the `getInterruptPri` and `trapNotifier` fields.

Function `getInterruptPriority` specifies the *priority* of the new local interrupts, with respect to standard interrupts and each other. This priority determines which interrupt is taken if multiple interrupts become pending at the same time. The function is defined by the `RISCV_GET_INTERRUPT_PRI_FN` macro in `riscvModelCallbacks.h`:

```
#define RISCV_GET_INTERRUPT_PRI_FN(_NAME) riscvExceptionPriority _NAME( \
    riscvP riscv, \
    Uns32 intNum, \
    void *clientData \
)
typedef RISCV_GET_INTERRUPT_PRI_FN((*riscvGetInterruptPriFn));
```

The interrupt priority callback is passed the current RISC-V processor, an interrupt number and a client data pointer (which is in fact the `vmiosObjectP` pointer for the current extension). It should return either 0 (if the interrupt is not recognized by this extension) or a priority based on the fixed priorities defined by the `riscvExceptionPriority` enumeration:

```
typedef enum riscvExceptionPriorityE {

    // this is the lowest architectural priority
    riscv_E_MinPriority      = -120,

    // low-priority local interrupt default priorities when AIA present
    riscv_E_Local32Priority  = riscv_E_MinPriority,
    riscv_E_Local16Priority  = -110,
    riscv_E_Local33Priority  = -100,
    riscv_E_Local34Priority  = -90,
    riscv_E_Local17Priority  = -80,
    riscv_E_Local35Priority  = -70,
    riscv_E_Local36Priority  = -60,
    riscv_E_Local18Priority  = -50,
    riscv_E_Local37Priority  = -40,
    riscv_E_Local38Priority  = -30,
    riscv_E_Local19Priority  = -20,
    riscv_E_Local39Priority  = -10,

    // standard major interrupts (0-15)
    riscv_E_UTimerPriority   = 10,
    riscv_E_USWPriority      = 20,
    riscv_E_UExternalPriority = 30,
    riscv_E_VSTimerPriority  = 40,
    riscv_E_VSSWPriority     = 50,
    riscv_E_VSEExternalPriority = 60,
    riscv_E_SGEIPriority    = 70,
    riscv_E_STimerPriority   = 80,
    riscv_E_SSWPriority      = 90,
    riscv_E_SEExternalPriority = 100,
    riscv_E_MTimerPriority   = 110,
    riscv_E_MSWSWPriority    = 120,
    riscv_E_MExternalPriority = 130,

    // high-priority local interrupt default priorities when AIA present
```

```
riscv_E_Local40Priority    = 140,
riscv_E_Local20Priority    = 150,
riscv_E_Local41Priority    = 160,
riscv_E_Local42Priority    = 170,
riscv_E_Local21Priority    = 180,
riscv_E_Local43Priority    = 190,
riscv_E_Local44Priority    = 200,
riscv_E_Local22Priority    = 210,
riscv_E_Local45Priority    = 220,
riscv_E_Local46Priority    = 230,
riscv_E_Local23Priority    = 240,
riscv_E_Local47Priority    = 250,

// local interrupt default priorities when AIA absent
riscv_E_LocalPriority      = 260

} riscvExceptionPriority;
```

In this case, the extension defines that both custom interrupts are of higher priority than all standard interrupts, with interrupt 22 being the highest priority of all:

```
static RISCV_GET_INTERRUPT_PRI_FN(getInterruptPriority) {

    riscvExceptionPriority result = 0;

    if(intNum==EXT_I_INT21) {
        result = riscv_E_LocalPriority;
    } else if(intNum==EXT_I_INT22) {
        result = riscv_E_LocalPriority+1;
    }

    return result;
}
```

Note that the exact value returned by `getInterruptPriority` is not significant: what matters is the *relative order* of different interrupt priority codes. The gaps in the specified priorities of standard interrupts mean that a local interrupt can be defined to have any intermediate priority between standard priorities: for example, a local interrupt of priority `riscv_E_MSWPriority+1` would be higher priority than a Machine Software Interrupt, but lower priority than a Machine External Interrupt.

Function `takeTrap` is called whenever the RISC-V processor takes a trap of any kind (interrupt or exception). It gives the extension object a chance to update state that is dependent on that trap. In this case, `takeTrap` simply reports whenever an interrupt is taken:

```
static RISCV_TRAP_NOTIFIER_FN(takeTrap) {

    // vmiosObjectP object = clientData;
    Uns64 pc = getPC(riscv);
    Bool isInt = RD_CSR_FIELD(riscv, mcause, Interrupt);

    if(isInt) {
        vmiMessage("I", CPU_PREFIX "_TRAP",
            SRCREF_FMT "TRAP:%u MODE:%u INTERRUPT:%u",
            SRCREF_ARGS(riscv, pc),
            RD_CSR_FIELD(riscv, mcause, ExceptionCode),
            mode,
            RD_CSR_FIELD(riscv, mcause, Interrupt)
        );
    }

}
```

See chapter 8 for a detailed explanation of trap notifier behavior.

9.5 Example Execution

When local interrupts are configured, the base model automatically modifies behavior of related CSRs to reflect their presence. For example, `mie` and `mideleg` bit fields corresponding to the new local interrupt positions become writable. To demonstrate this, the following test program shows writability of these two registers and also validates the behavior of the `sie` register, which is dependent upon `mideleg`.

This example can be found in

`Examples/Models/Processor/FeatureUsage/RISCV_ExtendedProcessor`

Using the assembler test program

`asmtest/test_ex_ch9.S`

The code below shows the main part of the simple assembler program that can be used to exercise the extension:

```
START_TEST:

    li        s0, -1

    // check sie (delegation disabled)
    csrwr     sie, zero
    csrwr     sie, s0

    // check mie & mideleg
    csrwr     mie, s0
    csrwr     mideleg, s0

    // check sie (delegation enabled)
    csrwr     sie, zero
    csrwr     sie, s0

EXIT_TEST
```

This can be run using the `iss.exe` simulator like this:

```
iss.exe \
--trace                \
--tracechange           \
--tracemode             \
--traceshowicount       \
--addressbits           32 \
--processorvendor       vendor.com \
--processorname         riscv \
--variant               RV64X \
--override              iss/cpu0/add_Extensions="S" \
--program               test.elf \
--extlib                iss/cpu0=riscv.ovpworld.org/intercept/customControl/1.0
```

Which produces this output:

```
Info 1: 'iss/cpu0', 0x0000000080000000(_start): Machine 4000206f j      80002400
Info 2: 'iss/cpu0', 0x0000000080002400(START_TEST): Machine 547d    li      s0,-1
Info   s0 0000000000000000 -> ffffffff
Info 3: 'iss/cpu0', 0x0000000080002402(START_TEST+2): Machine 10401073 csrwr   sie,zero
```

```
Info 4: 'iss/cpu0', 0x0000000080002406(START_TEST+6): Machine 10441073 csrwr sie,s0
Info 5: 'iss/cpu0', 0x000000008000240a(START_TEST+a): Machine 30441073 csrwr mie,s0
Info mie 0000000000000000 -> 0000000000600aaa
Info 6: 'iss/cpu0', 0x000000008000240e(START_TEST+e): Machine 30341073 csrwr mideleg,s0
Info sie 0000000000000000 -> 0000000000600222
Info mideleg 0000000000000000 -> 0000000000600222
Info 7: 'iss/cpu0', 0x0000000080002412(START_TEST+12): Machine 10401073 csrwr sie,zero
Info sie 0000000000600222 -> 0000000000000000
Info mie 0000000000600aaa -> 0000000000000888
Info 8: 'iss/cpu0', 0x0000000080002416(START_TEST+16): Machine 10441073 csrwr sie,s0
Info sie 0000000000000000 -> 0000000000600222
Info mie 0000000000000888 -> 0000000000600aaa
Info 9: 'iss/cpu0', 0x000000008000241a(START_TEST+1a): Machine 4501 li a0,0
Info 10: 'iss/cpu0', 0x000000008000241c(START_TEST+1c): Machine custom0
```

Instructions 3 and 4 attempt to write all-zeros and all-ones to `sie`. This has no effect since no interrupts are by default delegated to Supervisor mode:

```
Info 3: 'iss/cpu0', 0x0000000080002402(START_TEST+2): Machine 10401073 csrwr sie,zero
Info 4: 'iss/cpu0', 0x0000000080002406(START_TEST+6): Machine 10441073 csrwr sie,s0
```

Instructions 5 and 6 attempt to write all-ones to `mie` and `mideleg`. The writes update bits corresponding to the custom local interrupt positions (as well as standard interrupts):

```
Info 5: 'iss/cpu0', 0x000000008000240a(START_TEST+a): Machine 30441073 csrwr mie,s0
Info mie 0000000000000000 -> 0000000000600aaa
Info 6: 'iss/cpu0', 0x000000008000240e(START_TEST+e): Machine 30341073 csrwr mideleg,s0
Info sie 0000000000000000 -> 0000000000600222
Info mideleg 0000000000000000 -> 0000000000600222
```

Instructions 7 and 8 once more attempt to write all-zeros and all-ones to `sie`. This now has an effect for the interrupts that have been delegated to Supervisor mode:

```
Info 7: 'iss/cpu0', 0x0000000080002412(START_TEST+12): Machine 10401073 csrwr sie,zero
Info sie 0000000000600222 -> 0000000000000000
Info mie 0000000000600aaa -> 0000000000000888
Info 8: 'iss/cpu0', 0x0000000080002416(START_TEST+16): Machine 10441073 csrwr sie,s0
Info sie 0000000000000000 -> 0000000000600222
Info mie 0000000000000888 -> 0000000000600aaa
```

Although not shown by this simple example, net ports are now available for the new local interrupts so that they can be driven externally.

10 Adding Custom FIFOs (`fifoExtensions`)

The `fifoExtensions` extension object extends the basic RISC-V model by adding FIFO input and output ports and some new instructions to put data into a FIFO and get data from a FIFO. The instructions will block if the FIFO is full on a put or empty on a get.

All behavior of this extension object is implemented in file `fifoExtensions.c`. Sections will be discussed in turn below.

10.1 Intercept Attributes

The behavior of the library is defined using the standard `vmiosAttr` structure:

```
vmiosAttr modelAttrs = {  
  
    ///////////////////////////////////////////  
    // VERSION  
    ///////////////////////////////////////////  
  
    .versionString = VMI_VERSION,          // version string  
    .modelType     = VMI_INTERCEPT_LIBRARY, // type  
    .interceptType = VMI_IT_PROC_EXTENSION, // intercept type  
    .packageName  = "fifoExtensions",      // description  
    .objectSize   = sizeof(vmiosObject),    // size in bytes of OSS object  
  
    ///////////////////////////////////////////  
    // CONSTRUCTOR/DESTRUCTOR ROUTINES  
    ///////////////////////////////////////////  
  
    .constructorCB = fifoConstructor,      // object constructor  
    .docCB         = fifoDoc,              // documentation constructor  
  
    ///////////////////////////////////////////  
    // INSTRUCTION INTERCEPT ROUTINES  
    ///////////////////////////////////////////  
  
    .morphCB      = fifoMorph,             // instruction translation callback  
    .disCB        = fifoDisassemble,       // disassemble instruction  
  
    ///////////////////////////////////////////  
    // PORT ACCESS ROUTINES  
    ///////////////////////////////////////////  
  
    .fifoPortSpecsCB = fifoGetPortSpec,    // callback for next fifo port  
  
    ///////////////////////////////////////////  
    // PARAMETER CALLBACKS  
    ///////////////////////////////////////////  
  
    .paramSpecsCB      = fifoParamSpecs,    // iterate parameter declarations  
    .paramValueSizeCB = fifoParamTableSize, // get parameter table size  
  
    ///////////////////////////////////////////  
    // ADDRESS INTERCEPT DEFINITIONS  
    ///////////////////////////////////////////  
  
    .intercepts      = {{0}}  
};
```

In this library, there is a constructor, a documentation callback, a JIT translation (morpher) function and a disassembly function, as in the previous example. In addition,

there are functions to define the FIFO ports and to allow parameterization of the FIFO ports.

10.2 Object Type and Constructor

The object type is defined as follows:

```
typedef struct vmiosObjectS {  
  
    // Info for associated processor  
    riscvP      riscv;  
  
    // parameters  
    vmiParameterP  parameters;  
  
    // is this extension enabled?  
    Bool          enabled;  
  
    // temporary FIFO element  
    Uns64         FIFOTmp;  
  
    // configuration (including CSR reset values)  
    fifoConfig     config;  
  
    // extension CSR info  
    fifoCSRs       csr;                // FIFO extension CSR values  
    riscvCSRAttrs  csrs[XCSR_ID(LAST)]; // modified CSR definitions  
  
    // FIFO connections  
    vmiFifoPortP    fifoPorts;          // fifo port descriptions  
    vmiConnInputP   inputConn;          // input FIFO connection  
    vmiConnOutputP  outputConn;         // output FIFO connection  
  
    // extended instruction decode table  
    vmidDecodeTableP decode32;  
  
    // extension callbacks  
    riscvExtCB      extCB;  
  
} vmiosObject;
```

This structure contains the `riscv`, `decode32` and `extCB` fields that are always required when using the RISC-V extension support infrastructure documented here, together with a number of other extension-specific fields. The constructor initializes the fields as follows:

```
static VMIOS_CONSTRUCTOR_FN(fifoConstructor) {  
  
    riscvP      riscv = (riscvP)processor;  
    paramValuesP params = parameterValues;  
  
    object->riscv = riscv;  
  
    // prepare client data  
    object->extCB.clientData = object;  
  
    // register extension with base model  
    riscv->cb.registerExtCB(riscv, &object->extCB, EXTID_FIFO);  
  
    // copy configuration from template  
    object->config = *getExtConfig(riscv);  
  
    // override parameterized values  
    object->config.FIFO_bits = params->FIFO_bits;
```

```
// initialize CSRs
fifoCSRInit(object);

// is extension enabled in config info?
object->enabled = RD_FIFO_CSR_FIELD(object, fifo_cfg, fifoPresent);

// create fifoPorts
newFifoPorts(object);
}
```

In addition to the mandatory field setup, this constructor also defines extension specific CSRs and FIFO ports, as described in the next sections.

10.3 Extension CSRs

Extension `fifoExtensions` adds a single read-only CSR to the base model. The CSR is defined in file `fifoCSR.h`. An enumeration in that file first defines the set of additional CSRs:

```
typedef enum extCSRIdE {

    // FIFO configuration control and status registers
    XCSR_ID (fifo_cfg),      // 0xFF0

    // keep last (used to define size of the enumeration)
    XCSR_ID (LAST)

} extCSRId;
```

Then the fields in the `fifo_cfg` CSR are defined using a bitfield structure:

```
// -----
// fifo_cfg      (id 0xFF0)
// -----

// 32-bit view
typedef struct {
    Uns32 fifoPresent : 1;
    Uns32 _ul         : 31;
} CSR_REG_TYPE_32(fifo_cfg);

// define 32 bit type
CSR_REG_STRUCT_DECL_32(fifo_cfg);
```

A container structure is defined that holds all CSR values added by this extension:

```
typedef struct fifoCSRsS {
    CSR_REG_DECL(fifo_cfg);      // 0xFF0
} fifoCSRs;
```

The `vmiosObject` structure contains fields that hold CSR values and describe the CSRs:

```
typedef struct vmiosObjectS {

    . . . lines omitted . . .

    // extension CSR info
    fifoCSRs      csr;                // FIFO extension CSR values
    riscvCSRAttrs csrs[XCSR_ID(LAST)]; // modified CSR definitions

    . . . lines omitted . . .
```

```
} vmiosObject;
```

In file `fifoExtensions.c`, the set of CSRs to add is defined using an array of `extCSRAttrs` structures:

```
static const extCSRAttrs csrs[XCSR_ID(LAST)] = {

    XCSR_ATTR_T__(
        // name      num      arch extension attrs      description      rCB rwCB wCB
        fifo_cfg, 0xFF0, 0,    EXT_FIFO, 0,0,0,0,    "FIFO Configuration", 0, 0, 0
    )
};
```

Type `extCSRAttrs` is a structure type containing a standard base model CSR description structure (`riscvCSRAttrs`) together with an extension-specific identifier, `extension`. The purpose of the extension-specific identifier is to allow CSRs to be conditionally included based on parameter settings or other selection controls in the extension itself:

```
typedef struct extCSRAttrS {
    Uns32      extension;    // extension requirements
    riscvCSRAttrs baseAttrs; // base attributes
} extCSRAttrs;
```

The macro `XCSR_ATTR_T__` is defined in file `riscvModelCallbackTypes.h`:

```
#define XCSR_ATTR_T__( \
    _ID, _NUM, _ARCH, _EXT, _ENDB, _ENDRM, _NOTR, _TVMT, _DESC, _RCB, _RWCB, _WCB \
) [XCSR_ID(_ID)] = { \
    .extension = _EXT, \
    .baseAttrs = { \
        name      : #_ID, \
        desc      : _DESC, \
        csrNum     : _NUM, \
        arch      : _ARCH, \
        wEndBlock  : _ENDB, \
        wEndRM     : _ENDRM, \
        noTraceChange : _NOTR, \
        TVMT      : _TVMT, \
        readCB     : _RCB, \
        readWriteCB : _RWCB, \
        writeCB    : _WCB, \
        reg32      : XCSR_REG32_MT(_ID), \
        reg64      : XCSR_REG64_MT(_ID) \
    } \
}
```

The macro takes the following arguments:

1. The CSR *identifier*. This is used to construct both the CSR enumeration member name and the CSR name string (for reporting).
2. The CSR number, using standard RISC-V CSR numbering conventions.
3. Any architectural restrictions for the CSR, specified in the same way as architectural restrictions on instructions, discussed previously.
4. The extension identifier (see above).
5. *End-block* attribute: whether writes to the CSR should terminate a code block.
6. *End-rounding* attribute: whether writes to the CSR modify rounding mode.

7. *No-trace* attribute: whether changes to the CSR should not be reported when trace change is enabled.
8. *TVMT* attribute: whether accesses to the CSR are trapped by `mstatus.TVM`.
9. A CSR description string (used in documentation generation).
10. An optional read callback function.
11. An optional read-modify-write callback function (only for CSRs that have special behavior in this case).
12. An optional write callback function.

In this case, the CSR is implemented as a simple read-only register with a constant value. There are other macros available that allow specification of CSRs with a constant write mask (`XCSR_ATTR_TC_`), a variable write mask (`XCSR_ATTR_TV_`) and using callbacks only (`XCSR_ATTR_P_`).

Extension CSRs are initialized and registered with the base model by function `fifoCSRInit`, which is called by the constructor:

```
static void fifoCSRInit(vmiosObjectP object) {  
  
    riscvP    riscv = object->riscv;  
    extCSRId id;  
  
    // initialize CSR values that have configuration values defined  
    WR_FIFO_CSR(object, fifo_cfg, object->config.csr.fifo_cfg.u64.bits);  
  
    // register each CSR with the base model  
    for(id=0; id<XCSR_ID(LAST); id++) {  
  
        extCSRAttrsCP src = &csrs[id];  
        riscvCSRAttrs *dst = &object->csrs[id];  
  
        if(extensionPresent(object, src->extension)) {  
            riscv->cb.newCSR(dst, &src->baseAttrs, riscv, object);  
        }  
    }  
}
```

This function first sets the initial value of the `fifo_cfg` CSR from configuration defaults. It then iterates over all members of the CSR table, registering each CSR with the base model if the extension feature associated with that CSR is enabled (in fact, the FIFO feature is always enabled in this case, so `extensionPresent` always returns `True`). Registration with the base model is done by calling interface function `newCSR`, which takes these arguments:

1. A pointer to a destination `riscvCSRAttrs` object, which must be located in the current `vmiosObject` structure. This is filled with the source value passed as the second argument, augmented with a pointer back to the containing `vmiosObject` structure.
2. A source `riscvCSRAttrs` object, from the CSR table.
3. The RISC-V processor.
4. The containing `vmiosObject` structure.

Once the CSR is registered with the base model, reads and writes to it will be automatically performed by standard CSR access instructions with the relevant CSR index.

10.4 Extension FIFO Ports

Extension `fifoExtensions` adds two FIFO ports to the base model. The FIFO ports are defined by the `fifoPorts` array:

```
//  
// Return offset of the given field in the extension object  
//  
#define EXTENSION_FIELD_OFFSET(_F) ((void *)VMI_CPU_OFFSET(vmiosObjectP, _F))  
  
//  
// Template FIFO port list  
//  
static vmiFifoPort fifoPorts[] = {  
    {"fifoPortIn", vmi_FIFO_INPUT, 0, EXTENSION_FIELD_OFFSET(inputConn) },  
    {"fifoPortOut", vmi_FIFO_OUTPUT, 0, EXTENSION_FIELD_OFFSET(outputConn)}  
};
```

The macro `EXTENSION_FIELD_OFFSET` is used to initialize the handle field in each entry with the offset of the `inputConn` and `outputConn` fields in the `vmiosObject` structure:

```
typedef struct vmiosObjectS {  
  
    . . . fields omitted . . .  
  
    // FIFO connections  
    vmiFifoPortP    fifoPorts;           // fifo port descriptions  
    vmiConnInputP   inputConn;           // input FIFO connection  
    vmiConnOutputP  outputConn;          // output FIFO connection  
  
    . . . fields omitted . . .  
} vmiosObject;
```

Extension FIFOs are created by function `newFifoPorts`, which is called by the constructor:

```
static void newFifoPorts(vmiosObjectP object) {  
  
    Uns32 connBits = getConnBits(object);  
    Uns32 i;  
  
    object->fifoPorts = STYPE_CALLOC_N(vmiFifoPort, NUM_MEMBERS(fifoPorts));  
  
    for(i=0; i<NUM_MEMBERS(fifoPorts); i++) {  
  
        object->fifoPorts[i] = fifoPorts[i];  
  
        // correct FIFO port bit size  
        object->fifoPorts[i].bits = connBits;  
  
        // correct FIFO port handle  
        Uns8 *raw = (Uns8*)(object->fifoPorts[i].handle);  
        object->fifoPorts[i].handle = (void **)(raw + (UnsPS)object);  
    }  
}
```

This function first allocates an extension-specific array of FIFO port objects. It then fills each FIFO port object from the template array, adjusting the `bits` field to correspond to the value given in a model parameter and correcting the `handle` field to point to the field location within the current `vmiosObject` structure (by adding the offset in the template to the `vmiosObject` address). The bit size of each connection is extracted from the configuration using function `getConnBits`:

```
inline static Uns32 getConnBits(vmiosObjectP object) {
    return object->config.FIFO_bits;
}
```

A standard FIFO port iterator function, referenced in the `vmiosAttrs` structure, is used to indicate the presence of the new FIFO ports:

```
static VMIOS_FIFO_PORT_SPECS_FN(fifoGetPortSpec) {
    if (!object->enabled) {
        // Do not implement ports when not enabled
        return NULL;
    } else if (!prev) {
        // first port
        return object->fifoPorts;
    } else {
        // port other than the first
        Uns32 prevIndex = (prev-object->fifoPorts);
        Uns32 thisIndex = prevIndex+1;

        return (thisIndex<NUM_MEMBERS(fifoPorts)) ? &object->fifoPorts[thisIndex]:0;
    }
}
```

10.5 Instruction Decode

This extension object implements two new instructions, `pushb` and `popb`. These are defined by an enumeration and a table exactly as described in section 6.3:

```
typedef enum riscvExtITypeE {
    // extension instructions
    EXT_IT_PUSHB,
    EXT_IT_POPB,

    // KEEP LAST
    EXT_IT_LAST
} riscvExtIType;

const static riscvExtInstrAttrs attrsArray32[] = {
    EXT_INSTRUCTION(EXT_IT_PUSHB, "pushb", RVANY, RVIP_RD_RS1_RS2, FMT_R1,
        "|000000000000|00000|000|....|0001011|"),
    EXT_INSTRUCTION(EXT_IT_POPB, "popb", RVANY, RVIP_RD_RS1_RS2, FMT_R1,
        "|000000000000|00000|001|....|0001011|"),
};
```

In this case, the disassembly format is specified as `FMT_R1`, so that only one register (in the `Rd` position) is reported.

10.6 Instruction Disassembly

Instruction disassembly is implemented in exactly the same way as described in section 6.4.

10.7 Instruction Translation

Instruction translation uses a similar pattern to that previously described in section 6.5. In this example, the instruction translation table is specified like this:

```
const static riscvExtMorphAttr dispatchTable[] = {
    [EXT_IT_PUSHB] = {morph:emitPUSHB, variant:EXT_FIFO},
    [EXT_IT_POPB] = {morph:emitPOPB, variant:EXT_FIFO},
};
```

The JIT translation function is specified like this:

```
static VMIO_MORPH_FN(fifoMorph) {

    riscvP          riscv = (riscvP)processor;
    riscvExtMorphState state = {riscv:riscv, object:object};

    // get instruction and instruction type
    riscvExtIType type = decode(riscv, object, thisPC, &state.info);

    // action is only required if the instruction is implemented by this
    // extension
    if(type != EXT_IT_LAST) {

        riscvExtMorphAttrCP attrs = &dispatchTable[type];
        const char *reason = getDisableReason(object, attrs->variant);

        // fill translation attributes
        state.attrs = attrs;

        // translate instruction
        riscv->cb.morphExternal(&state, reason, opaque);
    }

    // no callback function is required
    return 0;
}
```

This is similar to the previous example, but includes an additional check for instruction validity, implemented by function `getDisableReason`:

```
static const char *getDisableReason(vmiosObjectP object, fifoVariant variant) {

    fifoVariant availableVariants = object->config.variant;
    const char *result = 0;

    // validate ext feature set
    if((availableVariants & variant) != variant) {
        result = "Unimplemented on this variant";
    }

    return result;
}
```

This function validates the feature set stated to be implemented by the extension configuration against the feature requirements of the instruction. If the instruction requires a feature set that is not implemented, the string "Unimplemented on this

variant" is returned. When this is passed to the `morphExternal` interface function, code will be emitted to take an Illegal Instruction exception instead of the normal instruction behavior.

In this example, the FIFO extension is always implemented so the Illegal Instruction behavior is never triggered, but this pattern is useful for an extension object that adds multiple extra instructions in different sets, enabled by parameters or feature registers.

In this example, `pushb` and `popb` are implemented by separate instruction callbacks. The implementation of `pushb` is this:

```
static EXT_MORPH_FN(emitPUSHB) {  
  
    vmiosObjectP object = state->object;  
    riscvP riscv = state->riscv;  
    riscvRegDesc rs = getRVReg(state, 0);  
    vmiReg rsa = getVMIReg(riscv, rs);  
    vmiReg conn = getOutputConn(object);  
    Uns32 bits = getRBits(rs);  
    Uns32 connBits = getConnBits(object);  
  
    // zero-extend source value to temporary if connection is wider than GPR  
    if(bits < connBits) {  
        vmiReg tmp = getFIFOTmp(object);  
        vmimtMoveExtendRR(connBits, tmp, bits, rsa, False);  
        rsa = tmp;  
    }  
  
    // put value  
    vmimtConnPutRB(connBits, conn, rsa, 0);  
}
```

This function obtains a `vmiReg` for register `rd` in the same way as the previous example. It also obtains a `vmiReg` for the output connection object using utility function `getOutputConn`:

```
//  
// Return VMI register for extension object field  
//  
inline static vmiReg getExtReg(vmiosObjectP object, void *field) {  
    return vmimtGetExtReg((vmiProcessorP)(object->riscv), field);  
}  
  
//  
// Return VMI register for output connection object  
//  
inline static vmiReg getOutputConn(vmiosObjectP object) {  
    return getExtReg(object, &object->outputConn);  
}
```

This uses the standard VMI Morph Time API function `vmimtGetExtReg` to get a `vmiReg` descriptor for a generic pointer.

The extension allows the size of the connection (FIFO) element in bits to be parameterized, up to `XLEN` bits in size. The next step is to get the size in bits of both the GPR and FIFO element:

```
Uns32 bits = getRBits(rs);
```

```
Uns32      connBits = getConnBits(object);
```

If the FIFO element is larger than `XLEN`, the value in `rd` is zero-extended to the full FIFO element width, using a temporary in the extension object to hold the extended value:

```
if(bits<connBits) {
    vmiReg tmp = getFIFOTmp(object);
    vmimtMoveExtendRR(connBits, tmp, bits, rsA, False);
    rsA = tmp;
}
```

Finally, the (possibly-extended) register value is written to the FIFO using a standard blocking put:

```
vmimtConnPutRB(connBits, conn, rsA, 0);
```

The implementation of `popb` is similar:

```
static EXT_MORPH_FN(emitPOPB) {
    vmiosObjectP object = state->object;
    riscvP      riscv   = state->riscv;
    riscvRegDesc rd      = getRVReg(state, 0);
    vmiReg      rdA      = getVMIReg(riscv, rd);
    vmiReg      conn     = getInputConn(object);
    Uns32       bits     = getRBits(rd);
    Uns32       connBits = getConnBits(object);
    Uns32       tmpBits  = (connBits<bits) ? connBits : bits;
    vmiReg      tmp      = getFIFOTmp(object);

    // get value into temporary
    vmimtConnGetRB(connBits, tmp, conn, False, 0);

    // commit value, zero-extending if necessary
    vmimtMoveExtendRR(bits, rdA, tmpBits, tmp, False);
}
```

Here, a blocking get is made from the input FIFO to a temporary, and then the temporary is zero-extended into the result register.

11 Adding Transactional Memory (`tmExtensions`)

The `tmExtensions` extension object extends the basic RISC-V model by adding custom transactional memory and some new instructions to manage the transactional memory state.

Transactional memory is likely to be highly implementation dependent. In this example extension, the extended processor operates in two modes:

1. *Normal mode*: when no transaction is active, loads and stores are performed to memory in the usual way.
2. *Transaction mode*: when a transaction is active, stores are accumulated in a cache. Dirty data is either committed atomically at the end of the transaction or discarded if the transaction is aborted for some reason (for example, a conflicting write by another processor, or too much data for the cache). If a transaction is aborted, all processor GPR and FPR values are reset to the state they had when the transaction was started, allowing the transaction to be retried easily if required.

The base processor model supports operating in normal and transaction mode using interface function `setTMode`, described later in this section. In normal mode, the base model behavior is unchanged; in transaction mode, all loads and stores are routed to functions in the extension object, which implement a cache model (or similar structure) to hold speculative data values. The extension object is responsible for implementing the cache model, performing memory reads to populate the model, and performing memory writes to drain the cache model when required.

This extension adds four instructions:

1. `xbegin`: executed to start a new transaction;
2. `xend`: executed to end a transaction;
3. `xabort`: executed to abort an active transaction; and
4. `wfe`: a *wait for event* pseudo-instruction, used to yield control to other harts in a multicore system.

All behavior of this extension object is implemented in file `tmExtensions.c`. Sections will be discussed in turn below.

11.1 Intercept Attributes

The behavior of the library is defined using the standard `vmiosAttr` structure:

```
vmiosAttr modelAttrs = {  
    ///////////////////////////////////////  
    // VERSION  
    ///////////////////////////////////////  
    .versionString = VMI_VERSION,          // version string  
    .modelType     = VMI_INTERCEPT_LIBRARY, // type  
    .interceptType = VMI_IT_PROC_EXTENSION, // intercept type  
}
```

```

.packageName = "transactionalMemory", // description
.objectSize  = sizeof(vmiosObject),  // size in bytes of OSS object

////////////////////////////////////////
// CONSTRUCTOR/DESTRUCTOR ROUTINES
////////////////////////////////////////

.constructorCB = tmConstructor,      // object constructor
.postConstructorCB = tmPostConstructor, // object post-constructor
.docCB         = tmDoc,              // documentation constructor

////////////////////////////////////////
// INSTRUCTION INTERCEPT ROUTINES
////////////////////////////////////////

.morphCB      = tmMorph,              // instruction morph callback
.disCB        = tmDisassemble,        // disassemble instruction

////////////////////////////////////////
// PARAMETER CALLBACKS
////////////////////////////////////////

.paramSpecsCB = tmParamSpecs,         // iterate parameter declarations
.paramValueSizeCB = tmParamTableSize, // get parameter table size

////////////////////////////////////////
// ADDRESS INTERCEPT DEFINITIONS
////////////////////////////////////////

.intercepts = {{0}}
};

```

In this library, there is a constructor and post-constructor, a documentation callback, a JIT translation (morpher) function, a disassembly function, and functions allowing the extension to specify parameters.

11.2 Intercept Parameters

This extension is parameterized as follows:

1. A parameter `diagnosticlevel` allows the verbosity of debug messages to be specified (0, 1, 2 or 3); and
2. A bit mask parameter `variant` allows the configured extensions to be defined: if bit 0 is set, the transactional instructions are present, and if bit 1 is set the WFE instruction is present).

The parameters are defined using the standard extension parameterization interface as follows:

```

typedef struct formalValuesS {
    VMI_UNSS32_PARAM(d diagnosticlevel);
    VMI_UNSS32_PARAM(features);
} formalValues, *formalValuesP;

// Parameter table
static vmiParameter parameters[] = {
    VMI_UNSS32_PARAM_SPEC(formalValues, diagnosticlevel, 0, 0, 3, "Override
the initial diagnostic level"),
    VMI_UNSS32_PARAM_SPEC(formalValues, features, EXT_ALL, 1, EXT_ALL, "Override
the configured variant features"),
    { 0 }
};

```

```
// Iterate formals
static VMIO_PARAM_SPEC_FN(tmParamSpecs) {
    if(!prev) {
        prev = parameters;
    } else {
        prev++;
    }
    return prev->name ? prev : 0;
}

// Return size of parameter structure
static VMIO_PARAM_TABLE_SIZE_FN(tmParamTableSize) {
    return sizeof(formalValues);
}
```

Functions `tmParamSpecs` and `tmParamTableSize` are referenced in the `vmiosAttr` structure for the extension (see section 11.1).

11.3 Object Type, Constructor and Post-Constructor

The object type is defined as follows:

```
typedef struct vmiosObjectS {

    // associated processor
    riscvP      riscv;

    // is this extension enabled?
    Bool        enabled;

    // configuration (including CSR reset values)
    tmConfig     config;

    // TM extension CSR registers info
    tmCSRs       csr;                // TM extension CSR values
    riscvCSRAttrs csrs[XCSR_ID(LAST)]; // modified CSR definitions

    // Transaction state info
    memDomainP   physicalMem;        // physical memory domain
    memRegionP   regionCache;        // cached physical memory region
    tmStatusE     tmStatus;          // current TM status
    Uns32        numPending;         // number of lines currently pending
    cacheLine     pending[CACHE_MAX_LINES]; // current transaction cached lines

    // Abort state info
    Uns32         xbeginReg;          // index of xbegin register
    Uns64         abortCode;          // xabort code
    Addr          abortPC;            // PC to jump to on abort
    Uns64         x[RISCV_GPR_NUM];   // GPR bank
    Uns64         f[RISCV_FPR_NUM];   // FPR bank

    // extended instruction decode table
    vmidDecodeTableP decode32;

    // Command argument values
    cmdArgValues  cmdArgs;

    // extension callbacks
    riscvExtCB    extCB;
} vmiosObject;
```

This structure contains the `riscv`, `decode32` and `extCB` fields that are always required when using the RISC-V extension support infrastructure documented here, together with

a number of other extension-specific fields. The constructor initializes the fields as follows:

```
static VMIO5_CONSTRUCTOR_FN(tmConstructor) {  
  
    riscvP      riscv = (riscvP)processor;  
    formalValuesP params = parameterValues;  
  
    object->riscv = riscv;  
  
    // prepare client data  
    object->extCB.clientData = object;  
  
    // Initialize diagnostic setting to value set for parameter  
    DIAG_LEVEL(object) = params->diagnosticlevel;  
  
    // Add commands  
    addCommands(object, &object->cmdArgs);  
  
    // add extension registers  
    addExtRegs(object);  
  
    // initialize base model callbacks  
    object->extCB.switchCB      = riscvSwitch;  
    object->extCB.tLoad         = riscvTLoad;  
    object->extCB.tStore        = riscvTStore;  
    object->extCB.trapNotifier  = riscvTrapNotifier;  
    object->extCB.ERETNotifier  = riscvERETNotifier;  
  
    // register extension with base model  
    riscv->cb.registerExtCB(riscv, &object->extCB, EXTID_TM);  
  
    // set status to not active to start  
    object->tmStatus = TM_NOTACTIVE;  
  
    // copy configuration from template  
    object->config = *getExtConfig(riscv);  
  
    // override configured variant  
    object->config.variant = params->features;  
  
    // initialize CSRs  
    tmCSRInit(object);  
  
    // is extension enabled in config info?  
    object->enabled = RD_XCSR_FIELD(object, tm_cfg, tmPresent);  
}
```

In addition to the mandatory field setup, this constructor also defines extension specific CSRs and installs notifier functions that are called when execution context switches to or from another processor in a multicore simulation (`riscvSwitch`), when loads and stores are performed in transaction mode (`riscvTLoad` and `riscvTStore`) and when taking or resuming from exceptions (`riscvTrapNotifier` and `riscvERETNotifier`). These are all described in following sections.

This extension also defines a *post-constructor* function, `tmPostConstructor`. The post-constructor is called *after all processor model and processor extension object constructors have been called but before simulation starts*:

```
static VMIO5_POST_CONSTRUCTOR_FN(tmPostConstructor) {  
  
    // record the processor physical memory domain  
    object->physicalMem = vmirtGetProcessorExternalDataDomain(processor);  
}
```

```
}
```

In this case, the post-constructor saves the processor *external memory data domain* object for later use. This memory domain object is the target for all loads and stores performed by this processor or others in a multicore simulation. It is required so that the extension object can implement cache line reads and writes (using VMI run time functions `vmirtReadNByteDomain` and `vmirtWriteNByteDomain`) and so that monitor callbacks can be installed on it to check for memory accesses by other processors that may invalidate an active transaction by this processor.

The call to `vmirtGetProcessorExternalDataDomain` must be placed in the *post-constructor* because it is *inspecting memory state*. The call cannot be made in the constructor because at that point there is no guarantee that all other processor and extension constructors have run, so any value returned by a VMI inspection function like this may return invalid state at that point. Any VMI function that references memory domains (such as `vmirtGetProcessorExternalDataDomain`) or processor state (such as `vmirtRegRead` or `vmirtRegWrite`) *must not be used in the constructor*.

The `diagnosticlevel` and `features` parameters are used to modify the initial extension configuration.

11.4 Extension Registers

Extension `tmExtensions` adds a single read-only extension register to the base model. This register is not a CSR: see the next section for information about adding CSRs.

The register added is called `TM` and holds the current value of the transaction mode state for the processor. The register is added by function `addExtRegs`, called in the constructor:

```
static void addExtRegs(vmiosObjectP object) {
    riscvP riscv = object->riscv;

    // create description of read-only transaction mode register
    vmiRegInfo tmReg = {
        name       : "TM",
        description : "Transaction mode state",
        bits       : 8,
        gdbIndex    : 0,
        access      : vmi_RA_R,
        raw         : vmimtGetExtReg((vmiProcessorP)riscv, &object->tmStatus)
    };

    // add register description
    riscv->cb.newExtReg(riscv, &tmReg);
}
```

The function first defines a `vmiRegInfo` object describing register `TM`. In this case, the register value is simply the value of field `tmStatus` in the extension object. More complex registers may have read and write callbacks to form and assign their value: see the *OVP Processor Modeling Guide* for more information about the use of `vmiRegInfo` objects to describe registers.

The register is added to the processor external register view using the `newExtReg` interface function:

```
riscv->cb.newExtReg(riscv, &tmReg);
```

This function can be called as many times as needed to add extension registers. Each added register should be given a unique name and index (`gdbIndex`). When added to the processor, the given index number will be modified using bitwise-or with `0x80000000` to ensure extension register indices do not conflict with base model register indices.

11.5 Extension CSRs

Extension `tmExtensions` adds a single read-only CSR to the base model. The CSR is defined in file `tmCSR.h`. An enumeration in that file first defines the set of additional CSRs:

```
typedef enum extCSRIdE {  
    // TM configuration control and status registers  
    XCSR_ID (tm_cfg),          // 0xFD0  
  
    // keep last (used to define size of the enumeration)  
    XCSR_ID (LAST)  
} extCSRId;
```

Then the fields in the `tm_cfg` CSR are defined using a bitfield structure:

```
// -----  
// tm_cfg      (id 0xFD0)  
// -----  
  
// 32-bit view  
typedef struct {  
    Uns32 tmPresent : 1;  
    Uns32 _ul       : 31;  
} CSR_REG_TYPE_32(tm_cfg);  
  
// define 32 bit type  
CSR_REG_STRUCT_DECL_32(tm_cfg);
```

A container structure is defined that holds all CSR values added by this extension:

```
typedef struct tmCSRss {  
    CSR_REG_DECL(tm_cfg);    // 0xFD0  
} tmCSR;
```

The `vmiosObject` structure contains fields that hold CSR values and describe the CSRs:

```
typedef struct vmiosObjectS {  
    . . . lines omitted . . .  
  
    // extension CSR info  
    tmCSR      csr;          // TM extension CSR values  
    riscvCSRAttrs csrs[XCSR_ID(LAST)]; // modified CSR definitions  
  
    . . . lines omitted . . .
```



```
} vmiosObject;
```

In file `tmExtensions.c`, the set of CSRs to add is defined using an array of `extCSRAttrs` structures:

```
static const extCSRAttrs csrs[XCSR_ID(LAST)] = {

    XCSR_ATTR_T__(
        // name      num      arch extension attrs      description      rCB rwCB wCB
        tm_cfg, 0xFD0, 0,      EXT_TM,      0,0,0,0, "TM Configuration", 0, 0, 0
    )
};
```

This field is used in function `tmCSRInit` to initialize CSR descriptions, exactly as previously described in section 10.3.

11.6 Context Switch Monitor (*riscvSwitch*)

Extension `tmExtensions` installs an *execution context switch monitor* function, `riscvSwitch`:

```
static VMIOS_CONSTRUCTOR_FN(tmConstructor) {

    . . . lines omitted . . .

    // initialize base model callbacks
    object->extCB.switchCB      = riscvSwitch;
    object->extCB.tLoad          = riscvTLoad;
    object->extCB.tStore         = riscvTStore;
    object->extCB.trapNotifier   = riscvTrapNotifier;
    object->extCB.ERETNotifier   = riscvERETNotifier;

    . . . lines omitted . . .
}
```

In a multiprocessor simulation, each processor is executed in turn for a number of instructions (the quantum). The execution context switch monitor function is called whenever execution context switches *to* this processor (it is about to start executing) or *away from* this processor (it has finished its quantum and another processor is about to run). When modeling transactional memory, one requirement is that the model is aware when conflicting reads or writes have been made to memory addresses by other processors that would cause an active transaction on this processor to be aborted.

`riscvSwitch` is defined like this:

```
static RISCV_IASSWITCH_FN(riscvSwitch) {

    vmiosObjectP object = clientData;

    . . . lines omitted . . .

    if(state==RS_SUSPEND) {
        installCacheMonitor(object);
    }
}
```

The function is of type `riscvIASSwitchFn`, defined in `riscvModelCallbacks.h` like this:

```
#define RISC_V_IASSWITCH_FN(_NAME) void _NAME( \
    riscvP      riscv,      \
    vmiIASRunState state,    \
    void        *clientData \
)
typedef RISC_V_IASSWITCH_FN((*riscvIASSwitchFn));
```

This function has effect if execution context is switching from this processor to another (state is RS_SUSPEND). In this case, function `installCacheMonitor` adds two kinds of memory callback to the physical memory domain cached by the post-constructor:

1. For any active line in the cache, a *write* callback is installed which is called if any other processor writes data to an address in that line;
2. For any dirty line in the cache, a *read* callback is installed which is called if any other processor reads data from an address in that line.

In both cases, a transaction abort is triggered because of a memory conflict.

11.7 Transactional Load and Store Functions (*riscvTLoad* and *riscvTStore*)

Extension `tmExtensions` installs transaction load and store functions, `riscvTLoad` and `riscvTStore`:

```
static VMIO_CONSTRUCTOR_FN(tmConstructor) {
    . . . lines omitted . . .

    // initialize base model callbacks
    object->extCB.switchCB      = riscvSwitch;
    object->extCB.tLoad          = riscvTLoad;
    object->extCB.tStore         = riscvTStore;
    object->extCB.trapNotifier   = riscvTrapNotifier;
    object->extCB.ERETNotifier   = riscvERETNotifier;

    . . . lines omitted . . .
}
```

When the processor is executing with transactional mode *enabled*, any load or store will cause these two functions to be executed instead of updating memory in the normal way. This allows the extension object to intervene to access data from another location (a cache model, in this case). Function `riscvTLoad` is of type `riscvTLoadFn`:

```
#define RISC_V_TLOAD_FN(_NAME) void _NAME( \
    riscvP riscv,      \
    void *buffer,       \
    Addr   VA,          \
    Uns32  bytes,       \
    void *clientData    \
)
typedef RISC_V_TLOAD_FN((*riscvTLoadFn));
```

This function must fill `buffer` with `bytes` bytes read from address `VA` to implement a load. Function `riscvTStore` is of type `riscvTStoreFn`:

```
#define RISC_V_TSTORE_FN(_NAME) void _NAME( \
    riscvP      riscv,      \
```

```
const void *buffer,      \
Addr          VA,        \
Uns32         bytes,     \
void          *clientData \
)
typedef RISC_V_TSTORE_FN((*riscvTStoreFn));
```

This function must take `bytes` bytes from `buffer` and save them in a data structure (in this case, representing cache lines). The implementation of the transactional memory will be implementation specific, so the details in this case will not be discussed here – refer to section 11.12 and the model source for a detailed example if required.

11.8 Trap and Exception Return Notifiers (*riscvTrapNotifier* and *riscvERETNotifier*)

Extension `tmExtensions` installs trap and exception return notifiers, `riscvTrapNotifier` and `riscvERETNotifier`:

```
static VMIO5_CONSTRUCTOR_FN(tmConstructor) {
    . . . lines omitted . . .

    // initialize base model callbacks
    object->extCB.switchCB    = riscvSwitch;
    object->extCB.tLoad       = riscvTLoad;
    object->extCB.tStore      = riscvTStore;
    object->extCB.trapNotifier = riscvTrapNotifier;
    object->extCB.ERETNotifier = riscvERETNotifier;

    . . . lines omitted . . .
}
```

`riscvTrapNotifier` is called whenever the processor takes a trap:

```
static RISC_V_TRAP_NOTIFIER_FN(riscvTrapNotifier) {
    vmiosObjectP object = clientData;

    if(object->tmStatus == TM_NOTACTIVE) {
        // ignore exceptions outside of transactions
    } else {
        . . . lines omitted . . .

        // update status to abort transaction
        object->tmStatus |= TM_ABORT_EXCEPTION;

        // clear transaction mode during exception so memory updates will occur
        // normally during exception
        setTMode(object->riscv, False);
    }
}
```

The notifier first detects whether the processor is operating in transaction mode. If it is, the current transaction is marked as aborted, for reason `TM_ABORT_EXCEPTION`. Then, transaction mode is disabled while the exception is handled by a call to `setTMode` (so that loads and stores in the exception routine behave normally). Function `setTMode` is a

wrapper round interface function `setTMode` in the base model. This function toggles the base model between normal and transaction mode:

```
static void setTMode(riscvP riscv, Bool enable) {
    riscv->cb.setTMode(riscv, enable);
}
```

The exception return notifier is similar except that its final stage is to *re-enable transaction mode* to cause a transaction abort on next transaction activity:

```
static RISC_V_TRAP_NOTIFIER_FN(riscvRETNotifier) {

    vmiosObjectP object = clientData;

    if(object->tmStatus == TM_NOTACTIVE) {

        // ignore exception returns outside of transactions

    } else {

        . . . lines omitted . . .

        // update status to abort transaction
        object->tmStatus |= TM_ABORT_EXCEPTION;

        // restore transaction mode after exception so transaction abort will
        // occur on next activity
        setTMode(object->riscv, True);
    }
}
```

11.9 Instruction Decode

This extension object implements four new instructions, `xbegin`, `xend`, `xabort` and `wfe`. These are defined by an enumeration and a table exactly as described in section 6.3:

```
typedef enum riscvExtITypeE {

    // extension instructions
    EXT_IT_XBEGIN,
    EXT_IT_XEND,
    EXT_IT_XABORT,
    EXT_IT_WFE,
    // KEEP LAST
    EXT_IT_LAST

} riscvExtIType;

const static riscvExtInstrAttrs attrsArray32[] = {
    EXT_INSTRUCTION(EXT_IT_XBEGIN, "xbegin", RVANY, RVIP_RD_RS1_RS2, FMT_R1,
        "|0000000|00000|00000|011|....|0001011|"),
    EXT_INSTRUCTION(EXT_IT_XEND, "xend", RVANY, RVIP_RD_RS1_RS2, FMT_NONE,
        "|0000000|00000|00000|010|00000|0001011|"),
    EXT_INSTRUCTION(EXT_IT_XABORT, "xabort", RVANY, RVIP_RD_RS1_RS2, FMT_R1,
        "|0000000|00000|00000|100|....|0001011|"),
    EXT_INSTRUCTION(EXT_IT_WFE, "wfe", RVANY, RVIP_RD_RS1_RS2, FMT_NONE,
        "|0000000|00000|00000|101|00000|0001011|")
};
```

In this case, the disassembly format is specified as `FMT_R1`, for `xbegin` and `xabort` so that only one register (in the `Rd` position) is reported, and as `FMT_NONE` for `xend` and `wfe`.

11.10 *Instruction Disassembly*

Instruction disassembly is implemented in exactly the same way as described in section 6.4.

11.11 *Instruction Translation*

Instruction translation uses a similar pattern to that previously described in section 6.5. In this example, the instruction translation table is specified like this:

```
const static riscvExtMorphAttr dispatchTable[] = {
    [EXT_IT_XBEGIN] = {morph:emitXBEGIN, variant:EXT_TM },
    [EXT_IT_XEND]   = {morph:emitXEND,   variant:EXT_TM },
    [EXT_IT_XABORT] = {morph:emitXABORT,  variant:EXT_TM },
    [EXT_IT_WFE]    = {morph:emitWFE,    variant:EXT_WFE},
};
```

This extension allows the transaction instructions and the `WFE` instruction to be enabled separately (so it is possible to configure a core with `WFE` only, for example). To handle this, the instructions are given different `variant` masks.

The JIT translation function follows the same pattern as used previously for the FIFO extension: refer to section 10.7 for more information.

The `xbegin` instruction is executed to start a new transaction. JIT code for this is created by function `emitXBEGIN`:

```
static EXT_MORPH_FN(emitXBEGIN) {
    // get abstract register operands
    riscvRegDesc rd = getRVReg(state, 0);

    // emit call implementing XBEGIN instruction
    vmimtArgNatAddress(state->object);
    vmimtArgUns32(getRIndex(rd));
    vmimtArgSimPC(64);
    vmimtCall((vmiCallFn)xBegin);

    // transaction mode change possible so end this code block
    vmimtEndBlock();
}
```

This emits an embedded call to function `xBegin`. Because `xBegin` could change transaction mode, a call to `vmimtEndBlock` is required to terminate the current code block (the next instruction could be executed in different transaction mode state, and a single code block must not contain instructions from different transaction mode states for correct behavior). The arguments to `xBegin` are the extension object, the register index for the one register in the instruction, and the current simulated program counter:

```
static void xBegin(vmiosObjectP object, Uns32 regIdx, Uns64 thisPC) {
    if(object->tmStatus != TM_NOTACTIVE) {
        // Nested transactions not supported
        object->tmStatus |= TM_ABORT_NESTED;
        doAbort(object);
    } else {
```

```
// save PC of next instruction (to be executed on abort)
object->abortPC = thisPC+4;
object->abortCode = 0;

// save xbegin instruction destination register index
object->xbeginReg = regIdx;

// save current values of registers for abort, if necessary
saveRegs(object);

// start a new transaction
object->tmStatus = TM_OK;
setTMode(object->riscv, True);

// set value in xbegin destination register
setXbeginReturnValue(object, object->tmStatus);
}
```

If there is already an active transaction when `xBegin` is called, that transaction is aborted. Otherwise, `xBegin` does this:

1. It saves the address of the instruction *after* the `xbegin` instruction. This is the *abort address*, to which control will be transferred in the transaction fails.
2. It saves the index of the register argument to `xbegin`. This register is assigned a *status code* when the transaction succeeds or fails, so that the initiator can react appropriately.
3. It saves the value of all GPRs and FPRs into a shadow block implemented in the extension object. This allows these registers to be restored if the transaction fails.
4. It enters transaction mode by calling `setTMode`, with initial state `TM_OK`.
5. It returns the initial state to the register argument to `xbegin`, so that the initiator can react appropriately.

The `xend` instruction is executed to terminate an active transaction. JIT code for this is created by function `emitXEND`:

```
static EXT_MORPH_FN(emitXEND) {

    // XEND instruction is a NOP when not in a transaction
    if(getTMode(state->riscv)) {

        // emit call implementing XEND instruction
        vmimtArgNatAddress(state->object);
        vmimtCall((vmiCallFn)xEnd);

        // transaction mode change possible so end this code block
        vmimtEndBlock();
    }
}
```

If the processor is not in transaction mode, this instruction behaves as a `NOP` and no code is emitted. Otherwise, an embedded call to `xEnd` is emitted, to terminate the transaction. Once again, a call to `vmimtEndBlock` is required to terminate the current code block (the next instruction could be executed in different transaction mode state). Function `xEnd` is defined like this:

```
static void xEnd(vmiosObjectP object) {  
    if(object->tmStatus != TM_OK) {  
        // abort is pending  
        doAbort(object);  
    } else {  
        // end current transaction  
        deactivate(object);  
    }  
}
```

The either aborts the current transaction (if status is not `TM_OK`) or commits results (if status is `TM_OK`).

The `xabort` instruction is executed to abort an active transaction. JIT code for this is created by function `emitXABORT`:

```
static EXT_MORPH_FN(emitXABORT) {  
    // XABORT instruction is a NOP when not in a transaction  
    if(getTMode(state->riscv)) {  
        // get abstract register operands  
        riscvRegDesc rs = getRVReg(state, 0);  
  
        // emit call implementing XABORT instruction  
        vmimtArgNatAddress(state->object);  
        vmimtArgUns32(getRIndex(rs));  
        vmimtCall((vmiCallFn)xAbort);  
  
        // transaction mode change possible so end this code block  
        vmimtEndBlock();  
    }  
}
```

If the processor is not in transaction mode, this instruction behaves as a NOP and no code is emitted. Otherwise, an embedded call to `xAbort` is emitted, to abort the transaction. Once again, a call to `vmimtEndBlock` is required to terminate the current code block (the next instruction could be executed in different transaction mode state). Function `xAbort` is defined like this:

```
static void xAbort(vmiosObjectP object, Uns32 regIdx) {  
    // get value from xabort source register  
    object->abortCode = object->riscv->x[regIdx];  
  
    // flag that this abort was from an instruction  
    object->tmStatus |= TM_ABORT_INST;  
  
    doAbort(object);  
}
```

Here, the abort is done for reason `TM_ABORT_INST`. Function `doAbort` is as follows:

```
static void doAbort(vmiosObjectP object) {  
    if (object->tmStatus == TM_NOTACTIVE) {  
        // Not active so nothing to abort - ignore
```

```
} else {  
    vmiProcessorP proc = (vmiProcessorP)object->riscv;  
  
    // restore values of RISC-V GPRs  
    restoreRegs(object);  
  
    // compute and set return value for xbegin  
    Uns64 returnValue = (  
        (object->tmStatus & 0xff) |  
        ((object->abortCode & 0xff) << 8)  
    );  
    setXbeginReturnValue(object, returnValue);  
  
    // clear current transaction  
    deactivate(object);  
  
    // set PC of next instruction (to be executed on abort)  
    vmirtSetPC(proc, object->abortPC);  
}  
}
```

To abort a transaction, the function does the following:

1. It restores GPR and FPR values to the state that was in effect *before* the transaction was started.
2. It constructs a return code by concatenating transaction status and the accumulated abort code, and then calls `setXbeginReturnValue` to assign that code to the GPR specified by the initiating `xbegin` operation.
3. It deactivates transaction mode by calling `deactivate`.
4. It uses `vmirtSetPC` to force a jump to the abort address, which is the address *after* the initiating `xbegin` instruction.

Note that in a linked model extension library it is legal to directly access fields of the base RISC-V model in cases where this can be done safely. As an example, function `restoreRegs` directly accesses the GPR and FPR values from the main model to restore them:

```
static void restoreRegs(vmiosObjectP object) {  
    riscvP riscv = object->riscv;  
    Uns32 i;  
  
    for(i=1; i<RISCV_GPR_NUM; i++) {  
        riscv->x[i] = object->x[i];  
    }  
    for(i=0; i<RISCV_FPR_NUM; i++) {  
        riscv->f[i] = object->f[i];  
    }  
}
```

Function `deactivate` is used to transition from transaction mode to normal mode:

```
static void deactivate(vmiosObjectP object) {  
    . . . lines omitted . . .  
    xCommit(object);  
}
```



```
setTMode(object->riscv, False);  
object->tmStatus = TM_NOTACTIVE;  
object->abortCode = 0;  
}
```

As part of the deactivation process, live data in the transaction mode cache model is drained to memory by calling function `xCommit`. Then, normal mode is enabled by calling `setTMode` with `enable` of `False`.

The `wfe` instruction is executed to suspend this processor until the end of its quantum to allow others to run (in real hardware, a similar instruction could cause a processor to stop executing and enter a low power state). JIT code for this is created by function `emitWFE`:

```
static EXT_MORPH_FN(emitWFE) {  
    vmimtIdle();  
}
```

11.12 *Memory Model Implementation Guidelines*

As previously stated, the transactional memory model is likely to be highly implementation dependent. This document will therefore not describe the chosen implementation in the `tmExtensions` example in detail, but instead will give some general guidelines on the approach to adopt.

1. Obtain handles to required memory domain objects in the post-constructor, as described above.
2. In the transactional load and store callback functions, use `vmirtReadNByteDomain` to read in cache line data.
3. When a transaction is being committed, use `vmirtWriteNByteDomain` to drain cache contents to physical memory.
4. When context switches *away* from the current processor, use `vmirtAddReadCallback` and `vmirtAddWriteCallback` to monitor address for which reads and writes by another processor should abort transactions on the current processor, respectively.
5. When a transaction is being committed or aborted, use `vmirtRemoveReadCallback` and `vmirtRemoveWriteCallback` to remove address monitors if required.

12 Implementing Custom PMA Behavior

The RISC-V architecture allows accesses to memory to be constrained by *physical memory attributes* (PMA). These PMA restrictions are not specified by the architecture but are machine-dependent. Typically, restrictions are either *statically defined* (e.g. by memory regions specified at build time) or *dynamic* (e.g. using custom CSRs in a similar way to the standard PMP CSRs).

The RISC-V model allows PMA constraints to be specified in several ways:

1. By static region definitions in the `riscvConfig` structure;
2. By commands callable when the simulation runs;
3. By extension model callbacks.

These three methods are each described below. The first two are appropriate for static PMA mappings, and the third for dynamic mappings.

12.1 Memory Domain Hierarchy

The RISC-V model uses OVP *memory domains* to efficiently implement access controls. A memory domain is essentially an address space with specified access permissions for subregions of that address space. Memory domains can implement physical memories, or peripherals (via callbacks on certain regions). They can also contain regions that are aliases of regions other memory domains, and so be used to model bus interconnect or MMU mappings very efficiently.

The RISC-V base model implements a static memory domain hierarchy, with a number of internally-created domain objects aliased to each other:

```
typedef struct riscvS {  
    . . . lines omitted . . .  
    memDomainP      extDomains[2];    // external domains (including CLIC)  
    memDomainP      pmaDomains[2];    // PMA domains  
    memDomainP      pmpDomains[2][2]; // PMP domains (M-mode and other modes)  
    . . . lines omitted . . .  
} riscv;
```

Furthest from the processor, there are the *external* domains (`extDomains`). These are provided by the instruction and data busses connected to the processor at instantiation time. In the `extDomains` array, index 0 holds the external *data* domain and index 1 the external *code* domain; often, these are identical.

Nearer the processor, there are the *PMA* domains (`pmaDomains`). These domains are aliases of the external domains. By default, these aliases are made with full RWX permissions and impose no other access constraints. If the external data and code domains are different objects, there will be distinct `pmaDomains` as well, otherwise they will be the same domain.

Nearer again to the processor, there are the *PMP* domains (`pmpDomains`), with separate versions for Machine mode and other modes, as well as for data and code accesses. These domains are aliases of the PMA domains, with additional constraints imposed by the

architectural *physical memory protection* (PMP) unit, if implemented. If PMA data and code domains are different objects, there will be distinct `pmpDomains` as well. If the PMA data and code domains are the same object, then whether `pmpDomains` code/data domains are also the same object is controlled by the `distinctPhysMem` extension object interface function (see section 15.30). Distinct domains are required if either address space view contains physical components that are not visible in the other (for example, closely-coupled memories).

If virtual memory is implemented, there are additional memory domains (not described in detail here), used to model the MMU. Regions in these domains are dynamically mapped to regions in the PMP domains to reflect MMU mappings as simulation runs.

Depending on the operating mode of the processor, the current data and code domains may either be MMU domains (if address translation is enabled) or PMP domains (if address translation is disabled).

The legality of an access is constrained by permission restrictions at *all levels* of the active domain hierarchy: a domain level can remove permissions for a type of access but not add permission. In order to model PMA, access permissions and other attributes need to be updated on the PMA domain objects.

12.2 Static PMA Region Definitions

The simplest way to model PMA constraints is by static PMA region definitions in the `riscvConfig` structure. This is done as follows for the defined configurations in the `vendor.com` template model:

```
//
// Static PMA mappings
//
static const riscvPMARegion pmaStaticRegions[] = {
    {.lo=0x0000000000, .hi=0x0000001FFF, .attrs="r-xa 1248"}, // BOOTROM
    {.lo=0x0000002000, .hi=0x0000FFFF, .attrs="rwx 1248"}, // RAM1
    {.lo=0x0080000000, .hi=0x00FFFFFF, .attrs="rwx- 1248"}, // RAM2
    {0} // KEEP LAST: terminator
};

//
// Defined configurations
//
static const riscvConfig configList[] = {

    {
        .name           = "RV32X",
        .arch            = ISA_U|RV32GC|ISA_X,
        .user_version    = RVUV_DEFAULT,
        .priv_version    = RVPV_DEFAULT,
        .tval_ii_code    = True,
        .ASID_bits       = 9,
        .local_int_num   = 7,           // enable local interrupts 16-22
        .unimp_int_mask  = 0x1f0000, // int16-int20 absent
        .extensionConfigs = allExtensions,
        .pmaStatic       = &pmaStaticRegions[0]
    },

    {
        .name           = "RV64X",
        .arch            = ISA_U|RV64GC|ISA_X,
```

```
.user_version      = RVUV_DEFAULT,  
.priv_version      = RVPV_DEFAULT,  
.tval_ii_code      = True,  
.ASID_bits         = 9,  
.local_int_num     = 7,           // enable local interrupts 16-22  
.unimp_int_mask    = 0x1f0000,   // int16-int20 absent  
.extensionConfigs  = allExtensions,  
.pmaStatic         = &pmaStaticRegions[0]  
},  
  
{0} // null terminator  
};
```

The `pmaStaticRegions` array is a NULL-terminated array of PMA region descriptions, each described by the following structure:

```
typedef struct riscvPMARegionS {  
    Uns64      lo;                // low region bound  
    Uns64      hi;                // high region bound  
    const char *attrs;           // region attributes  
} riscvPMARegion;
```

The bounds of a PMA region are specified by the `lo` and `hi` fields. Access constraints are specified by the `attrs` string, which can contain characters as follows:

r: read access allowed
w: write access allowed
x: execute access allowed
a: unaligned accesses disallowed
A: atomic instruction access disallowed (if `amo_constraint` specified)
L: `lr/sc` instruction access disallowed (if `lr_sc_constraint` specified)
P: `push/pop` instruction access disallowed (if `push_pop_constraint` specified)
V: vector instruction access disallowed (if `vector_constraint` specified)
M: `zicbom` instruction access disallowed
Z: `zicboz` instruction access disallowed
1: 1-byte accesses permitted
2: 2-byte accesses permitted
4: 4-byte accesses permitted
8: 8-byte accesses permitted
space: ignored, use for formatting if desired
-: ignored, use for formatting if desired

As an example, the string `"r-xa 1248"` allows 1, 2, 4 and 8 bytes read and fetch accesses, but not write accesses. Misaligned accesses are also disallowed. Any disallowed access that is not a misaligned access will cause an Access Fault exception of appropriate type. Misaligned accesses will by default cause Address Misaligned exceptions of appropriate type; this behavior can be modified using the `rdFaultCB` and `wrFaultCB` extension model interface functions (see sections 15.1 and 15.2). Instruction-type-specific accesses are controlled using instruction constraints, as described in section 5.8.1.

When a static PMA region definition is used, those definitions are applied once, when memory domains are created and initialized before simulation starts. Any areas of

memory without corresponding PMA region definitions will have all access denied. PMA region definitions are applied in sequence, replacing any previous definitions for the specified range: this means that it is possible to describe general permissions with a large region, and then restrict permissions for subranges of that region with subsequent region definitions if required.

12.3 PMA Region Definitions using Model Commands

PMA region definitions can also be specified by the `setPMA` command. For example, this command line applies some static PMA region definitions when the base RISC-V model is run using `iss.exe`:

```
iss.exe \  
--processorvendor riscv.ovpworld.org \  
--processorname riscv \  
--variant RV64GC \  
--program test.elf \  
--override iss/cpu0/simulateexceptions=T \  
--callcommand iss/cpu0/setPMA -lo 0x90000000 -hi 0x90000fff -attributes "r--" \  
--callcommand iss/cpu0/setPMA -lo 0x90001000 -hi 0x90001fff -attributes "-w-" \  
--callcommand iss/cpu0/setPMA -lo 0x90002000 -hi 0x90002fff -attributes "--x"
```

As for static region definitions, the `setPMA` commands are applied in sequence, so later calls can refine definitions specified by earlier calls. Unlike static definitions described in section 12.2, regions of the memory space not covered by `setPMA` commands have no PMA access constraints imposed.

It is possible to call `setPMA` commands at run time to modify PMA access constraints dynamically, in which case the specified region access constraints will replace any existing access constraints for the specified region.

12.4 PMA Control by Extension Models

Some RISC-V implementations allow PMA constraints to be configured dynamically, for example by using a custom CSR interface similar to the standard PMP interface. If such an interface is required, there are two places in which it can be activated:

1. When there are explicit updates to PMA control registers. For example, a CSR write that modifies range or access constraints of a programmable PMA region must modify memory domains in the `pmaDomains` array to reflect that change.
2. When there are accesses to memory (e.g. by a processor load or store) for which PMA access constraints are in effect. This is handled using the `PMACheck` extension interface function, which also results in modifications to memory domains in the `pmaDomains` array.

To implement PMA efficiently, it is best to use an approach in which PMA privileges are applied in a lazy fashion. Using this approach, any explicit updates to PMA control registers should remove all PMA permissions from a region range implied by the *original* value of the PMA control register as well as any region range implied by the *new* value of the PMA control register. The effect of this is that any subsequent memory access in

either range will cause the `PMACheck` extension interface function to be activated, which can update PMA privileges to either permit or deny the access.

12.4.1 Example PMA Control Implementation

The OpenHardware CV32E40X model implements PMA using regions specified as parameters, using the `PMACheck` extension interface function. It therefore provides a good template for modelling of lazy PMA mapping.

The `PMACheck` extension interface function is defined as:

```
static RISCVPMA_CHECK_FN(openhwPMACheck) {  
    vmiosObjectP object = clientData;  
  
    // apply PMA privileges  
    mapPMA(riscv, object, requiredPriv, lowPA, highPA);  
}
```

This function is passed the physical address range of a memory access using `lowPA` and `highPA` parameters. Function `mapPMA` then maps underlying PMA regions controlling the address range (there could be more than one region in the case of an access that straddles a PMA region boundary):

```
static void mapPMA(  
    riscvP      riscv,  
    vmiosObjectP object,  
    memPriv     requiredPriv,  
    Uns64       lowPA,  
    Uns64       highPA  
) {  
    Uns64 mask    = getAddressMask(riscv->extBits);  
    Uns64 highMap = lowPA-1;  
    Uns64 lowMap;  
  
    // iterate while unprocessed regions remain  
    do {  
        // get next region bounds to try  
        lowMap = highMap+1;  
        highMap = mask;  
  
        // attempt PMA privilege change for regions in implemented range  
        if(lowMap<=highMap) {  
            mapPMAInt(riscv, object, requiredPriv, lowPA, highPA, &lowMap, &highMap);  
        } else {  
            highMap = highPA;  
        }  
    } while(highMap<highPA);  
}
```

Function `mapPMAInt` is responsible for mapping a single region:

```
static void mapPMAInt(  
    riscvP      riscv,  
    vmiosObjectP object,  
    memPriv     requiredPriv,  
    Uns64       lowPA,  
    Uns64       highPA,  
    Uns64       *lowMapP,  
    Uns64       *highMapP
```

```
) {
    Uns64 PA = *lowMapP;
    memPriv priv = getDefaultPriv();
    Int32 i;

    // set widest possible range initially
    *lowMapP = 0;

    // handle all regions in lowest-to-highest priority order
    for(i=object->PMA_NUM_REGIONS-1; i>=0; i--) {
        refinePMARegionRange(object, lowMapP, highMapP, PA, i, &priv);
    }

    // get PMA region range
    Uns64 lowMap = *lowMapP;
    Uns64 highMap = *highMapP;

    // clamp physical range to maximum page size
    riscvClampPage(riscv, lowPA, highPA, &lowMap, &highMap);

    // update PMA privileges
    setPMAPriv(riscv, lowMap, highMap, priv, True);
}
```

This function traverses the set of configured PMA regions in lowest-to-highest priority order. After the traversal, it returns the bounds of the highest-priority region containing the accessed range and the access privileges required for that region. The returned bounds are possibly further restricted by the configured PMP/PMA maximum region size (see section 12.5) and then privileges on that range are updated by function `setPMAPriv`:

```
static void setDomainPriv(
    riscvP riscv,
    memDomainPP domains,
    Uns64 low,
    Uns64 high,
    memPriv priv,
    const char *name,
    Bool verbose
) {
    memDomainP dataDomain = domains[0];
    memDomainP codeDomain = domains[1];

    if(dataDomain==codeDomain) {

        // set permissions in unified domain
        vmirtProtectMemory(dataDomain, low, high, priv, MEM_PRIV_SET);

    } else {

        // get privileges for data and code domains
        memPriv privRW = priv&MEM_PRIV_RW ? priv&~MEM_PRIV_X : MEM_PRIV_NONE;
        memPriv privX = priv&MEM_PRIV_X ? priv&~MEM_PRIV_RW : MEM_PRIV_NONE;

        // set permissions in data domain
        vmirtProtectMemory(dataDomain, low, high, privRW, MEM_PRIV_SET);

        // set permissions in code domain
        vmirtProtectMemory(codeDomain, low, high, privX, MEM_PRIV_SET);

    }
}

static void setPMAPriv(
    riscvP riscv,
    Uns64 low,
    Uns64 high,
    memPriv priv,
    Bool verbose
)
```

```
) {  
    setDomainPriv(riscv, &riscv->pmaDomains[0], low, high, priv, "PMA", verbose);  
}
```

If the privileges set are insufficient for the access, the processor configuration then causes an appropriate Access Fault to be taken: this does not have to be explicitly handled by the PMACheck extension interface function.

12.5 PMA Mapped Page Size Restriction

By default, the example PMA control implementation in the previous section updates access permissions for the largest possible address range implied by the PMA settings; this could potentially be the entire address space range. This could cause a performance issue if the range being updated is highly fragmented, for example as a result of being divided into many separate subregions because of page descriptions implied by an active MMU.

To handle this potential performance issue, configuration parameter `PMP_max_page` can be used to limit the maximum size of PMP or PMA regions for which permission updates are made. If non-zero, the parameter specifies a power-of-two maximum page size that should be used for a single PMA or PMP mapping. The effect of this is to restrict the number of regions that are scanned when a large PMA region is updated.

The effect of `PMP_max_page` is applied by base model function `riscvClampPage`:

```
// get PMA region range  
Uns64 lowMap = *lowMapP;  
Uns64 highMap = *highMapP;  
  
// clamp physical range to maximum page size  
riscvClampPage(riscv, lowPA, highPA, &lowMap, &highMap);
```

This function is passed the address range of the memory access (`lowPA` and `highPA`) and the potentially large PMA region bounds (`lowMap` and `highMap`). It modifies `lowMap` and `highMap` to restrict them to page boundaries implied by `PMP_max_page` that include the required physical address range.

13 Appendix: Standard Instruction Patterns

This appendix describes the format of the standard instruction patterns specified by the `riscvExtInstrPattern` enumeration, including details of suitable disassembly format macros and fields set in the `riscvExtInstrInfo` structure during decode.

13.1 Pattern *RVIP_RD_RS1_RS2 (R-Type)*

Description: like `add x1, x2, x3`

Decode:

| | | | | |
|----------|-------|-------|-----|----------------|
| | rs2 | rs1 | rd | |
| vvvvvvvv | | | vvv | vvvvvvvv |

Format: `FMT_R1_R2_R3` (Or `FMT_R1_R2` or `FMT_R1`)

Fields set: `r[0]=rd, r[1]=rs1, r[2]=rs2`

13.2 Pattern *RVIP_RD_RS1_SI (I-Type)*

Description: like `addi x1, x2, imm` (immediate sign-extended to XLEN bits)

Decode:

| | | | |
|---------|-------|-----|----------------|
| imm11:0 | rs1 | rd | |
| | | vvv | vvvvvvvv |

Format: `FMT_R1_R2_SIMM, FMT_R1_R2_XIMM`

Fields set: `r[0]=rd, r[1]=rs1, c=sext(imm11:0)`

13.3 Pattern *RVIP_RD_RS1_SHIFT (I-Type - 5 or 6 bit shift)*

Description: like `slli x0, x1, shift`

Decode:

| | | | |
|--------------|-------|-----|------------------------|
| shift | rs1 | rd | |
| 0000000..... | | vvv | vvvvvvvv (RV32) |
| 0000000..... | | vvv | vvvvvvvv (RV64) |

Format: `FMT_R1_R2_SIMM, FMT_R1_R2_XIMM`

Fields set: `r[0]=rd, r[1]=rs1, c=shift`

13.4 Pattern *RVIP_BASE_RS2_OFFSET (S-Type)*

Description: like `sw x2, offset(x1)`

Decode:

| | | | | |
|---------|-------|-------|--------|----------------|
| imm11:5 | rs2 | rs1 | imm4:0 | |
| | | | vvv | vvvvvvvv |

Format: `FMT_R1_OFF_R2`

Fields set: `r[0]=rs2, r[1]=rs1, c=sext(imm11:0)`

13.5 Pattern *RVIP_RS1_RS2_OFFSET (B-Type)*

Description: like `beq x1, x2, offset`

Decode:

| | | | | |
|---------|-------|-------|--------|----------------|
| imm | rs2 | rs1 | imm | |
| 12,10:5 | | | 4:1,11 | vvvvvvvv |
| | | | vvv | vvvvvvvv |

Format: `FMT_R1_R2_TGT`

Fields set: `r[0]=rs1, r[1]=rs2, tgt=PC + sext(imm12:1 << 1)`

13.6 Pattern RVIP_RD_SI (U-Type)

Description: like `lui x1, imm`

Decode:

| | | | | | |
|--|----------|--|-------|--|----------|
| | imm31:12 | | rd | | vvvvvvvv |
| | | | | | |

Format: FMT_R1_UI

Fields set: `r[0]=rd, c=sext(imm31:12 << 12)`

13.7 Pattern RVIP_RD_OFFSET (J-Type)

Description: like `jal x1, offset`

Decode:

| | | | | | |
|--|---------------------|--|-------|--|----------|
| | imm20,10:1,11,19:12 | | rd | | vvvvvvvv |
| | | | | | |

Format: FMT_R1_TGT

Fields set: `r[0]=rd, tgt=PC + sext(imm20:1 << 1)`

13.8 Pattern RVIP_RD_RS1_RS2_RS3 (R4-Type)

Description: like `cmix x0, x1, x2, x3`

Decode:

| | | | | | | | | | |
|--|-------|--|-----|--|-------|--|-------|--|----------|
| | rs3 | | rs2 | | rs1 | | rd | | vvvvvvvv |
| | | | vv | | | | | | vvvvvvvv |

Format: FMT_R1_R2_R3_R4

Fields set: `r[0]=rd, r[1]=rs1, r[2]=rs2, r[3]=rs3`

13.9 Pattern RVIP_RD_RS1_RS3_SHIFT (Non-Standard)

Description: like `fsri x0, x1, x2, shift`

Decode:

| | | | | | | | | | |
|--|-------|--|----------|--|-------|--|-----|--|-----------------|
| | rs3 | | shift | | rs1 | | rd | | vvvvvvvv |
| | | | v 0..... | | | | vvv | | vvvvvvvv (RV32) |
| | | | v | | | | vvv | | vvvvvvvv (RV64) |

Format: FMT_R1_R2_R3_SIMM

Fields set: `r[0]=rd, r[1]=rs1, r[2]=rs3, c=shift`

13.10 Pattern RVIP_FD_FS1_FS2

Description: like `fmax.s f0, f1, f2`

Decode:

| | | | | | | | |
|--|-------|--|-------|--|-----|--|----------|
| | fs2 | | fs1 | | fd | | vvvvvvvW |
| | | | | | vvv | | vvvvvvvv |

Format: FMT_R1_R2_R3 (Or FMT_R1_R2 or FMT_R1)

Fields set: `r[0]=fd, r[1]=fs1, r[2]=fs2`

Width: `W=0:s, W=1:d`

13.11 Pattern RVIP_FD_FS1_FS2_RM

Description: like `fadd.s f0, f1, f2, rte`

Decode:

| | | | | | | | | | |
|--|-------|--|-------|--|-----|--|-------|--|----------|
| | fs2 | | fs1 | | rm | | fd | | vvvvvvvW |
| | | | | | ... | | | | vvvvvvvv |

Format: FMT_R1_R2_R3 (Or FMT_R1_R2 or FMT_R1)

Fields set: `r[0]=fd, r[1]=fs1, r[2]=fs2, rm=rm`

Width: `W=0:s, W=1:d`

13.12 Pattern RVIP_FD_FS1_FS2_FS3_RMDescription: like **fmadd.s f0, f1, f2, f3, rte**Decode:

| | | | | | | | | | | |
|-------|--|-----|--|-------|--|-------|--|-----|--|----------|
| fs3 | | fs2 | | fs1 | | rm | | fd | | vvvvvvvv |
| | | vW | | | | | | ... | | |

Format: FMT_R1_R2_R3_R4

Fields set: $r[0]=fd, r[1]=fs1, r[2]=fs2, r[3]=fs3, rm=rm$ Width: $W=0:s, W=1:d$ **13.13 Pattern RVIP_RD_FS1_FS2**Description: like **feq.s x0, f1, f2**Decode:

| | | | | | | |
|-------|--|-------|--|-----|--|----------|
| fs2 | | fs1 | | rd | | vvvvvvvW |
| | | | | vvv | | vvvvvvvv |

Format: FMT_R1_R2_R3

Fields set: $r[0]=rd, r[1]=fs1, r[2]=fs2$ Width: $W=0:s, W=1:d$ **13.14 Pattern RVIP_VD_VS1_VS2_M**Description: like **vadd.vv v1, v2, v3, v0.m**Decode:

| | | | | | | | | |
|---|--|-------|--|-------|--|-----|--|----------|
| m | | vs1 | | vs2 | | vd | | vvvvvvv |
| . | | | | | | vvv | | vvvvvvvv |

Format: FMT_R1_R2_R3_RM

Fields set: $r[0]=vd, r[1]=vs1, r[2]=vs2, mask=m?none:v0$ **13.15 Pattern RVIP_VD_VS1_SI_M**Description: like **vadd.vi v1, v2, imm, v0.m** (immediate sign-extended to SEW bits)Decode:

| | | | | | | | | |
|---|--|-------|--|-------|--|-----|--|----------|
| m | | vs1 | | imm | | vd | | vvvvvvv |
| . | | | | | | vvv | | vvvvvvvv |

Format: FMT_R1_R2_SIMM_RM

Fields set: $r[0]=vd, r[1]=vs1, c=imm, mask=m?none:v0$ **13.16 Pattern RVIP_VD_VS1_UI_M**Description: like **vsll.vi v1, v2, imm, v0.m** (immediate zero-extended to SEW bits)Decode:

| | | | | | | | | |
|---|--|-------|--|-------|--|-----|--|----------|
| m | | vs1 | | imm | | vd | | vvvvvvv |
| . | | | | | | vvv | | vvvvvvvv |

Format: FMT_R1_R2_SIMM_RM

Fields set: $r[0]=vd, r[1]=vs1, c=imm, mask=m?none:v0$ **13.17 Pattern RVIP_VD_VS1_RS2_M**Description: like **vadd.vx v1, v2, x3, v0.m**Decode:

| | | | | | | | | |
|---|--|-------|--|-------|--|-----|--|----------|
| m | | vs1 | | rs2 | | vd | | vvvvvvv |
| . | | | | | | vvv | | vvvvvvvv |

Format: FMT_R1_R2_R3_RM

Fields set: $r[0]=vd, r[1]=vs1, r[2]=rs2, mask=m?none:v0$

13.18 Pattern RVIP_VD_VS1_FS2_MDescription: like **vfadd.vf v1, v2, f3, v0.m**Decode:

| | | | | | | |
|---------|---|-------|-------|-----|-------|----------|
| | m | vs1 | fs2 | | vd | |
| vvvvvvv | . | | | vvv | | vvvvvvvv |

Format: FMT_R1_R2_R3_RM

Fields set: r[0]=vd, r[1]=vs1, r[2]=fs2, mask=m?none:v0

13.19 Pattern RVIP_RD_VS1_RS2Description: like **vext.v r1, v2, r3**Decode:

| | | | | | |
|---------|-------|-------|-----|-------|----------|
| | vs1 | rs2 | | rd | |
| vvvvvvv | | | vvv | | vvvvvvvv |

Format: FMT_R1_R2_R3

Fields set: r[0]=rd, r[1]=vs1, r[2]=rs2

13.20 Pattern RVIP_RD_VS1_MDescription: like **vpopc.m x1, v2, v0.m**Decode:

| | | | | | |
|---------|---|-------|----------|-------|----------|
| | m | vs1 | | rd | |
| vvvvvvv | . | | vvvvvvvv | | vvvvvvvv |

Format: FMT_R1_R2_RM

Fields set: r[0]=rd, r[1]=vs1, mask=m?none:v0

13.21 Pattern RVIP_VD_RS2Description: like **vmv.s.x v1, x2**Decode:

| | | | | |
|--------------|-------|-----|-------|----------|
| | rs2 | | vd | |
| vvvvvvvvvvvv | | vvv | | vvvvvvvv |

Format: FMT_R1_R2

Fields set: r[0]=vd, r[1]=rs2

13.22 Pattern RVIP_FD_VS1Description: like **vfmv.f.s f1, v2**Decode:

| | | | | |
|----------|-------|----------|-------|----------|
| | vs2 | | fd | |
| vvvvvvvv | | vvvvvvvv | | vvvvvvvv |

Format: FMT_R1_R2

Fields set: r[0]=fd, r[1]=vs2

13.23 Pattern RVIP_VD_FS2Description: like **vfmv.s.f v1, f2**Decode:

| | | | | |
|--------------|-------|-----|-------|----------|
| | fs2 | | vd | |
| vvvvvvvvvvvv | | vvv | | vvvvvvvv |

Format: FMT_R1_R2

Fields set: r[0]=vd, r[1]=fs2

14 Appendix: Base Model Interface Service Functions

This appendix describes *interface functions* implemented by the *base model* that are available to provide services for a *linked model extension library*. All such interface functions are installed in a structure of type `riscvModelCB` accessible via the `cb` field in the RISC-V processor structure. For example, an extension library can call the `takeException` interface function (which causes an exception to be immediately taken) like this:

```
riscv->cb.takeException(riscv, EXT_E_EXCEPT24, 0);
```

The `riscvModelCB` type is defined in file `riscvModelCallbacks.h` in the base model:

```
typedef struct riscvModelCBS {

    // from riscvUtils.h
    riscvRegisterExtCBFn      registerExtCB;
    riscvGetExtClientDataFn  getExtClientData;
    riscvGetExtConfigFn      getExtConfig;
    riscvGetXlenFn           getXlenMode;
    riscvGetXlenFn           getXlenArch;
    riscvGetRegNameFn        getXRegName;
    riscvGetRegNameFn        getFRegName;
    riscvGetRegNameFn        getVRegName;
    riscvSetTModeFn          setTMode;
    riscvGetTModeFn          getTMode;
    riscvGetDataEndianFn     getDataEndian;
    riscvReadCSRNumFn        readCSR;
    riscvWriteCSRNumFn       writeCSR;
    riscvReadBaseCSRFn       readBaseCSR;
    riscvWriteBaseCSRFn      writeBaseCSR;

    // from riscvExceptions.h
    riscvHaltRestartFn       halt;
    riscvHaltRestartFn       block;
    riscvHaltRestartFn       restart;
    riscvUpdateInterruptFn   updateInterrupt;
    riscvUpdateDisableFn     updateDisable;
    riscvUpdateDisableNMIFn  updateDisableNMI;
    riscvTestInterruptFn     testInterrupt;
    riscvResumeFromWFI       resumeFromWFI;
    riscvIllegalInstructionFn illegalInstruction;
    riscvIllegalVerboseFn    illegalVerbose;
    riscvIllegalInstructionFn virtualInstruction;
    riscvIllegalVerboseFn    virtualVerbose;
    riscvIllegalCustomFn     illegalCustom;
    riscvTakeExceptionFn     takeException;
    riscvPendFetchExceptionFn pendFetchException;
    riscvTakeResetFn         takeReset;
    riscvAcknowledgeCLICIntFn acknowledgeCLICInt;

    // from riscvDecode.h
    riscvFetchInstructionFn  fetchInstruction;

    // from riscvDisassemble.h
    riscvDisassInstructionFn disassInstruction;

    // from riscvMorph.h
    riscvInstructionEnabledFn instructionEnabled;
    riscvMorphExternalFn     morphExternal;
    riscvMorphIllegalFn      morphIllegal;
    riscvMorphIllegalFn      morphVirtual;
    riscvGetVMIRegFn         getVMIReg;
    riscvGetVMIRegFSFn       getVMIRegFS;
```

```
riscvWriteRegSizeFn    writeRegSize;
riscvWriteRegFn        writeReg;
riscvGetFPFlagsMtFn    getFPFlagsMt;
riscvGetDataEndianMtFn getDataEndianMt;
riscvLoadMtFn          loadMt;
riscvStoreMtFn         storeMt;
riscvRequireModeMtFn   requireModeMt;
riscvRequireNotVMtFn   requireNotVMt;
riscvCheckLegalRMMtFn  checkLegalRMMt;
riscvMorphTrapTVMFn    morphTrapTVM;
riscvMorphVOpFn        morphVOp;

// from riscvCSR.h
riscvNewCSRFn          newCSR;
riscvHPMAccessValidFn  hpmAccessValid;

// from riscvVM.h
riscvMapAddressFn       mapAddress;
riscvUnmapPMPRegionFn   unmapPMPRegion;
riscvUpdateLdStDomainFn updateLdStDomain;
riscvNewTLBEntryFn      newTLBEntry;
riscvFreeTLBEntryFn     freeTLBEntry;

// from riscvDebug.h
riscvNewExtRegFn        newExtReg;

} riscvModelCB;
```

Following subsections describe each interface function.

14.1 Function *registerExtCB*

```
#define RISC_V_REGISTER_EXT_CB_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvExtCBP extCB, \
    Uns32       id      \
)
typedef RISC_V_REGISTER_EXT_CB_FN((*riscvRegisterExtCBFn));

typedef struct riscvModelCBS {
    riscvRegisterExtCBFn    registerExtCB;
} riscvModelCB;
```

Description

This interface function is called from the extension object constructor to register that extension object with the base model. It requires the RISC-V processor, an object of type `riscvExtCBP` and an identifier as arguments. The identifier must be a unique index among all extensions added to the RISC-V processor.

The `riscvExtCBP` object should be pointer to a field of type `riscvExtCB` that is declared in the extension `vmiosObject` structure. The `riscvExtCB` is filled with callback functions and other information by the extension object constructor. These fields are used by the base model, primarily to notify the extension object of state changes in other events, and also to allow the extension object to modify some base model behavior.

Example

```
static VMIOS_CONSTRUCTOR_FN(addInstructionsConstructor) {
    riscvP riscv = (riscvP)processor;

    object->riscv = riscv;

    // prepare client data
    object->extCB.clientData = object;

    // register extension with base model using unique ID
    riscv->cb.registerExtCB(riscv, &object->extCB, EXTID_ADDINST);
}
```

Usage Context

Container or leaf level.

14.2 Function *getExtClientData*

```
#define RISC_V_GET_EXT_CLIENT_DATA_FN(_NAME) void *_NAME( \
    riscvP riscv,      \
    Uns32 id,          \
)
typedef RISC_V_GET_EXT_CLIENT_DATA_FN((*riscvGetExtClientDataFn));

typedef struct riscvModelCBS {
    riscvGetExtClientDataFn  getExtClientData;
} riscvModelCB;
```

Description

This interface function returns any extension object previously registered with the processor which has the given unique identifier.

Example

```
vmiosObjectP getExtObject(riscvP riscv) {
    vmiosObjectP object = riscv->cb.getExtClientData(riscv, EXT_ID);
    return object;
}
```

Usage Context

Container or leaf level.

14.3 Function *getExtConfig*

```
#define RISC_V_GET_EXT_CONFIG_FN(_NAME) riscvExtConfigCP _NAME( \
    riscvP riscv, \
    Uns32 id \
)
typedef RISC_V_GET_EXT_CONFIG_FN((*riscvGetExtConfigFn));

typedef struct riscvModelCBS {
    riscvGetExtConfigFn    getExtConfig;
} riscvModelCB;
```

Description

This interface function returns any extension configuration object for the processor, given an extension unique identifier. The extension configuration holds any variant-specific information that may modify the behavior of an extension. The extension configuration is held with the standard configuration information for a variant.

Example

With configuration list:

```
static riscvExtConfigCP allExtensions[] = {

    // example adding CSRs
    &(const riscvExtConfig){
        .id      = EXTID_ADDCSR,
        .userData = &(const addCSRsConfig){
            .csr = {
                .custom_rol = {u32 : {bits : 0x12345678}}
            }
        },
        0
    };
```

In extension object:

```
static addCSRsConfigCP getExtConfig(riscvP riscv) {
    riscvExtConfigCP cfg = riscv->cb.getExtConfig(riscv, EXTID_ADDCSR);
    return cfg->userData;
}
```

See section 7 for a complete example.

Usage Context

Container or leaf level.

14.4 Function *getXlenMode*

```
#define RISC_V_GET_XLEN_FN(_NAME) Uns32 _NAME(riscvP riscv)
typedef RISC_V_GET_XLEN_FN(*riscvGetXlenFn);

typedef struct riscvModelCBS {
    riscvGetXlenFn      getXlenMode;
} riscvModelCB;
```

Description

This interface function returns the currently-active XLEN.

Example

```
inline static Uns32 getXlenBits(riscvP riscv) {
    return riscv->cb.getXlenMode(riscv);
}
```

Usage Context

Leaf level only.

14.5 Function *getXlenArch*

```
#define RISC_V_GET_XLEN_FN(_NAME) Uns32 _NAME(riscvP riscv)
typedef RISC_V_GET_XLEN_FN(*riscvGetXlenFn);

typedef struct riscvModelCBS {
    riscvGetXlenFn      getXlenArch;
} riscvModelCB;
```

Description

This interface function returns the processor architectural XLEN.

Example

```
inline static Uns32 getXlenArch(riscvP riscv) {
    return riscv->cb.getXlenArch(riscv);
}
```

Usage Context

Leaf level only.

14.6 Function *getXRegName*

```
#define RISC_V_GET_REG_NAME_FN(_NAME) const char *_NAME(riscvP riscv, Uns32 index)
typedef RISC_V_GET_REG_NAME_FN((*riscvGetRegNameFn));

typedef struct riscvModelCBS {
    riscvGetRegNameFn      getXRegName;
} riscvModelCB;
```

Description

This interface function returns the name of a RISC-V GPR given its index.

Example

```
const char *getXRegName(riscvP riscv, Uns32 index) {
    return riscv->cb.getXRegName(riscv, index);
}
```

Usage Context

Leaf level only.

14.7 Function *getFRegName*

```
#define RISC_V_GET_REG_NAME_FN(_NAME) const char *_NAME(riscvP riscv, Uns32 index)
typedef RISC_V_GET_REG_NAME_FN((*riscvGetRegNameFn));

typedef struct riscvModelCBS {
    riscvGetRegNameFn      getFRegName;
} riscvModelCB;
```

Description

This interface function returns the name of a RISC-V FPR given its index.

Example

```
const char *getFRegName(riscvP riscv, Uns32 index) {
    return riscv->cb.getFRegName(riscv, index);
}
```

Usage Context

Leaf level only.

14.8 Function *getVRegName*

```
#define RISC_V_GET_REG_NAME_FN(_NAME) const char *_NAME(riscvP riscv, Uns32 index)
typedef RISC_V_GET_REG_NAME_FN((*riscvGetRegNameFn));

typedef struct riscvModelCBS {
    riscvGetRegNameFn      getVRegName;
} riscvModelCB;
```

Description

This interface function returns the name of a RISC-V vector register given its index.

Example

```
const char *getVRegName(riscvP riscv, Uns32 index) {
    return riscv->cb.getVRegName(riscv, index);
}
```

Usage Context

Leaf level only.

14.9 Function *setTMode*

```
#define RISC_V_SET_TMODE_FN(_NAME) void _NAME(riscvP riscv, Bool enable)
typedef RISC_V_SET_TMODE_FN(*riscvSetTModeFn);

typedef struct riscvModelCBS {
    riscvSetTModeFn      setTMode;
} riscvModelCB;
```

Description

This interface function enables or disabled *transactional memory mode* for the processor. When processors with transactional memory are being modeled, some instructions behave differently when the mode is enabled. For example, loads and stores typically accumulate information in cache lines or other structures without committing the values to memory in transactional mode.

Enabling transactional memory mode causes JIT code translations to be stored in and used from *a separate code dictionary* while that mode is active. This means that different behaviors for the same instructions can be modeled without the inefficiency inherent in transactional memory behavior affecting non-transaction mode.

Note that implementation of transactional memory requires standard load/store instructions (and others) to be reimplemented in the extension object in the case that transactional mode is active. Contact Imperas for further advice about modeling such features.

Example

```
static void setTMode(riscvP riscv, Bool enable) {
    riscv->cb.setTMode(riscv, enable);
}
```

Usage Context

Leaf level only.

14.10 *Function getTMode*

```
#define RISC_V_GET_TMODE_FN(_NAME) Bool _NAME(riscvP riscv)
typedef RISC_V_GET_TMODE_FN(*riscvGetTModeFn);

typedef struct riscvModelCBS {
    riscvGetTModeFn      getTMode;
} riscvModelCB;
```

Description

This interface function returns a Boolean indicating whether transactional memory mode is currently active. The function may be called at either run time or morph time (within the morph callback). See section 14.9 for more information.

Example

```
static Bool getTMode(riscvP riscv) {
    return riscv->cb.getTMode(riscv);
}
```

Usage Context

Leaf level only.

14.11 *Function `getDataEndian`*

```
#define RISC_V_GET_DATA_ENDIAN_FN(_NAME) memEndian _NAME( \
    riscvP    riscv,    \
    riscvMode mode      \
)
typedef RISC_V_GET_DATA_ENDIAN_FN((*riscvGetDataEndianFn));

typedef struct riscvModelCBS {
    riscvGetDataEndianFn    getDataEndian;
} riscvModelCB;
```

Description

This interface function returns the active endianness for loads and stores in the given processor mode. Usually, RISC-V processors use little-endian order for loads and stores, but this is configurable.

Example

```
static memEndian getDataEndian(riscvP riscv, riscvMode mode) {
    return riscv->cb.getDataEndian(riscv, mode);
}
```

Usage Context

Leaf level only.

14.12 *Function readCSR*

```
#define RISC_V_READ_CSR_NUM_FN(_NAME) Uns64 _NAME( \
    riscvP riscv, \
    Uns32 csrNum \
)
typedef RISC_V_READ_CSR_NUM_FN(*riscvReadCSRNumFn);

typedef struct riscvModelCBS {
    riscvReadCSRNumFn readCSR;
} riscvModelCB;
```

Description

This interface function returns the current value of a CSR, given its index number. The value returned can be either from the base model or from a CSR implemented in an extension object.

Example

```
static Uns64 readCSR(riscvP riscv, Uns32 csrNum) {
    return riscv->cb.readCSR(riscv, csrNum);
}
```

Usage Context

Leaf level only.

14.13 *Function writeCSR*

```
#define RISC_V_WRITE_CSR_NUM_FN(_NAME) Uns64 _NAME( \
    riscvP riscv, \
    Uns32  csrNum, \
    Uns64  newValue \
)
typedef RISC_V_WRITE_CSR_NUM_FN(*riscvWriteCSRNumFn);

typedef struct riscvModelCBS {
    riscvWriteCSRNumFn writeCSR;
} riscvModelCB;
```

Description

This interface function writes a new value to a CSR, given its index number. The CSR updated can be implemented either in the base model or in an extension object.

Example

```
static void writeCSR(riscvP riscv, Uns32 csrNum, Uns64 newValue) {
    riscv->cb.writeCSR(riscv, csrNum, newValue);
}
```

Usage Context

Leaf level only.

14.14 *Function readBaseCSR*

```
#define RISCV_READ_BASE_CSR_FN(_NAME) Uns64 _NAME( \
    riscvP    riscv, \
    riscvCSRId id \
)
typedef RISCV_READ_BASE_CSR_FN((*riscvReadBaseCSRFn));

typedef struct riscvModelCBS {
    riscvReadBaseCSRFn    readBaseCSR;
} riscvModelCB;
```

Description

This interface function returns the current value of a CSR, given its `riscvCSRId` identifier (*not* the CSR number). The CSR read will be that in the base model, disregarding any extension object redefinition. This function is useful when an extension object wishes to add fields into a standard CSR, retaining the behavior of the standard fields. In this case, the extension can install a replacement for the standard CSR, handle the new fields itself and merge in standard fields using `readBaseCSR` and `writeBaseCSR`.

Example

```
inline static Uns64 baseR(riscvP riscv, riscvCSRId id) {
    return riscv->cb.readBaseCSR(riscv, id);
}
```

Usage Context

Leaf level only.

14.15 *Function* `writeBaseCSR`

```
#define RISCV_WRITE_BASE_CSR_FN(_NAME) Uns64 _NAME( \
    riscvP    riscv,      \
    riscvCSRId id,        \
    Uns64     newValue    \
)
typedef RISCV_WRITE_BASE_CSR_FN((*riscvWriteBaseCSRFn));

typedef struct riscvModelCBS {
    riscvWriteBaseCSRFn    writeBaseCSR;
} riscvModelCB;
```

Description

This interface function writes the current value of a CSR, given its `riscvCSRId` identifier (*not* the CSR number). The CSR written will be that in the base model, disregarding any extension object redefinition. This function is useful when an extension object wishes to add fields into a standard CSR, retaining the behavior of the standard fields. In this case, the extension can install a replacement for the standard CSR, handle the new fields itself and merge in standard fields using `readBaseCSR` and `writeBaseCSR`.

Example

```
inline static void baseW(riscvP riscv, riscvCSRId id, Uns64 newValue) {
    riscv->cb.writeBaseCSR(riscv, id, newValue);
}
```

Usage Context

Leaf level only.

14.16 Function *halt*

```
#define RISC_V_HALT_RESTART_FN(_NAME) void _NAME( \
    riscvP          riscv, \
    riscvDisableReason reason \
)
typedef RISC_V_HALT_RESTART_FN((*riscvHaltRestartFn));

typedef struct riscvModelCBS {
    riscvHaltRestartFn      halt;
} riscvModelCB;
```

Description

This interface function causes the given hart to halt, for a reason given by the `riscvDisableReason` enumerated type:

```
typedef enum riscvDisableReasonE {
    RVD_ACTIVE      = 0x00, // running
    RVD_WFI         = 0x01, // waiting in WFI
    RVD_RESET       = 0x02, // waiting in reset
    RVD_DEBUG       = 0x04, // waiting for debug
    RVD_WRS         = 0x08, // waiting for reservation set (wrs.*)
    RVD_STO         = 0x10, // waiting for reservation set (wrs.sto)
    RVD_CUSTOM_WFI  = 0x20, // waiting for custom reason with WFI semantics
    RVD_CUSTOM_NMI  = 0x40, // waiting for custom reason with NMI semantics
};
```

Reason `RVD_WFI` is used by the base model to indicate the hart is stalled in a `wfi` instruction.

Reason `RVD_RESET` is used by the base model to indicate the hart is stalled because the reset signal is high.

Reason `RVD_DEBUG` is used only when debug mode is configured to halt the processor (`debug_mode` is `RVDM_HALT`) and indicates the hart is halted in debug state.

Reason `RVD_WRS` is used by the base model to indicate the hart is stalled in a `wrs.nto` or `wrs.sto` instruction (see the `Zawrs` extension).

Reason `RVD_STO` is used by the base model to indicate the hart is stalled in a `wrs.sto` instruction (see the `Zawrs` extension).

Reason `RVD_CUSTOM_WFI` is available for use in a custom extension. It indicates the processor is stalled for a reason with the same implicit wake-up semantics as the `wfi` instruction (for example, a locally-enabled interrupt becomes pending, even if globally disabled).

Reason `RVD_CUSTOM_NMI` is also available for use in a custom extension. It indicates the processor is stalled and should awaken on reset or if a fully-enabled interrupt is pending.

The hart will remain halted until a subsequent call to the `restart` function or a reset or other interrupt event, depending on the halt reason semantics.

Example

This example shows implementation of a wait-for-event (wfe) instruction with the same semantics as wfi for a processor with M-mode and U-mode.

```
static void doWFE(riscvP riscv) {  
  
    riscvMode mode = getCurrentMode5(riscv);  
  
    if((mode!=RISCV_MODE_M) && RD_CSR_FIELDC(riscv, mstatus, TW)) {  
  
        // WFE not available when mstatus.TW=1  
        riscv->cb.illegalVerbose(riscv, "mstatus.TW=1");  
  
    } else if(!riscv->cb.resumeFromWFI(riscv)) {  
  
        // stall for custom reason (WFE) with WFI wake semantics  
        riscv->cb.halt(riscv, RVD_CUSTOM_WFI);  
    }  
}
```

Usage Context

Leaf level only.

14.17 Function block

```
#define RISC_V_HALT_RESTART_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvDisableReason reason \
)
typedef RISC_V_HALT_RESTART_FN((*riscvHaltRestartFn));

typedef struct riscvModelCBS {
    riscvHaltRestartFn      block;
} riscvModelCB;
```

Description

This interface function causes the given hart to block during execution of an instruction, for a reason given by the `riscvDisableReason` enumerated type:

```
typedef enum riscvDisableReasonE {
    RVD_ACTIVE      = 0x00, // running
    RVD_WFI         = 0x01, // waiting in WFI
    RVD_RESET       = 0x02, // waiting in reset
    RVD_DEBUG       = 0x04, // waiting for debug
    RVD_WRS         = 0x08, // waiting for reservation set (wrs.*)
    RVD_STO         = 0x10, // waiting for reservation set (wrs.sto)
    RVD_CUSTOM_WFI  = 0x20, // waiting for custom reason with WFI semantics
    RVD_CUSTOM_NMI  = 0x40, // waiting for custom reason with NMI semantics
};
```

See section 14.16 for a description of the `riscvDisableReason` enumerated type.

The hart will remain blocked until a subsequent call to the `restart` function or a reset or other interrupt event, depending on the halt reason semantics. Upon restart, the currently-executing instruction will be *re-executed*, provided that the reason for restart was not to take a trap.

Typically, this function is used to allow simulation of load or store instructions with blocking semantics. The custom extension should add a net port allowing a signal to be raised while the load or store is active, to block execution if required. An external component that drives this signal needs to implement logic to determine whether a particular load or store should be blocked, taking into account the fact that the load or store will be restarted when the hart resumes execution from a blocked state.

Example

This example shows an example implementation of a block signal.

```
static VMI_NET_CHANGE_FN(blockCB) {
    vmiosObjectP object = userData;
    riscvP      riscv   = object->riscv;
    Bool        disable  = newValue & 1;

    if(disable) {
        riscv->cb.block(riscv, RVD_CUSTOM_NMI);
    } else {
        riscv->cb.restart(riscv, RVD_CUSTOM_NMI);
    }
}
```


Usage Context

Leaf level only.

14.18 *Function restart*

```
#define RISC_V_HALT_RESTART_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvDisableReason reason \
)
typedef RISC_V_HALT_RESTART_FN((*riscvHaltRestartFn));

typedef struct riscvModelCBS {
    riscvHaltRestartFn      restart;
} riscvModelCB;
```

Description

This interface function causes the given hart to restart, for the reason given by the `riscvDisableReason` enumerated type:

```
typedef enum riscvDisableReasonE {
    RVD_ACTIVE      = 0x00, // running
    RVD_WFI         = 0x01, // waiting in WFI
    RVD_RESET       = 0x02, // waiting in reset
    RVD_DEBUG       = 0x04, // waiting for debug
    RVD_WRS         = 0x08, // waiting for reservation set (wrs.*)
    RVD_STO         = 0x10, // waiting for reservation set (wrs.sto)
    RVD_CUSTOM_WFI  = 0x20, // waiting for custom reason with WFI semantics
    RVD_CUSTOM_NMI  = 0x40, // waiting for custom reason with NMI semantics
};
```

See section 14.16 for a description of the `riscvDisableReason` enumerated type.

The function has no effect if the hart is not halted for the given reason.

Example

This example shows implementation of a wait-for-event (`wfe`) instruction with the same semantics as `wfi`.

```
static VMI_NET_CHANGE_FN(wakeWFE) {
    vmiosObjectP object = userData;
    riscvP      riscv   = object->riscv;
    Bool        old      = object->wakeWFE;
    Bool        new      = newValue&1;

    // detect rising edge
    if(!old && new) {
        riscv->cb.restart(riscv, RVD_CUSTOM_WFI);
    }

    object->wakeWFE = new;
}
```

Usage Context

Leaf level only.

14.19 *Function updateInterrupt*

```
#define RISC_V_UPDATE_INTERRUPT_FN(_NAME) void _NAME( \
    riscvP riscv, \
    Uns32 index, \
    Bool newValue \
)
typedef RISC_V_UPDATE_INTERRUPT_FN((*riscvUpdateInterruptFn));

typedef struct riscvModelCBS {
    riscvUpdateInterruptFn updateInterrupt;
} riscvModelCB;
```

Description

This interface function updates the value of the indexed interrupt, as seen in the in the mip CSR. For example, to raise the M-mode software interrupt, index would be 3 and newValue 1.

Example

```
inline static void updateStandardInterrupt(
    vmiosObjectP object,
    Uns32 index,
    Bool newValue
) {
    riscvP riscv = object->riscv;

    riscv->cb.updateInterrupt(riscv, index, newValue);
}
```

Usage Context

Leaf level only.

14.20 *Function updateDisable*

```
#define RISC_V_UPDATE_DISABLE_FN(_NAME) void _NAME( \
    riscvP riscv, \
    Uns64 disableMask \
)
typedef RISC_V_UPDATE_DISABLE_FN((*riscvUpdateDisableFn));

typedef struct riscvModelCBS {
    riscvUpdateDisableFn updateDisable;
} riscvModelCB;
```

Description

This interface function allows a supplementary mask of disabled interrupts to be applied by a custom extension: any standard interrupt with an index number corresponding to a bit that is set in this mask will be disabled, overriding standard behavior specified by the mie CSR.

Example

```
static Uns64 ecsrW(vmiosObjectP object, Uns64 newValue, Uns64 mask) {

    riscvP riscv = object->riscv;
    Uns64 oldValue = RD_XCSR(object, mecsr);

    // update mecsr value
    WR_XCSR(object, mecsr, ((newValue & mask) | (oldValue & ~mask)));

    // get mask of externally-disabled interrupts
    Uns64 disableMask = RD_XCSR_FIELD(object, mecsr, WEDIS) ? (1<<CUST_MERROR) : 0;

    // update mask of externally-disabled interrupts
    riscv->cb.updateDisable(riscv, disableMask);

    // return composed value
    return RD_XCSR(object, mecsr) & mask;
}
```

Usage Context

Leaf level only.

14.21 *Function `updateDisableNMI`*

```
#define RISC_V_UPDATE_DISABLE_NMI_FN(_NAME) void _NAME( \
    riscvP riscv, \
    Bool disable \
)
typedef RISC_V_UPDATE_DISABLE_NMI_FN((*riscvUpdateDisableNMIFn));

typedef struct riscvModelCBS {
    riscvUpdateDisableNMIFn updateDisableNMI;
} riscvModelCB;
```

Description

This interface allows NMI interrupts to be enabled or disabled for a particular hart. This allows extensions to control which harts in a multi-hart implementation react to external NMI inputs. This is typically required in implementations that broadcast a single NMI to multiple harts.

Example

```
static void updateNMIPDEL(riscvP thread, Bool delegate) {
    thread->cb.updateDisableNMI(thread, !delegate);
}
```

Usage Context

Leaf level only.

14.22 *Function testInterrupt*

```
#define RISC_V_TEST_INTERRUPT_FN(_NAME) void _NAME(riscvP riscv)
typedef RISC_V_TEST_INTERRUPT_FN((*riscvTestInterruptFn));

typedef struct riscvModelCBS {
    riscvTestInterruptFn    testInterrupt;
} riscvModelCB;
```

Description

This interface function causes the base model to check for any pending interrupts. It should be called when an extension library has performed some action that might cause a pending interrupt to be activated (for example, pending it, or unmasking it when already pending).

Example

```
static void testInterrupt(riscvP riscv) {
    riscv->cb.testInterrupt(riscv);
}

static RISC_V_CSR_WRITEFN(slieW) {

    vmiosObjectP object    = attrs->object;
    Uns32          oldValue = RD_XCSR(object, slie);
    Uns32          mask     = slieMask(object);

    // update writeable bits
    newValue = (newValue & mask) | (oldValue & ~mask);

    // handle any change to interrupt state
    if(newValue != oldValue) {
        WR_XCSR(object, slie, newValue);
        testInterrupt(riscv);
    }

    return newValue;
}
```

Usage Context

Leaf level only.

14.23 *Function resumeFromWFI*

```
#define RISC_V_RESUME_FROM_WFI_FN(_NAME) Bool _NAME(riscvP riscv)
typedef RISC_V_RESUME_FROM_WFI_FN(*riscvResumeFromWFIFn);

typedef struct riscvModelCBS {
    riscvResumeFromWFIFn    resumeFromWFI;
} riscvModelCB;
```

Description

This interface function returns `True` if the base model is in a state where it should resume from a stalled state because of a `wfi` instruction (or an instruction with similar semantics). This is the case if there are pending interrupts that are locally enabled even if they are globally disabled.

Example

This example shows implementation of a wait-for-event (`wfe`) instruction with the same semantics as `wfi` for a processor with M-mode and U-mode.

```
static void doWFE(riscvP riscv) {
    riscvMode mode = getCurrentMode5(riscv);

    if((mode!=RISC_V_MODE_M) && RD_CSR_FIELDC(riscv, mstatus, TW)) {

        // WFE not available when mstatus.TW=1
        riscv->cb.illegalVerbose(riscv, "mstatus.TW=1");

    } else if(!riscv->cb.resumeFromWFI(riscv)) {

        // stall for custom reason (WFE) with WFI wake semantics
        riscv->cb.halt(riscv, RVD_CUSTOM_WFI);
    }
}
```

Usage Context

Leaf level only.

14.24 *Function illegalInstruction*

```
#define RISC_V_ILLEGAL_INSTRUCTION_FN(_NAME) void _NAME(riscvP riscv)
typedef RISC_V_ILLEGAL_INSTRUCTION_FN((*riscvIllegalInstructionFn));

typedef struct riscvModelCBS {
    riscvIllegalInstructionFn illegalInstruction;
} riscvModelCB;
```

Description

This interface function causes the base model to take an Illegal Instruction trap immediately.

Example

```
static void illegalInstruction(riscvP riscv) {
    riscv->cb.illegalInstruction(riscv);
}

static Bool accessUserCCTL(vmiosObjectP object, riscvP riscv) {

    if(riscv->artifactAccess) {

        // all artifact accesses are allowed
        return True;

    } else if(getCurrentMode3(riscv)==RISC_V_MODE_MACHINE) {

        // access always possible in Machine mode
        return True;

    } else if(RD_XCSR_FIELD(object, mcache_ctl, CCTL_SUEN)) {

        // access possible in Supervisor and User mode if mcache_ctl.CCTL_SUEN=1
        return True;

    } else {

        // take Illegal Instruction exception
        illegalInstruction(riscv);
        return False;

    }
}
```

Usage Context

Leaf level only.

14.25 *Function illegalVerbose*

```
#define RISC_V_ILLEGAL_VERBOSE_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    const char *reason \
)
typedef RISC_V_ILLEGAL_VERBOSE_FN((*riscvIllegalVerboseFn));

typedef struct riscvModelCBS {
    riscvIllegalVerboseFn    illegalVerbose;
} riscvModelCB;
```

Description

This interface function causes the base model to take an Illegal Instruction trap immediately, for a verbose reason indicated by the `reason` string.

Example

```
static void doWFE(riscvP riscv) {
    riscvMode mode = getCurrentMode5(riscv);

    if((mode!=RISC_V_MODE_M) && RD_CSR_FIELDC(riscv, mstatus, TW)) {
        // WFE not available when mstatus.TW=1
        riscv->cb.illegalVerbose(riscv, "mstatus.TW=1");
    } else if(!riscv->cb.resumeFromWFI(riscv)) {
        // stall for custom reason (WFE) with WFI wake semantics
        riscv->cb.halt(riscv, RVD_CUSTOM_WFI);
    }
}
```

Usage Context

Leaf level only.

14.26 *Function* `virtualInstruction`

```
#define RISC_V_ILLEGAL_INSTRUCTION_FN(_NAME) void _NAME(riscvP riscv)
typedef RISC_V_ILLEGAL_INSTRUCTION_FN(*riscvIllegalInstructionFn);

typedef struct riscvModelCBS {
    riscvIllegalInstructionFn virtualInstruction;
} riscvModelCB;
```

Description

This interface function causes the base model to take a Virtual Instruction trap immediately. It should be used only on processors that implement the Hypervisor extension.

Example

```
static void virtualInstruction(riscvP riscv) {
    riscv->cb.virtualInstruction(riscv);
}
```

Usage Context

Leaf level only.

14.27 *Function* ***virtualVerbose***

```
#define RISC_V_ILLEGAL_VERBOSE_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    const char *reason \
)
typedef RISC_V_ILLEGAL_VERBOSE_FN((*riscvIllegalVerboseFn));

typedef struct riscvModelCBS {
    riscvIllegalVerboseFn    virtualVerbose;
} riscvModelCB;
```

Description

This interface function causes the base model to take a Virtual Instruction trap immediately, for a verbose reason indicated by the `reason` string. It should be used only on processors that implement the Hypervisor extension.

Example

```
static void virtualVerbose(riscvP riscv) {
    riscv->cb.virtualVerbose(riscv);
}
```

Usage Context

Leaf level only.

14.28 *Function `illegalCustom`*

```
#define RISC_V_ILLEGAL_CUSTOM_FN(_NAME) void _NAME( \
    riscvP      riscv,      \
    riscvException exception, \
    const char   *reason     \
)
typedef RISC_V_ILLEGAL_CUSTOM_FN((*riscvIllegalCustomFn));

typedef struct riscvModelCBS {
    riscvIllegalCustomFn    illegalCustom;
} riscvModelCB;
```

Description

This interface function causes the base model to take a custom trap immediately, for a verbose reason indicated by the `reason` string. The exception passed as the `exception` argument can be either a standard exception or any other index number corresponding to a custom exception. The `xtval` CSR is updated in the same way as for a standard Illegal Instruction exception; that is, it will either be set to zero or to the opcode of the failing instruction, depending on the value of the `tval_ii_code` and `tval_zero` configuration options.

Example

```
static void illegalCustom(
    vmiosObjectP object,
    riscvException exception,
    const char   *reason
) {
    riscvP riscv = object->riscv;
    riscv->cb.illegalCustom(riscv, exception, reason);
}
```

Usage Context

Leaf level only.

14.29 Function *takeException*

```
#define RISC_V_TAKE_EXCEPTION_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvException exception, \
    Uns64      tval \
)
typedef RISC_V_TAKE_EXCEPTION_FN((*riscvTakeExceptionFn));

typedef struct riscvModelCBS {
    riscvTakeExceptionFn    takeException;
} riscvModelCB;
```

Description

This interface function causes the base model to take an exception immediately. The exception passed as the `exception` argument can be either a standard exception or any other index number corresponding to a custom exception. The `tval` argument provides a value that is reported in the corresponding `xtval` CSR.

Example

```
void andesTakeException(riscvP riscv, andesException exception, Uns64 tval) {
    riscv->cb.takeException(riscv, exception, tval);
}

static void doHSPEException(
    riscvP      riscv,
    vmiosObjectP object,
    andesException exception
) {
    // revert SP to its value prior to instruction execution
    riscv->x[RV_REG_X_SP] = object->oldSP;

    // clear stack protection enables
    WR_XCSR_FIELD(object, mhsp_ctl, OVF_EN, 0);
    WR_XCSR_FIELD(object, mhsp_ctl, UDF_EN, 0);

    // do standard exception actions
    andesTakeException(riscv, exception, 0);
}
```

Usage Context

Leaf level only.

14.30 *Function `pendFetchException`*

```
#define RISCVPEND_FETCH_EXCEPTION_FN(_NAME) void _NAME( \
    riscvp      riscv, \
    riscvException exception \
)
typedef RISCVPEND_FETCH_EXCEPTION_FN(*riscvpPendFetchExceptionFn);

typedef struct riscvModelCBS {
    riscvpPendFetchExceptionFn pendFetchException;
} riscvModelCB;
```

Description

This interface function causes the base model to schedule a pending exception that will be taken the next time the model attempts an instruction fetch. The exception passed as the exception argument will typically be an index number corresponding to a custom exception. The `xtval` CSR will be filled with the program counter corresponding to the fetch address.

If the exception is scheduled before an attempt to execute an instruction, then the exception will report that instruction address as the faulting address. If the exception is scheduled before an interrupt that is taken in CLIC mode, then the exception will report the CLIC handler address as the faulting address, and hart state will be updated to indicate a hardware vector table fetch fault.

Example

```
static VMI_NET_CHANGE_FN(instructionBusFault) {

    vmiosObjectP object = userData;
    riscvp      riscv   = object->riscv;
    Bool        old     = object->instructionBusFault;
    Bool        new     = newValue&1;

    // detect rising edge
    if(!old && new) {
        riscv->cb.pendFetchException(riscv, OHW_E_InstructionBusFault);
    }

    object->instructionBusFault = new;
}
```

Usage Context

Leaf level only.

14.31 *Function takeReset*

```
#define RISC_V_TAKE_RESET_FN(_NAME) void _NAME(riscvP riscv)
typedef RISC_V_TAKE_RESET_FN((*riscvTakeResetFn));

typedef struct riscvModelCBS {
    riscvTakeResetFn      takeReset;
} riscvModelCB;
```

Description

This interface function causes the base model to take a reset immediately.

Example

```
static void takeReset(riscvP riscv) {
    riscv->cb.takeReset(riscv);
}
```

Usage Context

Leaf level only.

14.32 *Function acknowledgeCLICInt*

```
#define RISC_V_ACKNOWLEDGE_CLIC_INT_FN(_NAME) void _NAME( \
    riscvP hart, \
    Uns32 intIndex \
)
typedef RISC_V_ACKNOWLEDGE_CLIC_INT_FN((*riscvAcknowledgeCLICIntFn));
```

Description

This callback clears the pending bit of an edge triggered interrupt in an internally implemented CLIC, and updates the pending-and-enabled state of the CLIC. The CLIC specification requires that the pending bit is cleared when taking an edge triggered interrupt configured to use selective hardware vectoring (shv).

When an extension implements a custom trap location using the `getHandlerPC` extension callback it must call this function before returning `True` from the `getHandlerPC` function, when taking a shv-configured interrupt. This is required because the call to this function is bypassed in the base model when a call to `getHandlerPC` returns `True`.

Example

```
//
// Use alternate base for CLIC vectored interrupts
//
static RISC_V_GET_HANDLER_PC_FN(vendorGetHandlerPC) {
    vmiosObjectP object = clientData;
    Bool custom = False;

    if (!isInterrupt(exception)) {
        // No custom behavior unless exception is an interrupt
    } else if (!(mode & riscv_int_CLIC)) {
        // No custom behavior unless in CLIC mode
    } else if (!CLICInternal(riscv)) {
        // No custom behavior if internal CLIC not being modeled
    } else if (!riscv->clic.sel.shv) {
        // No custom behavior if not a vectored interrupt
    } else {
        Addr tableAddr = object->local_mtv + (ecode*4);
        memEndian endian = riscv->cb.getDataEndian(riscv, mode);
        memDomainP domain = getCodeDomain(riscv);
        memAccessAttrs memAttrs = MEM_AA_TRUE|MEM_AA_FETCH;

        // read 4-byte table entry
        *handlerPCP = vmirtRead4ByteDomain(domain, tableAddr, endian, memAttrs);

        // SHV interrupts must be acknowledged when interrupt taken
        riscv->cb.acknowledgeCLICInt(riscv, ecode);

        // Use custom handler PC
        custom = True;
    }

    return custom;
}
```

Usage Context

Leaf level only.

14.33 *Function `fetchInstruction`*

```
#define RISC_V_FETCH_INSTRUCTION_FN(_NAME) Uns32 _NAME( \
    riscvP          riscv, \
    riscvAddr        thisPC, \
    riscvExtInstrInfoP info, \
    vmidDecodeTablePP tableP, \
    riscvExtInstrAttrsCP attrs, \
    Uns32            last, \
    Uns32            bits \
)
typedef RISC_V_FETCH_INSTRUCTION_FN((*riscvFetchInstructionFn));

typedef struct riscvModelCBS {
    riscvFetchInstructionFn  fetchInstruction;
} riscvModelCB;
```

Description

This interface function is used to decode a RISC-V instruction that conforms to a standard pattern, extracting information such as registers and constant values into the given `riscvExtMorphState` structure. See section 6 for a detailed example, and section 13 for more information about the available instruction patterns.

Example

```
static riscvExtIType decode(
    riscvP          riscv,
    vmiosObjectP     object,
    riscvAddr        thisPC,
    riscvExtInstrInfoP info
) {
    return riscv->cb.fetchInstruction(
        riscv, thisPC, info, &object->decode32, attrsArray32, EXT_IT_LAST, 32
    );
}
```

Usage Context

Leaf level only.

14.34 *Function `disassInstruction`*

```
#define RISC_V_DISASS_INSTRUCTION_FN(_NAME) const char *_NAME( \
    riscvP          riscv,          \
    riscvExtInstrInfoP instrInfo,    \
    vmiDisassAttrs   attrs          \
)
typedef RISC_V_DISASS_INSTRUCTION_FN((*riscvDisassInstructionFn));

typedef struct riscvModelCBS {
    riscvDisassInstructionFn  disassInstruction;
} riscvModelCB;
```

Description

This interface function is used to disassemble a RISC-V instruction using a standard format string, using information about registers and constant values previously extracted by interface function `fetchInstruction`. See section 6 for a detailed example, and section 13 for more information about the available instruction patterns.

Example

```
static VMIO_DISASSEMBLE_FN(addInstructionsDisassemble) {

    riscvP          riscv = (riscvP)processor;
    const char      *result = 0;
    riscvExtInstrInfo info;

    // action is only required if the instruction is implemented by this
    // extension
    if(decode(riscv, object, thisPC, &info) != EXT_IT_LAST) {
        result = riscv->cb.disassInstruction(riscv, &info, attrs);
    }

    return result;
}
```

Usage Context

Leaf level only.

14.35 *Function `instructionEnabled`*

```
#define RISC_V_INSTRUCTION_ENABLED_FN(_NAME) Bool _NAME( \
    riscvP      riscv, \
    riscvArchitecture requiredVariant \
)
typedef RISC_V_INSTRUCTION_ENABLED_FN((*riscvInstructionEnabledFn));

typedef struct riscvModelCBS {
    riscvInstructionEnabledFn instructionEnabled;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

This interface function is used to validate the legality of an instruction with respect to RISC-V feature letters (A-Z). The feature requirements for this instruction are passed as the `requiredVariant` argument; for example, if an instruction requires single-precision floating point to be enabled, the value `ISA_F` should be passed. Any required XLEN can also be specified; for example, `ISA_F|RV64` encodes a requirement for enabled floating point *and* XLEN of 64.

The function returns a Boolean indicating if the architectural constraint is satisfied. If it is not satisfied, the interface function emits code to cause an Illegal Instruction trap to be taken.

Example

```
static Bool instructionEnabled(riscvP riscv, riscvArchitecture riscvVariants) {
    return riscv->cb.instructionEnabled(riscv, riscvVariants);
}
```

Usage Context

Leaf level only, morph-time.

14.36 *Function morphExternal*

```
#define RISCVMORPH_EXTERNAL_FN(_NAME) void _NAME( \
    riscvExtMorphStateP state, \
    const char *disableReason, \
    Bool *opaque \
);
typedef RISCVMORPH_EXTERNAL_FN((*riscvMorphExternalFn));

typedef struct riscvModelCBS {
    riscvMorphExternalFn morphExternal;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object morphCB).

This interface function is used to emit translated code for an instruction implemented by an extension object, assuming that the instruction details have been decoded into a structure of type `riscvExtMorphState` by a previous call to the `fetchInstruction` interface function. The `opaque` argument should be passed down from the `morphCB` parameter of the same name. If the `disableReason` argument is non-NULL, code to cause an Illegal Instruction trap will be emitted and this reason printed in verbose mode. If the `disableReason` argument is NULL, the function `state.attrs->morph` (defined by the extension object) will be called to emit translated code for the instruction.

See section 6 for a detailed example.

Example

```
static VMIO_MORPH_FN(addInstructionsMorph) {
    riscvP riscv = (riscvP)processor;
    riscvExtMorphState state = {riscv:riscv, object:object};

    // decode instruction
    riscvExtIType type = decode(riscv, object, thisPC, &state.info);

    // action is only required if the instruction is implemented by this
    // extension
    if(type != EXT_IT_LAST) {
        // fill translation attributes
        state.attrs = &dispatchTable[type];

        // translate instruction
        riscv->cb.morphExternal(&state, 0, opaque);
    }

    // no callback function is required
    return 0;
}
```

Usage Context

Leaf level only, morph-time.

14.37 *Function `morphIllegal`*

```
#define RISC_V_MORPH_ILLEGAL_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    const char *reason \
)
typedef RISC_V_MORPH_ILLEGAL_FN((*riscvMorphIllegalFn));

typedef struct riscvModelCBS {
    riscvMorphIllegalFn      morphIllegal;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

This interface function is used to emit code to cause an Illegal Instruction exception, printing the given reason string in verbose mode.

Example

```
static void morphIllegal(riscvP riscv, const char *reason) {
    return riscv->cb.morphIllegal(riscv, reason);
}
```

Usage Context

Leaf level only, morph-time.

14.38 *Function morphVirtual*

```
#define RISC_V_MORPH_ILLEGAL_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    const char *reason \
)
typedef RISC_V_MORPH_ILLEGAL_FN((*riscvMorphIllegalFn));

typedef struct riscvModelCBS {
    riscvMorphIllegalFn      morphVirtual;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object morphCB).

This interface function is used to emit code to cause a Virtual Instruction exception, printing the given reason string in verbose mode. It should be used only on processors that implement the Hypervisor extension.

Example

```
static void morphVirtual(riscvP riscv, const char *reason) {
    return riscv->cb.morphVirtual(riscv, reason);
}
```

Usage Context

Leaf level only, morph-time.

14.39 *Function getVMIReg*

```
#define RISC_V_GET_VMI_REG_FN(_NAME) vmiReg _NAME(riscvP riscv, riscvRegDesc r)
typedef RISC_V_GET_VMI_REG_FN(*riscvGetVMIRegFn);

typedef struct riscvModelCBS {
    riscvGetVMIRegFn      getVMIReg;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

This interface function is used to convert a RISC-V register description (of type `riscvRegDesc`) to a VMI register description. The RISC-V register description will typically be extracted from a structure of type `riscvExtMorphState` filled by a previous call to the `fetchInstruction` interface function.

The VMI register description can then be used to specify instruction functional details using the VMI Morph Time Function API. See section 6 for a detailed example.

Example

```
inline static vmiReg getVMIReg(riscvP riscv, riscvRegDesc r) {
    return riscv->cb.getVMIReg(riscv, r);
}
```

Usage Context

Leaf level only, morph-time.

14.40 Function *getVMIRegFS*

```
#define RISC_V_GET_VMI_REG_FS_FN(_NAME) vmiReg _NAME( \
    riscvP      riscv,      \
    riscvRegDesc r,        \
    vmiReg      tmp        \
)
typedef RISC_V_GET_VMI_REG_FS_FN((*riscvGetVMIRegFSFn));

typedef struct riscvModelCBS {
    riscvGetVMIRegFSFn    getVMIRegFS;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object morphCB).

This interface function is used to convert a RISC-V register description (of type `riscvRegDesc`) to a VMI register description. The RISC-V register description will typically be extracted from a structure of type `riscvExtMorphState` filled by a previous call to the `fetchInstruction` interface function. The function must be used when the register being converted is potentially an FPR that is *narrower than FLEN* (for example, a single-precision source on a machine that supports double-precision). The function emits code to validate that the source value is correctly NaN-boxed, composing the resultant value in the `tmp` temporary, which is then returned. If the register does not require a NaN box test, a `vmiReg` object representing the register value in the processor structure is returned.

The VMI register description can then be used to specify instruction functional details using the VMI Morph Time Function API.

Example

```
inline static vmiReg getVMIRegFS(riscvP riscv, riscvRegDesc r, vmiReg tmp) {
    return riscv->cb.getVMIRegFS(riscv, r, tmp);
}
```

Usage Context

Leaf level only, morph-time.

14.41 *Function* `writeRegSize`

```
#define RISC_V_WRITE_REG_SIZE_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvRegDesc r, \
    Uns32       srcBits, \
    Bool        signExtend \
)
typedef RISC_V_WRITE_REG_SIZE_FN((*riscvWriteRegSizeFn));

typedef struct riscvModelCBS {
    riscvWriteRegSizeFn writeRegSize;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

When a result of size `srcBits` has been written into the VMI register equivalent to register `r`, this function is called to handle any required extension of that value from size `srcBits` to the architectural width of register `r`. Argument `signExtend` indicates whether the value is sign-extended (if `True`) or zero-extended (if `False`). See section 6 for a detailed example.

Example

```
inline static void writeRegSize(
    riscvP      riscv,
    riscvRegDesc r,
    Uns32       srcBits,
    Bool        signExtend
) {
    riscv->cb.writeRegSize(riscv, r, srcBits, signExtend);
}
```

Usage Context

Leaf level only, morph-time.

14.42 *Function writeReg*

```
#define RISC_V_WRITE_REG_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvRegDesc r, \
    Bool        signExtend \
)
typedef RISC_V_WRITE_REG_FN(*riscvWriteRegFn);

typedef struct riscvModelCBS {
    riscvWriteRegFn writeReg;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object morphCB).

When a result of the bit size encoded in the description of register `r` has been written into the VMI register equivalent to that register, this function is called to handle any required extension of that value from the encoded register size to the architectural width of register `r`. Argument `signExtend` indicates whether the value is sign-extended (if `True`) or zero-extended (if `False`).

This is equivalent to:

```
riscv->cb.writeRegSize(riscv, r, getRBits(r), signExtend);
```

Example

```
inline static void writeReg(riscvP riscv, riscvRegDesc r) {
    riscv->cb.writeReg(riscv, r, True);
}
```

Usage Context

Leaf level only, morph-time.

14.43 *Function `getFPFlagsMt`*

```
#define RISC_V_GET_FP_FLAGS_MT_FN(_NAME) vmiReg _NAME(riscvP riscv)
typedef RISC_V_GET_FP_FLAGS_MT_FN(*riscvGetFPFlagsMtFn);

typedef struct riscvModelCBS {
    riscvGetFPFlagsMtFn      getFPFlagsMt;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

When an extension object implements instructions that update floating point flag state, this function must be used to obtain a `vmiReg` descriptor for the standard RISC-V `fflags` CSR. The returned `vmiReg` register should then be used as the `flags` argument of any floating point VMI primitives used to implement that instruction.

Example

```
inline static vmiReg riscvGetFPFlagsMT(riscvP riscv) {
    return riscv->cb.getFPFlagsMt(riscv);
}
```

Usage Context

Leaf level only, morph-time.

14.44 *Function `getDataEndianMt`*

```
#define RISC_V_GET_DATA_ENDIAN_MT_FN(_NAME) memEndian _NAME(riscvP riscv)
typedef RISC_V_GET_DATA_ENDIAN_MT_FN((*riscvGetDataEndianMtFn));

typedef struct riscvModelCBS {
    riscvGetDataEndianMtFn    getDataEndianMt;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

This interface function returns the active endianness for loads and stores when translating an instruction. The function ensures that any JIT-compiled code generated is used only when the endianness matches, meaning that the returned endianness can be assumed to be constant by the JIT translation routine in the extension object.

Example

```
inline static memEndian getDataEndian(riscvP riscv) {
    return riscv->cb.getDataEndianMt(riscv);
}
```

Usage Context

Leaf level only, morph-time.

14.45 Function *loadMT*

```
#define RISC_V_LOAD_MT_FN(_NAME) void _NAME( \
    riscvP          riscv, \
    vmiReg          rd, \
    Uns32           rdBits, \
    vmiReg          ra, \
    Uns32           memBits, \
    Uns64           offset, \
    riscvExtLdStAttrs attrs \
)
typedef RISC_V_LOAD_MT_FN((*riscvLoadMtFn))

typedef struct riscvModelCBS {
    riscvLoadMtFn      loadMt;
} riscvModelCB;
```

Description

*This function must be called at morph time (from within the extension object *morphCB*).*

This interface function emits JIT code to perform a load of *memBits* wide data from the address *ra+offset*. The loaded value is extended to *rdBits* wide and written to register *rd*. Other attributes of the load are defined by the *attrs* parameter, which is defined in *riscvModelCallbackTypes.h* as follows:

```
typedef struct riscvExtLdStAttrsS {
    memConstraint constraint : 4; // access constraints
    Bool          sExtend    : 1; // sign-extension (load only)
    Bool          isVirtual   : 1; // whether HLV/HLVX/HSV
    Bool          isCode      : 1; // whether load as if fetch (HLVX)
} riscvExtLdStAttrs;
```

Fields in this structure are as follows:

constraint: this bitfield enumeration specifies constraints for the memory access. For the RISC-V model, these values are significant:

- MEM_CONSTRAINT_ALIGNED: whether the access must be aligned;
- MEM_CONSTRAINT_USER1: whether the access is atomic.

sExtend: this Boolean value indicates whether the loaded value must be sign-extended to *rdBits* width, if it is smaller. If *False*, the value is zero-extended.

isVirtual: this Boolean value indicates whether the load is a result of an instruction like *HLV* or *HLVX*, requiring access to guest virtual address space (only ever *True* when Hypervisor mode is implemented).

isCode: this Boolean value indicates whether the load is a result of an instruction like *HLVX*, where the load should be treated as if it was a fetch.

*When this function is used, any Trigger Module triggers sensitive to the specified access will fire if required, and the load will also comply with any transactional memory model installed by the extension - this will not be the case if more fundamental VMI primitives such as *vmimtLoadRRO* are used instead.*

Example

```
static void emitLoadCommon(
    andesMorphStateP state,
    vmiReg            rd,
    Uns32             rdBits,
    vmiReg            ra,
    memConstraint     constraint
) {
    riscvP riscv = state->riscv;
    Uns32 memBits = state->info.memBits;
    Uns64 offset = state->info.cl;

    riscvExtLdStAttrs attrs = {
        constraint : constraint,
        sExtend    : !state->info.unsExt
    };

    riscv->cb.loadMt(riscv, rd, rdBits, ra, memBits, offset, attrs);
}
```

Usage Context

Leaf level only, morph-time.

14.46 Function *storeMT*

```
#define RISC_V_STORE_MT_FN(_NAME) void _NAME( \
    riscvP          riscv, \
    vmiReg          rs,    \
    vmiReg          ra,    \
    Uns32           memBits, \
    Uns64           offset, \
    riscvExtLdStAttrs attrs \
)
typedef RISC_V_STORE_MT_FN((*riscvStoreMtFn))

typedef struct riscvModelCBS {
    riscvStoreMtFn      storeMt;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

This interface function emits JIT code to perform a store of `memBits` wide data to address `ra+offset`. The stored value is `memBits` wide and sourced from register `rs`. Other attributes of the load are defined by the `attrs` parameter, which is defined in `riscvModelCallbackTypes.h` as follows:

```
typedef struct riscvExtLdStAttrsS {
    memConstraint constraint : 4; // access constraints
    Bool          sExtend   : 1; // sign-extension (load only)
    Bool          isVirtual  : 1; // whether HLV/HLVX/HSV
    Bool          isCode     : 1; // whether load as if fetch (HLVX)
} riscvExtLdStAttrs;
```

Relevant fields in this structure for a store are as follows:

constraint: this bitfield enumeration specifies constraints for the memory access. For the RISC-V model, these values are significant:

`MEM_CONSTRAINT_ALIGNED`: whether the access must be aligned;

`MEM_CONSTRAINT_USER1`: whether the access is atomic.

isVirtual: this Boolean value indicates whether the store is a result of an instruction like `HSV`, requiring access to guest virtual address space (only ever `True` when Hypervisor mode is implemented).

When this function is used, any Trigger Module triggers sensitive to the specified access will fire if required, and the store will also comply with any transactional memory model installed by the extension - this will not be the case if more fundamental VMI primitives such as `vmimtStoreRRO` are used instead.

Example

```
static void emitStoreCommon(
    andesMorphStateP state,
    vmiReg          rs,
    vmiReg          ra,
    memConstraint    constraint
```

```
) {  
    riscvP riscv    = state->riscv;  
    Uns32 memBits = state->info.memBits;  
    Uns64 offset  = state->info.cl;  
  
    riscvExtLdStAttrs attrs = {constraint : constraint};  
  
    riscv->cb.storeMt(riscv, rs, ra, memBits, offset, attrs);  
}
```

Usage Context

Leaf level only, morph-time.

14.47 *Function checkLoadMT*

```
#define RISCVCHECKMEMMT_FN(_NAME) void _NAME( \
    riscvP riscv, \
    vmiReg base, \
    Int32 offset, \
    Uns32 bits, \
    void *clientData \
)
typedef RISCVCHECKMEMMT_FN((*riscvCheckMemMtFn));

typedef struct riscvModelCBS {
    riscvCheckMemMtFn checkLoadMt;
} riscvModelCB;
```

Description

This interface function is called by the base model when an instruction that loads from a location in memory is being translated. It gives the extended model the opportunity to emit JIT code to apply extra checks before the load is performed. This function can be used to implement extended features such as hardware stack protection where memory accesses using the stack pointer register (x2) as a base are bounds-checked.

The callback is passed the `vmiReg` object corresponding to the *base* register being used for the load and the offset from the base that must be added to form the full load address. It is also given the data width of the load, in bits. It will typically emit an embedded call to perform a run-time check on address validity.

Example

The following example shows a template for a hardware stack protection check where any access using the stack pointer (x2) is range-checked.

```
//
// Implement hardware stack protection check
//
static void doHSPCheckLimit(
    riscvP riscv,
    vmiosObjectP object,
    Int32 offset,
    Uns32 bytes,
    Uns64 loLimit,
    riscvException exception
) {
    Uns64 lo = riscv->x[RV_REG_X_SP] + offset;
    Uns64 hi = lo + bytes - 1;

    if(
        // highest byte address must be lower than the stack base address
        (hi < RD_XCSR(object, mspcba)) &&
        // lowest byte address equal to of higher than low limit
        (lo >= loLimit)
    ) {
        // valid stack access
    } else {
        // do standard exception actions
        takeException(riscv, exception, lo);
    }
}
```

```
//  
// Implement hardware stack protection check for load (low limit is SP)  
//  
static void doHSPCheckLoad(  
    riscvP      riscv,  
    vmiosObjectP object,  
    Int32       offset,  
    Uns32       bytes  
) {  
    doHSPCheckLimit(  
        riscv,  
        object,  
        offset,  
        bytes,  
        riscv->x[RV_REG_X_SP],  
        riscv_E_LoadAccessFault  
    );  
}  
  
//  
// Should HSP check be applied to the current instruction?  
//  
static Bool applyHSPMT(riscvP riscv, vmiosObjectP object, vmiReg base) {  
    return (  
        // HSP feature must be enabled  
        isHSPEnabledMT(object) &&  
        // base register must be the stack pointer  
        VMI_REG_EQUAL(base, RISCV_SP)  
    );  
}  
  
//  
// Emit code to check hardware stack protection on load if required  
//  
static RISCVCHECKMEMMTFN(emitCheckLoadHSP) {  
    vmiosObjectP object = clientData;  
  
    if(applyHSPMT(riscv, object, base)) {  
        vmimtArgProcessor();  
        vmimtArgNatAddress(object);  
        vmimtArgUns32(offset);  
        vmimtArgUns32(BITS_TO_BYTES(bits));  
        vmimtCallAttrs((vmiCallFn)doHSPCheckLoad, VMCA_NA);  
    }  
}
```

Usage Context

Leaf level only.

14.48 *Function checkStoreMt*

```
#define RISCVCHECKMEMMT_FN(_NAME) void _NAME( \
    riscvP riscv, \
    vmiReg base, \
    Int32 offset, \
    Uns32 bits, \
    void *clientData \
)
typedef RISCVCHECKMEMMT_FN((*riscvCheckMemMtFn));

typedef struct riscvModelCBS {
    riscvCheckMemMtFn checkStoreMt;
} riscvModelCB;
```

Description

This interface function is called by the base model when an instruction that stores to a location in memory is being translated. It gives the extended model the opportunity to emit JIT code to apply extra checks before the store is performed. This function can be used to implement extended features such as hardware stack protection where memory accesses using the stack pointer register (x2) as a base are bounds-checked.

The callback is passed the `vmiReg` object corresponding to the *base* register being used for the store and the offset from the base that must be added to form the full store address. It is also given the data width of the store, in bits. It will typically emit an embedded call to perform a run-time check on address validity.

Example

The following example shows a template for a hardware stack protection check where any access using the stack pointer (x2) is range-checked.

```
//
// Implement hardware stack protection check
//
static void doHSPCheckLimit(
    riscvP riscv,
    vmiosObjectP object,
    Int32 offset,
    Uns32 bytes,
    Uns64 loLimit,
    riscvException exception
) {
    Uns64 lo = riscv->x[RV_REG_X_SP] + offset;
    Uns64 hi = lo + bytes - 1;

    if(
        // highest byte address must be lower than the stack base address
        (hi < RD_XCSR(object, mspcba)) &&
        // lowest byte address equal to or higher than low limit
        (lo >= loLimit)
    ) {
        // valid stack access
    } else {
        // do standard exception actions
        takeException(riscv, exception, lo);
    }
}
```

```
//  
// Implement hardware stack protection check for store (low limit is mspcta)  
//  
static void doHSPCheckStore(  
    riscvP      riscv,  
    vmiosObjectP object,  
    Int32       offset,  
    Uns32       bytes  
) {  
    doHSPCheckLimit(  
        riscv,  
        object,  
        offset,  
        bytes,  
        RD_XCSR(object, mspcta),  
        riscv_E_StoreAMOAccessFault  
    );  
}  
  
//  
// Should HSP check be applied to the current instruction?  
//  
static Bool applyHSPMT(riscvP riscv, vmiosObjectP object, vmiReg base) {  
    return (  
        // HSP feature must be enabled  
        isHSPEnabledMT(object) &&  
        // base register must be the stack pointer  
        VMI_REG_EQUAL(base, RISCV_SP)  
    );  
}  
  
//  
// Emit code to check hardware stack protection on store if required  
//  
static RISCV_CHECK_MEM_MT_FN(emitCheckStoreHSP) {  
    vmiosObjectP object = clientData;  
  
    if(applyHSPMT(riscv, object, base)) {  
        vmimtArgProcessor();  
        vmimtArgNatAddress(object);  
        vmimtArgUns32(offset);  
        vmimtArgUns32(BITS_TO_BYTES(bits));  
        vmimtCallAttrs((vmiCallFn)doHSPCheckStore, VMCA_NA);  
    }  
}
```

Usage Context

Leaf level only.

14.49 *Function requireModeMt*

```
#define RISC_V_REQUIRE_MODE_MT_FN(_NAME) Bool _NAME( \
    riscvP    riscv,    \
    riscvMode mode      \
)
typedef RISC_V_REQUIRE_MODE_MT_FN((*riscvRequireModeMtFn));
```

Description

This function must be called at morph time (from within the extension object morphCB).

This interface function validates that the current processor mode is at least the given mode and emits code to generate either an Illegal Instruction or Virtual Instruction exception if not. If the required mode is Machine mode, or the processor is not currently in Virtual Supervisor or Virtual User mode, then an Illegal Instruction exception is taken; otherwise, a Virtual Instruction exception is taken.

The return value is `True` if the processor is executing in a sufficiently high privilege mode and `False` if not.

Example

```
inline static Bool requireModeMT(riscvP riscv, riscvMode required) {
    return riscv->cb.requireModeMt(riscv, required);
}
```

Usage Context

Leaf level only, morph-time.

14.50 *Function requireNotVMt*

```
#define RISC_V_REQUIRE_NOT_V_MT_FN(_NAME) Bool _NAME(riscvP riscv)
typedef RISC_V_REQUIRE_NOT_V_MT_FN(*riscvRequireNotVMtFn);
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

This interface function validates that the current processor mode is not a virtual mode; if it is, code to take a Virtual Instruction exception is emitted.

The return value is `True` if the processor is executing in a non-virtual mode and `False` if not.

Example

```
inline static Bool requireNonVirtual(riscvP riscv) {
    return riscv->cb.requireNotVMt(riscv);
}
```

Usage Context

Leaf level only, morph-time.

14.51 Function *checkLegalRMMt*

```
#define RISC_V_CHECK_LEGAL_RM_MT_FN(_NAME) Bool _NAME( \
    riscvP      riscv, \
    riscvRMDesc rm      \
)
typedef RISC_V_CHECK_LEGAL_RM_MT_FN((*riscvCheckLegalRMMtFn));

typedef struct riscvModelCBS {
    riscvCheckLegalRMMtFn    checkLegalRMMt;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

Given a rounding mode of type `riscvRMDesc` (typically extracted from a structure of type `riscvExtMorphState` filled by a previous call to the `fetchInstruction` interface function), this interface function inserts code to check legality of the rounding mode and take an Illegal Instruction trap if it is invalid. The Boolean return code indicates whether the rounding mode should be assumed to be valid by the calling extension object function.

Example

```
inline static Bool emitCheckLegalRM(riscvP riscv, riscvRMDesc rm) {
    return riscv->cb.checkLegalRMMt(riscv, rm);
}
```

Usage Context

Leaf level only, morph-time.

14.52 *Function morphTrapTVM*

```
#define RISC_V_MORPH_TRAP_TVM_FN(_NAME) void _NAME(riscvP riscv)
typedef RISC_V_MORPH_TRAP_TVM_FN((*riscvMorphTrapTVMFn));

typedef struct riscvModelCBS {
    riscvMorphTrapTVMFn      morphTrapTVM;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object `morphCB`).

This interface function is used to emit a trap when `mstatus.TVM=1` when executing in Supervisor mode.

Example

```
static void morphTrapTVM(riscvP riscv) {
    return riscv->cb.morphTrapTVM(riscv);
}
```

Usage Context

Leaf level only, morph-time.

14.53 Function *morphVOp*

```
#define RISC_V_MORPH_VOP_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    Uns64       thisPC, \
    riscvRegDesc r0,    \
    riscvRegDesc r1,    \
    riscvRegDesc r2,    \
    riscvRegDesc mask,  \
    riscvVShape  shape,  \
    riscvVExternalFn externalCB, \
    void         *userData \
)
typedef RISC_V_MORPH_VOP_FN((*riscvMorphVOpFn));

typedef struct riscvModelCBS {
    riscvMorphVOpFn      morphVOp;
} riscvModelCB;
```

Description

This function must be called at morph time (from within the extension object morphCB).

This interface function is used to create a *vector extension custom operation* in an extension object, according to the standard patterns implemented in the base model. The arguments are as follows:

1. `riscv`: the current processor.
2. `thisPC`: the current instruction address.
3. `r0`: first register operand.
4. `r1`: second register operand.
5. `r3`: third register operand.
6. `mask`: vector mask.
7. `shape`: vector operation shape (of type `riscvVShape`, described below).
8. `externalCB`: function of type `riscvVExternalFn`, implementing one operation.
9. `userData`: extension-specific context information.

Given an instruction previously decoded into a structure of type `riscvExtMorphState` filled by a previous call to the `fetchInstruction` interface function, most of these parameters can be extracted directly from fields in this structure, as shown in the example at the end of this section.

The `shape` argument describes the pattern of vector operation, as follows:

```
//
// Vector shape description. Each name is composed of:
// - prefix (RVVW)
// - two or three operand descriptors, each three letters:
//   1. argument type (V=vector, S=scalar, P=predicate);
//   2. VLMUL multiplier (1, 2 or 4);
//   3. element type (I=integer, F=float)
// - an optional generic suffix
//
typedef enum riscvVShapeE {

                                // INTEGER ARGUMENTS
    RVVW_V1I_V1I_V1I,          // SEW = SEW op SEW
```

```

RVVW_V1I_V1I_V1I_LD,    // vector load operations
RVVW_V1I_V1I_V1I_ST,    // vector store operations
RVVW_V1I_V1I_V1I_SAT,   // saturating result
RVVW_V1I_V1I_V1I_VXRM,  // uses vxrm
RVVW_V1I_S1I_V1I,       // src1 is scalar
RVVW_S1I_V1I_V1I,       // Vd is scalar
RVVW_P1I_V1I_V1I,       // Vd is predicate
RVVW_S1I_V1I_S1I,       // Vd and src2 are scalar
RVVW_V1I_V1I_V1I_CIN,   // mask is carry-in (VADC etc)
RVVW_P1I_V1I_V1I_CIN,   // Vd is predicate, mask is carry-in (VMADC etc)
RVVW_S2I_V1I_S2I,       // 2*SEW = SEW op 2*SEW, Vd and src2 are scalar
RVVW_V1I_V2I_V1I,       // SEW = 2*SEW op SEW
RVVW_V1I_V2I_V1I_SAT,   // SEW = 2*SEW op SEW, saturating result
RVVW_V2I_V1I_V1I_IW,    // 2*SEW = SEW op SEW, implicit widening
RVVW_V2I_V1I_V1I,       // 2*SEW = SEW op SEW
RVVW_V2I_V1I_V1I_SAT,   // 2*SEW = SEW op SEW, saturating result
RVVW_V4I_V1I_V1I,       // 4*SEW = SEW op SEW
RVVW_V2I_V2I_V1I,       // 2*SEW = 2*SEW op SEW
RVVW_V1I_V2I_FN,        // SEW = SEW/FN

                                // FLOATING POINT ARGUMENTS
RVVW_V1F_V1F_V1F,       // SEW = SEW op SEW
RVVW_V1F_S1F_V1F,       // src1 is scalar
RVVW_S1F_V1I_V1I,       // Vd is scalar
RVVW_P1I_V1F_V1F,       // Vd is predicate
RVVW_S1F_V1F_S1F,       // Vd and src2 are scalar
RVVW_S2F_V1F_S2F,       // 2*SEW = SEW op 2*SEW, Vd and src2 are scalar
RVVW_V1F_V2F_V1F_IW,    // SEW = 2*SEW op SEW, implicit widening
RVVW_V2F_V1F_V1F_IW,    // 2*SEW = SEW op SEW, implicit widening
RVVW_V2F_V1F_V1F,       // 2*SEW = SEW op SEW
RVVW_V2F_V2F_V1F,       // 2*SEW = 2*SEW op SEW
RVVW_V1F_V1F_V1I_UP,    // SEW, VFSLIDEUP instructions
RVVW_V1F_V1F_V1I_DN,    // SEW, VFSLIDEDOWN instructions

                                // CONVERSIONS
RVVW_V1F_V1I,           // SEW = SEW, Fd=Is
RVVW_V1I_V1F,           // SEW = SEW, Id=Fs
RVVW_V2F_V1I,           // 2*SEW = SEW, Fd=Is
RVVW_V2I_V1F,           // 2*SEW = SEW, Id=Fs
RVVW_V1F_V2I_IW,        // SEW = 2*SEW, Fd=Is, implicit widening
RVVW_V1I_V2F_IW,        // SEW = 2*SEW, Id=Fs, implicit widening

                                // MASK ARGUMENTS
RVVW_P1I_P1I_P1I,       // SEW = SEW op SEW
RVVW_P1I_P1I_P1I_VMSF,  // SEW = SEW op SEW, VMSBF/VMSOF/VMSIF instructions
RVVW_V1I_P1I_P1I_IOTA,  // SEW = SEW op SEW, VIOTA instruction
RVVW_V1I_P1I_P1I_ID,    // SEW = SEW op SEW, VID instruction

                                // SLIDING ARGUMENTS
RVVW_V1I_V1I_V1I_GR,    // SEW, VRGATHER instructions
RVVW_V1I_V1I_V1I_UP,    // SEW, VSLIDEUP instructions
RVVW_V1I_V1I_V1I_DN,    // SEW, VSLIDEDOWN instructions
RVVW_V1I_V1I_V1I_CMP,   // SEW, VCOMPRESS instruction

RVVW_LAST                // KEEP LAST: for sizing
} riscvVShape;

```

These shapes correspond to all standard vector instructions. Most extension instructions are likely to be of types `RVVW_V1I_V1I_V1I` or `RVVW_V1F_V1F_V1F` (i.e. same-width binary integer and floating point operations, respectively).

The `morphVOp` interface function automatically handles standard Vector Extension features such as element iteration, masking, and emitting Illegal Instruction exceptions if the Vector Extension is disabled. The extension object is required only to supply a

callback function of type `riscvVExternalFn` to implement the vector operation for a single element:

```
#define RISC_VEXTERNAL_FN(_NAME) void _NAME( \
    riscvP   riscv,      \
    void     *userData,   \
    vmiReg   *r,          \
    Uns32    SEW          \
)
typedef RISC_VEXTERNAL_FN((*riscvVExternalFn));
```

The callback function takes four arguments:

1. The current processor;
2. The `userData` pointer originally passed as the final argument to the `morphVOp` interface function;
3. An array of `vmiReg` objects for each register argument to the element operation;
4. The current selected element width (SEW).

Example

The following code snippet shows how a simple custom widening instruction that extends a `BFLOAT16` value to a single-precision value might be implemented:

```
//
// Per-element callback for VFWCVT.S.BF16
//
static RISC_VEXTERNAL_FN(emitVFWCVT_S_BF16CB) {

    vmiReg r0L = r[0];
    vmiReg r0H = VMI_REG_DELTA(r[0], 2);

    // move result to high part of register
    vmimtMoveRR(16, r0H, r[1]);
    vmimtMoveRC(16, r0L, 0);
}

//
// Emit VFWCVT.S.BF16
//
static void emitVFWCVT_S_BF16(
    riscvP   riscv,
    riscvExtMorphStateP state
) {
    riscv->cb.morphVOp(
        riscv,
        state->thisPC,
        state->r[0],
        state->r[1],
        state->r[2],
        state->mask,
        RVVW_V2F_V1I,
        emitVFWCVT_S_BF16CB,
        state
    );
}
```

Usage Context

Leaf level only, morph-time.

14.54 *Function newCSR*

```
#define RISC_V_NEW_CSR_FN(_NAME) void _NAME( \
    riscvCSRAttrsP  attrs,          \
    riscvCSRAttrsCP src,            \
    riscvP          riscv,          \
    vmiosObjectP    object          \
)
typedef RISC_V_NEW_CSR_FN((*riscvNewCSRFn));

typedef struct riscvModelCBS {
    riscvNewCSRFn      newCSR;
} riscvModelCB;
```

Description

Given a template CSR, this function registers that CSR with the base model. It should always be called from the extension object constructor. See section 7 for a detailed description and extended example.

Example

```
static void csrInit(vmiosObjectP object) {

    riscvP  riscv = object->riscv;
    extCSRId id;

    for(id=0; id<XCSR_ID(LAST); id++) {

        extCSRAttrsCP src = &csrs[id];
        riscvCSRAttrs *dst = &object->csrs[id];

        riscv->cb.newCSR(dst, &src->baseAttrs, riscv, object);
    }
}
```

Usage Context

Leaf level only, in constructor.

14.55 *Function hpmAccessValid*

```
#define HPM_ACCESS_VALID_FN(_NAME) Bool _NAME( \
    riscvCSRAttrsCP attrs,          \
    riscvP          riscv           \
)
typedef HPM_ACCESS_VALID_FN((*riscvHPMAccessValidFn));

typedef struct riscvModelCBS {
    riscvHPMAccessValidFn    hpmAccessValid;
} riscvModelCB;
```

Description

Access to performance counter CRSs is controlled by number of other registers (mcounteren, scounteren and, if Hypervisor mode is implemented, hcounteren). This function implements the logic of that control and returns a Boolean indicating whether access to a performance counter register defined by the `attrs` argument is legal in the current processor mode. It is useful when implementing custom performance counter CRSs in a derived model.

Example

```
static RISC_V_CSR_READFN(mtimeR) {
    vmiosObjectP object = attrs->object;
    Uns64         result = 0;

    if(riscv->artifactAccess) {
        // no action
    } else if(!riscv->cb.hpmAccessValid(attrs, riscv)) {
        // invalid access, standard exception
    } else {
        customTimeException(object);
    }

    return result;
}
```

Usage Context

Leaf level only.

14.56 Function *mapAddress*

```
#define RISC_V_MAP_ADDRESS_FN(_NAME) Bool _NAME( \
    riscvP      riscv, \
    memDomainP  domain, \
    memPriv     requiredPriv, \
    Uns64       address, \
    Uns32       bytes, \
    memAccessAttrs attrs \
)
typedef RISC_V_MAP_ADDRESS_FN((*riscvMapAddressFn));

typedef struct riscvModelCBS {
    riscvMapAddressFn    mapAddress;
} riscvModelCB;
```

Description

This function can be called by a derived model to attempt to establish a memory mapping for the given address and access size (*bytes*) in the given memory domain object, which must be one of the domains corresponding to an address space for the processor. The function returns *True* if the mapping was *unsuccessful* because a virtual memory address mapping failure and *False* otherwise. The function also establishes any permission restrictions implied by PMP/PMA regions (if implemented). The *attrs* parameter controls whether a mapping failure causes the processor to take an exception.

This function is usually called from within *rdSnapCB* or *wrSnapCB* callbacks to implement alignment checks.

Example

```
static memPriv getDomainPrivileges(
    riscvP      riscv,
    memDomainP  domain,
    Uns64       address,
    memPriv     priv
) {
    memPriv mappedPriv = vmirtGetDomainPrivileges(domain, address);

    // if no privilege is set, try mapping memory and get privilege again
    if(!mappedPriv) {
        riscv->cb.mapAddress(riscv, domain, priv, address, 1, MEM_AA_FALSE);
        mappedPriv = vmirtGetDomainPrivileges(domain, address);
    }

    return mappedPriv;
}
```

Usage Context

Leaf level only.

14.57 *Function unmapPMPRegion*

```
#define RISC_V_UNMAP_PMP_REGION_FN(_NAME) void _NAME( \
    riscvP riscv, \
    Uns32  regionIndex \
)
typedef RISC_V_UNMAP_PMP_REGION_FN(*riscvUnmapPMPRegionFn);

typedef struct riscvModelCBS {
    riscvUnmapPMPRegionFn  unmapPMPRegion;
} riscvModelCB;
```

Description

This function can be called by a derived model to force the indexed PMP region to be unmapped by the base model. This may be required if the derived model enhances the default permissions applied by the base model (for example, with supplementary custom CSRs).

Example

```
static void unmapPMPRegion (riscvP riscv, Uns32 regionIndex) {
    riscv->cb.unmapPMPRegion(riscv, regionIndex);
}
```

Usage Context

Leaf level only.

14.58 *Function updateLdStDomain*

```
#define RISC_V_UPDATE_LD_ST_DOMAIN_FN(_NAME) void _NAME(riscvP riscv)
typedef RISC_V_UPDATE_LD_ST_DOMAIN_FN(*riscvUpdateLdStDomainFn);

typedef struct riscvModelCBS {
    riscvUpdateLdStDomainFn    updateLdStDomain;
} riscvModelCB;
```

Description

Some CSR settings affect the access mode for load and store instructions in Machine mode. For example, `mstatus.MPRV=1` can cause load and store instructions to be executed in Supervisor or User modes instead of Machine mode.

This feature of the RISC-V architecture is implemented in the model by modifying the *memory domain* to which loads and stores are routed. This function causes the current memory domain to be refreshed so that it is valid for the current CSR settings. It should be called after any CSR update that affects Machine mode loads and stores in this way.

Example

```
static void updateLdStDomain(riscvP riscv) {
    riscv->cb.updateLdStDomain(riscv);
}
```

Usage Context

Leaf level only.

14.59 *Function newTLBEntry*

```
#define RISCV_NEW_TLB_ENTRY_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvTLBId  tlbId, \
    riscvExtVMMapping mapping \
)
typedef RISCV_NEW_TLB_ENTRY_FN(*riscvNewTLBEntryFn);

typedef struct riscvModelCBS {
    riscvNewTLBEntryFn    newTLBEntry;
} riscvModelCB;
```

Description

One RISC-V implementation choice is to implement virtual memory TLB updates using a trap to a Machine mode handler. In this case, memory mappings will typically be modified by writes to custom CSRs, or a similar mechanism.

This function is used to create a new mapping using a custom instruction. The TLB that is affected is specified by the `tlbId` enumeration:

```
typedef enum riscvTLBIdE {

    RISCV_TLB_HS,        // HS TLB
    RISCV_TLB_VS1,       // VS stage 1 virtual TLB
    RISCV_TLB_VS2,       // VS stage 2 virtual TLB

} riscvTLBId
```

(`RISCV_TLB_VS1` and `RISCV_TLB_VS2` TLBs should only be used for processors that implement the Hypervisor extension.)

The mapping to create is described by the `mapping` parameter, which is of type `riscvExtVMMapping`:

```
typedef struct riscvExtVMMappingsS {
    Uns64  lowVA;        // low VA
    Uns64  highVA;       // high VA
    Uns64  PA;           // low PA
    Uns16  entryId : 16; // custom unique identifier
    memPriv priv : 8;    // access privileges
    Bool   V       : 1;  // valid bit
    Bool   U       : 1;  // User-mode access
    Bool   G       : 1;  // global entry
    Bool   A       : 1;  // accessed
    Bool   D       : 1;  // dirty
} riscvExtVMMapping
```

Most entries correspond to fields of the same name in the RISC-V Privileged Architecture description. The `entryId` field is for application use, to identify a TLB entry uniquely in a TLB.

Example

```
static void installTLBEntry(riscvP riscv, custTLBEntryP entry) {
```

```
if(!entry->installed) {  
    riscv->cb.newTLBEntry(  
        riscv, entry->xatp, entry->mapping  
    );  
    entry->installed = True;  
}
```

Usage Context

Leaf level only.

14.60 Function *freeTLBEntry*

```
#define RISCV_FREE_TLB_ENTRY_FN(_NAME) void _NAME( \
    riscvP      riscv,          \
    riscvTLBId  tlbId,          \
    Uns16       entryId         \
)
typedef RISCV_FREE_TLB_ENTRY_FN(*riscvFreeTLBEntryFn);

typedef struct riscvModelCBS {
    riscvFreeTLBEntryFn    freeTLBEntry;
} riscvModelCB;
```

Description

One RISC-V implementation choice is to implement virtual memory TLB updates using a trap to a Machine mode handler. In this case, memory mappings will typically be modified by writes to custom CSRs, or a similar mechanism.

This function is used to invalidate a mapping using a custom instruction. The TLB that is affected is specified by the `tlbId` enumeration:

```
typedef enum riscvTLBIdE {
    RISCV_TLB_HS,          // HS TLB
    RISCV_TLB_VS1,         // VS stage 1 virtual TLB
    RISCV_TLB_VS2,         // VS stage 2 virtual TLB
} riscvTLBId
```

(`RISCV_TLB_VS1` and `RISCV_TLB_VS2` TLBs should only be used for processors that implement the Hypervisor extension.)

The mapping to invalidate is described by the `entryId` parameter, which is an index number corresponding to the field of the same name when the entry was created – see section 14.59 for more information.

Example

```
static void uninstallTLBEntry(riscvP riscv, custTLBEntryP entry) {
    if(entry->installed) {
        riscv->cb.freeTLBEntry(
            riscv, entry->xatp, entry->mapping.entryId
        );
        VMI_ASSERT(!entry->installed, "TLB entry not uninstalled");
    }
}
```

Usage Context

Leaf level only.

14.61 *Function newExtReg*

```
#define RISC_V_NEW_EXT_REG_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    vmiRegInfoCP extReg \
)
typedef RISC_V_NEW_EXT_REG_FN(*riscvNewExtRegFn);

typedef struct riscvModelCBS {
    riscvNewExtRegFn      newExtReg;
} riscvModelCB;
```

Description

Given a template generic extension register description, this function registers that extension register with the base model. It should always be called from the extension object constructor. See section 11.4 for a detailed description.

Example

```
static void addExtRegs(vmiosObjectP object) {

    riscvP riscv = object->riscv;

    // create description of read-only transaction mode register
    vmiRegInfo tmReg = {
        name       : "TM",
        description : "Transaction mode state",
        bits       : 8,
        gdbIndex    : 0,
        access      : vmi_RA_R,
        raw         : vmimtGetExtReg((vmiProcessorP)riscv, &object->tmStatus)
    };

    // add register description
    riscv->cb.newExtReg(riscv, &tmReg);
}
```

Usage Context

Leaf level only, in constructor.

15 Appendix: Extension Object Interface Functions

This appendix describes *interface functions* implemented by the *extension object* that allow the extension object to provide information to modify the behavior of the base model. All such interface functions are held in a structure of type `riscvExtCBS`, defined in file `riscvModelCallbacks.h` in the base model:

```
typedef struct riscvExtCBS {

    // link pointer and id (maintained by base model)
    riscvExtCBP      next;
    Uns32            id;

    // handle back to client data
    void             *clientData;

    // exception modification
    riscvRdWrFaultFn rdFaultCB;
    riscvRdWrFaultFn wrFaultCB;
    riscvRdWrSnapFn  rdSnapCB;
    riscvRdWrSnapFn  wrSnapCB;

    // exception actions
    riscvSuppressMemExceptFn suppressMemExcept;
    riscvCustomNMIFn        customNMI;
    riscvCustomIAssignFn    customIAssign;
    riscvTrapNotifierFn     trapNotifier;
    riscvTrapNotifierFn     trapPreNotifier;
    riscvTrapNotifierFn     ERETNotifier;
    riscvResetNotifierFn    preResetNotifier;
    riscvResetNotifierFn    resetNotifier;
    riscvFirstExceptionFn   firstException;
    riscvGetInterruptPriFn  getInterruptPri;
    riscvGetHandlerPCFn     getHandlerPC;
    riscvIntUpdateFn        intUpdate;

    // halt/restart actions
    riscvHRNotifierFn       haltRestartNotifier;
    riscvLRSCAbortFn        LRSCAbortFn;

    // code generation actions
    riscvDerivedMorphFn     preMorph;
    riscvDerivedMorphFn     postMorph;
    riscvDerivedMorphFn     AMOCheck;
    riscvDerivedMorphFn     AMOMorph;
    riscvEmitCSRCheckFn     emitCSRCheck;
    riscvUnitStrideCheckFn  unitStrideCheck;
    riscvVFREDSUMMorphFn    emitVFREDUSUM;
    riscvVFREDSUMMorphFn    emitVFWREDUSUM;

    // transaction support actions
    riscvIASSwitchFn        switchCB;
    riscvTLoadFn            tLoad;
    riscvTStoreFn           tStore;

    // memory access logging support actions
    riscvCheckMemMtFn       checkLoadMt;
    riscvCheckMemMtFn       checkStoreMt;

    // physical memory actions
    riscvDistinctPhysMemFn  distinctPhysMem;
    riscvPhysMemFn          installPhysMem;

    // PMP support actions
    riscvPMPPrivFn          PMPPriv;

    // PMA check actions
```

```

riscvPMAEnableFn      PMAEnable;
riscvPMACheckFn       PMACheck;

// virtual memory actions
riscvVMTrapFn         VMTrap;
riscvValidPTEFn       validPTE;
riscvSetDomainNotifierFn setDomainNotifier;
riscvFreeEntryNotifierFn freeEntryNotifier;

// CLIC actions
riscvClicCRdFn        clicCustomRd;
riscvClicCWrrFn       clicCustomWr;
riscvClicUpdatedFn    clicUpdated;

// documentation
riscvRestrictionsFn   restrictionsCB;
} riscvExtCBtype;

```

A structure of this type should be defined within the `vmiosObject` structure of the extension. Fields within the structure can either be initialized in the extension object constructor (to modify the standard base model behavior) or left as `NULL` (to use base model behavior unchanged).

Fields `next` and `id` are used by the base model to chain together multiple extensions to a single processor and should not be directly modified by the extension object. Field `clientData` should be initialized with the `vmiosObject` pointer of the extension object, as shown in all examples described in this document. Other fields may be modified by the extension object constructor and are described in following subsections.

15.1 Function *rdFaultCB*

```
#define RISC_V_RD_WR_FAULT_FN(_NAME) Bool _NAME( \
    riscvP      riscv,          \
    memDomainP  domain,         \
    Addr        address,        \
    Uns32        bytes,         \
    void        *clientData     \
)
typedef RISC_V_RD_WR_FAULT_FN((*riscvRdWrFaultFn))

typedef struct riscvExtCBS {
    riscvRdWrFaultFn      rdFaultCB;
} riscvExtCB;
```

Description

This interface function is called from the base model when an unaligned *load* access is detected. The access is of size *bytes* at address *address*. The function should return `True` if the access should cause a Load Access Fault exception (code 5) and `False` if the access should cause a Load Address Misaligned exception (code 4).

Example

This example shows how to force all misaligned loads to be reported as access faults:

```
static RISC_V_RD_WR_FAULT_FN(misalignedAccess) {
    return True;
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.rdFaultCB = misalignedAccess;
    . . . lines omitted . . .
}
```

15.2 Function *wrFaultCB*

```
#define RISC_V_RD_WR_FAULT_FN(_NAME) Bool _NAME( \
    riscvP      riscv,          \
    memDomainP  domain,         \
    Addr        address,        \
    Uns32        bytes,         \
    void        *clientData     \
)
typedef RISC_V_RD_WR_FAULT_FN((*riscvRdWrFaultFn))

typedef struct riscvExtCBS {
    riscvRdWrFaultFn wrFaultCB;
} riscvExtCB;
```

Description

This interface function is called from the base model when an unaligned *store* access is detected. The access is of size *bytes* at address *address*. The function should return *True* if the access should cause a Store/AMO Access Fault exception (code 7) and *False* if the access should cause a Store/AMO Address Misaligned exception (code 6).

Example

This example shows how to force all misaligned stores to be reported as access faults:

```
static RISC_V_RD_WR_FAULT_FN(misalignedAccess) {
    return True;
}

static VMIO_S_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.wrFaultCB = misalignedAccess;
    . . . lines omitted . . .
}
```


15.3 Function *rdSnapCB*

```
#define RISC_V_RD_WR_SNAP_FN(_NAME) Uns32 _NAME( \
    riscvP      riscv, \
    memDomainP  domain, \
    Addr        address, \
    Uns32       bytes, \
    atomicCode  atomic, \
    void        *clientData \
)
typedef RISC_V_RD_WR_SNAP_FN(*riscvRdWrSnapFn);

typedef struct riscvExtCBS {
    riscvRdWrSnapFn      rdSnapCB;
} riscvExtCB;
```

Description

This interface function is called from the base model when an unaligned *load* access is detected. The access is of size *bytes* at address *address*. The function can either allow the unaligned access to proceed normally (perhaps with data rotation) or return a code indicating that an exception should be taken. The required behavior is indicated by the Uns32 return code, whose value is constructed using the MEM_SNAP macro in *vmiTypes.h*. In practice, two return codes are likely to be required:

MEM_SNAP(1, 0): this indicates the unaligned access should proceed.

MEM_SNAP(0, 0): this indicates the unaligned access causes an exception.

Example

This example shows how unaligned accesses within a 64-byte cache line can be permitted, while unaligned accesses that straddle lines cause an exception:

```
//
// Get line index for address
//
inline static Uns32 getLine(Uns32 address) {
    return address/64;
}

//
// Unaligned accesses are allowed only within cache lines
//
static RISC_V_RD_WR_SNAP_FN(snapCB) {
    Uns32 snap = MEM_SNAP(1, 0);

    if(getLine(address) != getLine(address+bytes-1)) {
        snap = MEM_SNAP(0, 0);
    }

    return snap;
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.rdSnapCB = snapCB;
    . . . lines omitted . . .
}
```

15.4 Function *wrSnapCB*

```
#define RISC_V_RD_WR_SNAP_FN(_NAME) Uns32 _NAME( \
    riscvP      riscv, \
    memDomainP  domain, \
    Addr        address, \
    Uns32       bytes, \
    atomicCode  atomic, \
    void        *clientData \
)
typedef RISC_V_RD_WR_SNAP_FN(*riscvRdWrSnapFn);

typedef struct riscvExtCBS {
    riscvRdWrSnapFn      wrSnapCB;
} riscvExtCB;
```

Description

This interface function is called from the base model when an unaligned *store* access is detected. The access is of size *bytes* at address *address*. The function can either allow the unaligned access to proceed normally (perhaps with data rotation) or return a code indicating that an exception should be taken. The required behavior is indicated by the Uns32 return code, whose value is constructed using the MEM_SNAP macro in *vmiTypes.h*. In practice, two return codes are likely to be required:

MEM_SNAP(1, 0): this indicates the unaligned access should proceed.

MEM_SNAP(0, 0): this indicates the unaligned access causes an exception.

Example

This example shows how unaligned accesses within a 64-byte cache line can be permitted, while unaligned accesses that straddle lines cause an exception:

```
//
// Get line index for address
//
inline static Uns32 getLine(Uns32 address) {
    return address/64;
}

//
// Unaligned accesses are allowed only within cache lines
//
static RISC_V_RD_WR_SNAP_FN(snapCB) {

    Uns32 snap = MEM_SNAP(1, 0);

    if(getLine(address) != getLine(address+bytes-1)) {
        snap = MEM_SNAP(0, 0);
    }

    return snap;
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {

    . . . lines omitted . . .
    object->extCB.wrSnapCB = snapCB;
    . . . lines omitted . . .
}
```

15.5 Function *suppressMemExcept*

```
#define RISC_V.Suppress_Mem_Except_Fn(_Name) Bool _Name( \
    riscvP      riscv, \
    riscvException exception, \
    void        *clientData \
)
typedef RISC_V.Suppress_Mem_Except_Fn(*riscvSuppressMemExceptFn);

typedef struct riscvExtCBS {
    riscvSuppressMemExceptFn  suppressMemExcept;
} riscvExtCB;
```

Description

This interface function is called from the base model when a memory exception is about to be taken. It gives the extension library an opportunity to suppress that exception if required. Argument `exception` specifies the exception that is about to be taken.

Example

This example shows how memory exceptions can be suppressed, and a sticky bit set instead, in a custom extension CSR `mecsr`:

```
static RISC_V.Suppress_Mem_Except_Fn(suppressMemExcept) {

    vmiosObjectP object = clientData;
    Bool          suppress = False;

    switch(exception) {

        case riscv_E_LoadAddressMisaligned:
        case riscv_E_LoadAccessFault:
        case riscv_E_LoadPageFault:
            if(RD_XCSR_FIELD(object, mecsr, REDIS)) {
                suppress = True;
                WR_XCSR_FIELD(object, mecsr, RES, 1);
            }
            break;

        case riscv_E_StoreAMOAddressMisaligned:
        case riscv_E_StoreAMOAccessFault:
        case riscv_E_StoreAMOPageFault:
            if(RD_XCSR_FIELD(object, mecsr, WEDIS)) {
                suppress = True;
                WR_XCSR_FIELD(object, mecsr, WES, 1);
            }
            break;

        default:
            break;
    }

    return suppress;
}

static VMIO_Constructor_Fn(extConstructor) {

    . . . lines omitted . . .
    object->extCB.suppressMemExcept = suppressMemExcept;
    . . . lines omitted . . .
}
```

15.6 Function *customNMI*

```
#define RISC_V_CUSTOM_NMI_FN(_NAME) Bool _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_CUSTOM_NMI_FN(*riscvCustomNMIFn);

typedef struct riscvExtCBS {
    riscvCustomNMIFn      customNMI;
} riscvExtCB;
```

Description

This interface function is called from the base model when an NMI is to be taken, allowing the extension to define custom behavior for that exception. The return code indicates whether special NMI behavior has been performed; if `False`, the base model will handle the NMI in the normal way.

If the callback returns `False`, then depending on whether the RISC-V RNMI extension is configured, *the callback must set the full value of either `mcause` or `mncause` before returning*. The base model will then perform all other actions to indicate an M-mode exception has been taken, either as a standard exception or as a standard RNMI trap (depending on whether RNMI is configured). As part of this, any exception code specified by the `ecode_nmi` parameter or driven using the `nmi_cause` input signal will be combined with the value of `mcause` or `mncause` set by the extension callback using bitwise-or.

If the callback returns `True`, then no standard actions will be performed by the base model to handle the NMI: all required state changes must be made by the extension.

Example

This example shows how an extension could handle an NMI exception as a normal trap with custom cause:

```
static RISC_V_CUSTOM_NMI_FN(customNMI) {
    riscv->cb.takeException(riscv, riscv_E_Interrupt+EXT_NMI, 0);
    return True;
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.customNMI = customNMI;
    . . . lines omitted . . .
}
```

15.7 Function *customIAssign*

```
#define RISCVCUSTOM_IASSIGN_FN(_NAME) Uns64 _NAME( \
    riscvP          riscv, \
    Uns64           ip, \
    riscvBasicIntStateP intState, \
    void            *clientData \
)
typedef RISCVCUSTOM_IASSIGN_FN((*riscvCustomIAssignFn));

typedef struct riscvExtCBS {
    riscvCustomIAssignFn customIAssign;
} riscvExtCB;
```

Description

This interface function allows an extension to modify the allocation of standard interrupts to privilege levels, or to inject supplementary interrupts that are not visible using the standard mip CSR. It is called when the base model has determined the appropriate privilege levels for enabled standard interrupts but before the highest-priority interrupt has been determined.

The function is passed a pointer to an object of type `riscvBasicIntState`:

```
typedef struct riscvBasicIntStateS {
    Uns64 ip[RISCVMODE_LAST]; // pending and enabled interrupts per mode
    Bool hvictl;               // whether hvictl-injected interrupt
} riscvBasicIntState;
```

In this structure, the `ip` array holds bitmasks of pending and enabled interrupts that have been delegated to each privilege mode. The function can modify the values here to add or remove pending interrupts or change their priority assignment.

The function is also passed an `ip` parameter, which is the bitwise-or of pending and enabled interrupts at all privilege levels. It should return a value consistent with any changes it makes to interrupt assignments.

Example

Andes processors implement custom `slip`, `slie` and `mslideleg` CSRs allowing performance monitor interrupts to be delivered to S-mode even when those interrupts are disabled at M-mode. This is handled in the following custom interrupt assignment function:

```
RISCVCUSTOM_IASSIGN_FN(andesCustomIAssign) {
    vmiosObjectP object = clientData;
    Uns32 slip = RD_XCSR(object, slip) & RD_XCSR(object, slie);

    if(slip) {
        Uns32 mslideleg = RD_XCSR(object, mslideleg);

        // indicate slip interrupts are pending and enabled
        ip |= slip;

        // route interrupts to either M-mode or S-mode
        intState->ip[RISCVMODE_M] |= (slip & ~mslideleg);
    }
}
```

```
        intState->ip[RISCV_MODE_S] |= (slip & mslideleg);
    }

    return ip;
}

static VMIOS_CONSTRUCTOR_FN(constructor) {
    . . . lines omitted . . .
        object->extCB.customIAssign = andesCustomIAssign;
    . . . lines omitted . . .
}
```

15.8 Function *trapNotifier*

```
#define RISC_V_TRAP_NOTIFIER_FN(_NAME) void _NAME( \
    riscvP      riscv,      \
    riscvMode    mode,      \
    riscvERETType eretType,  \
    void        *clientData \
)
typedef RISC_V_TRAP_NOTIFIER_FN((*riscvTrapNotifierFn));

typedef struct riscvExtCBS {
    riscvTrapNotifierFn trapNotifier;
} riscvExtCB;
```

Description

This interface function is called from the base model when a trap is taken. It gives the extension library the opportunity to modify trap behavior (perhaps by recording extra information about certain trap types). The notifier is called when all hart state has been modified to handle the trap: section 15.9 describes an alternative notifier that is called *before* hart state is modified.

The `mode` argument specifies the mode to which the trap is taken.

The `eretType` argument is significant only when the callback is used in the context of an exception return notifier (see section 15.10). Here, it is always passed the value `ERT_NA`.

Example

This example shows how an extension could record extra data in field `subCause` of a custom CSR `mcustcause` when a custom interrupt `EXT_INT1` is taken:

```
static RISC_V_TRAP_NOTIFIER_FN(takeTrap) {
    vmiosObjectP object = clientData;

    if(
        (mode==RISC_V_MODE_MACHINE) &&
        RD_CSR_FIELD(riscv, mcause, Interrupt) &&
        (RD_CSR_FIELD(riscv, mcause, ExceptionCode)==EXT_INT1)
    ) {
        WR_XCSR_FIELD(object, mcustcause, subCause, object->subCause);
    }
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.trapNotifier = takeTrap;
    . . . lines omitted . . .
}
```

15.9 Function *trapPreNotifier*

```
#define RISC_V_TRAP_NOTIFIER_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvMode    mode, \
    riscvERETType eretType, \
    void        *clientData \
)
typedef RISC_V_TRAP_NOTIFIER_FN((*riscvTrapNotifierFn));

typedef struct riscvExtCBS {
    riscvTrapNotifierFn    trapPreNotifier;
} riscvExtCB;
```

Description

This interface function is called from the base model when a trap is about to be taken. It gives the extension library the opportunity to save current model state before it is modified in the process of taking the trap. Section 15.8 describes an alternative notifier that is called *after* hart state is modified.

The `mode` argument specifies the mode to which the trap is taken.

The `eretType` argument is significant only when the callback is used in the context of an exception return notifier (see section 15.10). Here, it is always passed the value `ERT_NA`.

Example

This example shows how an extension could save the current value of `mstatus.MPP` in a custom CSR field `mcuststatus.SMPP` before a custom interrupt `EXT_INT1` is taken. This field is modified as part of the standard process of taking a trap.

```
static RISC_V_TRAP_NOTIFIER_FN(preTrap) {
    vmiosObjectP object = clientData;

    if(
        RD_CSR_FIELD(riscv, mcause, Interrupt) &&
        (RD_CSR_FIELD(riscv, mcause, ExceptionCode)==EXT_INT1)
    ) {
        Uns32 MPP = RD_CSR_FIELD_M(riscv, mstatus, MPP);
        WR_XCSR_FIELD(object, mcuststatus, SMPP, MPP);
    }
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.trapPreNotifier = preTrap;
    . . . lines omitted . . .
}
```


15.10 Function *ERETNotifier*

```
#define RISC_V_TRAP_NOTIFIER_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    riscvMode    mode, \
    riscvERETType eretType, \
    void        *clientData \
)
typedef RISC_V_TRAP_NOTIFIER_FN((*riscvTrapNotifierFn));

typedef struct riscvExtCBS {
    riscvTrapNotifierFn    ERETNotifier;
} riscvExtCB;
```

Description

This interface function is called from the base model when a return from exception instruction is executed. It gives the extension library the opportunity to modify exception return behavior (perhaps by restoring extra custom information).

The `mode` argument specifies the mode from which the trap return is being made.

The `eretType` argument indicates the exact type of exception return instruction that has been executed, described by the `riscvERETType` enumeration:

```
typedef enum riscvERETTypeE {
    ERT_NA,    // not an exception return instruction
    ERT_M,     // mret instruction
    ERT_D,     // dret instruction
    ERT_MN,    // mnret instruction (from resumable NMI)
    ERT_S,     // sret (from S/HS mode)
    ERT_VS,    // sret (from VS mode)
    ERT_U,     // uret (from U/HU mode)
    ERT_VU,    // uret (from VU mode)
} riscvERETType;
```

Example

This example shows how the `andes.ovpworld.org` model uses this function to implement the Andes-specific CRASHSAVE extension and restore custom fields in the `mxstatus` CSR when executing an MRET instruction:

```
static RISC_V_TRAP_NOTIFIER_FN(ERETNotifier) {

    vmiosObjectP object = clientData;

    // restore CRASHSAVE extension information if required
    if(RD_CSR_FIELD(object->riscv, mcause, Interrupt)) {
        // no action if last cause was an interrupt
    } else if(eretType!=ERT_M) {
        // no action unless an mret instruction
    } else if(!RD_XCSR_FIELD(object, mmisc_cfg, CRASHSAVE)) {
        // CRASHSAVE feature is absent
    } else {
        WR_CSR( object->riscv, mepc,    RD_XCSR(object, msaveepc1));
        WR_CSR( object->riscv, mcause,  RD_XCSR(object, msavecause1));
        WR_XCSR(object,          mdcause, RD_XCSR(object, msavedcause1));
    }

    // restore mxstatus fields when MRET is executed
    if(eretType==ERT_M) {
        COPY_FIELD(object, mxstatus, PFT_EN, PPFT_EN);
    }
}
```

```
        COPY_FIELD(object, mxstatus, IME,    PIME);
        COPY_FIELD(object, mxstatus, DME,    PDME);
        COPY_FIELD(object, mxstatus, TYP,    PTYP);
    }

    // refresh all counter objects
    refreshCounters(object);
}

void andesCSRInit(vmiosObjectP object) {
    . . . lines omitted . . .
    object->extCB.ERETNotifier = ERETNotifier;
    . . . lines omitted . . .
}
```

15.11 Function *preResetNotifier*

```
#define RISC_V_RESET_NOTIFIER_FN(_NAME) void _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_RESET_NOTIFIER_FN((*riscvResetNotifierFn));

typedef struct riscvExtCBS {
    riscvResetNotifierFn    preResetNotifier;
} riscvExtCB;
```

Description

This interface function is called from the base model when a reset is executed, *before* any standard reset actions have been performed. It allows the extension to implement custom features such as crash state logging, which require saving current processor state before it is modified as part of the reset process.

Example

This example shows a template for crash state logging based on a feature defined for Andes RISC-V cores:

```
//
// Fill a field in mcrash_statesave
//
#define FILL_STATESAVE(_O, _F, _V) WR_XCSR_FIELD(_O, mcrash_statesave, _F, _V)

//
// Perform CSR changes prior to reset
//
static RISC_V_RESET_NOTIFIER_FN(CSRPreReset) {

    vmiosObjectP object = clientData;
    riscvMode mode = getCurrentMode3(riscv);

    // fill mcrash_statesave
    FILL_STATESAVE(object, MIE, RD_CSR_FIELD(riscv, mstatus, MIE));
    FILL_STATESAVE(object, CP, mode);
    FILL_STATESAVE(object, PPFT_EN, RD_XCSR_FIELD(object, mxstatus, PPFT_EN));
    FILL_STATESAVE(object, PIME, RD_XCSR_FIELD(object, mxstatus, PIME));
    FILL_STATESAVE(object, PDME, RD_XCSR_FIELD(object, mxstatus, PDME));
    FILL_STATESAVE(object, PTYP, RD_XCSR_FIELD(object, mxstatus, PTYP));

    . . . lines omitted . . .
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {

    . . . lines omitted . . .
    object->extCB.preResetNotifier = CSRPreReset;
    . . . lines omitted . . .
}
```

15.12 *Function resetNotifier*

```
#define RISC_V_RESET_NOTIFIER_FN(_NAME) void _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_RESET_NOTIFIER_FN((*riscvResetNotifierFn));

typedef struct riscvExtCBS {
    riscvResetNotifierFn    resetNotifier;
} riscvExtCB;
```

Description

This interface function is called from the base model when a reset is executed, *after* all standard reset actions have been performed. It gives the extension the opportunity to perform custom reset actions and also set standard CSR fields to implementation-defined values.

Example

This example shows how both a custom CSR and a standard CSR field can be reset:

```
static RISC_V_RESET_NOTIFIER_FN(CSRReset) {

    vmiosObjectP object = clientData;

    // reset custom CSR
    WR_XCSR(object, custom1, 0);

    // reset standard CSR fields
    WR_CSR_FIELD(riscv, mstatus, TW, 1);
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {

    . . . lines omitted . . .
    object->extCB.resetNotifier = CSRReset;
    . . . lines omitted . . .
}
```

15.13 *Function firstException*

```
#define RISC_V_FIRST_EXCEPTION_FN(_NAME) vmiExceptionInfoCP _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_FIRST_EXCEPTION_FN((*riscvFirstExceptionFn));

typedef struct riscvExtCBS {
    riscvFirstExceptionFn    firstException;
} riscvExtCB;
```

Description

This interface function returns the first member of a NULL-terminated list of custom exceptions implemented by an extension. The base model adds all exception descriptions in the list to the visible exceptions of the derived model.

Example

This example shows addition of a custom exception (EXCEPT24) to the base model:

```
#define EXT_EXCEPTION(_NAME, _DESC) { \
    name: #_NAME, code: EXT_E_##_NAME, description: _DESC, \
}

static const vmiExceptionInfo exceptions[] = {
    EXT_EXCEPTION (EXCEPT24, "Custom Exception 24"),
    {0}
};

static RISC_V_FIRST_EXCEPTION_FN(firstException) {
    return exceptions;
}

static VMIO5_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.firstException = firstException;
    . . . lines omitted . . .
}
```

15.14 *Function `getInterruptPri`*

```
#define RISC_V_GET_INTERRUPT_PRI_FN(_NAME) riscvExceptionPriority _NAME( \
    riscvP riscv, \
    Uns32 intNum, \
    void *clientData \
)
typedef RISC_V_GET_INTERRUPT_PRI_FN((*riscvGetInterruptPriFn));

typedef struct riscvExtCBS {
    riscvGetInterruptPriFn    getInterruptPri;
} riscvExtCB;
```

Description

This interface function returns the relative priority of an interrupt, or 0 if the default priority for that interrupt should be used. If non-zero, priorities are expressed relative to specified priorities of standard interrupts, as define by the `riscvExceptionPriority` enumeration.

Example

This example shows a function returning priorities for two custom interrupts, `EXT_I_INT21` and `EXT_I_INT22`:

```
static RISC_V_GET_INTERRUPT_PRI_FN(getInterruptPriority) {

    riscvExceptionPriority result = 0;

    if(intNum==EXT_I_INT21) {
        result = riscv_E_LocalPriority;
    } else if(intNum==EXT_I_INT22) {
        result = riscv_E_LocalPriority+1;
    }

    return result;
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {

    . . . lines omitted . . .
    object->extCB.getInterruptPri = getInterruptPriority;
    . . . lines omitted . . .
}
```

15.15 Function *getHandlerPC*

```
#define RISC_V_GET_HANDLER_PC_FN(_NAME) Bool _NAME( \
    riscvP      riscv, \
    Uns64       tvec, \
    riscvException exception, \
    Uns32       ecode, \
    Uns64       *handlerPCP, \
    void        *clientData \
)
typedef RISC_V_GET_HANDLER_PC_FN((*riscvGetHandlerPCFn));

typedef struct riscvExtCBS {
    riscvGetHandlerPCFn    getHandlerPC;
} riscvExtCB;
```

Description

This interface function allows an extension model to return a custom interrupt or exception handler address. It is called when a trap occurs and is passed the `xtvec` address for the target mode, the exception that is being taken, the exception code (usually derived from the exception, but can differ for interrupts when the external interrupt ID signals are driven) and a pointer to an `Uns64` handler PC to be filled with the desired handler address if required.

If the custom handler address function is active, it should fill the `handlerPCP` result and return `True`; otherwise, it should return `False`.

Example

This example shows how this function is used in the `andes.ovpworld.org` model. In this model, a custom vectored trap mode is enabled when `mmisc_ctl.VEC_PLIC=1`:

```
RISC_V_GET_HANDLER_PC_FN(andesGetHandlerPC) {
    vmiosObjectP object = clientData;
    Bool         custom = RD_XCSR_FIELD(object, mmisc_ctl, VEC_PLIC);

    if(custom) {
        riscvMode      mode      = getCurrentMode5(riscv);
        memEndian      endian    = riscv->cb.getDataEndian(riscv, mode);
        memDomainP     domain    = getDataDomain(riscv);
        memAccessAttrs memAttrs = MEM_AA_TRUE;
        Uns32          offset    = 0;
        Uns64          handlerPC;

        // get table offset for this exception type
        switch(exception) {
            case riscv_E_UExternalInterrupt:
            case riscv_E_SExternalInterrupt:
            case riscv_E_MExternalInterrupt:
                offset = ecode*4;
                break;

            default:
                break;
        }

        // read 4-byte table entry
        handlerPC = vmirtRead4ByteDomain(domain, tvec+offset, endian, memAttrs);
    }
}
```

```
        // mask off LSB from result
        *handlerPCP = handlerPC &= -2;
    }

    return custom;
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.getHandlerPC = andesGetHandlerPC;
    . . . lines omitted . . .
}
```


15.16 Function *intUpdate*

```
#define RISC_V_INTERRUPT_UPDATE_FN(_NAME) Bool _NAME( \
    riscvP      riscv, \
    Uns32       intNum, \
    Uns64       *newValue, \
    void        *clientData \
)
typedef RISC_V_INTERRUPT_UPDATE_FN((*riscvIntUpdateFn));
```

Description

This notifier function informs the extension when a write has occurred on an interrupt input pin of a processor.

intNum is the interrupt number where the write occurred.

**newValue* is a pass by reference of the value that was written. The value may be changed in the notifier.

If the function returns true then a re-evaluation of the interrupt state will be forced, regardless of the value that was written (otherwise, re-evaluation will occur only if the new value causes a change in the interrupt state.) This should be done if the interrupt state is modified in the callback,

Example

This example shows how an extension can keep state info on the values on the first 32 interrupt inputs:

```
static RISC_V_INTERRUPT_UPDATE_FN(intUpdate) {
    vmiosObjectP object = clientData;

    if (intNum < 32) {
        Uns32 mask      = 1 << intNum;
        Bool  intIn     = *newValue != 0;

        if (intIn) {
            object->interruptLevels |= mask;
        } else {
            object->interruptLevels &= ~mask;
        }
    }

    return False;
}
```

15.17 Function *haltRestartNotifier*

```
#define RISC_V_HR_NOTIFIER_FN(_NAME) void _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_HR_NOTIFIER_FN((*riscvHRNotifierFn));

typedef struct riscvExtCBS {
    riscvHRNotifierFn      haltRestartNotifier;
} riscvExtCB;
```

Description

This interface function is called when a processor transitions from running to halted state, or from halted to running state. It allows the extension object to perform any state changes required at this point (typically, to components like instruction counters).

Example

This example shows how this function is used in the `andes.ovpworld.org` model. In this model, when the processor transitions between running and halted states, counters need to be started and stopped:

```
static void refreshCounter(andesCounterP counter) {

    riscvP      riscv = object->riscv;
    andesCounterMode cmode = ACM_INACTIVE;

    // get raw counter mode
    if(counter->TYPE) {
        // no action
    } else if(counter->SEL==1) {
        cmode = ACM_CY;
    } else if(counter->SEL==2) {
        cmode = riscv->disable ? ACM_INACTIVE : ACM_IR;
    }

    . . . counter state modified based in cmode here . . .
}

static void refreshCounters(vmiosObjectP object) {

    andesCounterID id;

    for(id=0; id<AT_LAST; id++) {

        andesCounterP counter = &object->counters[id];

        if(counter->vmi) {
            refreshCounter(counter);
        }
    }
}

static RISC_V_HR_NOTIFIER_FN(haltRestartNotifier) {
    refreshCounters(clientData);
}

static VMIO_CONSTRUCTOR_FN(constructor) {

    . . . lines omitted . . .
    object->extCB.haltRestartNotifier = haltRestartNotifier;
    . . . lines omitted . . .
}
```

15.18 *Function LRSCAbortFn*

```
#define RISC_V_LRSC_ABORT_FN(_NAME) void _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_LRSC_ABORT_FN((*riscvLRSCAbortFn));

typedef struct riscvExtCBS {
    riscvLRSCAbortFn      LRSCAbortFn;
} riscvExtCB;
```

Description

This interface function is called when an active LR/SC sequence is aborted (for example, because of a conflicting store by another processor). It allows the extension object to perform any state changes required at this point.

Example

This example shows how this function could be used to restart a processor that is halted for an implementation-defined reason when another processor in a multicore simulation causes an active LR/SC sequence to be aborted:

```
RISC_V_LRSC_ABORT_FN(custLRSCAbort) {
    if(riscv->disable & RVD_CUSTOM_WFI) {
        riscv->cb.restart(riscv, RVD_CUSTOM_WFI);
    }
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {
    . . . lines omitted . . .
    object->extCB.LRSCAbortFn = custLRSCAbort;
    . . . lines omitted . . .
}
```

15.19 Function *preMorph*

```
#define RISC_V_DERIVED_MORPH_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    void        *clientData, \
    riscvInstrInfoP instrInfo, \
    Uns64       thisPC \
)
typedef RISC_V_DERIVED_MORPH_FN((*riscvDerivedMorphFn));

typedef struct riscvExtCBS {
    riscvDerivedMorphFn    preMorph;
} riscvExtCB;
```

Description

This interface function is called before an instruction is translated by the base model. It allows the extension object to emit code to perform an action that should be done before any standard instruction. The function is passed the address of the instruction (*thisPC*) and a pointer to a structure giving information about the decoded instruction (*instrInfo*); see section 16 for more information about this structure.

Example

This example shows how this function is used in the `andes.ovpworld.org` model. In this model, modeling hardware stack protection (HSP) requires that the current stack pointer is saved before each instruction executes so that it can be compared with the new value after the instruction. The callback is also used to specify the alignment constraints for any load/store instructions (there is a custom register that enables or disables unaligned access for some instruction types).

```
void andesRecordSP(riscvP riscv, vmiosObjectP object) {
    if(!isHSPEnabledMT(riscv, object)) {
        // no action if feature is currently disabled
    } else {
        Uns32 bits = andesGetXlenArch(riscv);
        vmiReg oldSPReg = andesObjectReg(object, RISC_V_OBJ_REG(oldSP));

        // move current stack pointer to a temporary
        vmintMoveRR(bits, oldSPReg, RISC_V_GPR(RV_REG_X_SP));
    }
}

RISC_V_DERIVED_MORPH_FN(andesPreMorph) {
    vmiosObjectP object = clientData;
    andesMorphState state;

    // handle hardware stack protection if required
    if(RD_XCSR_FIELD(object, mmisc_cfg, HSP)) {
        andesRecordSP(riscv, object);
    }

    // get instruction and instruction type
    andesDecode(riscv, object, thisPC, &state.info);

    if(state.info.type == AN_IT_MSA_UNA) {
        // base model instructions controlled by mmisc_ctl.MSA_UNA
    }
}
```

```
    riscv->configInfo.unaligned = RD_XCSR_FIELD(object, mmisc_ctl, MSA_UNA);

    } else if(state.info.type == AN_IT_LAST) {

        // other base model instructions always disallow unaligned access
        riscv->configInfo.unaligned = False;
    }
}

static VMIOS_CONSTRUCTOR_FN(constructor) {

    . . . lines omitted . . .
    object->extCB.preMorph      = andesPreMorph;
    . . . lines omitted . . .
}
```

15.20 Function *postMorph*

```
#define RISC_V_DERIVED_MORPH_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    void        *clientData, \
    riscvInstrInfoP instrInfo, \
    Uns64       thisPC \
)
typedef RISC_V_DERIVED_MORPH_FN((*riscvDerivedMorphFn));

typedef struct riscvExtCBS {
    riscvDerivedMorphFn    postMorph;
} riscvExtCB;
```

Description

This interface function is called after an instruction is translated by the base model. It allows the extension object to emit code to perform an action that should be done after any standard instruction. The function is passed the address of the instruction (*thisPC*) and a pointer to a structure giving information about the decoded instruction (*instrInfo*); see section 16 for more information about this structure.

Example

This example shows how this function is used in the `andes.ovpworld.org` model. In this model, modeling hardware stack protection (HSP) requires that the current stack pointer is checked after each instruction executes to detect illegal stack pointer changes. There is also a requirement to revert unaligned access behavior to a default state:

```
static void doHSPCheck(riscvP riscv, vmiosObjectP object) {
    . . . check SP against base and limit, maybe take exception . . .
}

void andesCheckHSP(riscvP riscv, vmiosObjectP object) {

    if(!isHSPEnabledMT(riscv, object)) {

        // no action if feature is currently disabled

    } else if(riscv->writtenXMask & (1<<RV_REG_X_SP)) {

        // check is required only if the instruction updates SP
        vmimtArgProcessor();
        vmimtArgNatAddress(object);
        vmimtCallAttrs((vmiCallFn)doHSPCheck, VMCA_NA);
    }
}

RISC_V_DERIVED_MORPH_FN(andesPostMorph) {

    vmiosObjectP object = clientData;

    // handle hardware stack protection if required
    if(RD_ACSR_FIELD(object, mmisc_cfg, HSP)) {
        andesCheckHSP(riscv, object);
    }

    // reset default unaligned access behavior (required for PMP updates to
    // be correctly handled)
    riscv->configInfo.unaligned = True;
}

static VMIOS_CONSTRUCTOR_FN(constructor) {
```

```
. . . lines omitted . . .  
object->extCB.postMorph      = andesPostMorph;  
. . . lines omitted . . .  
}
```

15.21 Function AMOCheck

```
#define RISC_V_DERIVED_MORPH_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    void        *clientData, \
    riscvInstrInfoP instrInfo, \
    Uns64       thisPC \
)
typedef RISC_V_DERIVED_MORPH_FN((*riscvDerivedMorphFn));

typedef struct riscvExtCBS {
    riscvDerivedMorphFn    AMOCheck;
} riscvExtCB;
```

Description

This interface function is called *before* any code is emitted for an atomic memory access (AMO) instruction. It allows the extension object to emit code to validate the legality of the AMO instruction in a custom manner if required. The custom behavior will be performed before any other behavior of that instruction. The function is passed the address of the instruction (`thisPC`) and a pointer to a structure giving information about the decoded instruction (`instrInfo`); see section 16 for more information about this structure.

See also interface function `AMOMorph`, which allows custom behavior to be specified at a later stage in execution of AMO instructions.

Example

This example shows how this function can be used to cause all AMO instructions to take a custom exception if a custom CSR enable bit `custctl.AMOEN` is zero.

```
static void illegalCustom(
    vmiosObjectP object,
    riscvException exception,
    const char *reason
) {
    riscvP riscv = object->riscv;
    riscv->cb.illegalCustom(riscv, exception, reason);
}

void customAMOCheck(vmiosObjectP object) {
    if(!RD_XCSR_FIELD(object, custctl, AMOEN)) {
        illegalCustom(object, custom_E_IllegalInstruction, "Custom AMO disable");
    }
}

RISC_V_DERIVED_MORPH_FN(customAMOCheck) {
    vmimtArgNatAddress(clientData);
    vmimtCallAttrs((vmiCallFn)customAMOCheck, VMCA_NA);
}

static VMIOS_CONSTRUCTOR_FN(constructor) {
    . . . lines omitted . . .
    object->extCB.AMOCheck = customAMOCheck;
    . . . lines omitted . . .
}
```


15.22 Function AMOMorph

```
#define RISC_V_DERIVED_MORPH_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    void        *clientData, \
    riscvInstrInfoP instrInfo, \
    Uns64        thisPC \
)
typedef RISC_V_DERIVED_MORPH_FN((*riscvDerivedMorphFn));

typedef struct riscvExtCBS {
    riscvDerivedMorphFn      AMOMorph;
} riscvExtCB;
```

Description

This interface function allows an extension object to emit code that is executed *after* permission checks have been performed for an atomic memory access (AMO) instruction, but *before* any of the accesses specified for that instruction have been performed. It allows the extension object to modify the behavior of the AMO instruction in a custom manner if required. The function is passed the address of the instruction (`thisPC`) and a pointer to a structure giving information about the decoded instruction (`instrInfo`); see section 16 for more information about this structure.

See also interface function `AMOCheek`, which allows custom behavior to be specified at an earlier stage in execution of AMO instructions.

Example

This example shows how this function can be used to cause all AMO instructions to take a custom exception if the permission checks succeed. The effect will be that AMO instructions preferentially take any access exceptions implied by their behavior, and then, if there are no exceptions generated as a result, a custom trap is taken.

```
static void illegalCustom(
    vmiosObjectP  object,
    riscvException exception,
    const char    *reason
) {
    riscvP riscv = object->riscv;
    riscv->cb.illegalCustom(riscv, exception, reason);
}

void customAMOException(vmiosObjectP object) {
    illegalCustom(object, custom_E_IllegalInstruction, "Custom AMO emulation");
}

RISC_V_DERIVED_MORPH_FN(customAMOMorph) {

    vmimtArgNatAddress(clientData);
    vmimtCallAttrs((vmiCallFn)customAMOException, VMCA_EXCEPTION);
}

static VMIOES_CONSTRUCTOR_FN(constructor) {

    . . . lines omitted . . .
    object->extCB.AMOMorph      = customAMOMorph;
    . . . lines omitted . . .
}
```

15.23 Function *emitCSRCheck*

```
#define RISC_V_EMIT_CSR_CHECK_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    void        *clientData, \
    riscvInstrInfoP instrInfo, \
    Bool        isRead, \
    Bool        isWrite, \
    Bool        useRS1 \
)
typedef RISC_V_EMIT_CSR_CHECK_FN((*riscvEmitCSRCheckFn));

typedef struct riscvExtCBS {
    riscvEmitCSRCheckFn    emitCSRCheck;
} riscvExtCB;
```

Description

Usually, CSR access constraints can be cleanly handled when a CSR is defined (as shown in section 7). Sometimes, however, CSRs have custom access constraints that cannot be described in the standard definition. For example, a CSR may permit access using `csrrw` instructions but deny access using `csrrs` or `csrrc` instructions. In such cases, the custom *CSR check* interface function can be used. This function is called before any behavioral code is emitted to implement the CSR and allows the extension to emit higher-priority checks that cause an exception to be taken if access constraints are not satisfied.

The function is passed the following arguments, in addition to the `riscv` and `clientData`:

1. `instrInfo`: this is a pointer to a structure giving information about the decoded CSR access instruction - see section 16 for more information;
2. `isRead`: this indicates whether the instruction is reading the CSR.
3. `isWrite`: this indicates whether the instruction is writing the CSR.
4. `useRS1`: if `True`, this indicates the instruction is writing the CSR and taking the value to write from the GPR specified in the `rs1` field of the instruction.

CSR-related fields from the `instrInfo` structure can be extracted for use in this function as follows:

```
// CSR index number
Uns32 csr = instrInfo->csr;

// CSR update semantics:
//   RV_CSR_RW : read/write
//   RV_CSR_RS : read/set
//   RV_CSR_RC : read/clear
riscvCSRUDesc csrUpdate = instrInfo->csrUpdate;

// whether constant is used (CSRRI)
Bool useC = (instrInfo->type==RV_IT_CSRRI_I);

// constant value (if CSRRI)
Uns32 c = instrInfo->c;
```

To clarify parameter and `instrInfo` field use with different classes of instruction, the following table shows values that will be used for different types of CSR access instruction.

| Disassembly | Alias | csrUpdate | isRead | isWrite | useRS1 | useC |
|-------------------------------|---------------------------|------------------------|--------|---------|--------|-------|
| <code>csrrw x0,satp,x0</code> | <code>csrw satp,x0</code> | <code>RV_CSR_RW</code> | False | True | False | False |
| <code>csrrw x1,satp,x0</code> | | <code>RV_CSR_RW</code> | True | True | False | False |
| <code>csrrw x0,satp,x1</code> | <code>csrw satp,x1</code> | <code>RV_CSR_RW</code> | False | True | True | False |
| <code>csrrw x1,satp,x2</code> | | <code>RV_CSR_RW</code> | True | True | True | False |
| <code>csrrs x0,satp,x0</code> | <code>csrr x0,satp</code> | <code>RV_CSR_RS</code> | True | False | False | False |
| <code>csrrs x1,satp,x0</code> | <code>csrr x1,satp</code> | <code>RV_CSR_RS</code> | True | False | False | False |
| <code>csrrs x0,satp,x1</code> | <code>csrs satp,x1</code> | <code>RV_CSR_RS</code> | True | True | True | False |
| <code>csrrs x1,satp,x2</code> | | <code>RV_CSR_RS</code> | True | True | True | False |
| <code>csrrc x0,satp,x0</code> | <code>csrc satp,x0</code> | <code>RV_CSR_RC</code> | True | False | False | False |
| <code>csrrc x1,satp,x0</code> | | <code>RV_CSR_RC</code> | True | False | False | False |
| <code>csrrc x0,satp,x1</code> | <code>csrc satp,x1</code> | <code>RV_CSR_RC</code> | True | True | True | False |
| <code>csrrc x1,satp,x2</code> | | <code>RV_CSR_RC</code> | True | True | True | False |
| <code>csrrwi x0,satp,0</code> | <code>csrwi satp,0</code> | <code>RV_CSR_RW</code> | False | True | False | True |
| <code>csrrwi x1,satp,0</code> | | <code>RV_CSR_RW</code> | True | True | False | True |
| <code>csrrwi x0,satp,1</code> | <code>csrwi satp,1</code> | <code>RV_CSR_RW</code> | False | True | False | True |
| <code>csrrwi x1,satp,1</code> | | <code>RV_CSR_RW</code> | True | True | False | True |
| <code>csrrsi x0,satp,0</code> | <code>csrsi satp,0</code> | <code>RV_CSR_RS</code> | True | False | False | True |
| <code>csrrsi x1,satp,0</code> | | <code>RV_CSR_RS</code> | True | False | False | True |
| <code>csrrsi x0,satp,1</code> | <code>csrsi satp,1</code> | <code>RV_CSR_RS</code> | True | True | False | True |
| <code>csrrsi x1,satp,1</code> | | <code>RV_CSR_RS</code> | True | True | False | True |
| <code>csrrci x0,satp,0</code> | <code>csrci satp,0</code> | <code>RV_CSR_RC</code> | True | False | False | True |
| <code>csrrci x1,satp,0</code> | | <code>RV_CSR_RC</code> | True | False | False | True |
| <code>csrrci x0,satp,1</code> | <code>csrci satp,1</code> | <code>RV_CSR_RC</code> | True | True | False | True |
| <code>csrrci x1,satp,1</code> | | <code>RV_CSR_RC</code> | True | True | False | True |

Example

This example shows how this function can be used to deny all accesses to a set of CSRs unless those accesses use the `csrrw` instruction variant with `rs1` that is not `x0`.

```
static RISC_V_EMIT_CSR_CHECK_FN(emitCSRCheck) {
    Uns32      csr      = instrInfo->csr;
    riscvCSRUDesc csrUpdate = instrInfo->csrUpdate;

    if(!((csr==0xBF9) || (csr==0xBFA) || (csr==0xBFC))) {
        // no action
    } else if(isWrite && useRS1 && (csrUpdate==RV_CSR_RW)) {
        // valid csrrw write access using RS1
    } else {
        riscv->cb.morphIllegal(riscv, "illegal access");
    }
}

static VMIO_S_CONSTRUCTOR_FN(constructor) {
    . . . lines omitted . . .
    object->extCB.emitCSRCheck = emitCSRCheck;
    . . . lines omitted . . .
}
```

15.24 Function `unitStrideCheck`

```
#define RISC_V_UNIT_STRIDE_CHECK_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    void        *clientData, \
    riscvInstrInfoP instrInfo, \
    vmiReg      baseAddr, \
    Uns32       memBits, \
    Bool        isLoad \
)
typedef RISC_V_UNIT_STRIDE_CHECK_FN((*riscvUnitStrideCheckFn));

typedef struct riscvExtCBS {
    riscvUnitStrideCheckFn    unitStrideCheck;
} riscvExtCB;
```

Description

This interface function is called for Vector extension unit-stride load/store instructions. It allows the extension object to emit code that is executed *after* the validity of the load/store instruction has been verified, but *before* any of the accesses specified for that instruction have been performed. It allows the extension to insert extra custom checks on unit-stride load/store instructions if required.

Argument `instrInfo` gives information about the decoded vector instruction (see section 16 for more information about this structure), argument `baseAddr` is the base address of the load/store, argument `memBits` indicates the element size, and argument `isLoad` is `True` if the instruction is a load and `False` if it is a store.

Example

This example shows how this function can be used to perform an additional custom check that the entire range potentially accessed by a unit-stride load/store is accessible before any accesses are performed.

```
//
// Check vector unit-stride load
//
static void checkVLd(vmiosObjectP object, Uns64 addr, Uns32 memBits) {

    riscvP riscv = object->riscv;
    Uns32 vl      = RD_CSR32(riscv, vl);

    if(vl) {

        memDomainP domain = vmirtGetProcessorDataDomain((vmiProcessorP)riscv);
        Uns32 bytes = (vl*memBits)/8;

        // do try-load of required address
        vmirtReadNByteDomain(
            domain, addr, 0, bytes, &object->rcache, MEM_AA_TRUE
        );
    }
}

//
// Check vector unit-stride store
//
static void checkVSt(vmiosObjectP object, Uns64 addr, Uns32 memBits) {

    riscvP riscv = object->riscv;
    Uns32 vl      = RD_CSR32(riscv, vl);
```

```
if(vl) {

    memDomainP domain = vmirtGetProcessorDataDomain((vmiProcessorP)riscv);
    Uns32      bytes  = (vl*memBits)/8;

    // do try-store of required address
    vmirtWriteNByteDomain(
        domain, addr, 0, bytes, &object->rcache, MEM_AA_TRUE
    );
}

static RISC_V_UNIT_STRIDE_CHECK_FN(customUnitStrideCheck) {

    vmimtArgNatAddress(clientData);
    vmimtArgReg(VPRRAT_64, baseAddr);
    vmimtArgUns32(memBits);

    if(isLoad) {
        vmimtCall((vmiCallFn)checkVLd);
    } else {
        vmimtCall((vmiCallFn)checkVSt);
    }
}

static VMIO_S_CONSTRUCTOR_FN(constructor) {

    . . . lines omitted . . .
    object->extCB.unitStrideCheck = customUnitStrideCheck;
    . . . lines omitted . . .
}
```

15.25 Function *emitVFREDUSUM*

```
#define RISC_VFREDSUM_MORPH_FN(_NAME) void _NAME( \
    riscvP riscv, \
    void *clientData, \
    vmiReg vdAcc, \
    void *vs2, \
    void *vm, \
    Uns32 SEW \
)
typedef RISC_VFREDSUM_MORPH_FN((*riscvVFREDSUMMorphFn));

typedef struct riscvExtCBS {
    riscvVFREDSUMMorphFn    emitVFREDUSUM;
} riscvExtCB;
```

Description

This interface function is called for Vector extension floating point single-width reduction instructions (*vfredusum.vs*). The specification does not define precisely how these are implemented, so this enables a custom implementation to be provided by code emitted by an extended model. The base model performs all instruction validity checks, so these do not have to be done by the custom implementation. If no custom implementation is provided, the default behavior matches *vfredosum.vs*.

When the instruction is executed, argument *vdAcc* contains the scalar input value initially and should be written with the reduction result on completion. *vs2* is a pointer to the vector input argument and *vm* is a pointer to the mask argument, or *NULL* if the operation is unmasked. *SEW* indicates the operation width.

Example

This example shows how this function could be used to reimplement the default single-width floating point reduction operation using floating point emulation functions provided by the *VMI Run Time Function* API. Note that this is an example only, as this behavior is the default if the interface function is omitted.

```
//
// Return current rounding mode
//
static vmiFPRC getRC(riscvP riscv) {

    static const vmiFPRC map[8] = {
        [0] = vmi_FPR_NEAREST,
        [1] = vmi_FPR_ZERO,
        [2] = vmi_FPR_NEG_INF,
        [3] = vmi_FPR_POS_INF,
        [4] = vmi_FPR_AWAY,
    };

    return map[RD_CSR_FIELDC(riscv, fcsr, frm)];
}

//
// Return floating point type for the given bits
//
static vmiFType getFType(Uns32 fBits) {

    vmiFType result = 0;

    if(fBits==16) {
```

```

        result = vmi_FT_16_IEEE_754;
    } else if(fBits==32) {
        result = vmi_FT_32_IEEE_754;
    } else if(fBits==64) {
        result = vmi_FT_64_IEEE_754;
    }
}

return result;
}

//
// If the operation is masked, indicate whether the indexed element is enabled
//
static Bool elementEnabled(Uns8 *vm, Uns32 i) {

    Bool enabled = True;

    if(vm) {

        Uns32 byte = i/8;
        Uns32 mask = 1<<(i&7);

        enabled = vm[byte] & mask;
    }

    return enabled;
}

//
// Get floating point operation descriptor for add of the passed width
//
static vmiFPBinopDescCP getFAddDesc(riscvP riscv, Uns32 fBits, Uns32 num) {

    vmiProcessorP processor = (vmiProcessorP)riscv;
    vmiFType      fType      = getFType(fBits);
    vmiFPRC       rc         = getRC(riscv);

    return vmirtGetFBinopRRRDesc(processor, fType, num, vmi_FADD, 0, rc, False);
}

//
// Execute floating point binary operation
//
static void doFBinop(
    riscvP      riscv,
    vmiFPBinopDescCP opDesc,
    void        *rd,
    void        *rs1,
    void        *rs2
) {
    Uns8 flags = 0;

    vmirtFBinopSimdRRR((vmiProcessorP)riscv, opDesc, rd, rs1, rs2, &flags);

    riscv->fpFlagsMT |= flags;
}

//
// Custom VFREDUSUM implementation
//
static Uns64 doVFREDUSUM(
    vmiosObjectP object,
    Uns64        vdAcc,
    Uns8         *vs2,
    Uns8         *vm,
    Uns32        SEW,
    Uns32        SEWMul
) {
    riscvP riscv = object->riscv;
    Uns32  vl     = RD_CSRC(riscv, vl);
    Uns32  fBits  = SEW*SEWMul;

```

```

Uns32 fBytes = fBits/8;
Uns32 i;

// get operation description for single floating point add
vmiFPBinopDescCP opDescX1 = getFAddDesc(riscv, fBits, 1);

// handle all enabled elements
for(i=0; i<vl; i++) {

    if(elementEnabled(vm, i)) {

        Uns8 *vs2Elem = vs2+(i*fBytes);

        doFBinop(riscv, opDescX1, &vdAcc, &vdAcc, vs2Elem);

    }

}

return vdAcc;
}

//
// Common routine to emit callback for VFREDUSUM and VFWREDUSUM
//
static void emitVFREDUSUMCommon(
    vmiosObjectP object,
    vmiCallFn opCB,
    vmiReg vdAcc,
    void *vs2,
    void *vm,
    Uns32 SEW,
    Uns32 SEWMul
) {
    // extend temporary accumulator to 64 bits
    vmimtMoveExtendRR(64, vdAcc, SEW*SEWMul, vdAcc, False);

    // call custom implementation
    vmimtArgNatAddress(object);
    vmimtArgReg(64, vdAcc);
    vmimtArgNatAddress(vs2);
    vmimtArgNatAddress(vm);
    vmimtArgUns32(SEW);
    vmimtArgUns32(SEWMul);
    vmimtCallResultAttrs(opCB, 64, vdAcc, VMCA_NO_INVALIDATE);
}

static RISC_VFREDSUM_MORPH_FN(customVFREDUSUM) {

    vmiCallFn opCB = (vmiCallFn)doVFREDUSUM;

    emitVFREDUSUMCommon(clientData, opCB, vdAcc, vs2, vm, SEW, 1);

}

static VMIO_CONSTRUCTOR_FN(constructor) {

    . . . lines omitted . . .
    object->extCB.emitVFREDUSUM = customVFREDUSUM;
    . . . lines omitted . . .
}

```


15.26 Function *emitVFWREDUSUM*

```
#define RISC_VFREDSUM_MORPH_FN(_NAME) void _NAME( \
    riscvP riscv, \
    void *clientData, \
    vmiReg vdAcc, \
    void *vs2, \
    void *vm, \
    Uns32 SEW \
)
typedef RISC_VFREDSUM_MORPH_FN((*riscvVFREDSUMMorphFn));

typedef struct riscvExtCBS {
    riscvVFREDSUMMorphFn    emitVFWREDUSUM;
} riscvExtCB;
```

Description

This interface function is called for Vector extension floating point widening reduction instructions (*vfwredusum.vs*). The specification does not define precisely how these are implemented, so this enables a custom implementation to be provided by code emitted by an extended model. The base model performs all instruction validity checks, so these do not have to be done by the custom implementation. If no custom implementation is provided, the default behavior matches *vfredusum.vs*.

When the instruction is executed, argument *vdAcc* contains the scalar input value initially and should be written with the reduction result on completion. *vs2* is a pointer to the *unwidened* vector input argument and *vm* is a pointer to the mask argument, or *NULL* if the operation is unmasked. *SEW* indicates the operation width.

Example

This example shows how this function could be used to reimplement the default widening floating point reduction operation using floating point emulation functions provided by the *VMI Run Time Function* API. Note that this is an example only, as this behavior is the default if the interface function is omitted. Some of the implementation is shared with the *vfredusum.vs* example (section 15.25) so common functions are omitted here.

```
//
// Get floating point operation descriptor for widening conversion
//
static vmiFPConvertDescCP getFCvtDesc(
    riscvP riscv,
    Uns32 srcBits,
    Uns32 dstBits
) {
    vmiProcessorP processor = (vmiProcessorP)riscv;
    vmiFType srcType = getFType(srcBits);
    vmiFType dstType = getFType(dstBits);
    vmiFPRC rc = vmi_FPR_NEAREST;

    return vmirtGetFConvertRRDesc(processor, 1, dstType, srcType, 0, rc, False);
}

//
// Execute floating point widening conversion
//
static void doFConvert(
    riscvP riscv,
    vmiFPConvertDescCP opDesc,
```

```

        void          *rd,
        void          *rs
    ) {
        Uns8 flags = 0;

        vmirtFConvertSimdRR((vmiProcessorP)riscv, opDesc, rd, rs, &flags);

        riscv->fpFlagsMT |= flags;
    }

//
// Custom VFWREDUSUM implementation
//
static Uns64 doVFWREDUSUM(
    vmiosObjectP object,
    Uns64        vdAcc,
    Uns8         *vs2,
    Uns8         *vm,
    Uns32        SEW,
    Uns32        SEWMul
) {
    riscvP riscv = object->riscv;
    Uns32 vl = RD_CSRC(riscv, vl);
    Uns32 srcBits = SEW;
    Uns32 dstBits = SEW*SEWMul;
    Uns32 srcBytes = srcBits/8;
    Uns32 dstBytes = dstBits/8;
    Uns8 vs2Tmp[vl*dstBytes];
    Uns32 i;

    // get operation description for single widening conversion
    vmiFPConvertDescCP opDescXl = getFCvtDesc(riscv, srcBits, dstBits);

    // widen all enabled elements
    for(i=0; i<vl; i++) {

        if(elementEnabled(vm, i)) {

            Uns8 *src = vs2+(i*srcBytes);
            Uns8 *dst = &vs2Tmp[i*dstBytes];

            doFConvert(riscv, opDescXl, dst, src);

        }
    }

    // use common algorithm with widened source
    return doVFREDUSUM(object, vdAcc, vs2Tmp, vm, SEW, SEWMul);
}

static RISC_V_FREDSUM_MORPH_FN(customVFWREDUSUM) {

    vmiCallFn opCB = (vmiCallFn)doVFWREDUSUM;

    emitVFREDUSUMCommon(clientData, opCB, vdAcc, vs2, vm, SEW, 2);
}

static VMIO_CONSTRUCTOR_FN(constructor) {

    . . . lines omitted . . .
    object->extCB.emitVFWREDUSUM = customVFWREDUSUM;
    . . . lines omitted . . .
}

```

15.27 *Function switchCB*

```
#define RISC_V_IASSWITCH_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    vmiIASRunState state, \
    void        *clientData \
)
typedef RISC_V_IASSWITCH_FN(*riscvIASSwitchFn);

typedef struct riscvExtCBS {
    riscvIASSwitchFn      switchCB;
} riscvExtCB;
```

Description

This interface function is called when a processor has been scheduled and is about to be run, or has been descheduled and is about to stop running, in a multiprocessor simulation. It allows the extension object to make any state changes required at those points.

Example

This example shows how this function is used in the transactional memory extension (tmExtensions) in the vendor.com template model. In this model, modeling transactional memory requires that memory watchpoints are placed on all active cache lines when a processor is descheduled so that conflicting writes to those lines by another processor can be detected:

```
static void installCacheMonitor(vmiosObjectP object) {

    if (object->tmStatus != TM_OK) {

        // Ignore if not in transaction or there
        // has already been an abort event

    } else {

        memDomainP domain = object->physicalMem;
        cacheLineP this;
        Uns32      i;

        for(i = 0; i < object->numPending; i++) {

            this = &object->pending[i];

            Addr lo = getLineLowPA(this);
            Addr hi = getLineHighPA(this);

            // register for a callback if any current line is written
            if (!this->readCallbackInstalled) {
                vmirtAddWriteCallback(domain, 0, lo, hi, memoryConflict, object);
                this->readCallbackInstalled = True;
            }

            // register for a callback if any dirty line is read
            if (this->dirty && !this->writeCallbackInstalled) {
                vmirtAddReadCallback(domain, 0, lo, hi, memoryConflict, object);
                this->writeCallbackInstalled = True;
            }

        }

    }

}

static RISC_V_IASSWITCH_FN(riscvSwitch) {
```

```
vmiosObjectP object = clientData;

if(state==RS_SUSPEND) {
    installCacheMonitor(object);
}

static VMIO_CONSTRUCTOR_FN(tmConstructor) {

    . . . lines omitted . . .
    object->extCB.switchCB    = riscvSwitch;
    . . . lines omitted . . .
}
```

15.28 Function *tLoad*

```
#define RISC_V_TLOAD_FN(_NAME) void _NAME( \
    riscvP riscv, \
    void *buffer, \
    Addr VA, \
    Uns32 bytes, \
    void *clientData \
)
typedef RISC_V_TLOAD_FN((*riscvTLoadFn));

typedef struct riscvExtCBS {
    riscvTLoadFn tLoad;
} riscvExtCB;
```

Description

For an extension implementing transactional memory, this function is called whenever a load is performed and transactional memory mode is enabled (base model interface function `setTMode` in the base model has been called with `enable` of `True` – see section 14.9). The function should implement a transactional load, typically by modeling active cache lines.

Example

This example shows how this function is used in the transactional memory extension (`tmExtensions`) in the `vendor.com` template model:

```
static RISC_V_TLOAD_FN(riscvTLoad) {
    vmiosObjectP object = clientData;

    if (object->tmStatus != TM_OK) {
        // abort is pending
        doAbort(object);
    } else {
        Uns8 *buffer8 = buffer;
        Addr PA;
        Uns32 thisBytes;
        Uns32 i;

        while ((thisBytes = bytes) > 0) {
            if(mapVAToPA(object, VA, &PA)) {
                // get pointer to data in cache (loads cache line if required)
                Uns8 *data = getCacheAddress(object, PA, &thisBytes, False);

                if (!data) {
                    object->tmStatus |= TM_ABORT_OVERFLOW;
                    doAbort(object);
                } else {
                    for(i=0; i<thisBytes; i++) {
                        buffer8[i] = data[i];
                    }
                }
            }
            // reduce byte count by count of bytes on this line
        }
    }
}
```

```
        bytes -= thisBytes;
        VA    += thisBytes;
    }
}

static VMIO_CONSTRUCTOR_FN(tmConstructor) {
    . . . lines omitted . . .
    object->extCB.tLoad      = riscvTLoad;
    . . . lines omitted . . .
}
```

15.29 Function *tStore*

```
#define RISC_V_TSTORE_FN(_NAME) void _NAME( \  
    riscvP      riscv,      \  
    const void *buffer,      \  
    Addr        VA,          \  
    Uns32       bytes,       \  
    void        *clientData  \  
    ) \  
typedef RISC_V_TSTORE_FN((*riscvTStoreFn)); \  
 \  
typedef struct riscvExtCBS { \  
    riscvTStoreFn      tStore; \  
} riscvExtCB;
```

Description

For an extension implementing transactional memory, this function is called whenever a store is performed and transactional memory mode is enabled (base model interface function `setTMode` in the base model has been called with `enable` of `True` – see section 14.9). The function should implement a transactional store, typically by modeling active cache lines.

Example

This example shows how this function is used in the transactional memory extension (`tmExtensions`) in the `vendor.com` template model:

```
static RISC_V_TSTORE_FN(riscvTStore) { \  
    vmiosObjectP object = clientData; \  
    if (object->tmStatus != TM_OK) { \  
        // abort is pending \  
        doAbort(object); \  
    } else { \  
        const Uns8 *buffer8 = buffer; \  
        Addr        PA; \  
        Uns32       thisBytes; \  
        Uns32       i; \  
        while ((thisBytes = bytes) > 0) { \  
            if(mapVAToPA(object, VA, &PA)) { \  
                // get pointer to data in cache (loads cache line if required) \  
                Uns8 *data = getCacheAddress(object, PA, &thisBytes, True); \  
                if (!data) { \  
                    object->tmStatus |= TM_ABORT_OVERFLOW; \  
                    doAbort(object); \  
                } else { \  
                    for(i=0; i<thisBytes; i++) { \  
                        data[i] = buffer8[i]; \  
                    } \  
                } \  
            } \  
            // reduce byte count by count of bytes on this line
```

```
        bytes -= thisBytes;
        VA     += thisBytes;
    }
}

static VMIOS_CONSTRUCTOR_FN(tmConstructor) {
    . . . lines omitted . . .
    object->extCB.tStore      = riscvTStore;
    . . . lines omitted . . .
}
```


15.30 *Function `distinctPhysMem`*

```
#define RISC_V_DISTINCT_PHYS_MEM_FN(_NAME) Bool _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_DISTINCT_PHYS_MEM_FN((*riscvDistinctPhysMemFn));

typedef struct riscvExtCBS {
    riscvDistinctPhysMemFn    distinctPhysMem;
} riscvExtCB;
```

Description

This interface function is required when an extension implements physical memory map modifications (for example, by installing custom local memories into the physical memory hierarchy – see section 15.31 for more details). The function should return a Boolean indicating whether, as a result of these custom physical memory changes, code and data domain views of memory are different. For example, the function should return `True` if a model implements Code and Data local memories where the Data local memory is not visible in the instruction address space.

Example

The Andes extended model implements Instruction and Data local memories in the physical address space, and the Data local memory is not visible in the instruction address space. The function below is used to indicate that separate code and data domain views are required if either local memory is present.

```
//
// Are separate code and data physical memory view required?
//
RISC_V_DISTINCT_PHYS_MEM_FN(andesDistinctPhysicalMem) {

    vmiosObjectP object = clientData;

    // determine whether either local memory is present
    Uns32 ILMSize = createLocalMemory(object, object->csr.micm_cfg, True);
    Uns32 DLMSize = createLocalMemory(object, object->csr.mdcn_cfg, False);

    // separate view is required if either LM is present
    return ILMSize || DLMSize;
}
```

15.31 *Function installPhysMem*

```
#define RISC_V_PHYS_MEM_FN(_NAME) void _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_PHYS_MEM_FN((*riscvPhysMemFn));

typedef struct riscvExtCBS {
    riscvPhysMemFn      installPhysMem;
} riscvExtCB;
```

Description

This interface function allows an extension to modify the physical memory map of a processor. It is called once the default physical memory constructor has been called. Typically, the function will create new memory domain (`memDomainP`) objects and replace the default physical memory domains with those (see lines in bold in the following example 1). This is typically required if, for example, the extended model has components such as local memories that reside at fixed locations in the physical memory map.

If a processor implements custom physical memory attributes (PMA), this function can also be used to specify any initial PMA constraints (the default is that no constraints are imposed initially).

Example 1

The Andes extended model implements Instruction and Data local memories in the physical address space. The details of the operation of these are quite complex; code below shows how the standard physical memory domains are replaced with Andes-specific domains by the `installPhysMem` callback. Refer to the source of the Andes model for more information on the implementation of local memories.

```
//
// Create local memory physical alias domain for the given mode and access type
//
static void newLMDomain(vmiosObjectP object, riscvMode mode, Bool isCode) {

    riscvP      riscv = object->riscv;
    memDomainP base  = riscv->physDomains[mode][isCode];
    Uns32      bits  = vmirtGetDomainAddressBits(base);

    // create local memory domain
    memDomainP result = createLMDomain(mode, bits, isCode);

    // initially make it an alias of the physical domain
    vmirtAliasMemory(base, result, 0, getAddressMask(bits), 0, 0);

    // replace physical domain
    riscv->physDomains[mode][isCode] = result;
}

//
// Install local memory domains in the domain hierarchy
//
static void installLocalMemoryDomains(vmiosObjectP object) {

    riscvP      riscv = object->riscv;
    riscvMode mode;
```

```

// create physical domain aliases for all non-User modes
for(mode=RISCV_MODE_S; mode<=RISCV_MODE_M; mode++) {

    memDomainP dataDomain = riscv->physDomains[mode][0];
    memDomainP codeDomain = riscv->physDomains[mode][1];

    // save current physical domain to allow aliasing to it later
    object->physDomains[mode][0] = dataDomain;
    object->physDomains[mode][1] = codeDomain;

    if(dataDomain) {
        newLMDomain(object, mode, False);
        newLMDomain(object, mode, True);
    }
}

// use Supervisor aliases for User mode
riscv->physDomains[RISCV_MODE_U][0] = riscv->physDomains[RISCV_MODE_S][0];
riscv->physDomains[RISCV_MODE_U][1] = riscv->physDomains[RISCV_MODE_S][1];
}

//
// Install ILM/DLM domains if required
//
RISCV_PHYS_MEM_FN(andesInstallPhysicalMem) {

    vmiosObjectP object = clientData;

    // create local memories if required
    Uns32 ILMSize = createLocalMemory(object, object->csr.micm_cfg, True);
    Uns32 DLMSize = createLocalMemory(object, object->csr.mdcn_cfg, False);

    // if either local memory is present, insert ILM domains into the memory
    // domain hierarchy
    if(ILMSize || DLMSize) {
        installLocalMemoryDomains(object);
    }

    // enable ILM if required
    if(RD_XCSR_FIELD(object, milmb, EN)) {
        updateLocalMemory(object, True, True);
    }

    // enable DLM if required
    if(RD_XCSR_FIELD(object, mdlmb, EN)) {
        updateLocalMemory(object, False, True);
    }
}

```

Example 2

The OpenHardware CV32E40X extended model requires that all PMA access is initially denied (the model is structured so that PMA mappings are created on demand using configured PMA range definitions). This is implemented as follows:

```

//
// Set privileges in the given domain set
//
static void setDomainPriv(
    riscvP      riscv,
    memDomainPP domains,
    Uns64       low,
    Uns64       high,
    memPriv     priv,
    const char *name,
    Bool        verbose
) {
    memDomainP dataDomain = domains[0];

```

```
memDomainP codeDomain = domains[1];

if(dataDomain==codeDomain) {

    // set permissions in unified domain
    vmirtProtectMemory(dataDomain, low, high, priv, MEM_PRIV_SET);

} else {

    // get privileges for data and code domains
    memPriv privRW = priv&MEM_PRIV_RW ? priv&~MEM_PRIV_X : MEM_PRIV_NONE;
    memPriv privX  = priv&MEM_PRIV_X  ? priv&~MEM_PRIV_RW : MEM_PRIV_NONE;

    // set permissions in data domain
    vmirtProtectMemory(dataDomain, low, high, privRW, MEM_PRIV_SET);

    // set permissions in code domain
    vmirtProtectMemory(codeDomain, low, high, privX, MEM_PRIV_SET);
}
}

//
// Set privileges in PMA domain
//
static void setPMAPriv(
    riscvP riscv,
    Uns64   low,
    Uns64   high,
    memPriv priv,
    Bool    verbose
) {
    setDomainPriv(riscv, &riscv->pmaDomains[0], low, high, priv, "PMA", verbose);
}

//
// Handle PMA initialization
//
static RISCVP_PHYS_MEM_FN(openhwInstallPhysicalMem) {

    // remove all PMA permissions
    setPMAPriv(riscv, 0, -1, MEM_PRIV_NONE, False);
}
```

15.32 *Function PMPPriv*

```
#define RISCVPMP_PRIV_FN(_NAME) memPriv _NAME( \
    riscvP riscv, \
    memPriv priv, \
    Uns32 regionIndex, \
    void *clientData \
)
typedef RISCVPMP_PRIV_FN(*riscvPMPPrivFn);

typedef struct riscvExtCBS {
    riscvPMPPrivFn      PMPPriv;
} riscvExtCB;
```

Description

This interface function allows an extension to modify the memory privileges of a PMP region in a custom manner. When a new PMP mapping is being established, this function is called with the region index and default access permissions; it returns the final access permissions after applying custom behavior.

Example

This example shows how alignment might be forced using extension CSRs with a one-to-one mapping to the base model PMP CSRs:

```
static RISCVPMP_PRIV_FN(updatePMAPriv) {
    vmiosObjectP object = clientData;
    Uns32 CCA = object->pmacfg.u8[regionIndex] & 7;

    // if CCA!=0, accesses in this region must be aligned
    if(CCA) {
        priv |= MEM_PRIV_ALIGN;
    }

    return priv;
}

static VMIO_CONSTRUCTOR_FN(constructor) {
    . . . lines omitted . . .
    object->extCB.PMAPriv = updatePMAPriv;
    . . . lines omitted . . .
}
```

15.33 *Function PMAEnable*

```
#define RISCVPMA_ENABLE_FN(_NAME) Bool _NAME( \
    riscvp riscv, \
    void *clientData \
)
typedef RISCVPMA_ENABLE_FN(*riscvpMAEnableFn);

typedef struct riscvExtCBS {
    riscvpMAEnableFn      PMAEnable;
} riscvExtCB;
```

Description

For an extension implementing physical memory attributes (PMA) when maximum PMA page size is restricted by configuration parameter `PMP_max_page`, this interface function allows the extension to indicate whether PMA is enabled. If so, the `PMACheck` interface function is used to perform any memory mappings needed to model PMA for a given address range.

Example

The `andes.ovpworld.org` model uses this function to specify whether PMA is enabled:

```
RISCVPMA_ENABLE_FN(andesPMAEnable) {
    vmiosObjectP object = clientData;

    return RD_XCSR_FIELD(object, mmio_cfg, DPMA);
}

static VMIO_CONSTRUCTOR_FN(constructor) {
    . . . lines omitted . . .
    object->extCB.PMAEnable = andesPMAEnable;
    . . . lines omitted . . .
}
```

See section 12 for further details about this function.

15.34 *Function PMACheck*

```
#define RISCVPMA_CHECK_FN(_NAME) void _NAME( \
    riscvP riscv, \
    memPriv requiredPriv, \
    Uns64 lowPA, \
    Uns64 highPA, \
    void *clientData \
)
typedef RISCVPMA_CHECK_FN((*riscvPMACheckFn));

typedef struct riscvExtCBS {
    riscvPMACheckFn PMACheck;
} riscvExtCB;
```

Description

For an extension implementing physical memory attributes (PMA), this interface function allows the extension to perform any memory mappings needed to model PMA for an address range. The function is called with required privileges and an address range. If the address range can be accessed legally, permissions on the memory domain objects referenced in the `pmaDomains` field of the base model should be updated to enable the access; otherwise, privileges in these domains should be left unaltered.

Example

See section 12 for a detailed example using this function.

15.35 *Function `validPTE`*

```
#define RISC_V_VALID_PTE_FN(_NAME) Bool _NAME( \
    riscvP      riscv,      \
    riscvTLBId  id,         \
    riscvVAMode  vaMode,     \
    Uns64        PTE,        \
    void         *clientData \
)
typedef RISC_V_VALID_PTE_FN((*riscvValidPTEFn));

typedef struct riscvExtCBS {
    riscvValidPTEFn    validPTE;
} riscvExtCB;
```

Description

When virtual memory is active, this function allows an extension to perform custom checks for PTE validity. It should return `True` if the given page table entry is valid and `False` otherwise. This validity check is performed in addition to the standard checks performed by the base model and described in the *Privileged Architecture Specification*.

15.36 *Function VMTrap*

```
#define RISCVM_TRAP_FN(_NAME) riscvException _NAME( \
    riscvP      riscv, \
    riscvTLBId id, \
    memPriv     requiredPriv, \
    Uns64       VA, \
    void        *clientData \
)
typedef RISCVM_TRAP_FN((*riscvVMTrapFn));

typedef struct riscvExtCBS {
    riscvVMTrapFn      VMTrap;
} riscvExtCB;
```

Description

One RISC-V implementation choice is to implement virtual memory TLB updates using a trap to a Machine mode handler. In this case, memory mappings will typically be modified by writes to custom CSRs, or a similar mechanism.

When virtual memory is enabled, this notifier is called to indicate an address lookup for an address for which there is currently no valid mapping in the TLB. It gives the extension the opportunity to initiate an implementation-specific trap to handle the address. The TLB that is affected is specified by the `tlbId` enumeration:

```
typedef enum riscvTLBIdE {
    RISCVM_TLB_HS,      // HS TLB
    RISCVM_TLB_VS1,     // VS stage 1 virtual TLB
    RISCVM_TLB_VS2,     // VS stage 2 virtual TLB
} riscvTLBId
```

(`RISCVM_TLB_VS1` and `RISCVM_TLB_VS2` TLBs are only used for processors that implement the Hypervisor extension.)

Example

This example shows how a failing address lookup can cause one of six custom exceptions:

```
typedef enum custExceptionE {
    cust_InstTLBMiss      = 24,
    cust_LoadTLBMiss      = 25,
    cust_StoreTLBMiss     = 27,
    cust_GuestInstTLBMiss = 28,
    cust_GuestLoadTLBMiss = 29,
    cust_GuestStoreTLBMiss = 31,
} custException;

RISCVM_TRAP_FN(custVMTrap) {
    vmiosObjectP object = clientData;
    Bool          S2     = (id==RISCVM_TLB_VS2);
    riscvException result = 0;

    if(!object->config.software_table_walk) {
        // use standard hardware page table walk
    } else if(requiredPriv & MEM_PRIV_R) {
        result = S2 ? cust_GuestLoadTLBMiss : cust_LoadTLBMiss;
    }
```

```
    } else if(requiredPriv & MEM_PRIV_W) {  
        result = S2 ? cus _GuestStoreTLBMiss : cust_StoreTLBMiss;  
    } else {  
        result = S2 ? cust_GuestInstTLBMiss : cust_InstTLBMiss;  
    }  
  
    return result;  
}  
  
static VMIO_CONSTRUCTOR_FN(extConstructor) {  
  
    . . . lines omitted . . .  
    object->extCB.VMTrap          = custVMTrap;  
    . . . lines omitted . . .  
}
```

15.37 Function *setDomainNotifier*

```
#define RISC_V_SET_DOMAIN_NOTIFIER_FN(_NAME) void _NAME( \
    riscvP      riscv, \
    memDomainPP domainP, \
    void        *clientData \
)
typedef RISC_V_SET_DOMAIN_NOTIFIER_FN((*riscvSetDomainNotifierFn));

typedef struct riscvExtCBS {
    riscvSetDomainNotifierFn setDomainNotifier;
} riscvExtCB;
```

Description

Some CSR settings affect the access mode for load and store instructions in Machine mode. For example, `mstatus.MPRV=1` can cause load and store instructions to be executed in Supervisor or User modes instead of Machine mode. This feature of the RISC-V architecture is implemented in the model by modifying the *memory domain* to which loads and stores are routed.

This notifier is called after the base model has selected the active load/store memory domain for the current processor mode. It gives the extension the opportunity to modify the choice of domain if required. The choice of domain is indicated by updating the `domainP` by-reference argument, typically with one of the PMP domains or the guest page table walk domain from the base model.

Example

This example shows how the load/store domain might be affected in a processor that implements TLB updates using a Machine mode trap. In this processor, when `cstatus.MTW=1`, all loads and stores should be performed with table walk privilege:

```
RISC_V_SET_DOMAIN_NOTIFIER_FN(custSetDomain) {

    riscvMode    mode    = getCurrentMode5(riscv);
    vmiosObjectP object = clientData;

    if(mode!=RISC_V_MODE_M) {
        // no action unless in Machine mode
    } else if(!RD_XCSR_FIELD(object, cstatus, MTW)) {
        // Machine Table Walk not enabled
    } else if(!RD_CSR_FIELD64(riscv, mstatus, GVA)) {
        *domainP = riscv->pmpDomains[RISC_V_MODE_S][False];
    } else if(!RD_CSR_FIELD_S(riscv, hgatp, MODE)) {
        *domainP = riscv->pmpDomains[RISC_V_MODE_S][False];
    } else {
        *domainP = riscv->guestPTWDomain;
    }
}

static VMIOS_CONSTRUCTOR_FN(extConstructor) {

    . . . lines omitted . . .
    object->extCB.setDomainNotifier = custSetDomain;
    . . . lines omitted . . .
}
```

15.38 *Function freeEntryNotifier*

```
#define RISC_V_FREE_ENTRY_NOTIFIER_FN(_NAME) void _NAME( \
    riscvP      riscv,          \
    riscvTLBId id,              \
    Uns16       entryId,        \
    void        *clientData     \
)
typedef RISC_V_FREE_ENTRY_NOTIFIER_FN((*riscvFreeEntryNotifierFn));

typedef struct riscvExtCBS {
    riscvFreeEntryNotifierFn freeEntryNotifier;
} riscvExtCB;
```

Description

One RISC-V implementation choice is to implement virtual memory TLB updates using a trap to a Machine mode handler. In this case, memory mappings will typically be modified by writes to custom CSRs, or a similar mechanism.

This notifier is called to indicate that the base model is deleting a cached entry in a TLB. It gives the extension the opportunity to update data structures to make them consistent with removal of that entry. The entry that is being removed is indicated by the unique `entryId` parameter, which was supplied when the TLB entry was created (see section 14.59). The TLB that is affected is specified by the `tlbId` enumeration:

```
typedef enum riscvTLBIdE {

    RISC_V_TLB_HS,          // HS TLB
    RISC_V_TLB_VS1,         // VS stage 1 virtual TLB
    RISC_V_TLB_VS2,         // VS stage 2 virtual TLB

} riscvTLBId
```

(RISC_V_TLB_VS1 and RISC_V_TLB_VS2 TLBs are only used for processors that implement the Hypervisor extension.)

Example

This template shows how the notifier might be used to mark an entry in an array of implementation-specific entries as uninstalled so that it is available for reuse:

```
RISC_V_FREE_ENTRY_NOTIFIER_FN(custFreeEntry) {
    vmiosObjectP object = clientData;
    object->tlb[entryId] = False;
}

static VMIO_CONSTRUCTOR_FN(extConstructor) {

    . . . lines omitted . . .
    object->extCB.freeEntryNotifier = custFreeEntry;
    . . . lines omitted . . .
}
```

15.39 Function *cllcCustomRd*

```
#define RISCVC_CUSTOM_READ_FN(_NAME) Bool _NAME( \
    riscvP riscv, \
    Uns32 offset, \
    Uns8 *val, \
    void *clientData \
)
typedef RISCVC_CUSTOM_READ_FN((*riscvClicCRdFn));
```

Description

This notifier is called when the custom region of the CLIC MMIO has been read. Reads are handled one byte at a time, so a 4 byte read will result in 4 separate calls to the notifier.

offset is the offset from the *mclicbase* location.

If the extension implements a register at the offset location it should copy the value to be read for that byte in **val* and return *True* indicating the register is present.

If there is no register implemented by the extension at the offset then it should return *False*.

Example

```
RISCVC_CUSTOM_READ_FN(customClicRd) {
    vmiosObjectP object = clientData;

    if ((offset & ~0x3) == 0x800) { // Custom register: 4 bytes at 0x800
        Uns32 byte = offset & 0x3;

        *val = (object->custom >> (byte*8)) & 0xff;
        return True;
    }

    return False
}
```

Notes

In a multi hart processor *riscv* is the root processor of the cluster

Once any callback has returned *True* no additional callbacks will be done.

15.40 Function *cllcCustomWr*

```
#define RISCVC_CLIC_CUSTOM_WRITE_FN(_NAME) Bool _NAME( \
    riscvP riscv, \
    Uns32 offset, \
    Uns8 val, \
    void *clientData \
)
typedef RISCVC_CLIC_CUSTOM_WRITE_FN((*riscvClicCWrFn));
```

Description

This notifier is called when the custom region of the CLIC MMIO has been written. Writes are handled one byte at a time, so a 4 byte write will result in 4 separate calls to the notifier.

offset is the offset from the *mclicbase* location.

If the extension implements a register at the offset location it should copy *val* to the appropriate byte of the register and return *True* indicating the register is present.

If there is no register implemented by the extension at the offset then it should return *False*.

Example

```
RISCVC_CLIC_CUSTOM_WRITE_FN(customClicWr) {
    vmiosObjectP object = clientData;

    if ((offset & ~0x3) == 0x800) { // Custom register: 4 bytes at 0x800
        Uns32 byte = offset & 0x3;
        Uns32 newByte = ((Uns32)val) << (byte*8);
        Uns32 byteMask = 0xff << (byte*8);

        object->custom &= ~byteMask;
        object->custom |= newByte;

        return True;
    }

    return False;
}
```

Notes

In a multi hart processor *riscv* is the root processor of the cluster

Once any callback has returned *True* no additional callbacks will be done.

15.41 *Function clicUpdated*

```
# define RISC_V_CLIC_UPDATED_FN(_NAME) Bool _NAME( \
    riscvP riscv, \
    void *clientData \
)
typedef RISC_V_CLIC_UPDATED_FN((*riscvClicUpdatedFn));
```

Description

This notifier is called when the CLIC configuration has been updated by software through a write to a CLIC MMIO configuration register.

It gives an extension an opportunity to update any CLIC related state after a configuration change. If the CLIC state is updated in the callback then the function should return True, which will force a re-evaluation of the CLIC status by the base model.

Example

```
static RISC_V_CLIC_UPDATED_FN(cirrusCLICUpdateNotify) {
    vmiosObjectP object = clientData;

    return updateExtensionClicState(riscv, object);
}
```

Notes

In a multi hart processor *riscv* is the root processor of the cluster

15.42 *Function restrictionsCB*

```
#define RISC_V_RESTRICTIONS_FN(_NAME) void _NAME( \
    riscvP      riscv,      \
    vmiDocNodeP node,      \
    void        *clientData \
)
typedef RISC_V_RESTRICTIONS_FN((*riscvRestrictionsFn));

typedef struct riscvExtCBS {
    riscvRestrictionsFn restrictionsCB;
} riscvExtCB;
```

Description

This function is called to add documentation to a derived model listing any restrictions of that model. Documentation should be added beneath the given `node` (of type `vmiDocNodeP`) using the `vmidocAddText` function.

16 Appendix: riscvInstrInfo Structure

Some extension object interface functions described in section 15 are passed a pointer to a structure of type `riscvInstrInfo`. This structure gives comprehensive information obtained by decoding the instruction that is currently being translated, which can be useful when generating custom behavior. For example, the structure can be used to implement custom behavior to cause Illegal Instructions to be generated for some instruction types or argument combinations. This section describes this structure and indicates when particular fields within it may be useful.

The structure is defined in file `riscvDecodeTypes.h` as follows:

```
typedef struct riscvInstrInfoS {

    // COMMON FIELDS FOR ALL INSTRUCTION TYPES
    riscvIType      type;           // instruction type
    riscvAddr       thisPC;         // instruction address
    Uns64           instruction;    // instruction word
    Uns8            bytes;         // instruction size in bytes
    riscvRegDesc    r[RV_MAX_AREGS]; // argument registers
    riscvAddr       tgt;           // constant target address
    Uns64           c;             // constant value
    Int32           memBits;       // load/store size
    Bool            unsExt;        // whether to extend unsigned

    // CSR UPDATE INSTRUCTIONS
    riscvCSRUDesc   csrUpdate;     // CSR update semantics
    Uns32           csr;           // CSR index

    // FLOATING POINT INSTRUCTIONS
    riscvRMDesc     rm;            // rounding mode
    Bool            useF;          // whether F registers used

    // VECTOR EXTENSION INSTRUCTIONS
    riscvRegDesc    mask;          // vector mask register
    riscvWholeDesc  isWhole;       // is this a whole-register instruction?
    riscvVType      vtype;        // vector type information
    Uns32           eew;          // explicit EEW encoding
    Uns8            eewDiv;       // explicit EEW divisor
    Uns8            eewIndex;     // number of EEW index operand
    Uns8            nf;           // nf value
    Bool            isFF;         // is this a first-fault instruction?

    // CODE SIZE REDUCTION INSTRUCTIONS
    Uns32           rlist;        // register list (Zcea push/pop)
    Uns32           alist;        // argument register list (Zcea push)
    riscvRetValDesc retVal;       // return value (Zcea pop)
    riscvCompressSet Zc;          // compressed extension
    Bool            doRet;        // do return (Zcea pop)
    Bool            embedded;     // whether embedded modifier required

    // INTEGER DIVISION INSTRUCTIONS
    Bool            Zmmul;        // whether affected by Zmmul

    // SIMD EXTENSION INSTRUCTIONS
    riscvCrossOpDesc crossOp;     // cross operation
    riscvHalfDesc   half;         // top/bottom half operation
    riscvPackDesc   pack;         // byte packing
    Uns8            elemSize;     // element size
    Bool            doDouble;     // whether additional doubling step
    Bool            round;        // whether additional rounding step

    // INSTRUCTION SUBSET VALIDATION
    riscvArchitecture arch;       // architecture requirements
}
```

```

// INSTRUCTION DISASSEMBLY
const char *opcode;           // opcode name
const char *format;          // disassembly format string
riscvVIType VIType;          // vector instruction type
riscvAQRLEDesc aqrl;         // acquire/release specifier
riscvFenceDesc pred;         // predecessor fence
riscvFenceDesc succ;         // successor fence
riscvXPERMDesc xperm;        // XPERM descriptor
riscvCSufDesc cSuffix;       // compressed suffix
Uns8 explicitType;           // whether types are explicit in opcode
Bool explicitW;              // whether 'w' explicit in opcode
Bool explicitRM;             // whether rounding explicit in opcode
Bool explicitDot;            // whether explicit type has dot
Bool unsPfx;                 // show unsigned as z/s prefix not u suffix
Bool csrInOp;                // whether to emit CSR as part of opcode
Bool cPrefix;                // whether compressed prefix enabled
Uns8 shN;                    // shN prefix
} riscvInstrInfo;

```

Fields in the structure are grouped according to their applicability. These groups are described below.

16.1 Common Fields

These fields will have useful values for a wide range of instruction types:

```

riscvIType type;             // instruction type
riscvAddr thisPC;            // instruction address
Uns64 instruction;           // instruction word
Uns8 bytes;                  // instruction size in bytes
riscvRegDesc r[RV_MAX_REGS]; // argument registers
riscvAddr tgt;               // constant target address
Uns64 c;                     // constant value
Int32 memBits;               // load/store size
Bool unsExt;                  // whether to extend unsigned

```

16.1.1 riscvIType type

This field is an enumeration member describing the abstract type of the decoded instruction; there are a very large number of these. The value is typically useful in contexts where particular instructions have custom behavior or are unimplemented and should cause an Illegal Instruction trap. For example, in an atomic instruction context, the particular type of atomic instruction that has been encountered can be found by comparing against these values:

```

//
// This enumerates generic instructions
//
typedef enum riscvITypeE {

    . . .values omitted . . .

    // A-extension R-type instructions
    RV_IT_AMOADD_R,
    RV_IT_AMOAND_R,
    RV_IT_AMOMAX_R,
    RV_IT_AMOMAXU_R,
    RV_IT_AMOMIN_R,
    RV_IT_AMOMINU_R,
    RV_IT_AMOOR_R,
    RV_IT_AMOSWAP_R,
    RV_IT_AMOXOR_R,

```

```
RV_IT_LR_R,
RV_IT_SC_R

. . .values omitted . . .
};
```

16.1.2 riscvAddr thisPC

This field gives the simulated address of the current instruction.

16.1.3 Uns64 instruction

This field gives the byte pattern of the current instruction.

16.1.4 Uns8 bytes

This field gives the size in bytes of the current instruction.

16.1.5 riscvRegDesc r[RV_MAX_AREGS]

This array field gives the registers used by the current instruction in left-to-right order in an encoded form, using the `riscvRegDesc` type, defined in file `riscvRegisterTypes.h`:

```
typedef enum riscvRegDescE {

    RV_RD_NA          = 0x0,
                                // REGISTER INDEX
    RV_RD_INDEX_MASK = 0x01f,
                                // mask to select register index

                                // REGISTER SIZE (in bits)
    RV_RD_4           = RV_RD_INDEX_MASK+1, // 4-bit
    RV_RD_8           = RV_RD_4 *2,         // 8-bit
    RV_RD_16          = RV_RD_8 *2,         // 16-bit
    RV_RD_32          = RV_RD_16*2,         // 32-bit
    RV_RD_64          = RV_RD_32*2,         // 64-bit
    RV_RD_128         = RV_RD_64*2,         // 128-bit
    RV_RD_BITS_MASK   = (RV_RD_4|RV_RD_8|RV_RD_16|RV_RD_32|RV_RD_64|RV_RD_128),

                                // REGISTER TYPE
    RV_RD_X           = RV_RD_128*2,        // integer (X) register
    RV_RD_F           = RV_RD_X*2,          // floating point register
    RV_RD_V           = RV_RD_F*2,          // vector register
    RV_RD_TYPE_MASK   = (RV_RD_X|RV_RD_F|RV_RD_V|RV_RD_BITS_MASK),

                                // REGISTER TYPE MODIFIERS
    RV_RD_ZFINX       = RV_RD_V*2,          // Zfinx-modified register
    RV_RD_BF16        = RV_RD_ZFINX*2,      // explicit BFLOAT16 register

                                // DISASSEMBLY CONTROL
    RV_RD_Q           = RV_RD_BF16*2,       // quiet (don't show type)
    RV_RD_WL          = RV_RD_Q*2,          // explicit w/l type name
    RV_RD_FX          = RV_RD_WL*2,         // explicit x type name
    RV_RD_U           = RV_RD_FX*2,         // explicit u type name

} riscvRegDesc;
```

File `riscvRegisterTypes.h` also implements a set of inline helper functions that can be used to extract information from a `riscvRegDesc` value. The most useful of these are summarized below.

```
// Is the register an X register?
inline static Bool isXReg(riscvRegDesc r);
```

```
// Is the register an F register?
inline static Bool isFReg(riscvRegDesc r);

// Is the register a V register?
inline static Bool isVReg(riscvRegDesc r);

// Is the register a BF16 register?
inline static Bool isBF16Reg(riscvRegDesc r);

// Is the register an X register holding a floating point value (Zfinx)?
inline static Bool isZfinxReg(riscvRegDesc r);

// Return register index
inline static Uns32 getRIndex(riscvRegDesc r);

// Return register size in bits
inline static Uns32 getRBits(riscvRegDesc r);
```

16.1.6 riscvAddr tgt

For direct branch instructions, this field gives the target address.

16.1.7 Uns64 c

For instructions with encoded constants, this gives the constant value, sign-extended to 64 bits if required.

16.1.8 Int32 memBits

For instructions that access memory, this gives the width of the memory access in bits. The special value -1 is used for polymorphic vector instructions that are accessing memory using the current SEW.

16.1.9 Int32 Bool unsExt

For instructions that extend narrow values to wider ones (either memory reads or register-to-register operations), this indicates whether the extension is unsigned.

16.2 CSR Instruction Fields

These fields will have useful values only for CSR access instructions (types RV_IT_CSRR_I or RV_IT_CSRRI_I):

```
riscvCSRUDesc    csrUpdate;    // CSR update semantics
Uns32            csr;          // CSR index
```

16.2.1 riscvCSRUDesc csrUpdate

This field indicates the update semantics of the CSR access instruction. Type `riscvCSRUDesc` is defined in file `riscvTypes.h` as follows:

```
typedef enum riscvCSRUDescE {
    RV_CSR_NA,    // no update semantics
    RV_CSR_RW,    // read/write
    RV_CSR_RS,    // read/set
    RV_CSR_RC,    // read/clear
} riscvCSRUDesc;
```

16.2.2 Uns32 csr

This field gives the index number of the CSR being accessed:

16.3 Floating Point Instruction Fields

These fields provide information about floating point instructions:

| | | |
|-------------|-------|-----------------------------|
| riscvRMDesc | rm; | // rounding mode |
| Bool | useF; | // whether F registers used |

16.3.1 riscvRMDesc rm

This field indicates the rounding mode encoded in the instruction. Type `riscvRMDesc` is defined in file `riscvTypes.h` as follows:

```
typedef enum riscvRMDescE {
    RV_RM_NONE,        // no rounding mode
    RV_RM_CURRENT,     // round using current rounding mode
    RV_RM_RNE,         // round to nearest, ties to even
    RV_RM_RTZ,         // round towards zero
    RV_RM_RDN,         // round towards -infinity
    RV_RM_RUP,         // round towards +infinity
    RV_RM_RMM,         // round to nearest, ties away
    RV_RM_ROD,         // round to odd (jamming)
    RV_RM_BAD5,        // illegal rounding mode 5
    RV_RM_BAD6,        // illegal rounding mode 6
} riscvRMDesc;
```

16.3.2 Bool useF

This field indicates whether the current instruction accesses any FPR:

16.4 Vector Extension Instruction Fields

These fields provide information about vector extension instructions:

| | | |
|----------------|-----------|--|
| riscvRegDesc | mask; | // vector mask register |
| riscvWholeDesc | isWhole; | // is this a whole-register instruction? |
| riscvVType | vtype; | // vector type information |
| Uns32 | eew; | // explicit EEW encoding |
| Uns8 | eewDiv; | // explicit EEW divisor |
| Uns8 | eewIndex; | // number of EEW index operand |
| Uns8 | nf; | // nf value |
| Bool | isFF; | // is this a first-fault instruction? |

16.4.1 riscvRegDesc mask

This field specifies the mask register to use with the current vector instruction; if zero, the instruction is unmasked. For vector extension versions up to 1.0, only register `v0` may be used as a mask.

16.4.2 riscvWholeDesc isWhole

For whole-register vector instructions, this specifies the class of instruction, as follows:

```
typedef enum riscvWholeDescE {
    RV_WD_NA,          // not a whole register instruction
    RV_WD_LD_ST,       // whole register load/store
    RV_WD_MV,          // whole register move
} riscvWholeDesc;
```

16.4.3 riscvVType vtype

This field is used to categorize a vector instruction shape. It is mainly useful for disassembly:

```
typedef enum riscvVTypeE {
    RV_VIT_NA,      // not a vector instruction
    RV_VIT_V,       // instruction type .v
    RV_VIT_W,       // instruction type .w
    RV_VIT_VV,      // instruction type .vv
    RV_VIT_VI,      // instruction type .vi
    RV_VIT_VX,      // instruction type .vx
    RV_VIT_WV,      // instruction type .wv
    RV_VIT_WI,      // instruction type .wi
    RV_VIT_WX,      // instruction type .wx
    RV_VIT_VF,      // instruction type .vf
    RV_VIT_WF,      // instruction type .wf
    RV_VIT_VS,      // instruction type .vs
    RV_VIT_M,       // instruction type .m
    RV_VIT_MM,      // instruction type .mm
    RV_VIT_VM,      // instruction type .vm
    RV_VIT_VVM,     // instruction type .vvm
    RV_VIT_VXM,     // instruction type .vxm
    RV_VIT_VIM,     // instruction type .vim
    RV_VIT_VFM,     // instruction type .vfm
    RV_VIT_VN,      // instruction type .v/.w (version-dependent)
    RV_VIT_VVN,     // instruction type .vv/.wv (version-dependent)
    RV_VIT_VIN,     // instruction type .vi/.wi (version-dependent)
    RV_VIT_VXN,     // instruction type .vx/.wx (version-dependent)
    RV_VIT_V_V,     // instruction type .v.v
    RV_VIT_LAST     // KEEP LAST: for sizing
} riscvVType;
```

16.4.4 Uns32 eew

This field gives any **EEW** value explicitly encoded in the instruction. If zero, there is no explicit **EEW** encoding.

16.4.5 Uns8 eewDiv

This field gives any **EEW** divisor explicitly encoded in the instruction. If zero, there is no explicit **EEW** divisor.

16.4.6 Uns8 eewIndex

For instructions that have **EEW** specified in a register operand, this gives the index number of that operand.

16.4.7 Uns8 nf

This field gives any **nf** field number explicitly encoded in the instruction. If zero, there is no explicit **nf** field number.

16.4.8 Bool isFF

This field indicates whether the current instruction has *first fault* behavior.

16.5 Code Size Reduction Extension Instruction Fields

These fields provide information about code size reduction extension instructions:

```

Uns32      rlist;           // register list (Zcea push/pop)
Uns32      alist;          // argument register list (Zcea push)
riscvRetValDesc retval;    // return value (Zcea pop)
riscvCompressSet Zc;       // compressed extension
Bool       doRet;          // do return (Zcea pop)
Bool       embedded;       // whether embedded modifier required

```

16.5.1 Uns32 rlist

This field is a bitmask of registers read or written by Zcea subset push/pop instructions. Register xN is read or written if bit N is set in the mask.

16.5.2 Uns32 alist

This field is a bitmask of argument registers used by Zcea subset push instructions. Register xN is an argument if bit N is set in the mask.

16.5.3 riscvRetValDesc retval

This field indicates the return value for the Zcea subset pop instruction:

```

typedef enum riscvRetValDescE {
    RV_RV_NA,      // {}
    RV_RV_0,       // {0}
    RV_RV_P1,      // {1}
    RV_RV_M1,      // {-1}
    RV_RV_Z,       // explicit Z in opcode
    RV_RV_LAST     // KEEP LAST: for sizing
} riscvRetValDesc;

```

16.5.4 riscvCompressSet Zc

This field indicates any extension subset requirements for the instruction. Depending on the extension version, either legacy or new values are indicated. The value is a bitfield; if more than one bit is set, then the instruction is configured if any of the subsets is present:

```

typedef enum riscvCompressSetE {
    // legacy values
    RVCS_Zcea = (1<<0),      // Zcea subset
    RVCS_Zceb = (1<<1),      // Zceb subset
    RVCS_Zcee = (1<<2),      // Zcee subset
    // new values
    RVCS_Zca = (1<<3),        // Zca subset
    RVCS_Zcb = (1<<4),        // Zcb subset
    RVCS_Zcd = (1<<5),        // Zcd subset (implicit)
    RVCS_Zcf = (1<<6),        // Zcf subset
    RVCS_Zcmb = (1<<7),       // Zcmb subset
    RVCS_Zcmp = (1<<8),       // Zcmp subset
    RVCS_Zcmpe = (1<<9),      // Zcmpe subset
    RVCS_Zcmt = (1<<10),      // Zcmt subset
} riscvCompressSet;

```

16.5.5 Bool doRet

For the Zcea subset pop instruction, this indicates whether the instruction does a return.

16.5.6 Bool embedded

This indicates whether the embedded extension (E) is active. If it is, the registers stored by the Zcea subset push instruction are modified.

16.6 *Zmmul* Instruction Fields

One field is relevant for the *Zmmul* extension, which describes a subset of the standard *M* extension:

```
Bool      Zmmul;      // whether affected by Zmmul
```

16.6.1 Bool *Zmmul*

This field indicates whether the current instruction is absent if extension *Zmmul* is active (currently, this is integer divide and modulus instructions only).

16.7 *Packed SIMD* Extension Instruction Fields

These fields provide information about packed SIMD extension instructions:

```
riscvCrossOpDesc  crossOp;      // cross operation
riscvHalfDesc     half;         // top/bottom half operation
riscvPackDesc     pack;        // byte packing
Uns8              elemSize;     // element size
Bool              doDouble;     // whether additional doubling step
Bool              round;        // whether additional rounding step
```

This extension has not been ratified and the fields are not discussed further here.

16.8 *Other Fields*

Remaining fields in the structure are concerned with standard extension subset validation and disassembly. They are unlikely to be useful in intended models and not discussed further here.

17 Appendix: Custom CSR Description

This appendix provides detailed information about the `riscvCSRAttrs` type, used to define custom CSRs in an extended processor model. It should be read in conjunction with chapter 7, which describes common use cases.

17.1 *riscvCSRAttrs* Fields

The `riscvCSRAttrs` type is defined as follows:

```
typedef struct riscvCSRAttrsS {
    const char      *name;           // register name
    const char      *desc;           // register description
    vmiosObjectP    object;          // custom extension
    Uns32           csrNum;          // CSR number (includes privilege and r/w access)
    riscvArchitecture arch;         // required architecture (presence)
    riscvArchitecture access;       // required architecture (access)
    riscvPrivVer    version;         // minimum specification version
    riscvCSRPresentFn presentCB;     // CSR present callback
    riscvCSRReadFn  readCB;         // read callback
    riscvCSRReadFn  readWriteCB;    // read callback (in r/w context)
    riscvCSRWriteFn writeCB;        // write callback
    riscvCSRWStateFn wstateCB;      // adjust JIT code generator state
    vmiReg          reg;            // register
    vmiReg          writeMaskV;     // configuration-dependent write mask
    Uns32           writeMaskC32;   // constant 32-bit write mask
    Uns64           writeMaskC64;   // constant 64-bit write mask

    riscvCSRStateenBit Smstateen :8; // whether xstateen-controlled access
    riscvCSRTrap       trap      :2; // whether trapped
    riscvCSRTrace      noTraceChange:2; // trace mode
    Bool              wEndBlock  :1; // whether write terminates this block
    Bool              wEndRM     :1; // whether write invalidates RM assumption
    Bool              noSaveRestore:1; // whether to exclude from save/restore
    Bool              writeRd    :1; // whether write updates Rd
    Bool              aliasV     :1; // whether CSR has virtual alias
    Bool              undefined  :1; // whether CSR is undefined
    Bool              forceRO    :1; // type is RO even if write CB defined
} riscvCSRAttrs;
```

The fields in this structure are as follows:

17.1.1 Field `name`

This is the name of the CSR, used in tracing and for named access from a harness.

17.1.2 Field `desc`

This is a CSR description, used by documentation generation.

17.1.3 Field `object`

This field should be zeroed in the CSR template structure. It is initialized by the RISC-V model with a pointer to the extension CSR object.

17.1.4 Field `csrNum`

This field should be set to the CSR number, in the range 0-4095.

17.1.5 Field arch

This field gives special architectural constraints which determine whether the CSR is *present* or not. For example, if `arch` is set to `ISA_S`, then the CSR will only be present for variants that implement *Supervisor* mode. If the field is zero, then whether the CSR is present is determined by the standard RISC-V rules applying to the CSR number, specified by the `csrNum` field (see the RISC-V Privileged Architecture Specification).

17.1.6 Field access

This field gives special architectural constraints which determine whether the CSR is *accessible* or not. For example, if `arch` is set to `ISA_S`, then the CSR will only be accessible if the hart is operating in *Supervisor* mode (or any higher privilege level, Machine mode in this case). If the field is zero, then whether the CSR is accessible is determined by the standard RISC-V rules applying to the CSR number, specified by the `csrNum` field (see the RISC-V Privileged Architecture Specification).

17.1.7 Field version

This field indicates the minimum Privileged Architecture version for which the CSR is supported. In extension objects this is likely always to be zero, implying to restriction.

17.1.8 Field presentCB

This field specifies an optional function that implements custom rules to determine whether a CSR is present. The function is of type `riscvCSRPresentFn`, defined as follows:

```
#define RISC_V_CSR_PRESENTFN(_NAME) Bool _NAME( \
    riscvCSRAttrsCP attrs,      \
    riscvP          riscv       \
)
typedef RISC_V_CSR_PRESENTFN((*riscvCSRPresentFn));
```

The function should return `True` if the CSR is present and `False` otherwise; typically, it will examine some state within the extension object to determine this. For example:

```
static RISC_V_CSR_PRESENTFN(csrPresent) {
    return attrs->object->csrPresent; // set by an extension parameter?
}
```

Note that the `object` field is used here to obtain a pointer to the extension object.

17.1.9 Field readCB

This field specifies an optional function that is called implement a read of the CSR. The function is of type `riscvCSRReadFn`, defined as follows:

```
#define RISC_V_CSR_READFN(_NAME) Uns64 _NAME( \
    riscvCSRAttrsCP attrs,      \
    riscvP          riscv       \
)
typedef RISC_V_CSR_READFN((*riscvCSRReadFn));
```

The function should return the CSR value as an `Uns64`. For example:

```
static RISC_V_CSR_READFN(custom_rw3_32R) {  
    vmiosObjectP object = attrs->object;  
    Int32 result = RD_XCSR(object, custom_rw3_32);  
    return result;  
}
```

Note that the `object` field is used here to obtain a pointer to the extension object.

If the `readCB` callback is not specified, then a read of the CSR will return the raw value from any register specified by the `reg` field. If that field is `VMI_NOREG`, then a read of the register will return zero.

17.1.10 Field `readWriteCB`

This field specifies an optional function that is called implement a read of the CSR in a read-write context (i.e. using instruction `csrrw` or an alias of it). A few standard CSRs such as `mip` have special behavior in such cases. If the `readWriteCB` is not specified, then the `readCB` is used for CSR reads in such cases as well.

17.1.11 Field `writeCB`

This field specifies an optional function that is called implement a write of the CSR. The function is of type `riscvCSRWriteFn`, defined as follows:

```
#define RISC_V_CSR_WRITEFN(_NAME) Uns64 _NAME( \  
    riscvCSRAttrsCP attrs,    \  
    riscvP riscv,             \  
    Uns64 newValue           \  
)  
typedef RISC_V_CSR_WRITEFN((*riscvCSRWriteFn));
```

The function takes the value to be written and should return the resulting CSR value as an `Uns64` (this may be different to the given value if CSR bits are read-only or if special WARL restrictions are implemented). For example:

```
static RISC_V_CSR_WRITEFN(custom_rw3_32W) {  
    vmiosObjectP object = attrs->object;  
    WR_XCSR(object, custom_rw3_32, newValue);  
    return newValue;  
}
```

Note that the `object` field is used here to obtain a pointer to the extension object.

If the `writeCB` callback is not specified, then whether the CSR is writable is determined by `csrNum[11:10]` using standard RISC-V rules (see the RISC-V Privileged Architecture Specification). However, if the `writeCB` callback is specified for a CSR with a number that would *not* normally be writable, that CSR is writable (overriding the standard rules).

If the `writeCB` callback is absent for a CSR that is defined to be writable by standard RISC-V CSR number rules, then a write of the CSR will update the raw value in any register specified by the `reg` field, using the masking constraints described for that field below. If the `reg` field is `VMI_NOREG`, then writes of the register will be ignored.

17.1.12 Field `wstateCB`

This field specifies an optional function that is called update code generator state after a write of the CSR. It is used in the RISC-V base model but is not intended for use in extensions.

17.1.13 Field `reg`

This `vmiReg` field specifies an optional description of the register value in the extension object. If the register value is not held as a field in the extension object, `reg` should be set to `VMI_NOREG`; otherwise, it should be initialized using the `XCSR_REG_MT` macro. For example, given this extension object definition:

```
typedef struct addCSRsCSRss {
    CSR_REG_DECL (custom_rwl_32);
} addCSRsCSRss;

typedef struct vmiosObjectS {
    addCSRsCSRss csr;
} vmiosObject;
```

Then the `custom_rwl_32` CSR value in the extension object can be referenced like this:

```
XCSR_REG_MT(custom_rwl_32)
```

If read callbacks are not specified for the CSR, then the value indicated by the `reg` field will be returned; if this is `VMI_NOREG`, then zero will be returned.

If a write callback is not specified for the CSR and the CSR number indicates it has write access, then the value indicated by the `reg` field will be written; if this is `VMI_NOREG`, then the write will be ignored. The written value will be masked according to these rules:

1. If the `writeMaskV` field is not `VMI_NOREG`, then this specifies a configurable mask register in the extension object. Only bits that are *non-zero* in this mask will be updated in the written register; bits that are *zero* in the mask will be preserved in the target register.
2. If the `writeMaskV` field is `VMI_NOREG` and a *32-bit* CSR is written, then field `writeMaskC32` specifies a constant 32-bit mask of writable bits in the target register. Only bits that are *non-zero* in this mask will be updated in the written register; bits that are *zero* in the mask will be preserved in the target register. Set this field to all-ones to allow unrestricted access to the target register.
3. If the `writeMaskV` field is `VMI_NOREG` and a *64-bit* CSR is written, then field `writeMaskC64` specifies a constant 64-bit mask of writable bits in the target register. Only bits that are *non-zero* in this mask will be updated in the written

register; bits that are *zero* in the mask will be preserved in the target register. Set this field to all-ones to allow unrestricted access to the target register.

17.1.14 Field `writeMaskV`

This field specifies a configurable write mask for the CSR. See section 17.1.13.

17.1.15 Field `writeMaskC32`

This field specifies a constant 32-bit write mask for the CSR. See section 17.1.13.

17.1.16 Field `writeMaskC64`

This field specifies a constant 64-bit write mask for the CSR. See section 17.1.13.

17.1.17 Field `Smstateen`

This field specifies any `Smstateen` bit that controls access to this CSR. If it is non-zero then `Smstateen` control register state will be added to control access to the register using the indicated bit. Some basic control bit aliases are defined by the `riscvCSRStateenBit` enumeration:

```
typedef enum riscvCSRStateenBitE {
    bit_stateen_NA      = 0,
    bit_stateen_Zfinx    = 1,
    bit_stateen_Zcmt     = 2,
    bit_stateen_xcse     = 57,
    bit_stateen_IMSIC    = 58,
    bit_stateen_AIA      = 59,
    bit_stateen_sireg    = 60,
    bit_stateen_xenvcfg   = 62,
    bit_stateen_xstateen = 63,
} riscvCSRStateenBit;
```

Extensions may specify these values or any other values in the range 1-255.

17.1.18 Field `trap`

This field specifies whether CSR accesses should be trapped by either `mstatus.TVM` or `hvipctl.VTI`, according to the `riscvCSRTrap` enumeration:

```
typedef enum riscvCSRTrapE {
    CSRT_NA,           // CSR not trapped
    CSRT_TVM = 1<<0,  // trapped by mstatus.TVM=1 (e.g. satp register)
    CSRT_VTI = 1<<1,  // trapped by hvipctl.VTI=1 (e.g. sip, sie registers)
} riscvCSRTrap;
```

17.1.19 Field `noTraceChange`

This field specifies whether the CSR value should be shown during simulation when register tracing is enabled. Normally register values should be traced, but for some volatile registers, like instruction counters, this can produce a lot of noise. Trace options are defined by the `riscvCSRTrace` enumeration:

```
typedef enum riscvCSRTraceE {
    RCSRT_YES,          // always trace CSR
    RCSRT_NO,           // never trace CSR
    RCSRT_VOLATILE,     // trace only in volatile mode
}
```

```
} riscvCSRTrace;
```

Value `RCSRT_YES` indicates the CSR value should always be shown in the trace.

Value `RCSRT_NO` indicates the CSR value should never be shown in the trace.

Value `RCSRT_VOLATILE` indicates the CSR value should be shown in the trace only when parameter `traceVolatile` on the RISC-V model is `True`.

17.1.20 Field `wEndBlock`

This field specifies whether, when the CSR is written by a CSR update instruction, that instruction must be the last in a JIT-translated code block. This is necessary if writes to the CSR can modify constraints enforced using VMI block masks or similar constructs. See the *OVP Processor Modeling Guide* for more information about the use of block masks in processor models.

17.1.21 Field `wEndRM`

This field specifies whether, when the CSR is written by a CSR update instruction, any assumptions about floating point rounding mode must be invalidated. This is intended for use in the base model and not likely to be required in an extension.

17.1.22 Field `noSaveRestore`

The simulation environment supports save and restore of processor model state. To implement save and restore, CSRs are usually simply read and written using the defined read/write callbacks for that CSR, or using the raw register value if no read/write callbacks are specified. This field indicates that the CSR should not be saved and restored in this way.

17.1.23 Field `writerd`

This field indicates that, when a CSR is written and has a write callback, the value returned by the write callback should be assigned to the `rd` field in the `csrrs/csrrc` instruction that accesses the CSR. This is not standard CSR behavior, but is required to implement CSRs like `mnxti` in the CLIC extension.

17.1.24 Field `aliasv`

This field indicates that a Supervisor mode CSR has a virtual alias, with index `csrNum+0x100`. When a hart is executing in VS mode, apparent references to the CSR will instead access the virtual alias CSR.

17.1.25 Field `undefined`

This field indicates that a CSR is undefined, and that accesses to it should always generate an Illegal Instruction trap. This is useful for extensions that remove standard CSRs for any reason.

17.1.26 Field `forceRO`

This field indicates that a CSR is read-only, regardless of:

- whether a write callback is defined for it, which would normally make it be considered read-write even when the CSR index is in the read-only region.
- the index is in a CSR region defined as a read-write region

This may be used when defining a write callback that supports artifact writes for a read-only CSR.

17.2 CSR Definition Macros

CSRs can be defined using C structure declarations using named-field syntax if required. However, there are several macros defined in file `riscvModelCallbackTypes.h` that cover the common types of declaration in a more concise form. These are described below.

17.2.1 Macro `XCSR_ATTR_UIP`

This macro declares an undefined CSR. This can be used where an extension model does not implement a standard CSR for some reason. It is defined as:

```
#define XCSR_ATTR_UIP(_ID, _NUM, _ARCH, _EXT) [XCSR_ID(_ID)] = { \
    .extension = _EXT, \
    .baseAttrs = { \
        name      : #_ID, \
        csrNum     : _NUM, \
        arch       : _ARCH, \
        undefined  : True \
    } \
}
```

17.2.2 Macro `XCSR_ATTR_NIP`

This macro declares a defined but unimplemented CSR. This can be used where an extension model is under development and a temporary placeholder is needed for future implementation. It is defined as:

```
#define XCSR_ATTR_NIP( \
    _ID, _NUM, _ARCH, _EXT, _ENDB, _ENDRM, _NOTR, _TRAP, _DESC \
) [XCSR_ID(_ID)] = { \
    .extension = _EXT, \
    .baseAttrs = { \
        name      : #_ID, \
        desc      : _DESC " (not implemented)", \
        csrNum     : _NUM, \
        arch       : _ARCH, \
        wEndBlock  : _ENDB, \
        wEndRM     : _ENDRM, \
        noTraceChange : _NOTR, \
        trap       : _TRAP, \
    } \
}
```

17.2.3 Macro `XCSR_ATTR_T__`

This macro declares a CSR that has a value implemented by a field in the extension object with no write mask constraints. There can optionally be read or write callbacks. It is defined as:

```
#define XCSR_ATTR_T__( \
    _ID, _NUM, _ARCH, _EXT, _ENDB, _ENDRM, _NOTR, _TRAP, _DESC, _RCB, _RWC, _WCB \
) [XCSR_ID(_ID)] = { \
    .extension = _EXT, \
    .baseAttrs = { \
        name      : #_ID, \
        desc      : _DESC, \
        csrNum     : _NUM, \
        arch       : _ARCH, \
        wEndBlock  : _ENDB, \
        wEndRM     : _ENDRM, \
        noTraceChange : _NOTR, \
        trap       : _TRAP, \
        readCb     : _RCB, \
        writeCb    : _RWC, \
        writeMaskCb : _WCB \
    } \
}
```



```

        wEndRM      : _ENDRM,
        noTraceChange : _NOTR,
        trap        : _TRAP,
        readCB      : _RCB,
        readWriteCB  : _RWCb,
        writeCB     : _WCB,
        reg         : XCSR_REG_MT(_ID),
        writeMaskC32 : -1,
        writeMaskC64 : -1
    }
}

```

17.2.4 Macro XCSR_ATTR_TC_

This macro declares a CSR that has a value implemented by a field in the extension object and *constant* write masks. There can optionally be read or write callbacks. It is defined as:

```

#define XCSR_ATTR_TC_( \
    _ID, _NUM, _ARCH, _EXT, _ENDB, _ENDRM, _NOTR, _TRAP, _DESC, _RCB, _RWCb, _WCB \
) [XCSR_ID(_ID)] = { \
    .extension = _EXT,
    .baseAttrs = {
        name      : #_ID,
        desc      : _DESC,
        csrNum     : _NUM,
        arch      : _ARCH,
        wEndBlock  : _ENDB,
        wEndRM     : _ENDRM,
        noTraceChange : _NOTR,
        trap      : _TRAP,
        readCB     : _RCB,
        readWriteCB : _RWCb,
        writeCB    : _WCB,
        reg        : XCSR_REG_MT(_ID),
        writeMaskC32 : WM32_##_ID,
        writeMaskC64 : WM64_##_ID
    }
}

```

17.2.5 Macro XCSR_ATTR_TV_

This macro declares a CSR that has a value implemented by a field in the extension object and *configurable* write mask. There can optionally be read or write callbacks. It is defined as:

```

#define XCSR_ATTR_TV_( \
    _ID, _NUM, _ARCH, _EXT, _ENDB, _ENDRM, _NOTR, _TRAP, _DESC, _RCB, _RWCb, _WCB \
) [XCSR_ID(_ID)] = { \
    .extension = _EXT,
    .baseAttrs = {
        name      : #_ID,
        desc      : _DESC,
        csrNum     : _NUM,
        arch      : _ARCH,
        wEndBlock  : _ENDB,
        wEndRM     : _ENDRM,
        noTraceChange : _NOTR,
        trap      : _TRAP,
        readCB     : _RCB,
        readWriteCB : _RWCb,
        writeCB    : _WCB,
        reg        : XCSR_REG_MT(_ID),
        writeMaskV  : XCSR_MASK_MT(_ID)
    }
}

```

17.2.6 Macro XCSR_ATTR_P__

This macro declares a CSR that is implemented by callbacks only and has no corresponding value field in the extension object. It is defined as:

```
#define XCSR_ATTR_P__( \
    _ID, _NUM, _ARCH, _EXT, _ENDB, _ENDRM, _NOTR, _TRAP, _DESC, _RCB, _RWC, _WCB \
) [XCSR_ID(_ID)] = { \
    .extension = _EXT, \
    .baseAttrs = { \
        name      : #_ID, \
        desc      : _DESC, \
        csrNum    : _NUM, \
        arch      : _ARCH, \
        wEndBlock : _ENDB, \
        wEndRM    : _ENDRM, \
        noTraceChange : _NOTR, \
        trap      : _TRAP, \
        readCB    : _RCB, \
        readWriteCB : _RWC, \
        writeCB   : _WCB, \
        writeMaskC32 : -1, \
        writeMaskC64 : -1 \
    } \
}
```

If both read and write callbacks are absent, then this specifies a CSR that reads as zero and ignores writes.

18 Appendix: Configuring Standard Extensions

This appendix provides detailed information about configuring the model to implement standard extensions using parameters.

18.1 Ratified Extension Support

The following ratified extensions are supported in the model using appropriate parameterization, with the exception of the `Sscofpmf` extension which is not currently supported. See subsections on following pages for information about how these extensions can be configured in the OVP RISC-V model. Ratified extensions are listed at [Ratified Extensions - Home - RISC-V International \(riscv.org\)](https://riscv.org).

| | |
|--|---|
| <i>RISC-V Quality-of-Service (QoS) Identifiers</i> | <code>Ssqosid</code> |
| <i>Obviating Memory-Management Instructions after Marking PTEs Valid</i> | <code>Svvptc</code> |
| <i>Resumable Non-Maskable Interrupts</i> | <code>Smrnmi</code> |
| <i>Shadow Stacks and Landing Pads</i> | <code>Zicfiss, Zicfilp</code> |
| <i>BF16 Extensions</i> | <code>Zfbfmin, Zvfbfmin, Zvfbfwma</code> |
| <i>Zaamo and Zalrsc Extensions</i> | <code>Zaamo, Zalrsc</code> |
| <i>B Standard Extension for Bit Manipulation Instructions</i> | <code>B</code> |
| <i>Byte and Halfword Atomic Memory Operations (Zabha)</i> | <code>Zabha</code> |
| <i>RISC-V Supervisor Counter Delegation</i> | <code>Smcdeleg, Ssccfg</code> |
| <i>May-Be-Operations</i> | <code>Zimop, Zcmop</code> |
| <i>RISC-V Indirect CSR Access (Smcsrind/Sscsrind)</i> | <code>Smcsrind, Sscsrind</code> |
| <i>RISC-V Pointer Masking Extensions</i> | <code>Ssnpm, Smnpm, Smmnpm</code> |
| <i>RISC-V Integer Conditional (Zicond) operations extension</i> | <code>Zicond</code> |
| <i>Hardware Updating of PTE A/D Bits (Svadu)</i> | <code>Svadu</code> |
| <i>RISC-V Cycle and Instret Privilege Mode Filtering (Smcntrpmf)</i> | <code>Smcntrpmf</code> |
| <i>Atomic Compare-and-Swap (CAS) Instructions (Zacas)</i> | <code>Zacas</code> |
| <i>RISC-V Cryptography Extensions Volume II: Vector Instructions</i> | <code>Zvbb, Zvbc, Zvkb, Zvkg, Zvkn, Zvknk, Zvkned, Zvkng, Zvknha, Zvknhb, Zvks, Zvksc, Zvkse, Zvksg, Zvksh, Zvkt</code> |
| <i>"Zfa" Standard Extension for Additional Floating-Point Instructions</i> | <code>Zfa</code> |
| <i>RISC-V Advanced Interrupt Architecture</i> | <code>Smaia, Ssaia</code> |
| <i>"Zvfh/Zvfhmin:" Vector Extension for Half-Precision Floating-Point Arithmetic/Vector Extension for Minimal Half-Precision Floating-Point Arithmetic</i> | <code>Zvfh, Zvfhmin</code> |
| <i>"Zihintntl" Non-Temporal Locality Hints</i> | <code>Zihintntl</code> |
| <i>RISC-V Code Size Reduction</i> | <code>Zca, Zcb, Zcd, Zce, Zcf, Zcmp, Zcmt</code> |
| <i>RISC-V Profiles</i> | <code>RVA20, RVI20, RVA22</code> |
| <i>"Zicntr" and "Zihpm" Counters</i> | <code>Zicntr, Zihpm</code> |
| <i>RV32E and RV64E Base Integer Instruction Sets</i> | <code>RV32E/RV64E</code> |

| | |
|---|---|
| <i>"Ztso" Standard Extension for Total Store Ordering</i> | Ztso |
| <i>RISC-V Wait-on-Reservation-Set (Zawrs) extension</i> | Zawrs |
| <i>Zmmul Extension</i> | Zmmul |
| <i>PMP Enhancements for memory access and execution prevention on Machine mode (Smepmp)</i> | Smepmp |
| <i>RISC-V Base Cache Management Operation ISA Extensions</i> | Zicbom, Zicbop, Zicboz |
| <i>RISC-V Bit-Manipulation ISA-extensions</i> | Zba, Zbb, Zbc, Zbs |
| <i>RISC-V Count Overflow and Mode-Based Filtering Extension</i> | Sscofpmf |
| <i>RISC-V Cryptography Extensions Volume I: Scalar & Entropy Source Instructions</i> | Zbkb, Zbkc, Zbkx, Zknd, Zkne, Zknh, Zksed, Zksh, Zkn, Zks, Zkt, Zk, Zkr |
| <i>RISC-V State Enable Extension</i> | Smstateen |
| <i>RISC-V "stimecmp / vstimecmp" Extension</i> | Sstc |
| <i>RISC-V Vector Extension</i> | Zve32x, Zve32f, Zve64x, Zve64f, Zve64d, Zve, Zvl32b, Zvl64b, Zvl128b, Zvl256b, Zvl512b, Zvl1024b, Zvl, Zv |
| <i>The RISC-V Instruction Set Manual Volume II: Privileged Architecture</i> | Smlpl2, Sslpl2, Sv57, Hypervisor, Svinval, Svnapot, Svpbmt |
| <i>"Zfh" and "Zfhmin" Standard Extensions for Half-Precision Floating-Point</i> | Zfh, Zfhmin |
| <i>"Zfinx", "Zdinx", "Zhinx", "Zhinxmin": Standard Extensions for Floating-Point in Integer Registers</i> | Zfinx, Zdinx, Zhinx, Zhinxmin |
| <i>"Zhintpause" Pause Hint</i> | Zhintpause |
| <i>"Zicsr", Control and Status Register (CSR) Instructions</i> | Zicsr |
| <i>"Zifencei" Instruction-Fetch Fence</i> | Zifencei |

18.1.1 RISC-V Quality-of-Service (QoS) Identifiers

Enabled using parameter `Ssqosid`.

18.1.2 Obviating Memory-Management Instructions after Marking PTEs Valid

The model always behaves as if this extension is enabled.

18.1.3 Resumable Non-Maskable Interrupts

Enabled using parameter `rnmi_version` with version 0.4.

18.1.4 Shadow Stacks and Landing Pads

The following table shows how to configure the model for each sub-extension listed here.

| | |
|---------|-----------------------------|
| Zicfiss | Set parameter Zicfiss to T. |
| Zicfilp | Set parameter Zicfilp to T. |

18.1.5 *BF16 Extensions*

The following table shows how to configure the model for each sub-extension listed here.

| | |
|----------|---|
| zfbfmin | Set parameter zfbfmin to T on a variant with the F extension. |
| zvfbfmin | Set parameter zvfbfmin to T on a variant with the F and V extensions. |
| zvfbfwma | Set parameter zvfbfwma to T on a variant with the F and V extensions. |

18.1.6 *Zaamo and Zalrsc Extensions*

The following table shows how to configure the model for each sub-extension listed here when the A extension is configured.

| | |
|--------|----------------------------|
| Zaamo | Set parameter zaamo to T. |
| Zalrsc | Set parameter zalrsc to T. |

18.1.7 *B Standard Extension for Bit Manipulation Instructions*

Enabled using parameter misa_B_Zba_Zbb_Zbs when the B extension is configured.

18.1.8 *Byte and Halfword Atomic memory Operations (Zabha)*

Enabled using parameter zabha when the A extension is configured.

18.1.9 *RISC-V Supervisor Counter Delegation (Smcdeleg/Ssccfg)*

Enabled using parameter smcdeleg.

18.1.10 *May-Be-Operations*

The following table shows how to configure the model for each sub-extension listed here.

| | |
|-------|---------------------------|
| Zimop | Set parameter zimop to T. |
| Zcmop | Set parameter zcmop to T. |

18.1.11 *RISC-V Indirect CSR Access (Smcsrind/Sscsrind)*

Enabled using parameter smcsrind. (Sscsrind is enabled if this parameter is set and Supervisor mode is present).

18.1.12 *RISC-V Pointer Masking Extensions*

The following table shows how to configure the model for each sub-extension listed here.

| | |
|--------|---------------------------|
| Ssnpm | Set parameter ssnpm to T. |
| Snnpm | Set parameter snnpm to T. |
| Smmnpm | Set parameter smnpm to T. |

18.1.13 *RISC-V Integer Conditional (Zicond) operations extension*

Enabled using parameter zicond.

18.1.14 Hardware Updating of PTE A/D Bits (Svadu)

Enabled using parameter Svadu.

18.1.15 RISC-V Cycle and Instret Privilege Mode Filtering (Smcntrpmf)

Enabled using parameter Smcntrpmf.

18.1.16 Atomic Compare-and-Swap (CAS) Instructions (Zacas)

Enabled using parameter Zacas.

18.1.17 RISC-V Cryptography Extensions Volume II: Vector Instructions

For all listed extensions, the V and K extensions must be enabled by setting parameter add_Extensions to include v and k. The following table shows how to further configure the model for each sub-extension listed here.

| | |
|---------|--|
| Zvbb | Set parameter Zvbb to T. |
| Zvbc | Set parameter Zvbc to T. |
| Zvkb | Set parameter Zvkb to T. |
| Zvkg | Set parameter Zvkg to T. |
| Zvkn | Set parameter Zvkn to T. |
| Zvknc | Set parameters Zvkn and Zvbc to T. |
| Zvkned | Set parameter Zvkned to T. |
| Zvkng | Set parameters Zvkn and Zvkg to T. |
| Zvknha | Set parameter Zvknha to T. |
| Zvknhb | Set parameter Zvknhb to T. |
| Zvks | Set parameter Zvks to T. |
| Zvksc | Set parameters Zvks and Zvbc to T. |
| Zvksg | Set parameters Zvks and Zvkg to T. |
| Zvksh | Set parameter Zvksh to T. |
| Zvk sed | Set parameter Zvk sed to T. |
| Zvkt | <i>Data-independent execution latency, not relevant for this model</i> |

18.1.18 "Zfa" Standard Extension for Additional Floating-Point Instructions

Enabled using parameter zfa on variant with floating point configured.

18.1.19 RISC-V Advanced Interrupt Architecture Extension

The following table shows how to configure the model for each sub-extension listed here. See section 5.7 for details of further parameters to refine the Advanced Interrupt Architecture configuration.

| | |
|-------|--|
| Smaia | Set parameter Smaia to T. |
| Ssaia | Set parameter Smaia to T, use variant with S extension |

18.1.20 “Zvfh/Zvfhmin:” Vector Extension for Half-Precision Floating-Point Arithmetic/Vector Extension for Minimal Half-Precision Floating-Point Arithmetic

The following table shows how to configure the model for each sub-extension listed here.

| | |
|---------|---|
| zvfh | Set parameter <code>zvfh</code> to T, use variant with V extension and floating point configured |
| zvfhmin | Set parameter <code>zvfhmin</code> to T, use variant with V extension and floating point configured |

18.1.21 “Zihintntl” Non-Temporal Locality Hints

Enabled using parameter `zihintntl`. The hint instructions are executed but otherwise have no effect on the model.

18.1.22 RISC-V Code Size Reduction

For all listed extensions, the C extension must be enabled by either setting parameter `add_Extensions` to include C, or instantiating a base variant with C extension (e.g. RV32GC). The following table shows how to further configure the model for each sub-extension listed here.

| | |
|------|--|
| zca | Set parameter <code>zca</code> to T |
| zcb | Set parameter <code>zcb</code> to T |
| zcd | Implied if C and D extensions are enabled and all <code>zcm*</code> extensions are absent. |
| zce | (Collective name for <code>zca</code> , <code>zcb</code> , <code>zcmp</code> , <code>zcmt</code> and optionally <code>zcf</code> extensions) |
| zcf | Set parameter <code>zcf</code> to T |
| zcmp | Set parameter <code>zcmp</code> to T |
| zcmt | Set parameter <code>zcmt</code> to T |

18.1.23 RISC-V Profiles

Enabled using more fundamental parameters documented here.

18.1.24 "Zicntr" and "Zihpm" Counters

These extensions are both enabled by default. The CSRs that they affect can be selectively disabled as follows:

| | |
|------------|---|
| cycle | Set parameter <code>cycle_undefined</code> to T |
| mcycle | Set parameter <code>mcycle_undefined</code> to T |
| time | Set parameter <code>time_undefined</code> to T |
| instret | Set parameter <code>instret_undefined</code> to T |
| minstret | Set parameter <code>minstret_undefined</code> to T |
| hpmcounter | Set parameter <code>hpmcounter_undefined</code> to T. Note that these registers are implementation-defined and always read as zero in the OVP RISC- |

| | |
|-------------|--|
| | V base model. |
| mhpmcounter | Set parameter mhpmcounter_undefined to T. Note that these registers are implementation-defined and always read as zero in the OVP RISC-V base model. |

18.1.25 RV32E and RV64E Base Integer Instruction Sets

Set parameter `add_Extensions` to include E, or instantiate base variant with E extension (e.g. RV32EC).

18.1.26 “Ztso” Standard Extension for Total Store Ordering

The OVP RISC-V model behaves as if this extension is always enabled.

18.1.27 RISC-V Wait-on-Reservation-Set (Zawrs) extension

Enabled using parameter `zawrs`.

18.1.28 Zmmul Extension

Enabled using parameter `zmmul`.

18.1.29 RISC-V PMP Enhancements for memory access and execution prevention on Machine mode (Smepmp)

Enabled using parameter `smepmp`.

18.1.30 RISC-V Base Cache Management Operation ISA Extensions

The RISC-V model does not implement caches, but supports the instructions specified in this extension. The following table shows how to configure the model for each sub-extension listed here. See section 5.8 for information about other parameters that affect these extensions.

| | |
|--------|--|
| Zicbom | Set parameter <code>zicbom</code> to T |
| Zicbop | Set parameter <code>zicbop</code> to T |
| Zicboz | Set parameter <code>zicboz</code> to T |

18.1.31 RISC-V Bit-Manipulation ISA-extensions

For all listed extensions, the B extension must be enabled by either setting parameter `add_Extensions` to include B, or instantiating a base variant with B extension (e.g. RV32GCB). The following table shows how to further configure the model for each sub-extension listed here.

| | |
|-----|-------------------------------------|
| Zba | Set parameter <code>zba</code> to T |
| Zbb | Set parameter <code>zbb</code> to T |
| Zbc | Set parameter <code>zbc</code> to T |
| Zbs | Set parameter <code>zbs</code> to T |

18.1.32 **RISC-V Count Overflow and Mode-Based Filtering Extension**

Enabled using parameter `Sscofpmf`. Note that performance monitors are implementation-defined, so this model implements only the CSR state defined by the `Sscofpmf` extension and not the counters themselves (all counters are hard-wired to zero).

18.1.33 **RISC-V Cryptography Extensions Volume I: Scalar & Entropy Source Instructions**

For all listed extensions, the K extension must be enabled by either setting parameter `add_Extensions` to include `K`, or instantiating a base variant with K extension (e.g. `RV32GCK`). The following table shows how to further configure the model for each sub-extension listed here.

| | |
|-------|---|
| Zbkb | Set parameter <code>Zbkb</code> to T |
| Zbkc | Set parameter <code>Zbkc</code> to T |
| Zbkx | Set parameter <code>Zbkx</code> to T |
| Zknd | Set parameter <code>Zknd</code> to T |
| Zkne | Set parameter <code>Zkne</code> to T |
| Zknh | Set parameter <code>Zknh</code> to T |
| Zksed | Set parameter <code>Zksed</code> to T |
| Zksh | Set parameter <code>Zksh</code> to T |
| Zkn | (Collective name for <code>Zbk</code> -prefixed and <code>Zkn</code> -prefixed extensions above) |
| Zks | (Collective name for <code>Zbk</code> -prefixed and <code>Zks</code> -prefixed extensions above) |
| Zkt | <i>Data-independent execution latency, not relevant for this model</i> |
| Zkr | Set parameter <code>Zkr</code> to T |
| Zk | (Collective name for <code>Zkn</code> , <code>Zkr</code> , and <code>Zkt</code> extensions above) |

18.1.34 **RISC-V State Enable Extension**

Enabled using parameter `Smstateen`.

18.1.35 **RISC-V "stimecmp / vstimecmp" Extension**

Enabled using parameter `Sstc`.

18.1.36 **RISC-V Vector Extension**

For all listed extensions, the V extension must be enabled by either setting parameter `add_Extensions` to include `V`, or instantiating a base variant with V extension (e.g. `RV32GCV`). The following table shows how to further configure the model for each sub-extension listed here.

| | |
|--------|--|
| Zve32x | Set parameter <code>Zve32x</code> to T |
| Zve32f | Set parameter <code>Zve32f</code> to T |
| Zve64x | Set parameter <code>Zve64x</code> to T |
| Zve64f | Set parameter <code>Zve64f</code> to T |
| Zve64d | Set parameter <code>Zve64d</code> to T |

| | |
|----------|---|
| Zve | (Family name for zve-prefixed extensions above) |
| Zvl32b | Set parameter VLEN=32. |
| Zvl64b | Set parameter VLEN=64. |
| Zvl128b | Set parameter VLEN=128. |
| Zvl256b | Set parameter VLEN=256. |
| Zvl512b | Set parameter VLEN=512. |
| Zvl1024b | Set parameter VLEN=1024. |
| Zvl | (Family name for zvl-prefixed extensions above) |
| Zv | Instantiate either RV32GCV or RV64GCV variants |

18.1.37 *The RISC-V Instruction Set Manual Volume II: Privileged Architecture*

The following table shows how to configure the model for each sub-extension listed here.

| | |
|------------|--|
| Smlp12 | Set parameter <code>priv_version</code> to 1.12 |
| Sslp12 | Set parameter <code>priv_version</code> to 1.12 |
| Sv57 | Set parameter <code>Sv_modes</code> to include bit 10 (0x400). Other bits in this mask can be set to specify Sv32, Sv39, Sv48 and Sv64. |
| Hypervisor | Set parameter <code>add_Extensions</code> to include H, or instantiate base variant with H extension (e.g. RV64GCH). |
| Svinval | Set parameter <code>Svinval</code> to T |
| Svnapot | Set parameter <code>Svnapot_page_mask</code> to indicate implemented intermediate page sizes (e.g. 1<<16 means 64KiB contiguous regions are supported) |
| Svpbmt | Set parameter <code>Svpbmt</code> to T |

18.1.38 *"Zfh" and "Zfhmin" Standard Extensions for Half-Precision Floating-Point*

The following table shows how to configure the model for each sub-extension listed here.

| | |
|--------|--|
| Zfh | Set parameter <code>add_Extensions</code> to include F and set parameter <code>zfh</code> to T. |
| Zfhmin | Set parameter <code>add_Extensions</code> to include F and set parameter <code>zfhmin</code> to T. |

18.1.39 *"Zfinx", "Zdinx", "Zhinx", "Zhinxmin": Standard Extensions for Floating-Point in Integer Registers*

The following table shows how to configure the model for each sub-extension listed here.

| | |
|-------|--|
| Zfinx | Set parameter <code>add_Extensions</code> to include F and set parameter <code>zfinx_version</code> to 1.0. |
| Zdinx | Set parameter <code>add_Extensions</code> to include F and D, and set parameter <code>zfinx_version</code> to 1.0. |
| Zhinx | Set parameter <code>add_Extensions</code> to include F, set parameter |

| | |
|----------|--|
| | zfinx_version to 1.0 and set parameter zfh to T. |
| Zhinxmin | Set parameter add_Extensions to include F, set parameter zfinx_version to 1.0 and set parameter zfhmin to T. |

18.1.40 “Zhintpause” Pause Hint

The PAUSE hint is always treated as a NOP by the model, so there is no option to explicitly enable this extension.

18.1.41 “Zicsr”, Control and Status Register (CSR) Instructions

Enabled using parameter zicsr. Note that disabling zicsr is not useful unless the model is extended to support a custom replacement for standard CSRs.

18.1.42 “Zifencei” Instruction-Fetch Fence

Enabled using parameter zifencei.

18.2 RISC-V Profile Extensions

The RISC-V Profile definition ([riscv-profiles/profiles.adoc at main · riscv/riscv-profiles · GitHub](https://github.com/riscv/riscv-profiles/blob/main/riscv-profiles.adoc)) gives new names for a number of existing RISC-V features. The following table shows how to configure the model for each extension listed there.

| | |
|--------------|--|
| Ziccif | <i>Always enabled</i> |
| Ziccrse | <i>Always enabled</i> |
| Ziccamoa | <i>Always enabled</i> |
| Zicclsm | Set parameter <code>unaligned</code> to T |
| Za64rs | Set parameter <code>lr_sc_grain</code> to 64 |
| Za128rs | Set parameter <code>lr_sc_grain</code> to 128 |
| Zic64b | Set parameters <code>cmomp_bytes</code> and <code>cmoz_bytes</code> to 64 |
| Svbare | Set parameter <code>Sv_modes</code> to include bit 0 (0x1). |
| Svade | Set parameters <code>updatePTEA</code> and <code>updatePTED</code> to F |
| Sccptr | <i>Always enabled</i> |
| Sscounterenw | <i>Not applicable: base model does not implement HPM counters</i> |
| Sstvecd | <i>Always enabled</i> |
| Sstvala | Set parameter <code>tval_zero</code> to F and parameter <code>tval_ii_code</code> to T |
| Ssu64xl | Run variant with <code>XLEN=64</code> |
| Ssstateen | Set parameter <code>Smstateen</code> to T |
| Shcounterenw | <i>Not applicable: base model does not implement HPM counters</i> |
| Shvstvala | Set parameter <code>tval_zero</code> to F and parameter <code>tval_ii_code</code> to T |
| Shtvala | Set parameter <code>tval_zero</code> to F and parameter <code>tval_ii_code</code> to T |
| Shvsatpa | <i>Always enabled</i> |
| SvNNx4 | <i>Always enabled</i> |

18.3 Other Extensions

The following extensions are supported in the model using appropriate parameterization, but are not mentioned on the *Recently Ratified Extensions* page. Some are not yet ratified, in which case note that behavior may change prior to ratification.

| | |
|--|--|
| “Sdext” ISA Extension | Sdext |
| “Sdtrig” ISA Extension | Sdtrig |
| Core-Local Interrupt Controller (CLIC) RISC-V Privileged Architecture Extensions | smcllic, ssclic, suclic, smcllicshv, smcllicconfig |

18.3.1 “Sdext” ISA Extension

This is enabled by setting parameter `debug_mode` to a value other than `none`. See section 5.17 for details of further parameters to refine Debug mode.

18.3.2 “Sdtrig” ISA Extension

This is enabled by setting parameter `trigger_num` to a non-zero value. See section 5.18 for details of further parameters to refine the trigger module.

18.3.3 Core-Local Interrupt Controller (CLIC) RISC-V Privileged Architecture Extensions

The CLIC extension can be configured to either use an internal model or an external CLIC model – see section 5.5.

| | |
|---------------|--|
| smcllic | Set parameter <code>CLICLEVELS</code> to a non-zero value |
| ssclic | Set parameter <code>CLICLEVELS</code> to a non-zero value on a variant with S mode. |
| suclic | Set parameter <code>CLICLEVELS</code> to a non-zero value on a variant with U mode and N extension. |
| smcllicshv | Set parameter <code>CLICLEVELS</code> to a non-zero value and set parameter <code>CLICSELHVEC</code> to T. |
| smcllicconfig | Set parameter <code>CLICLEVELS</code> to a non-zero value |